EECS 498-008 / 598-008
Deep Learning for Computer Vision

Lecture 1: Introduction

Logistics:

- Remote for first two weeks (Lectures 1 – 3)

- After that, in-person lecture in Chrysler 220

Deep Learning for Computer Vision

Deep Learning for Computer Vision

Building artificial systems that process, perceive, and reason about visual data

Computer Vision is everywhere!




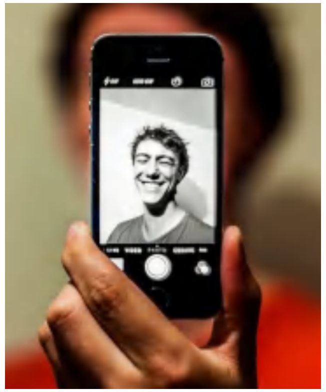

Left to right:
Image by Roger H Goun is licensed under CC BY 2.0
Image is CC0 1.0 public domain
Image is CC0 1.0 public domain
Image is CC0 1.0 public domain




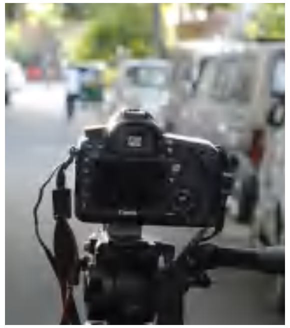

Left to right:
Image is free to use
Image is CC0 1.0 public domain
Image by NASA is licensed under CC BY 2.0
Image is CC0 1.0 public domain




Bottom row, left to right
Image is CC0 1.0 public domain
Image by Derek Keats is licensed under CC BY 2.0; changes made
Image is public domain
Image is licensed under CC-BY 2.0; changes made

Deep Learning for Computer Vision

Building artificial systems that
learn from data and experience

<u>Deep Learning</u> for Computer Vision

Hierarchical learning algorithms with many “layers”, (very) loosely inspired by the brain

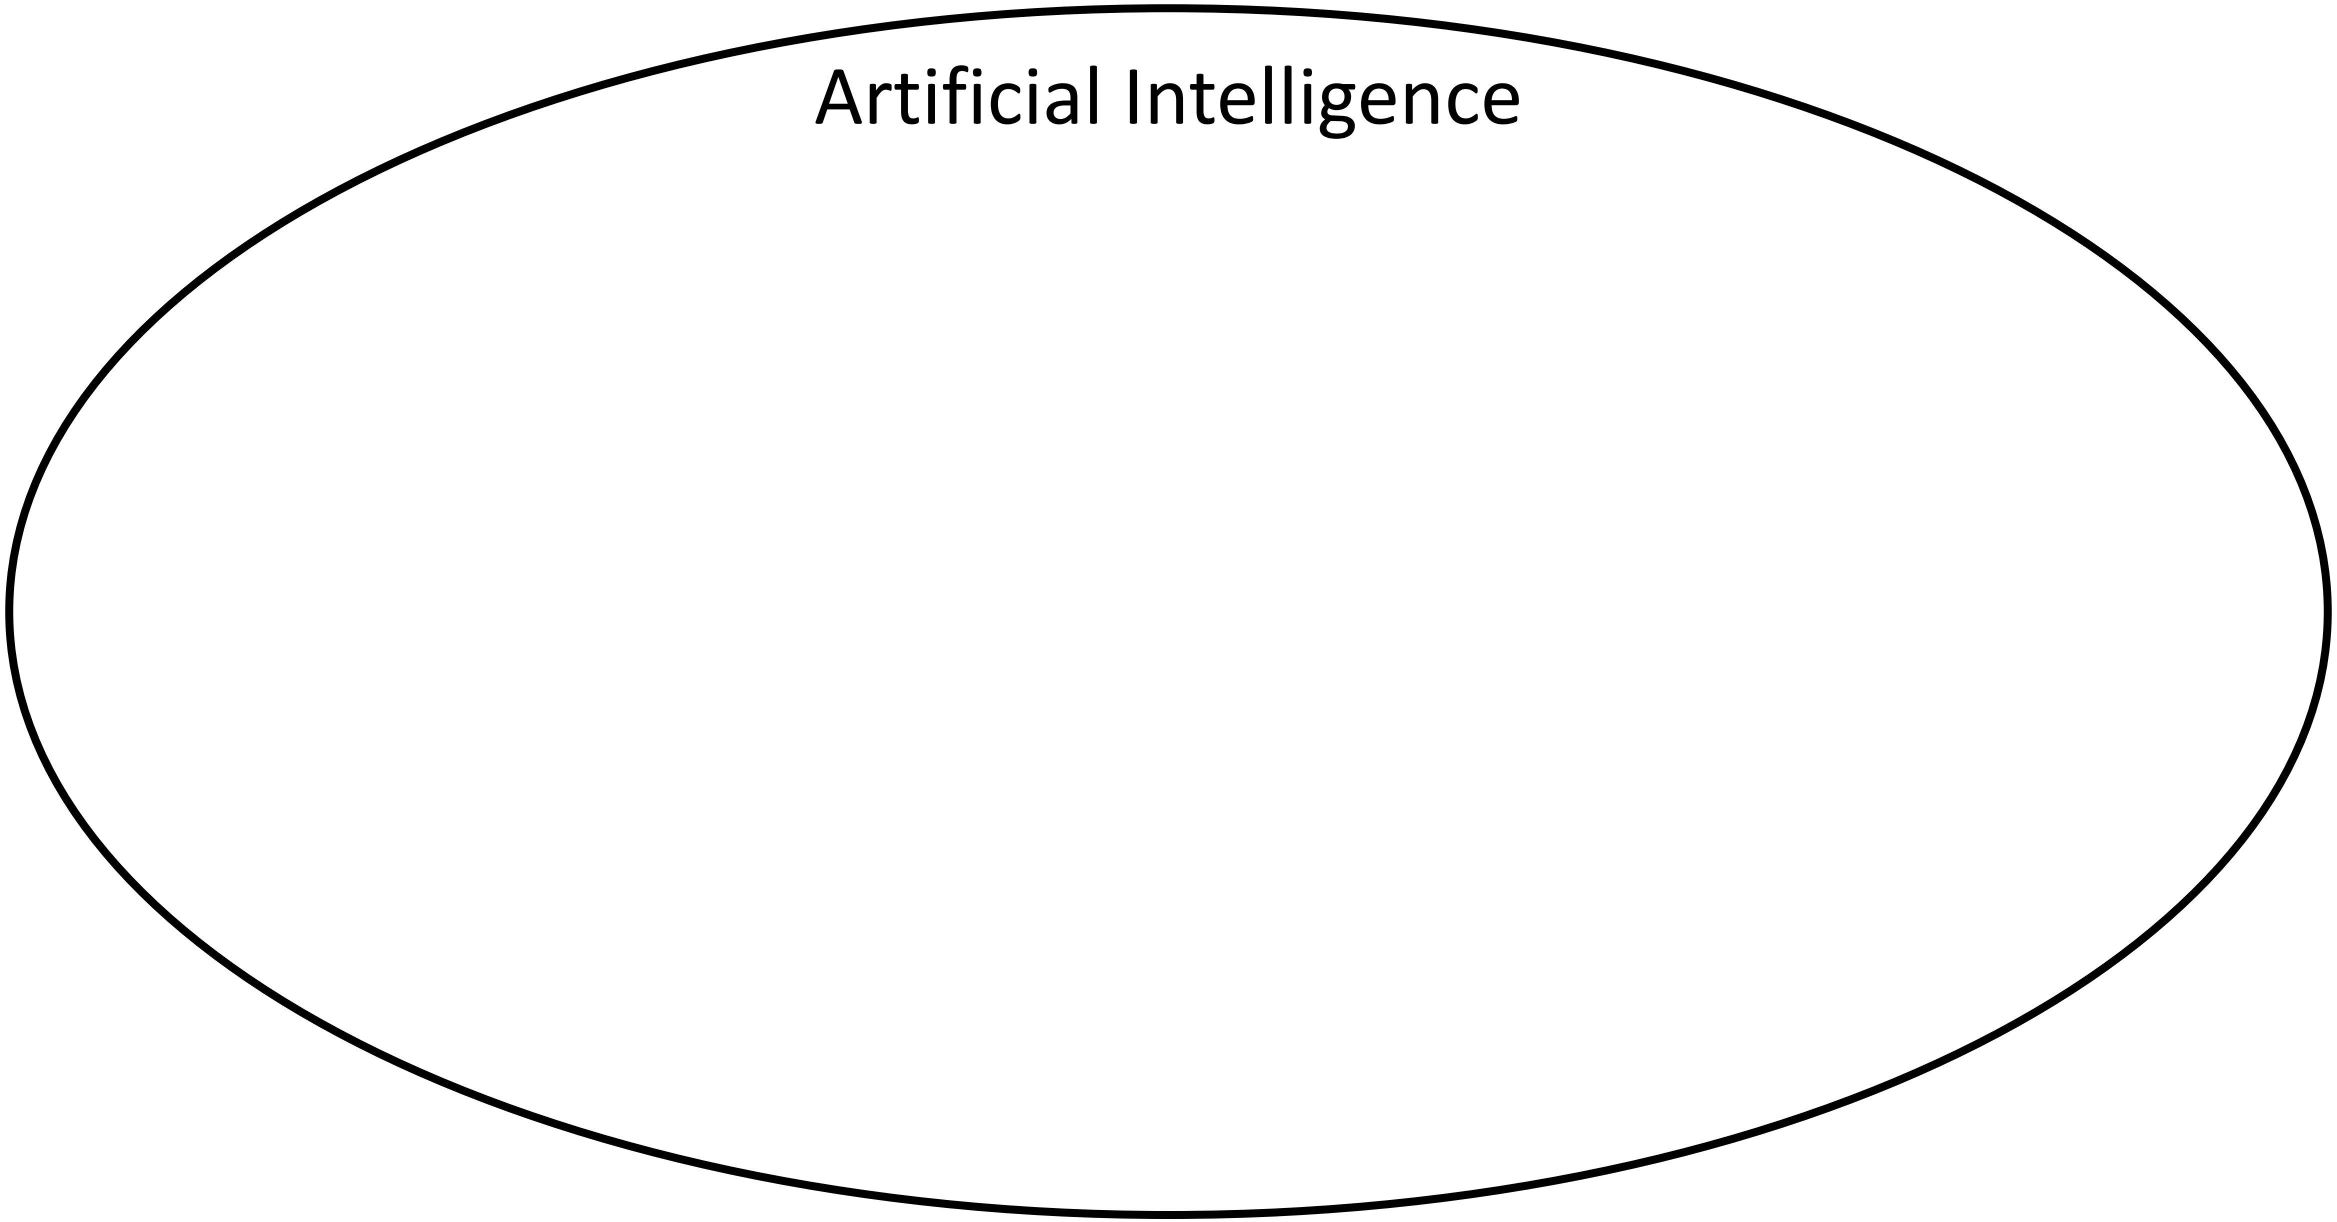
Artificial Intelligence

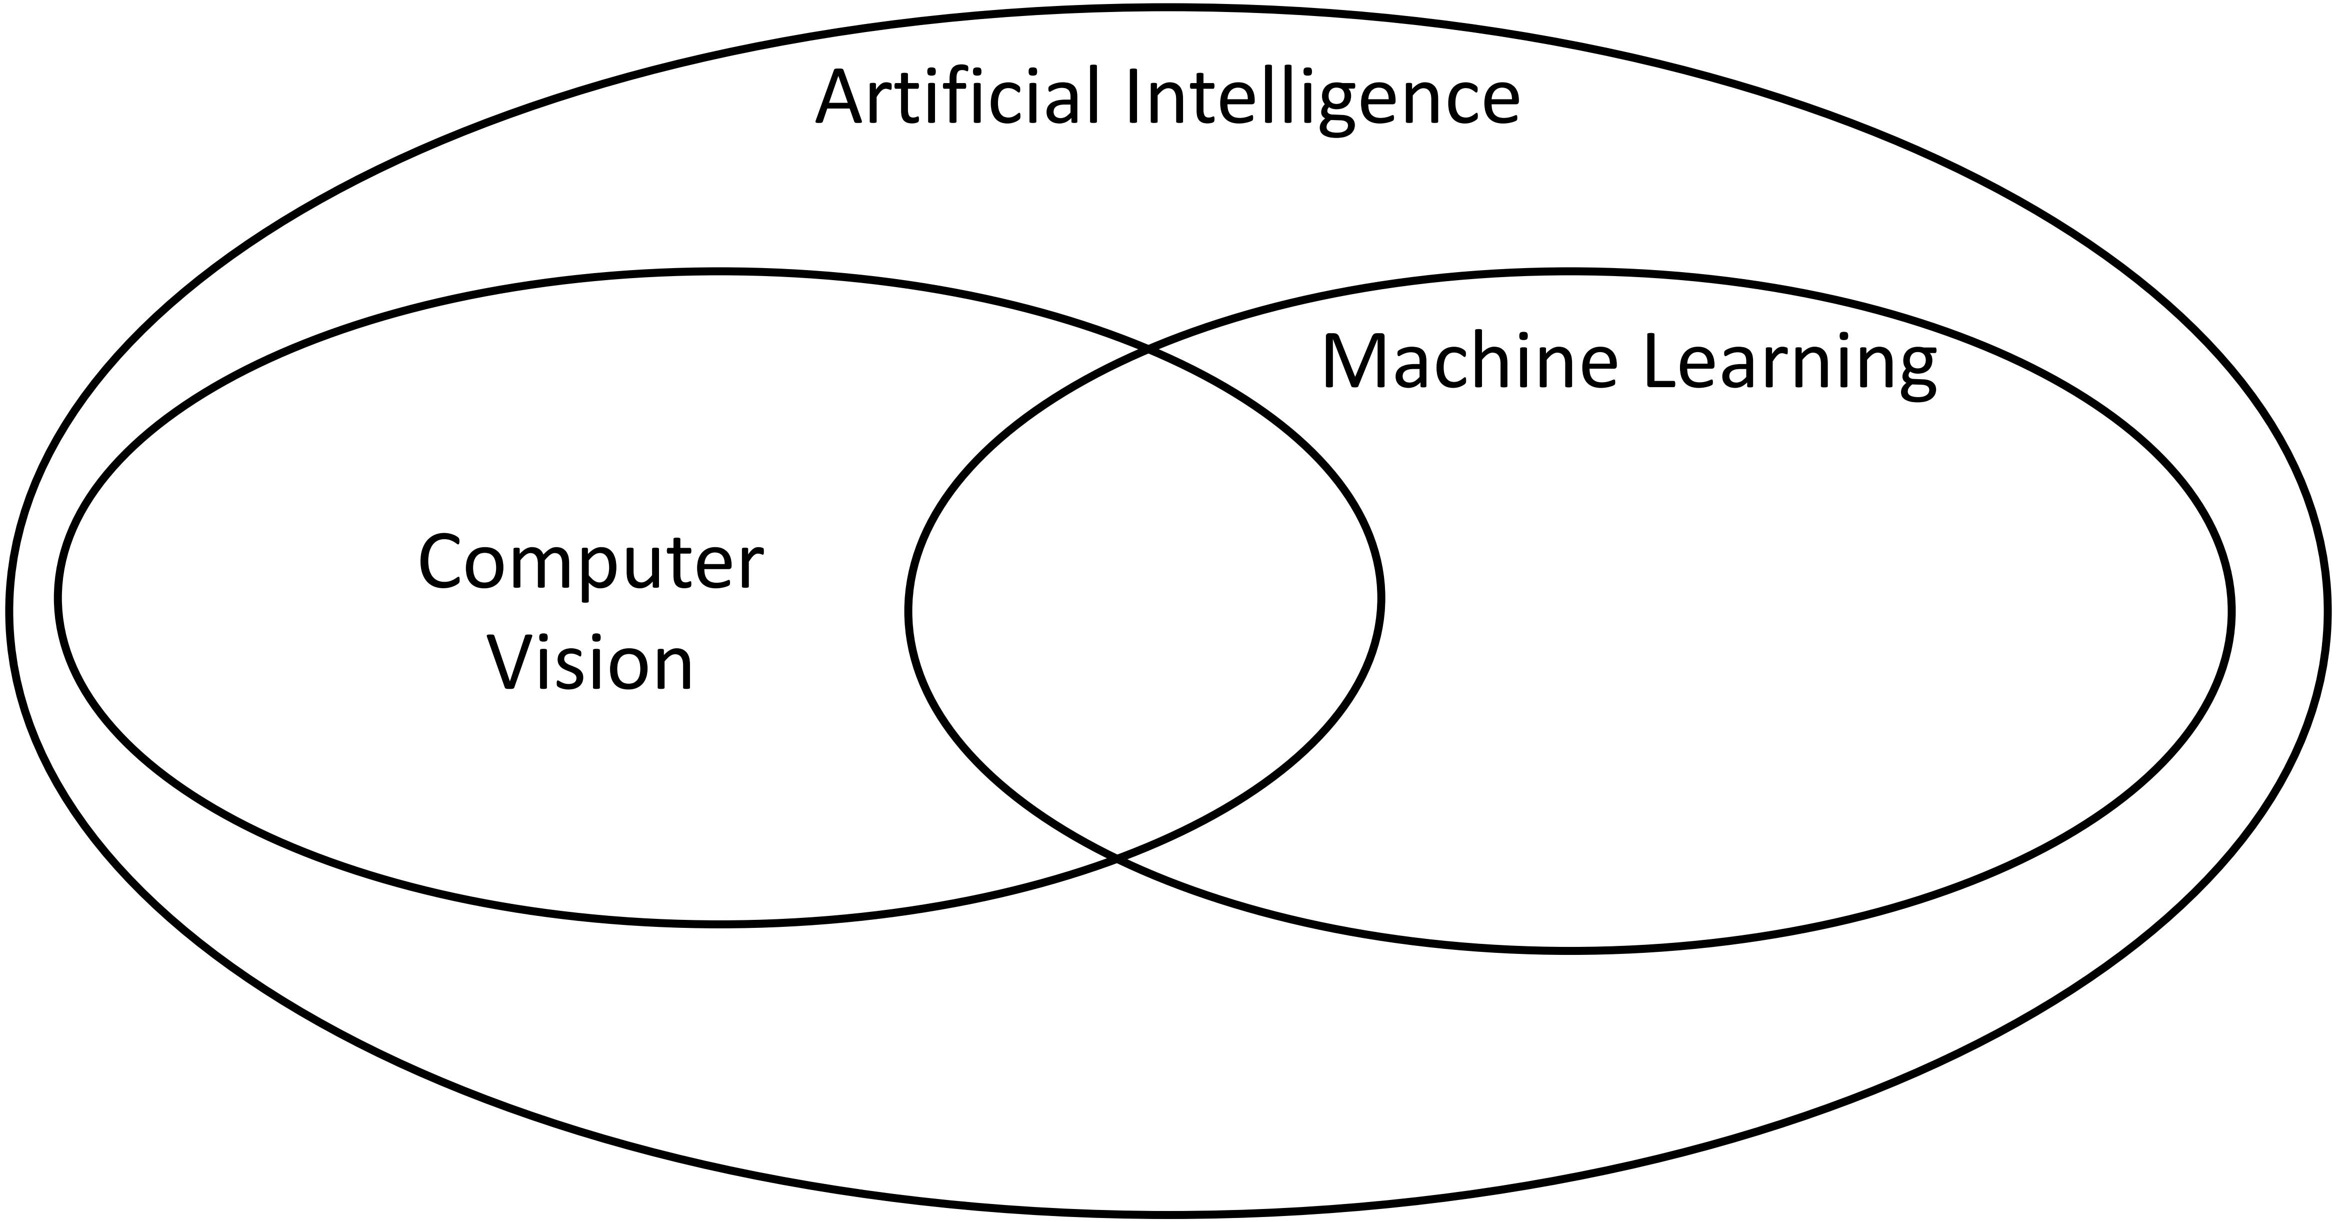
Artificial Intelligence
Machine Learning
Computer
Vision

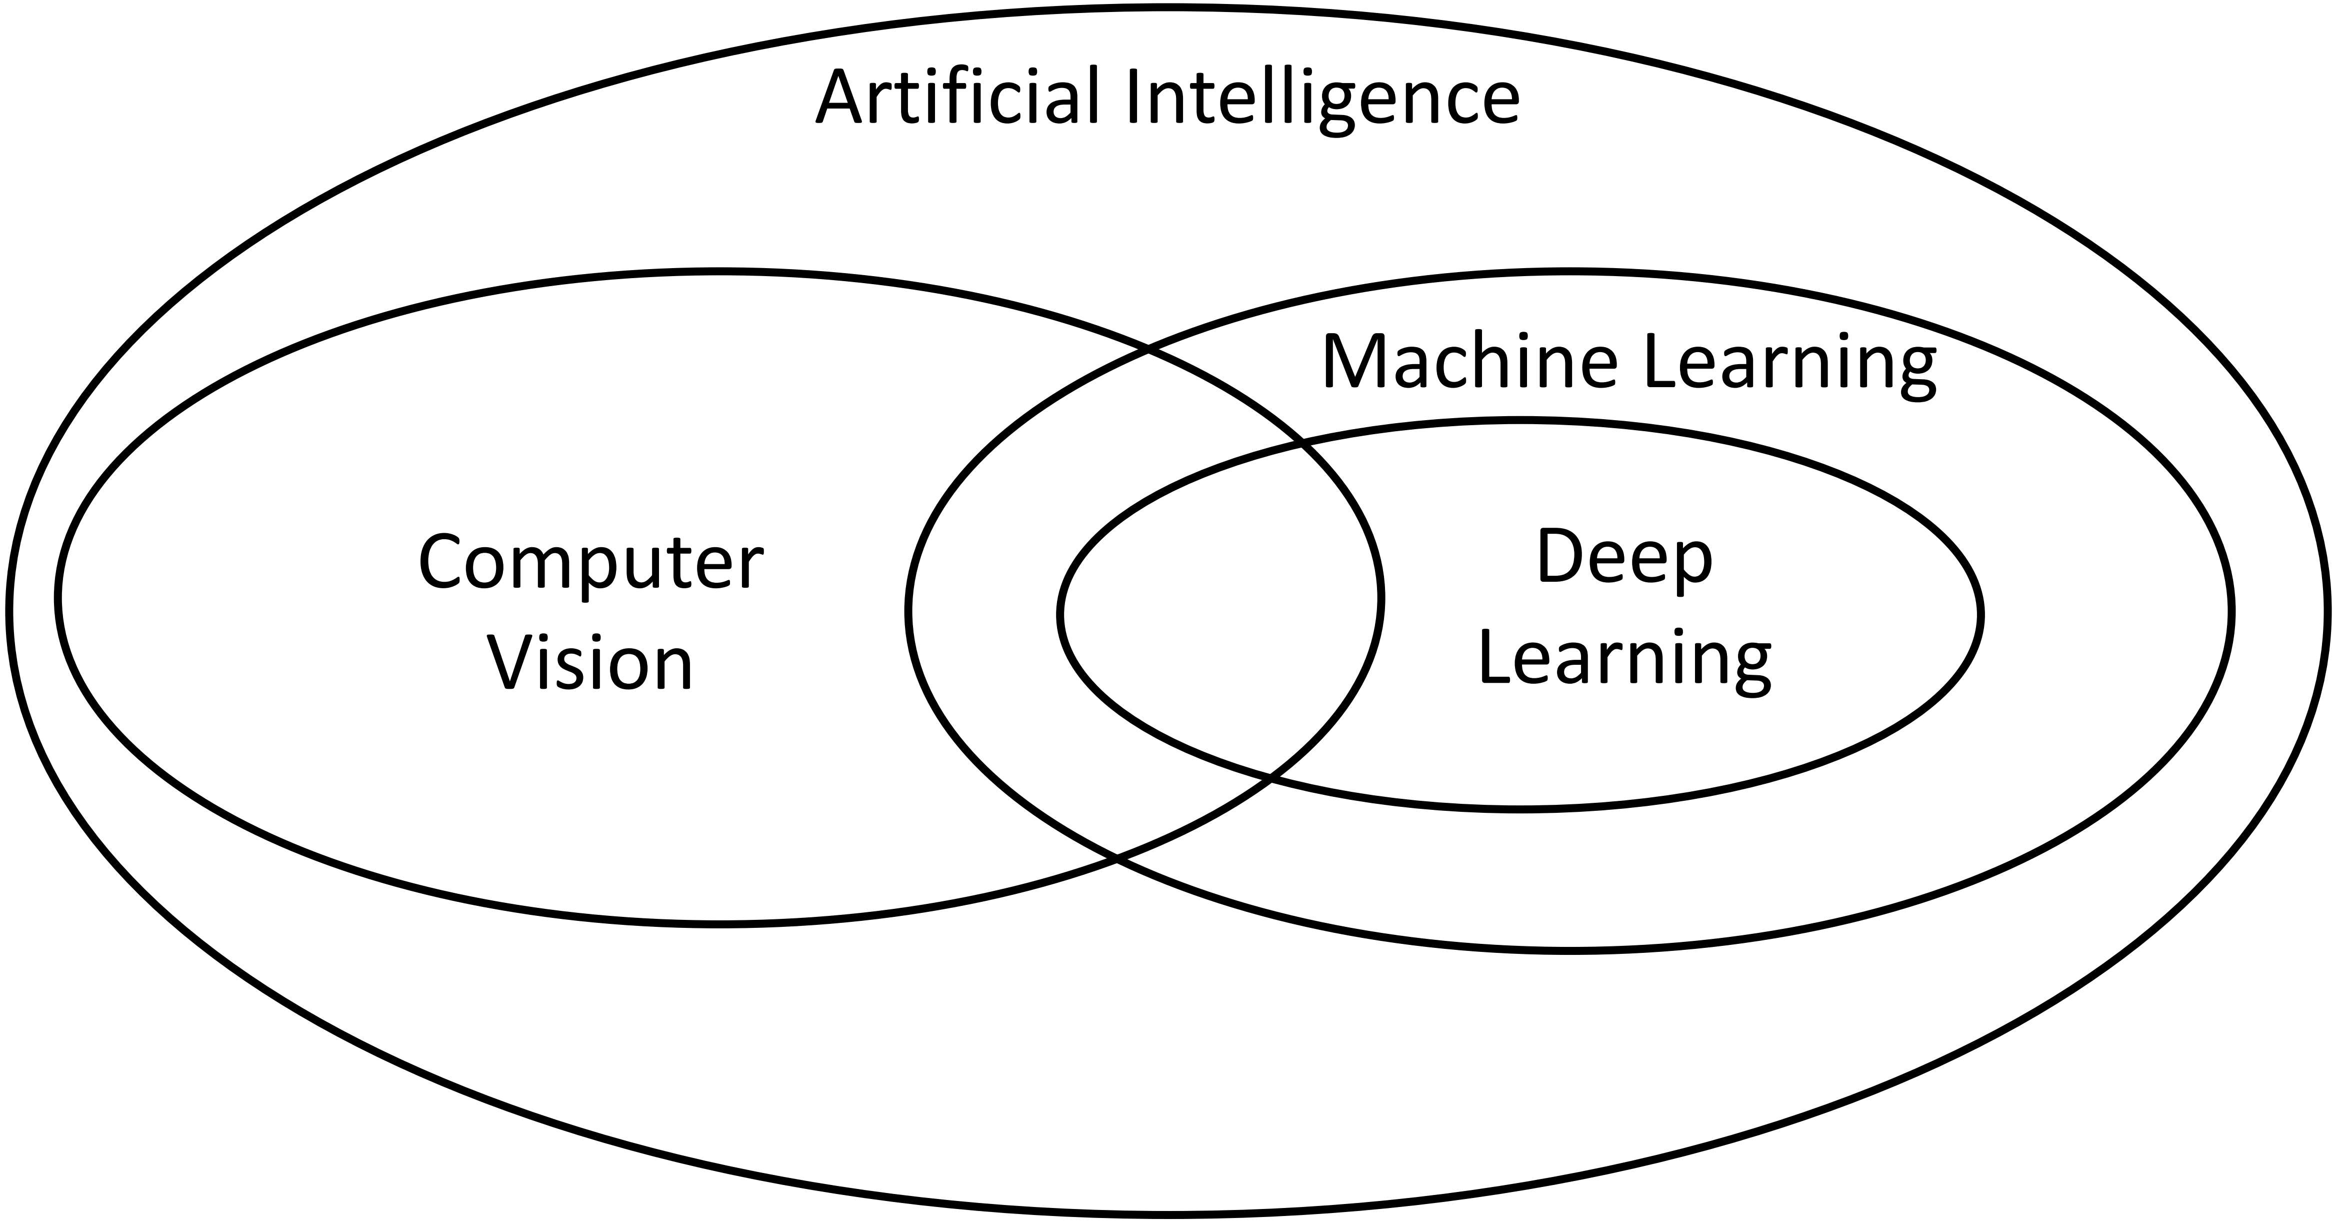
Artificial Intelligence
Machine Learning
Computer Vision
Deep Learning

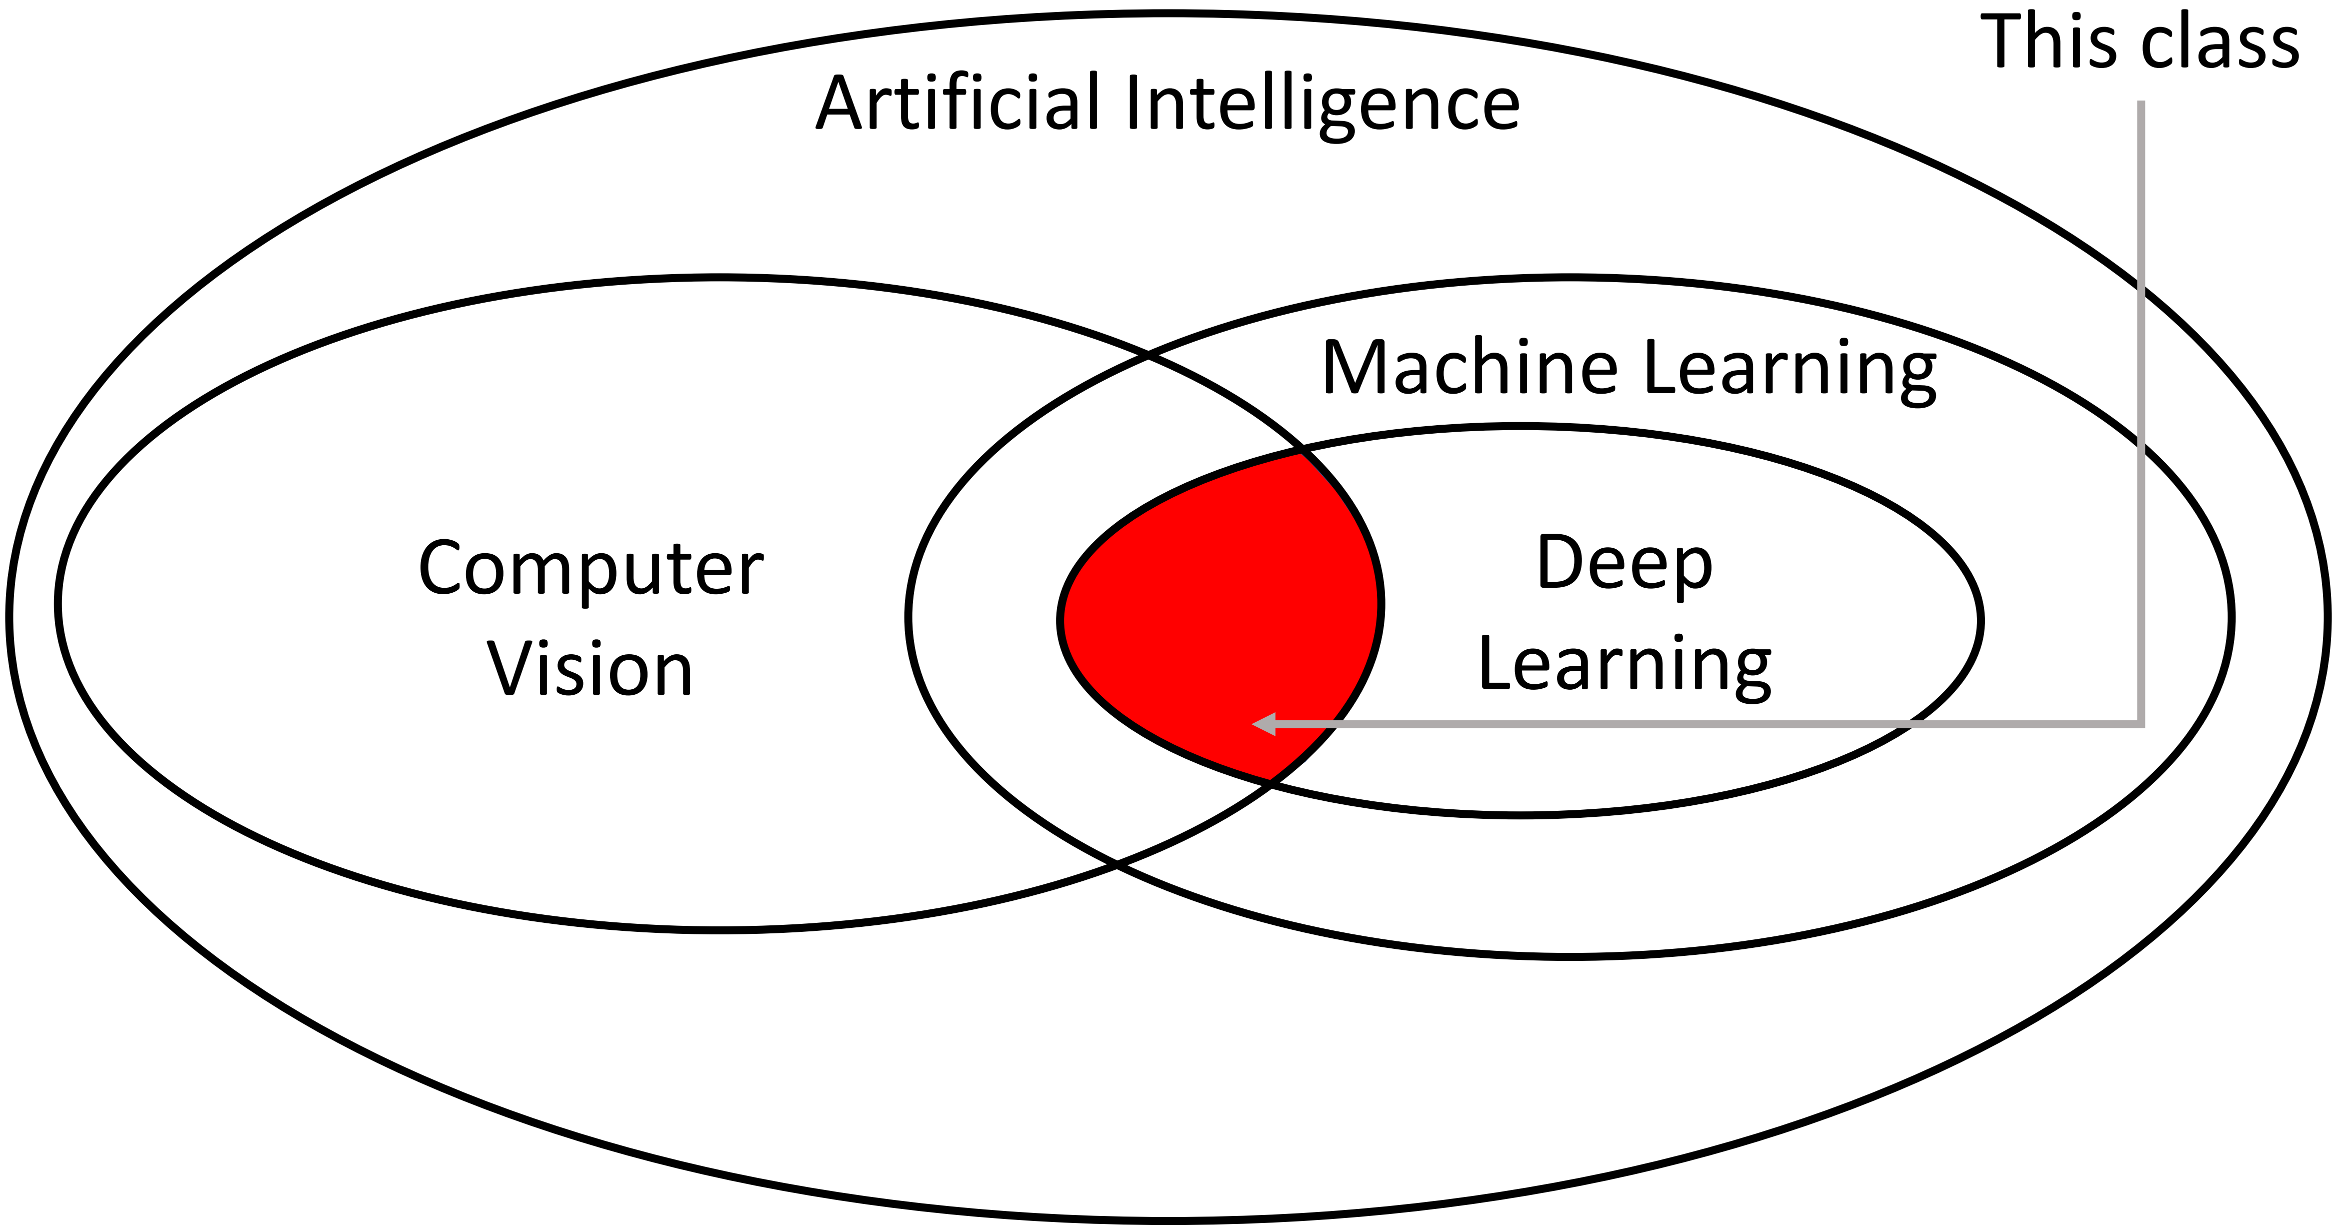
This class
Artificial Intelligence
Machine Learning
Computer
Vision
Deep
Learning

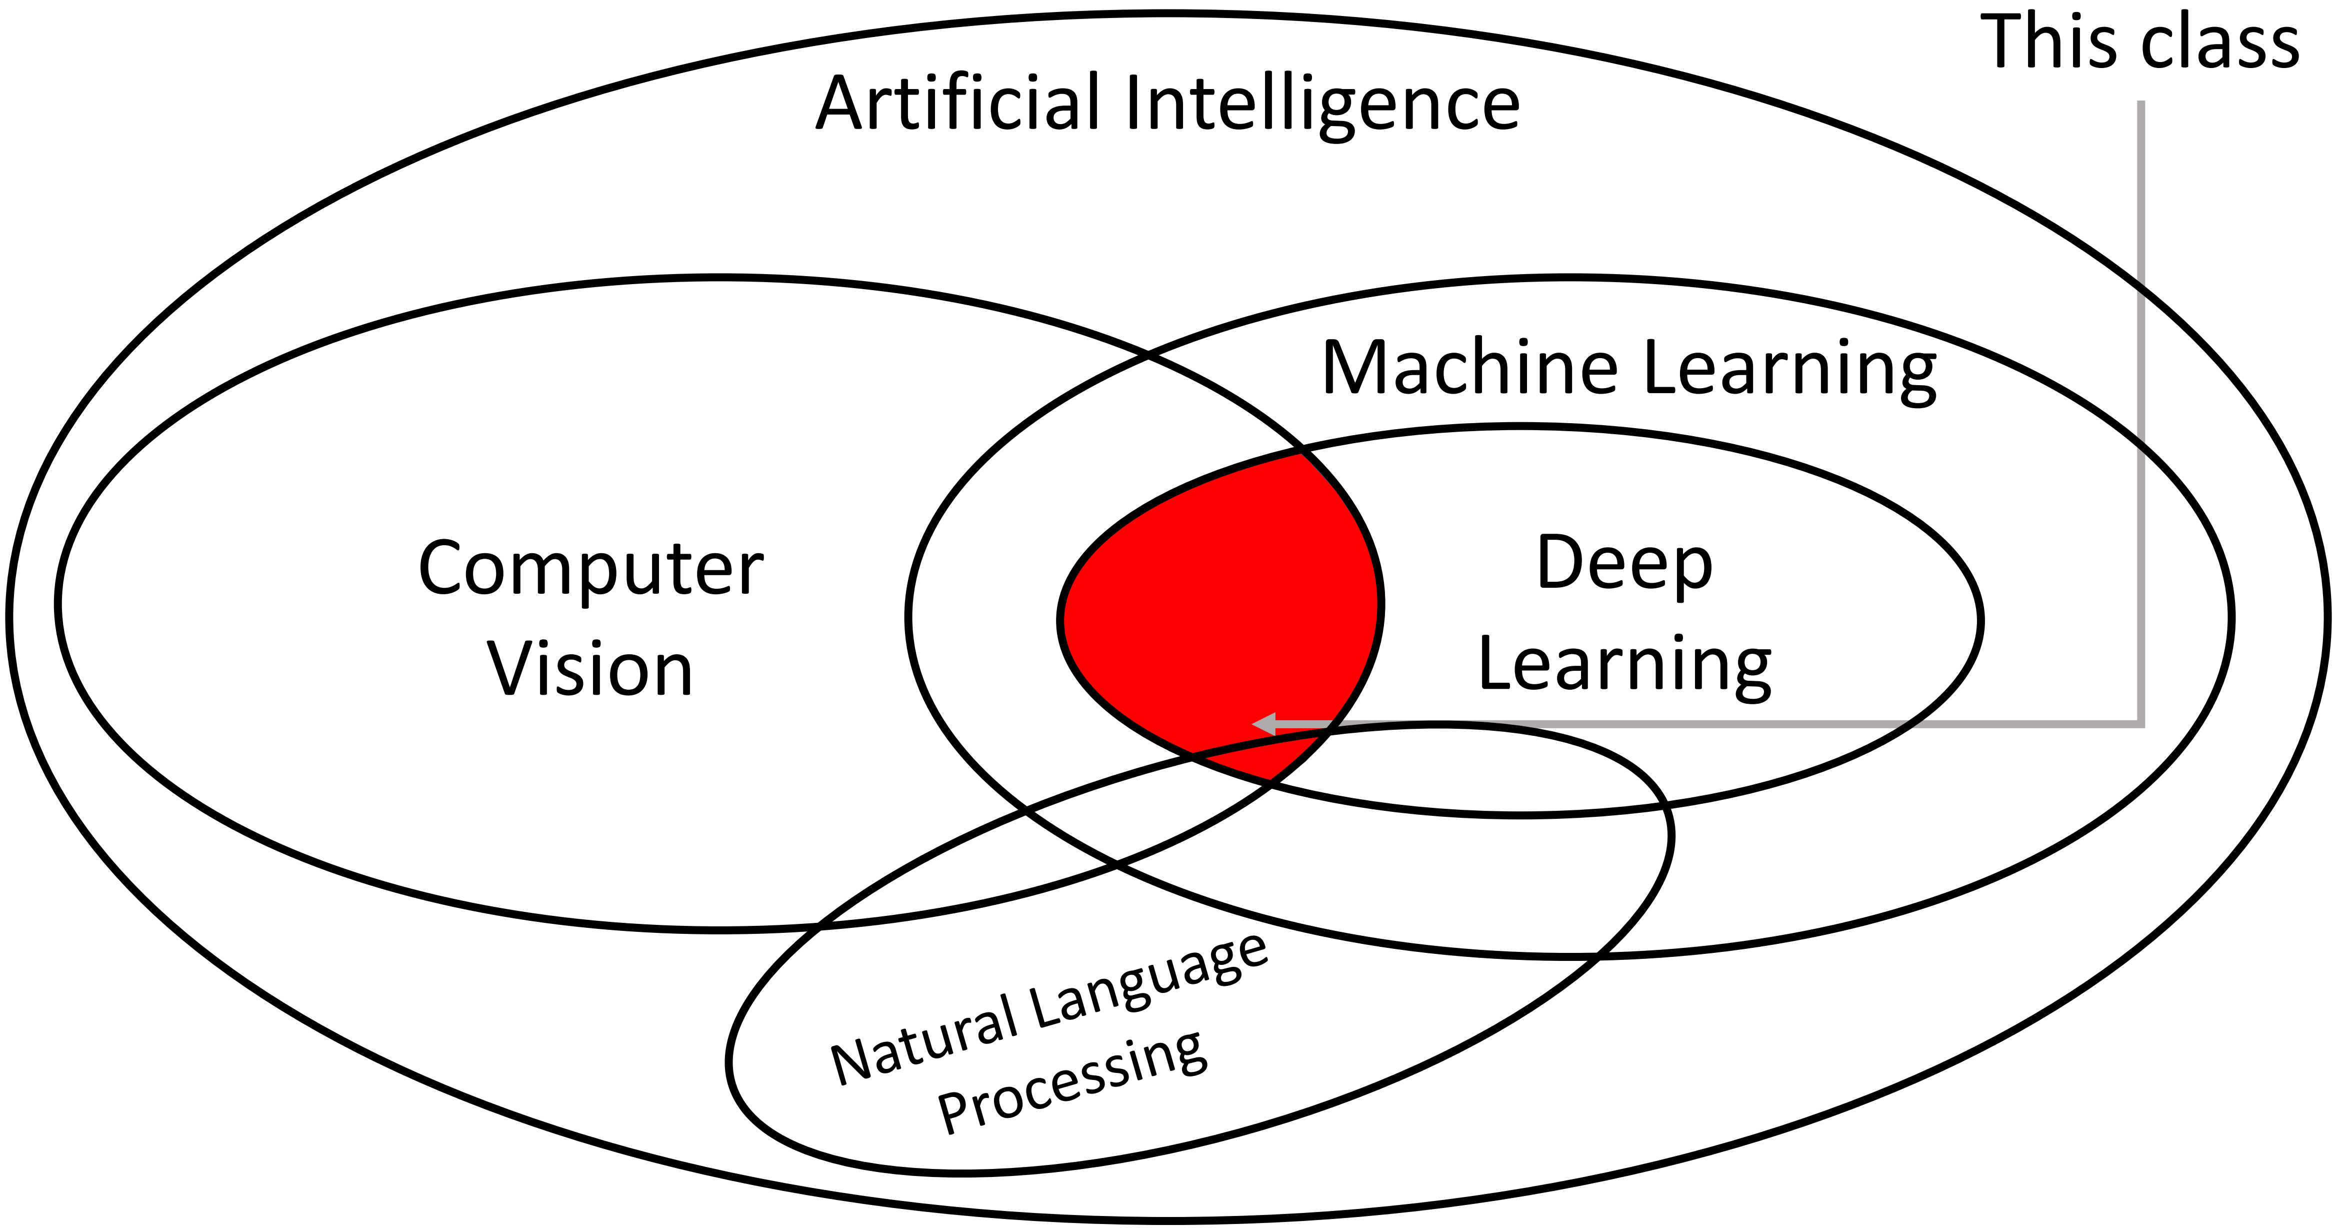
This class
Artificial Intelligence
Machine Learning
Computer Vision
Deep Learning
Natural Language Processing

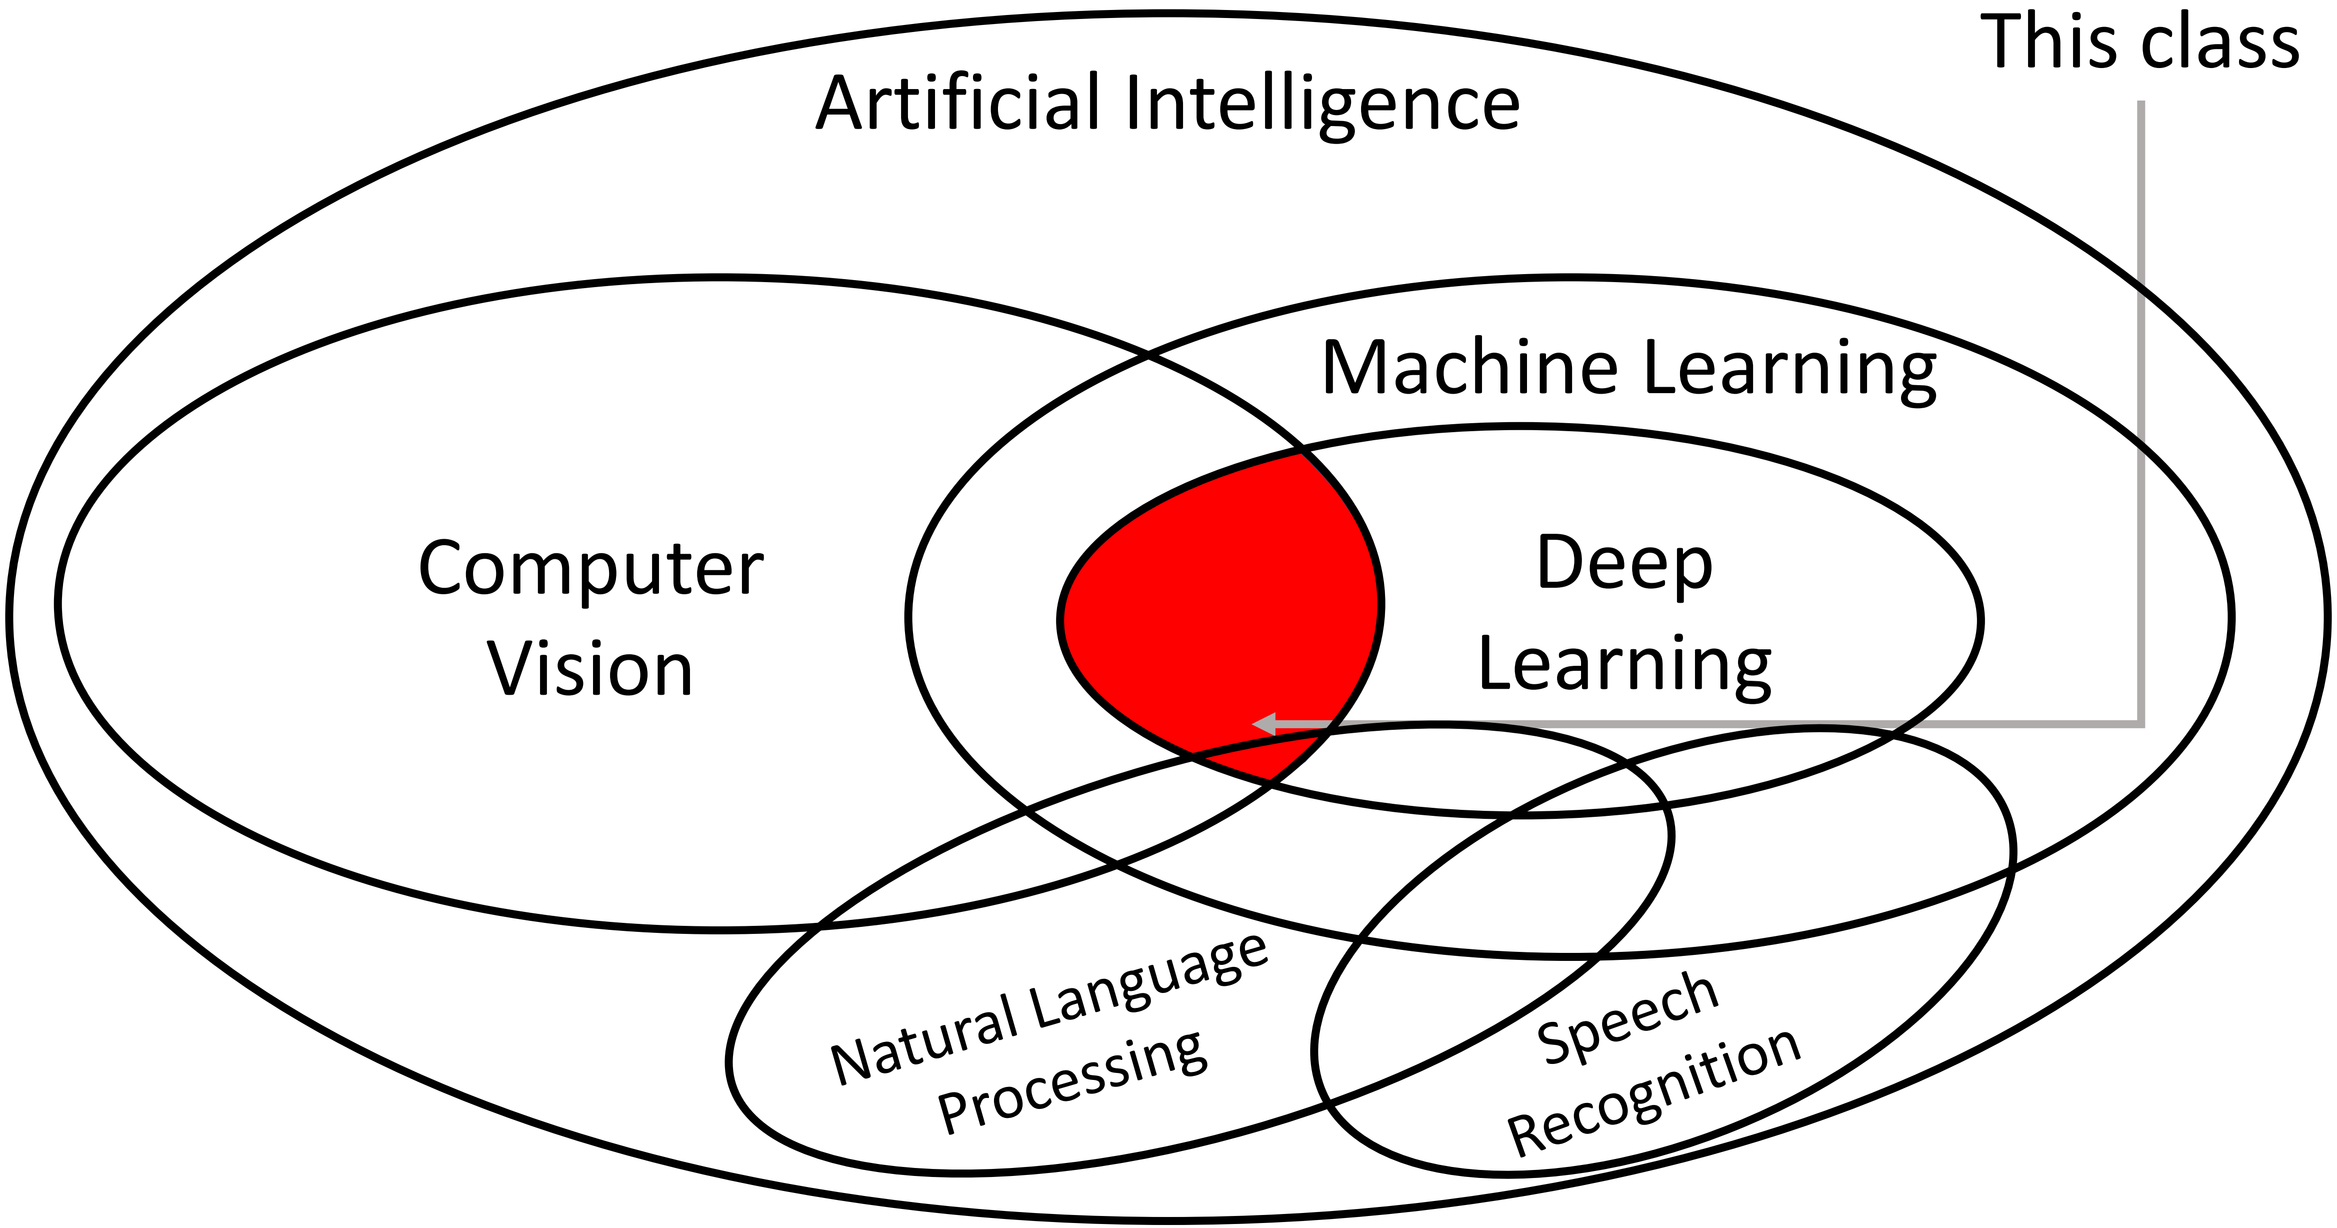
This class
Artificial Intelligence
Machine Learning
Computer Vision
Deep Learning
Natural Language Processing
Speech Recognition

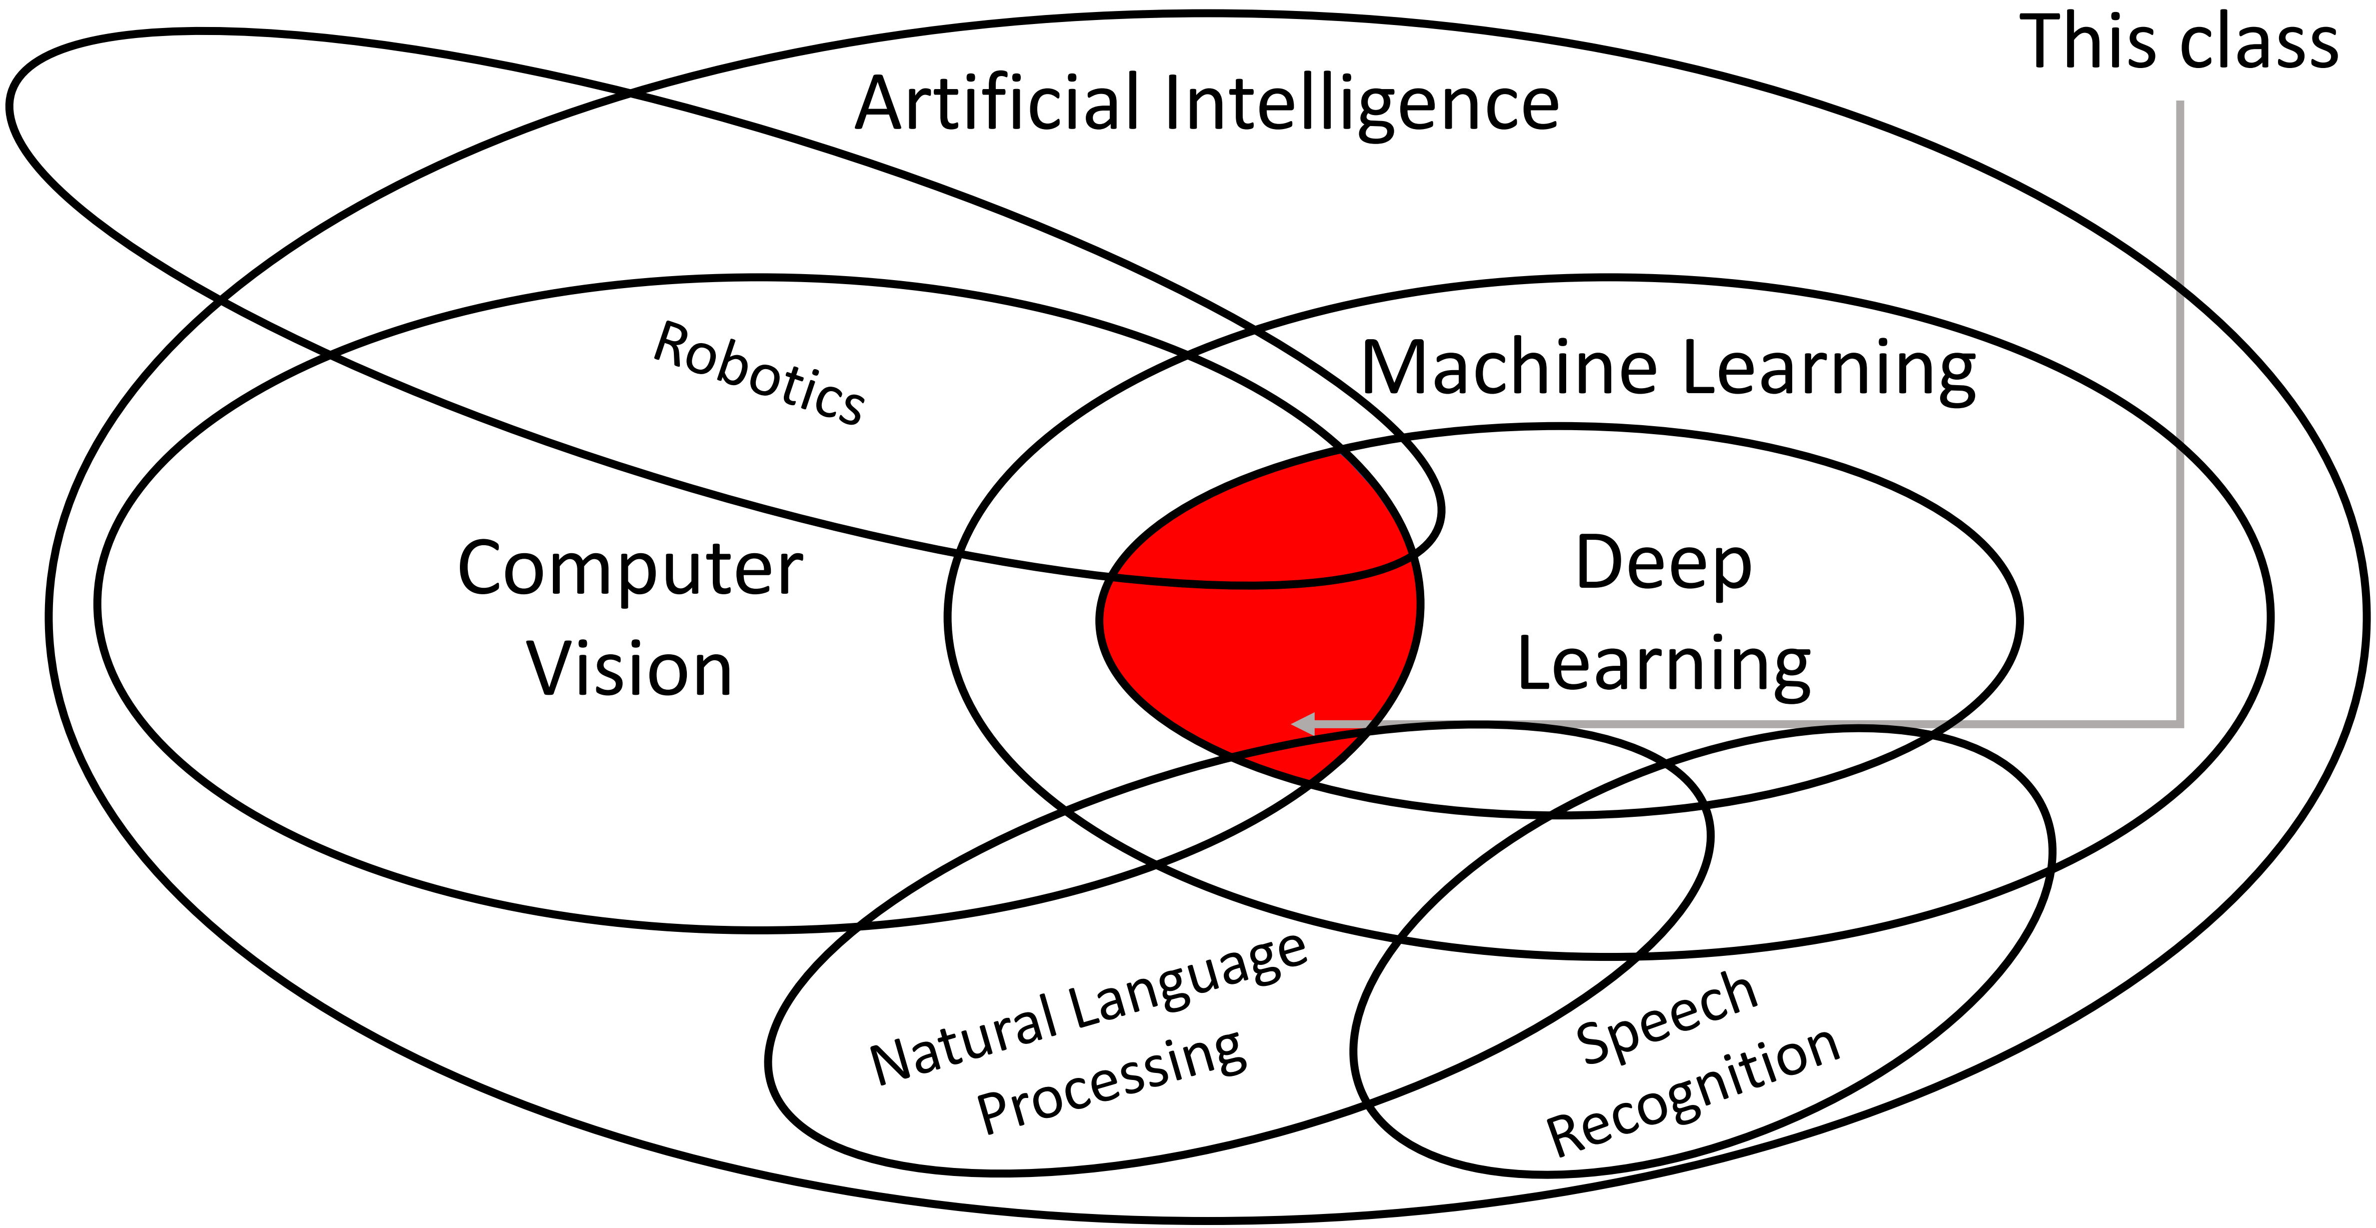
This class
Artificial Intelligence
Robotics
Machine Learning
Computer Vision
Deep Learning
Natural Language Processing
Speech Recognition

Today's Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics

Hubel and Wiesel, 1959

Measure
brain activity

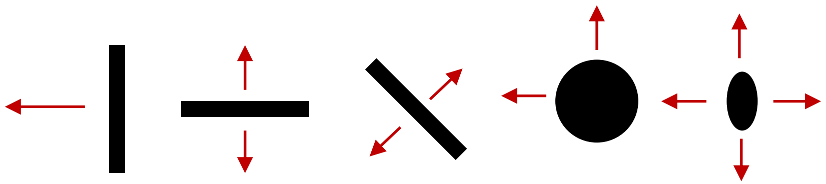

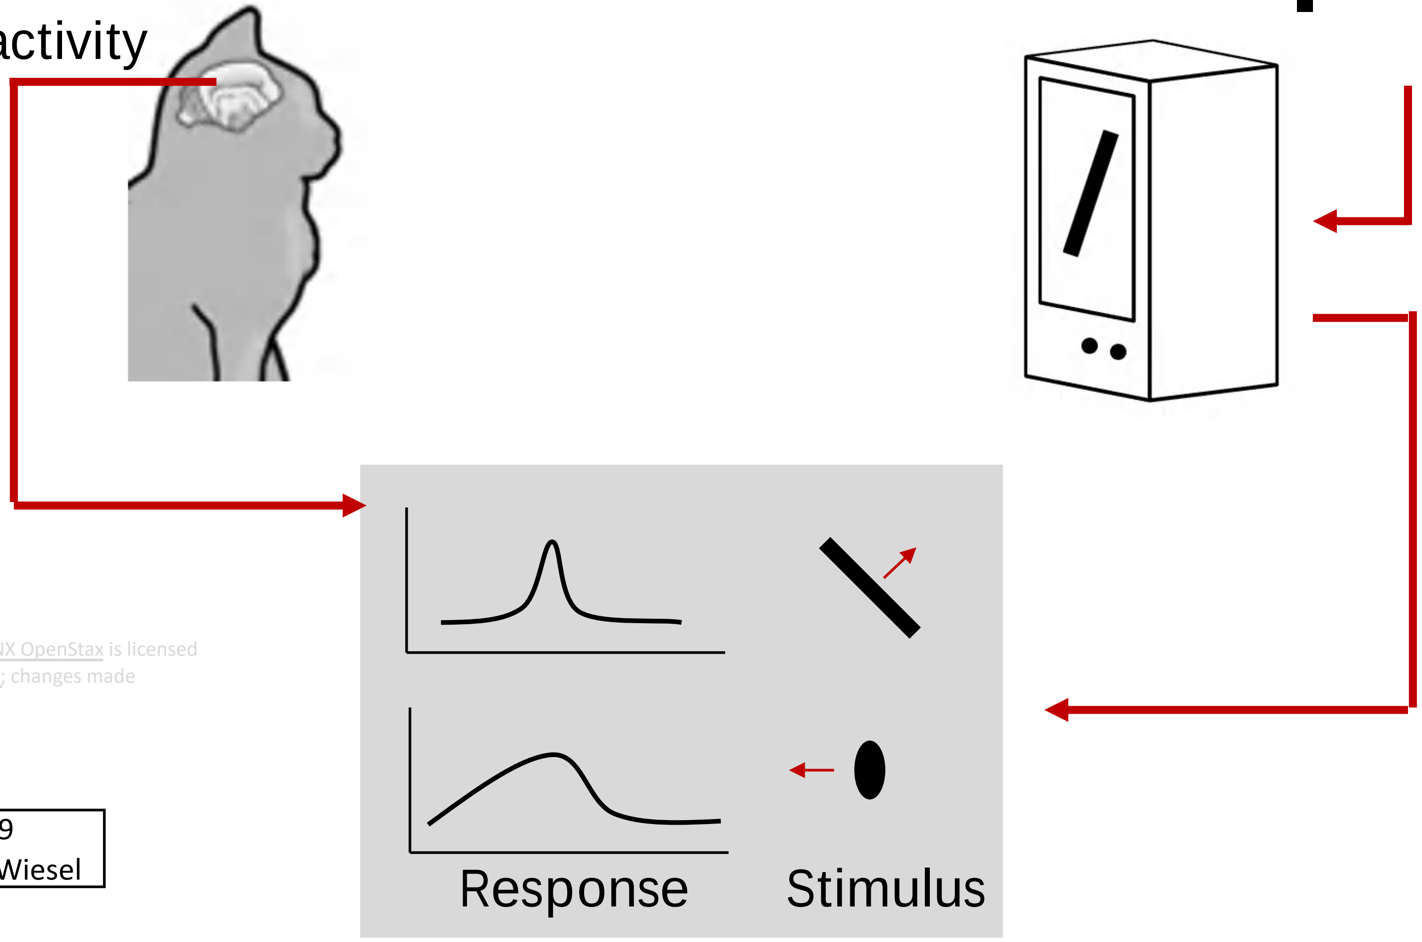

Simple cells:
Response to light
orientation

Cat image by CNX OpenStax is licensed
under CC BY 4.0; changes made

1959
Hubel & Wiesel

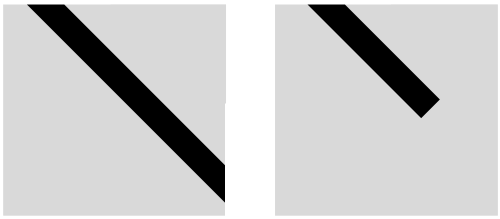

No response

Hubel and Wiesel, 1959

Measure brain activity

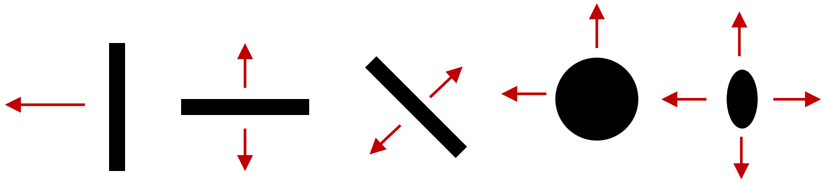

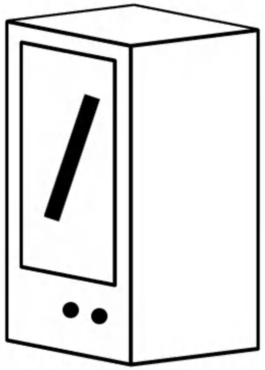

Simple cells:
Response to light orientation

Complex cells:
Response to light orientation and movement

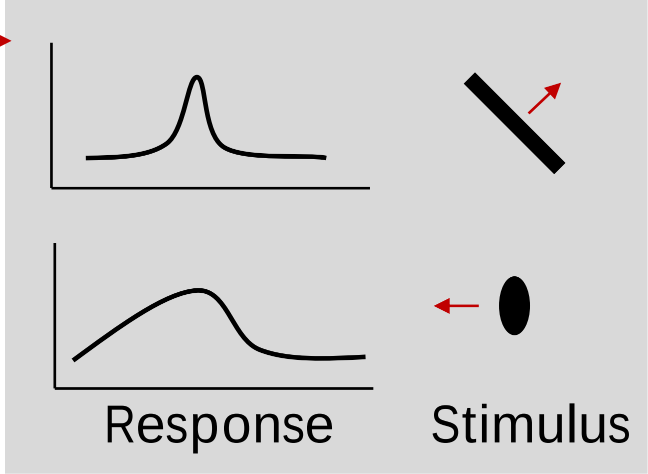

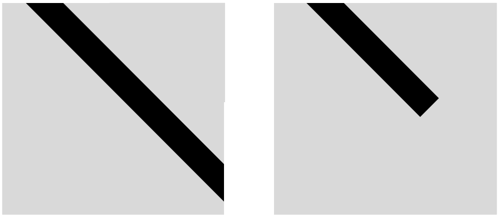

No response

Cat image by CNX OpenStax is licensed under CC BY 4.0; changes made

1959
Hubel & Wiesel

Larry Roberts, 1963

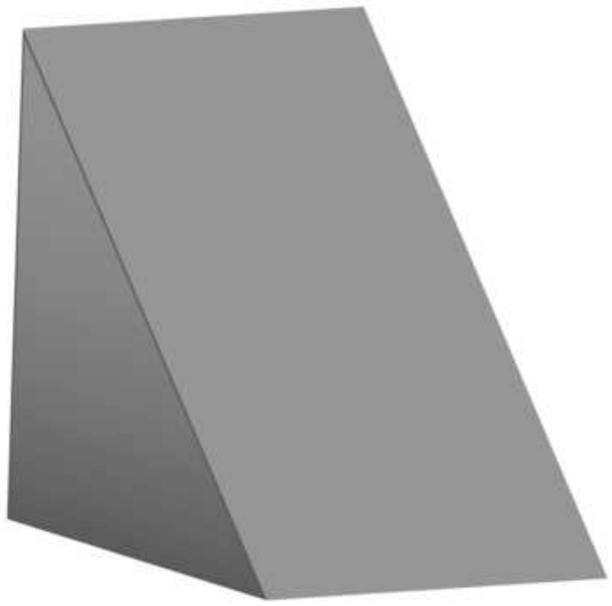

(a) Original picture

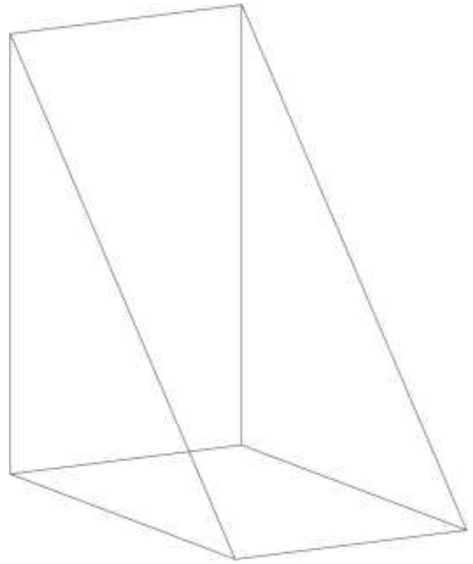

(b) Differentiated picture

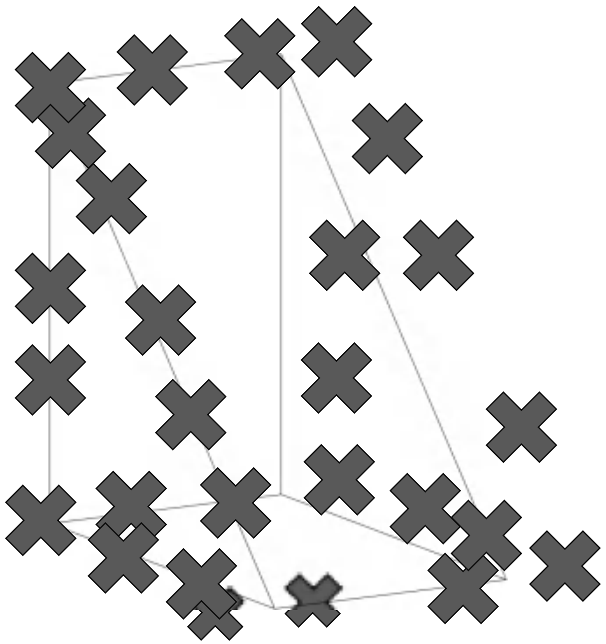

(c) Feature points selected

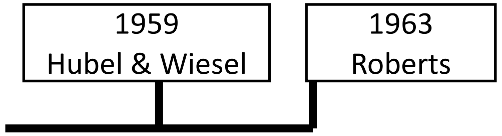

Lawrence Gilman Roberts, “Machine Perception of Three-Dimensional Solids”, 1963

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

PROJECT MAC

Artificial Intelligence Group July 7, 1966

Vision Memo. No. 100.

THE SUMMER VISION PROJECT

Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

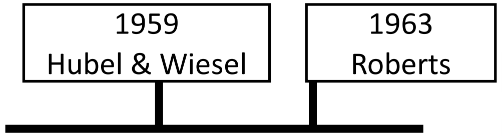

https://dspace.mit.edu/handle/1721.1/6125

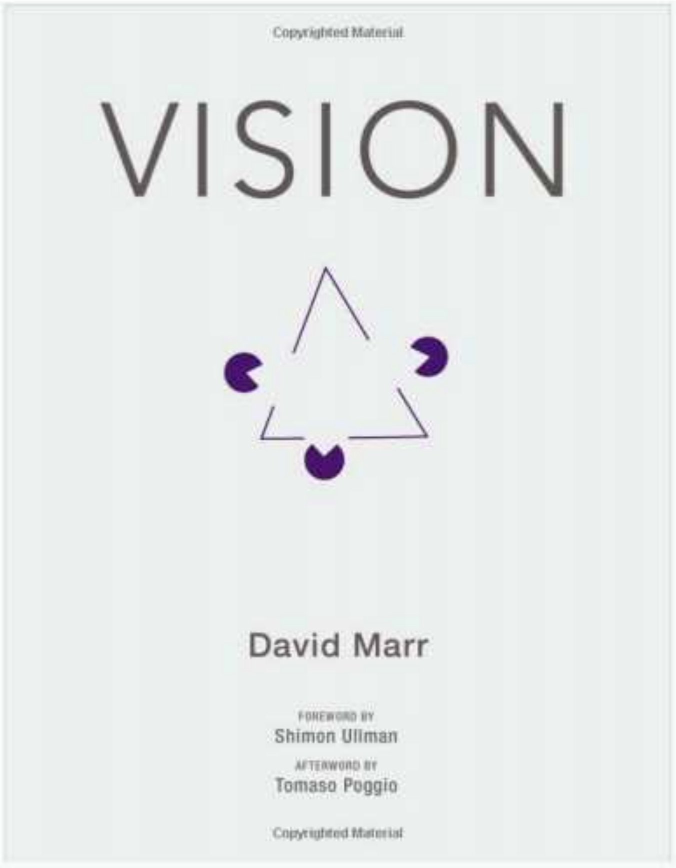

1959 Hubel & Wiesel	1963 Roberts	1970s David Marr

Input image

This image is CC0 1.0 public domain

Edge image

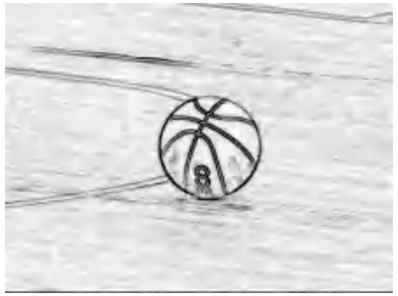

2 ½-D sketch

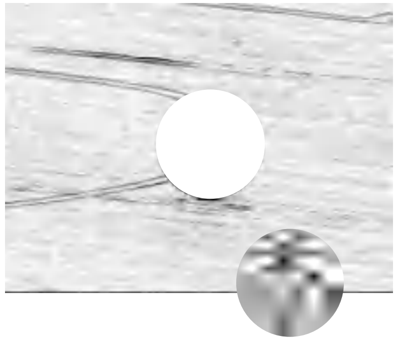

3-D model

This image is CC0 1.0 public domain

Input Image	Primal Sketch	2 ½-D Sketch	3-D Model Representation
Perceived intensities	Zero crossings, blobs, edges, bars, ends, virtual lines, groups, curves boundaries	Local surface orientation and discontinuities in depth and in surface orientation	3-D models hierarchically organized in terms of surface and volumetric primitives

Stages of Visual Representation, David Marr, 1970s

Recognition via Parts (1970s)

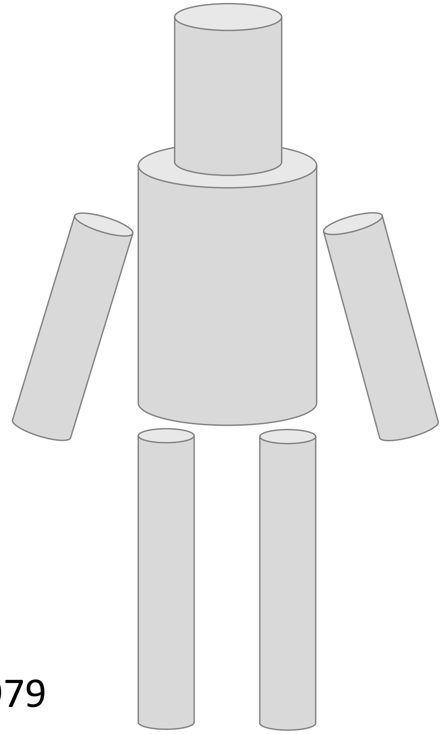

Generalized Cylinders,
Brooks and Binford, 1979

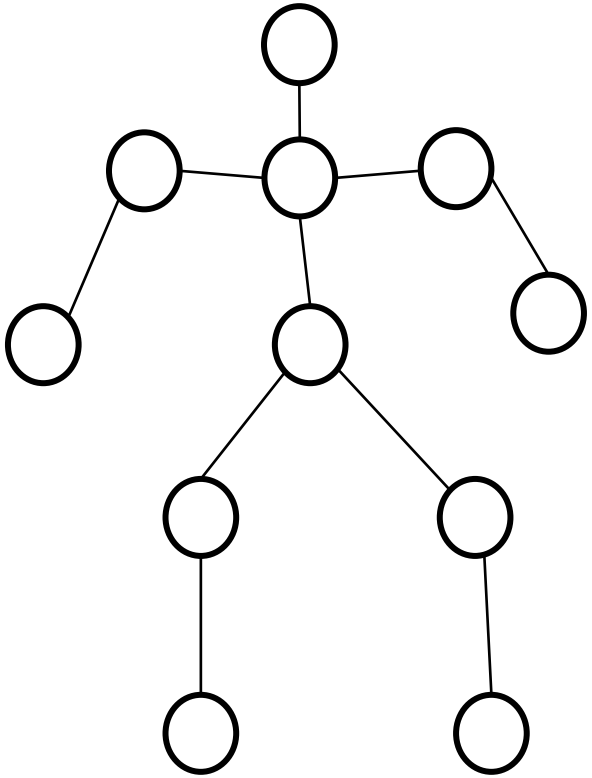

Pictorial Structures,
Fischler and Elshlager, 1973

1959	1963	1970s	1979
Hubel & Wiesel	Roberts	David Marr	Gen. Cylinders

Recognition via Edge Detection (1980s)

1959 Hubel & Wiesel	1963 Roberts	1970s David Marr	1979 Gen. Cylinders	1986 Canny

John Canny, 1986
David Lowe, 1987

Image is CC0 1.0 public domain

Recognition via Grouping (1990s)

1959 Hubel & Wiesel | 1963 Roberts | 1970s David Marr | 1979 Gen. Cylinders | 1986 Canny | AI Winter | 1997 Norm. Cuts

Normalized Cuts, Shi and Malik, 1997

Left Image is CC BY 3.0   Middl Image is public domain   Right Image is CC-BY 2.0; changes made

Recognition via Matching (2000s)

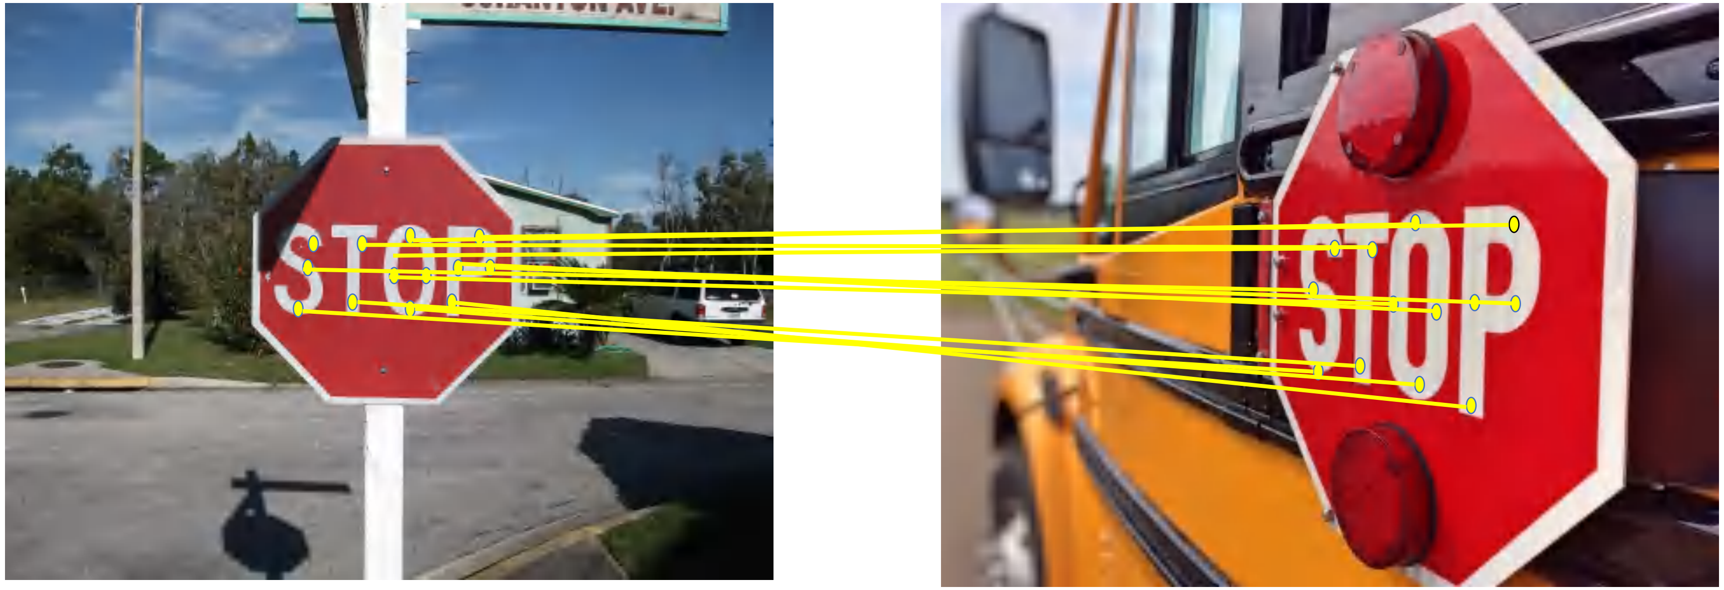

Image is public domain

Image is public domain

1959 Hubel & Wiesel | 1963 Roberts | 1970s David Marr | 1979 Gen. Cylinders | 1986 Canny | 1997 Norm. Cuts | 1999 SIFT

AI Winter

SIFT, David Lowe, 1999

Face Detection

Viola and Jones, 2001

One of the first successful applications of machine learning to vision

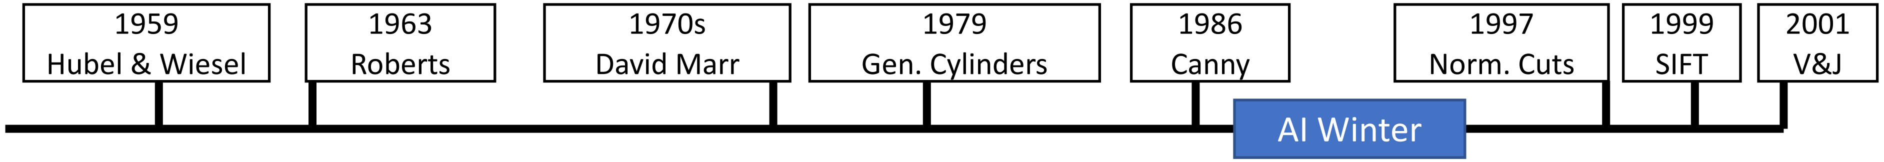

PASCAL Visual Object Challenge

Image is CC0 1.0 public domain

Image is CC0 1.0 public domain

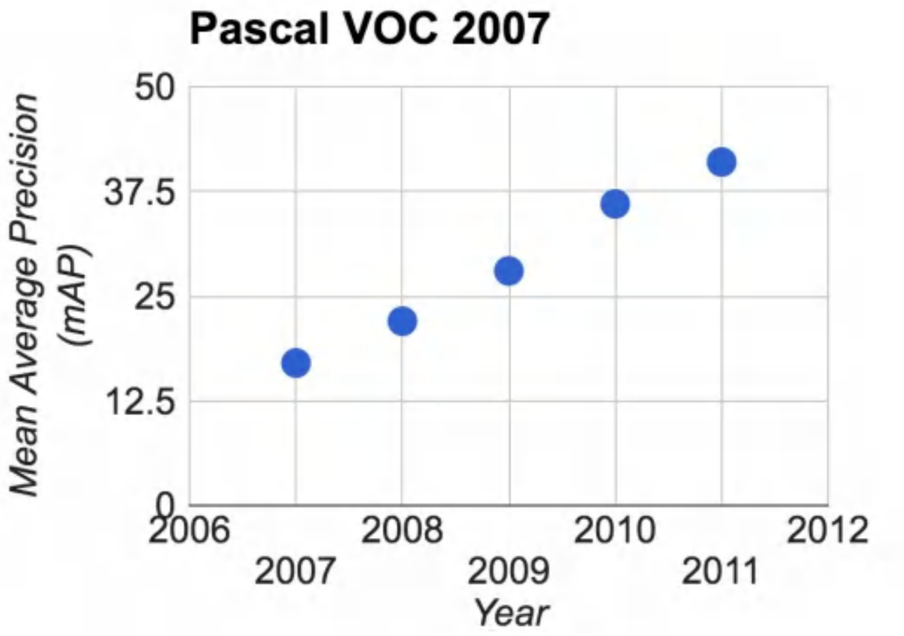

1959 Hubel & Wiesel | 1963 Roberts | 1970s David Marr | 1979 Gen. Cylinders | 1986 Canny | AI Winter | 1997 Norm. Cuts | 1999 SIFT | 2001 V&J | 2007 PASCAL

IMAGENET Large Scale Visual Recognition Challenge

The Image Classification Challenge:

1,000 object classes

1,431,167 images

Output:
Scale
T-shirt
<u>Steel drum</u>
Drumstick
Mud turtle

Deng et al, 2009
Russakovsky et al. IJCV 2015

1959 Hubel & Wiesel | 1963 Roberts | 1970s David Marr | 1979 Gen. Cylinders | 1986 Canny | AI Winter | 1997 Norm. Cuts | 1999 SIFT | 2001 V&J | 2007 PASCAL | 2009 ImageNet

IMAGENET Large Scale Visual Recognition Challenge

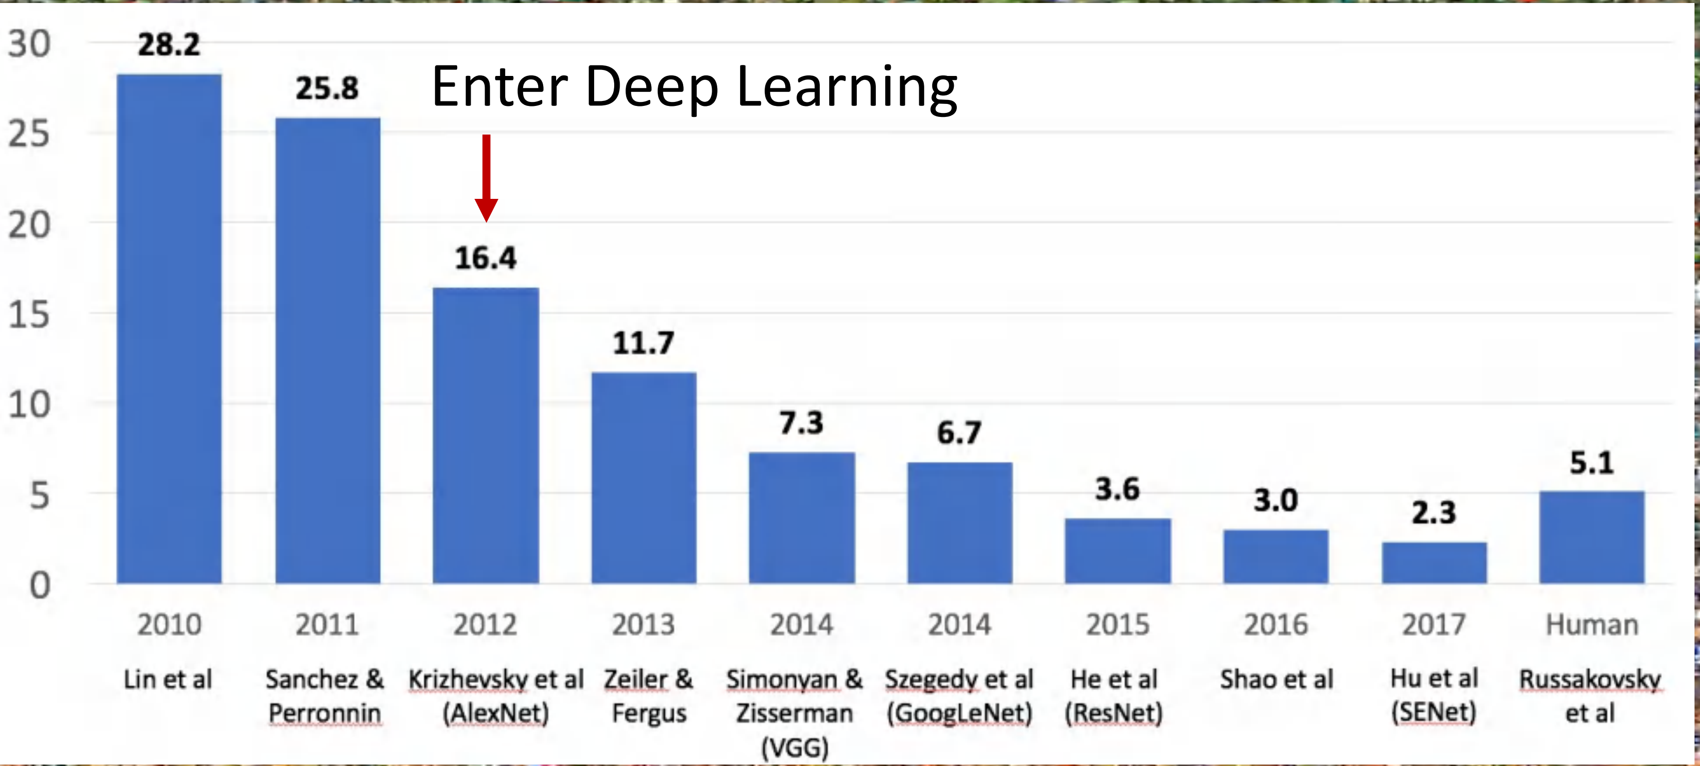

1959 Hubel & Wiesel

1963 Roberts

1970s David Marr

1979 Gen. Cylinders

1986 Canny

AI Winter

1997 Norm. Cuts

1999 SIFT

2001 V&J

2007 PASCAL

2009 ImageNet

2012 AlexNet

AlexNet: Deep Learning Goes Mainstream

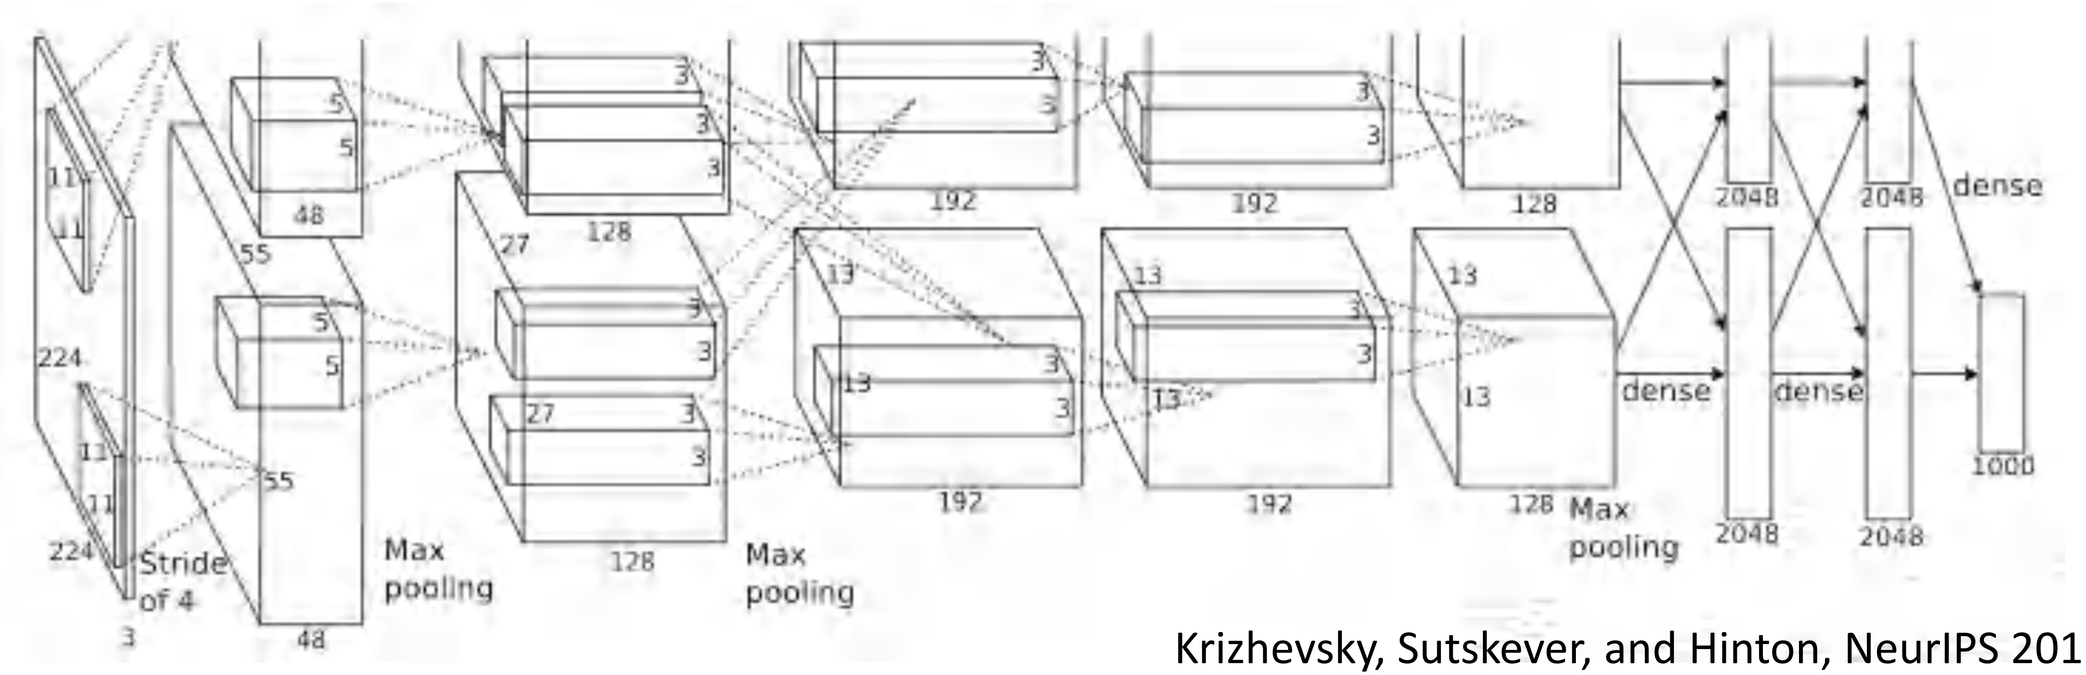

Krizhevsky, Sutskever, and Hinton, NeurIPS 2012

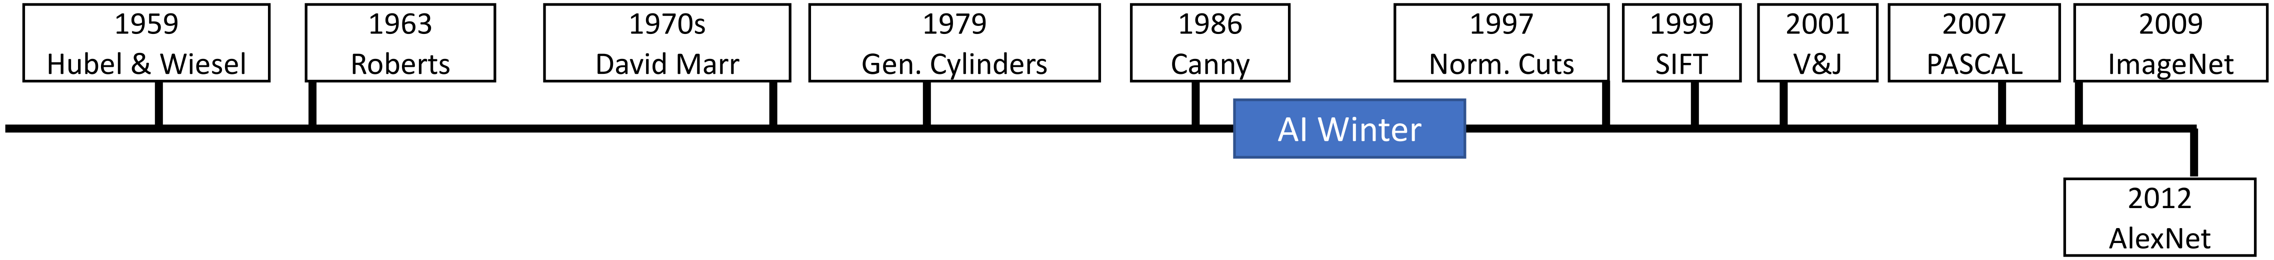

Perceptron

One of the earliest algorithms that could learn from data

Implemented in hardware! Weights stored in potentiometers, updated with electric motors during learning

Connected to a camera that used 20x20 cadmium sulfide photocells to make a 400-pixel image

Could learn to recognize letters of the alphabet

Today we would recognize it as a **linear classifier**

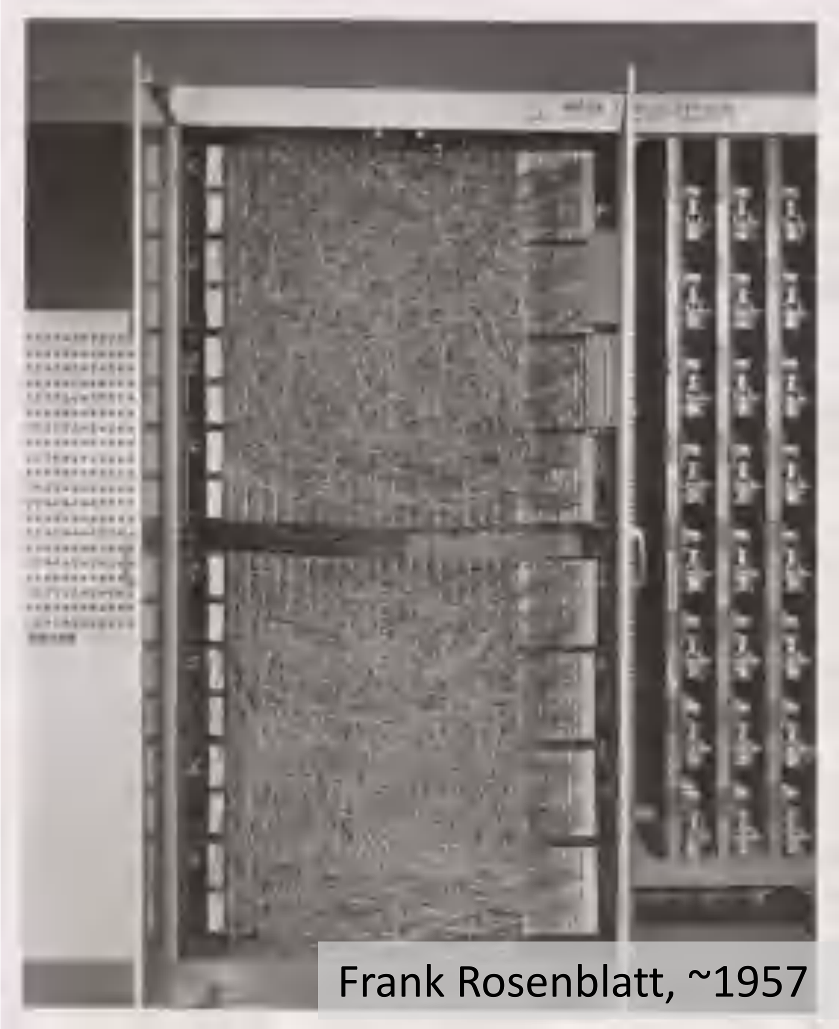

Frank Rosenblatt, ~1957

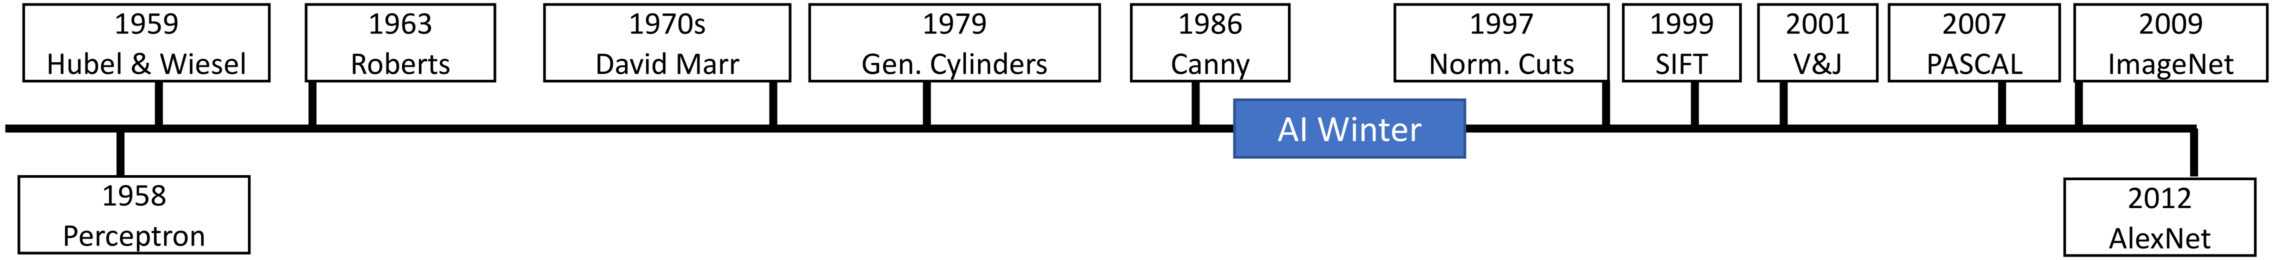

Minsky and Papert, 1969

X	Y	F(x,y)
0	0	0
0	1	1
1	0	1
1	1	0

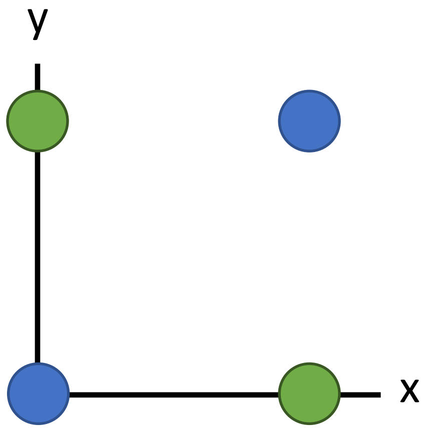

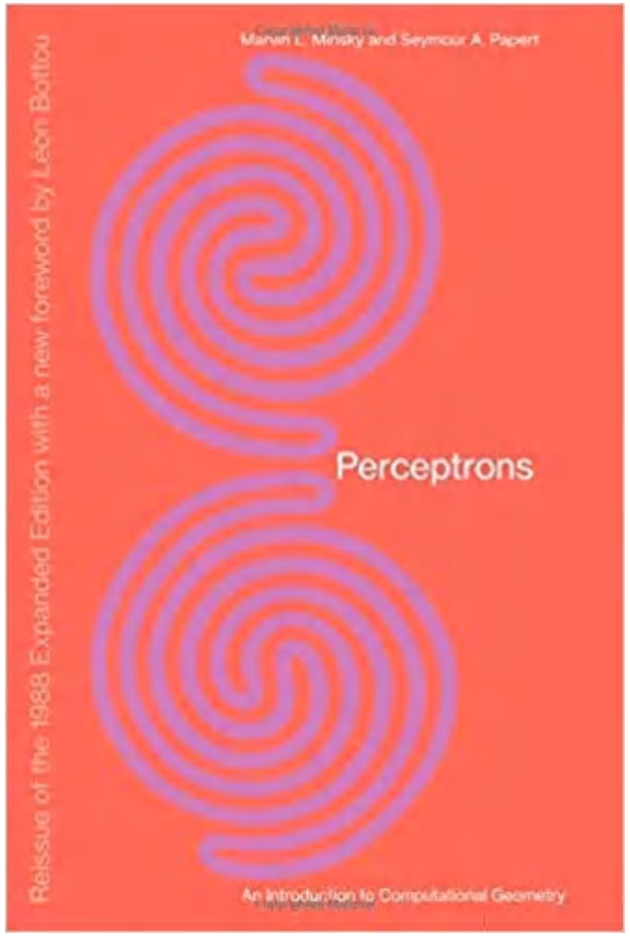

Showed that Perceptrons could not learn the XOR function
Caused a lot of disillusionment in the field

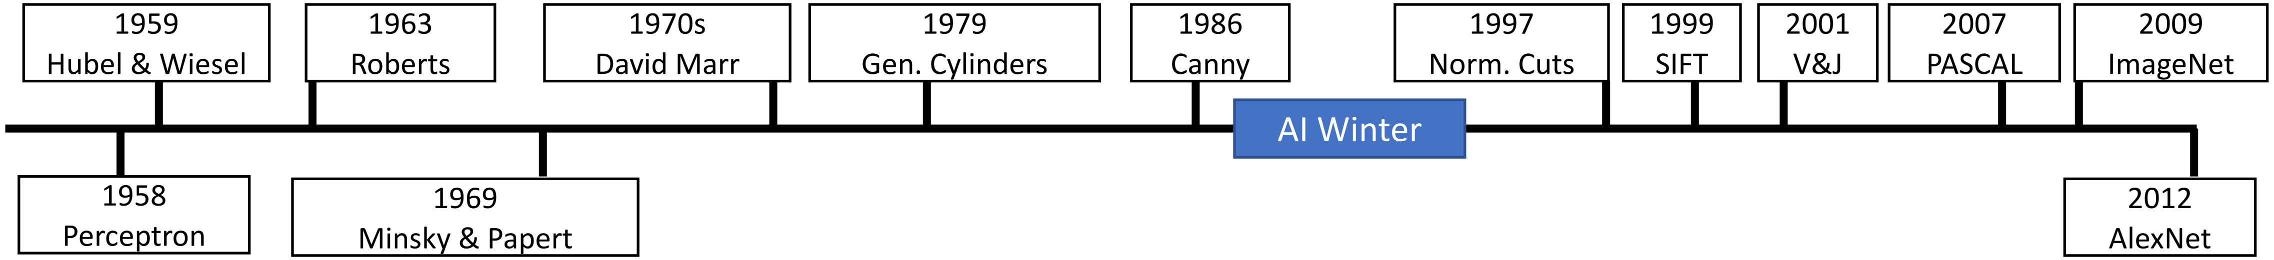

Neocognitron: Fukushima, 1980

Computational model the visual system, directly inspired by Hubel and Wiesel's hierarchy of complex and simple cells

Interleaved simple cells (convolution) and complex cells (pooling)

No practical training algorithm

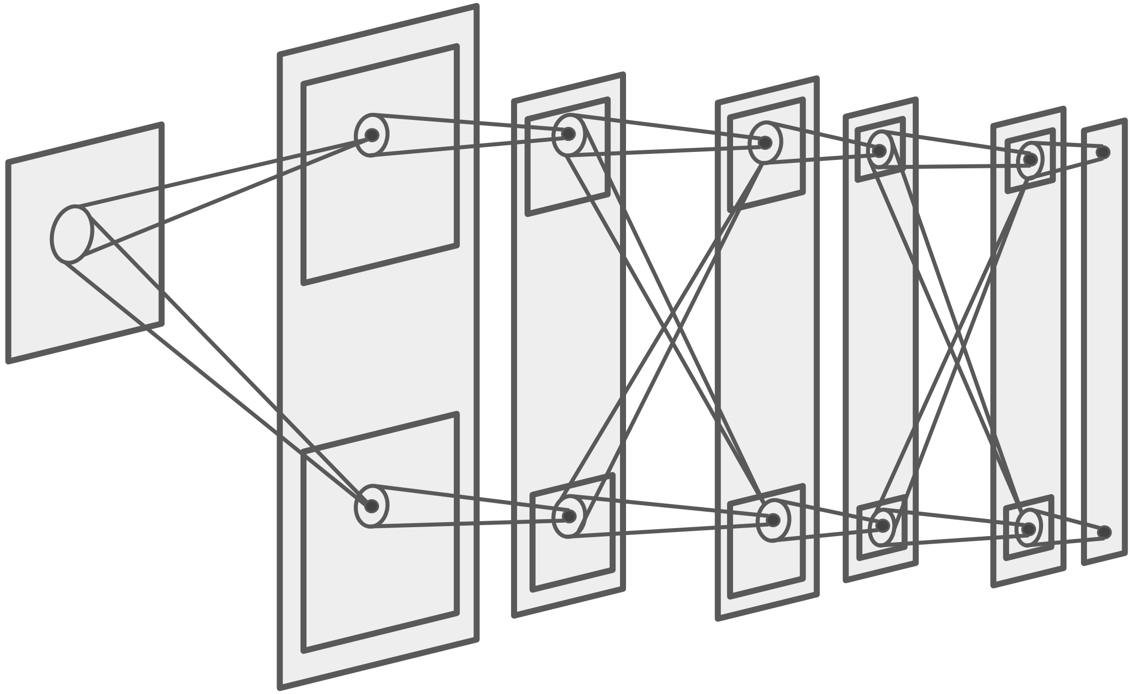

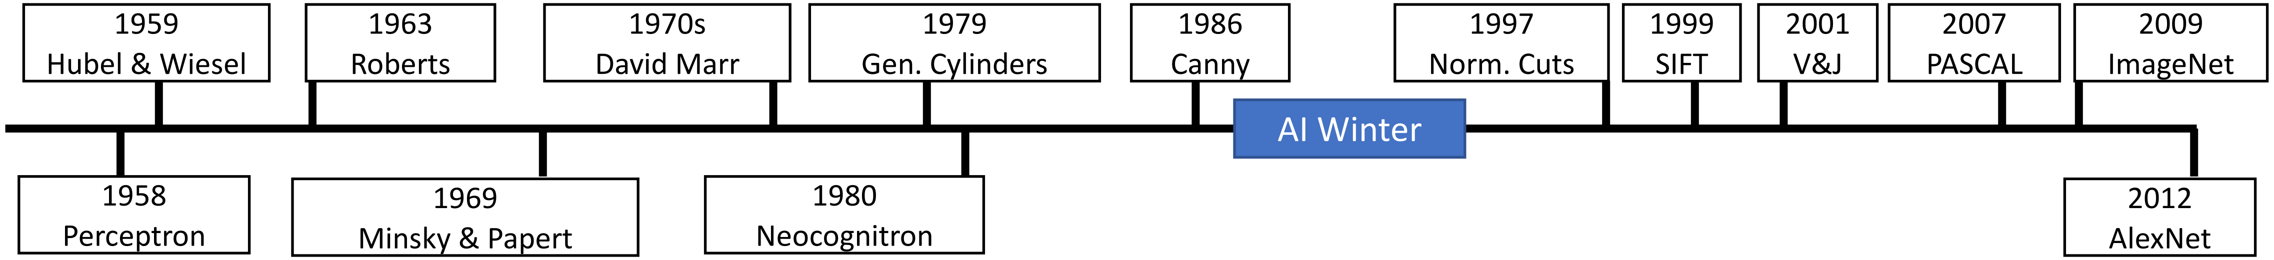

Neocognitron: Fukushima, 1980

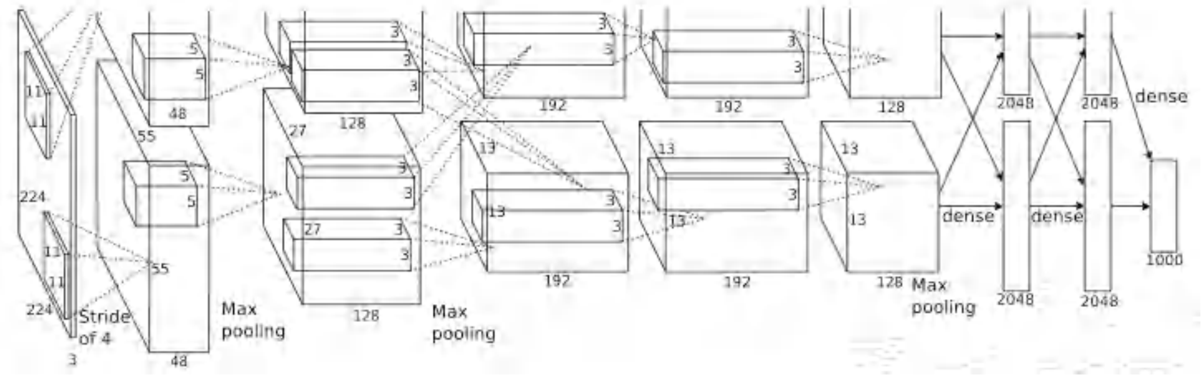

Looks a lot like AlexNet
more than 32 years later!

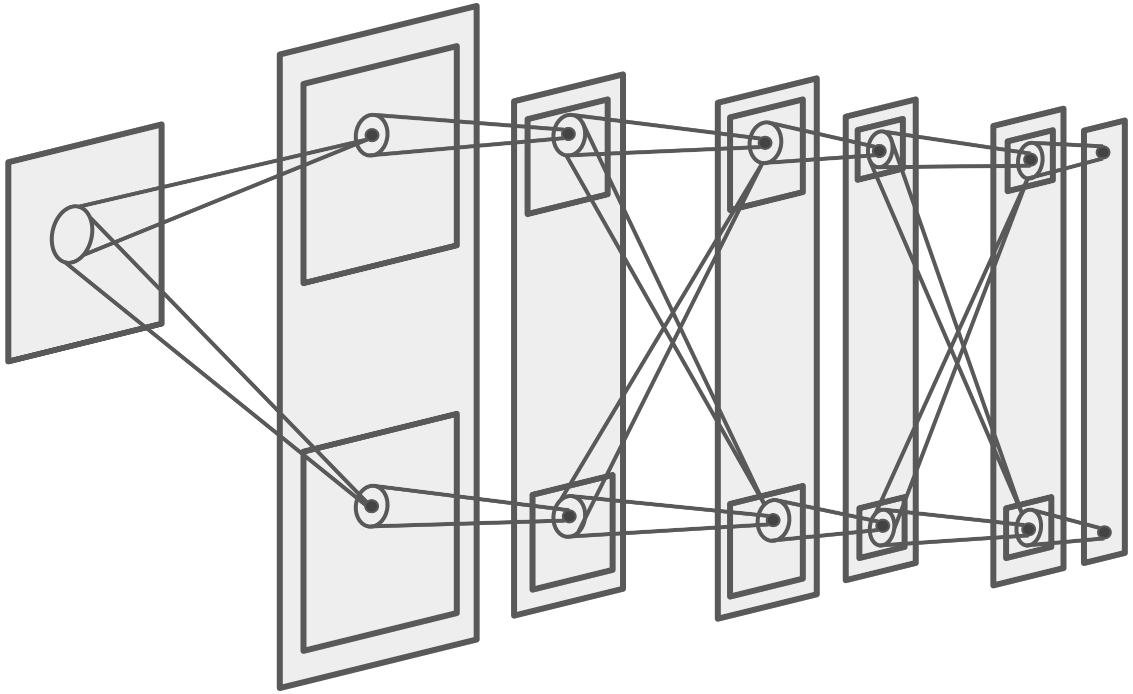

1959 Hubel & Wiesel

1963 Roberts

1970s David Marr

1979 Gen. Cylinders

1986 Canny

1997 Norm. Cuts

1999 SIFT

2001 V&J

2007 PASCAL

2009 ImageNet

AI Winter

1958 Perceptron

1969 Minsky & Papert

1980 Neocognitron

2012 AlexNet

Backprop: Rumelhart, Hinton, and Williams, 1986

Introduced backpropagation for computing gradients in neural networks

Successfully trained perceptrons with multiple layers

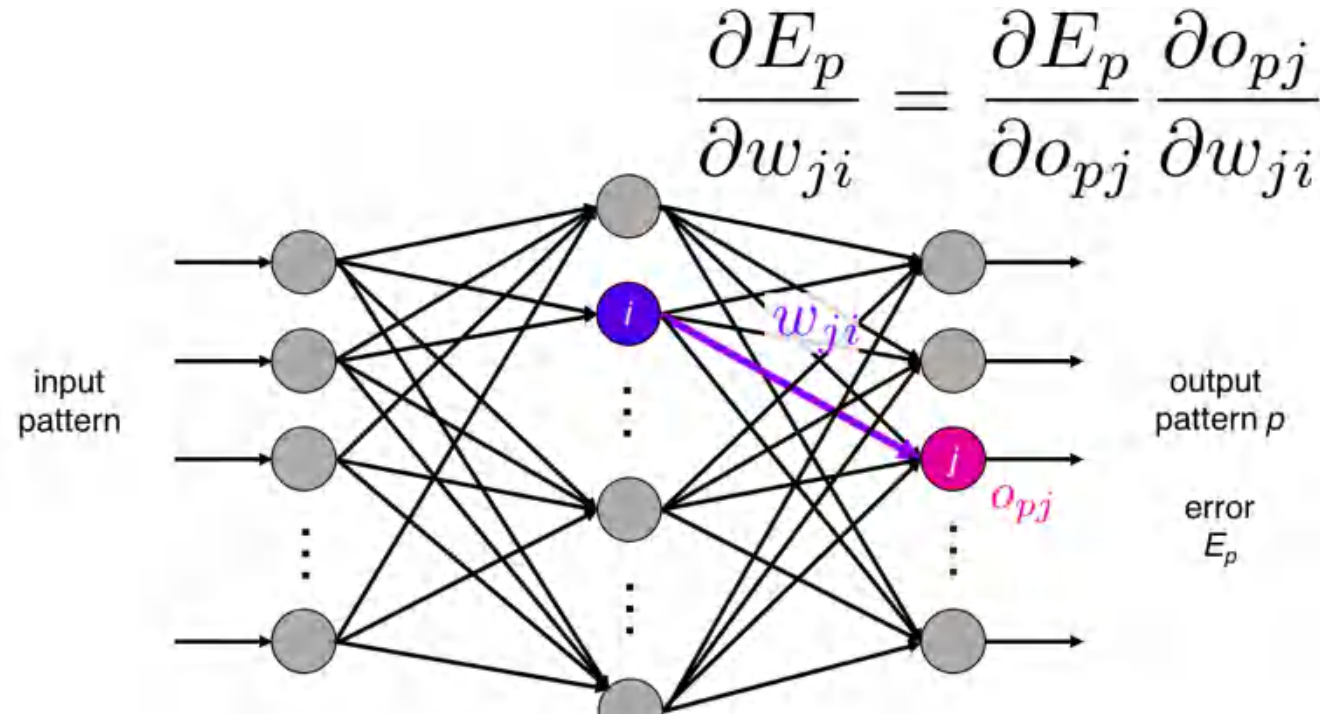

recognizable math

Illustration of Rumelhart et al., 1986 by Lane McIntosh, copyright CS231n 2017

1959 Hubel & Wiesel
1963 Roberts
1970s David Marr
1979 Gen. Cylinders
1986 Canny
1997 Norm. Cuts
1999 SIFT
2001 V&J
2007 PASCAL
2009 ImageNet

AI Winter

1958 Perceptron
1969 Minsky & Papert
1980 Neocognitron
1985 Backprop
2012 AlexNet

Convolutional Networks: LeCun et al, 1998

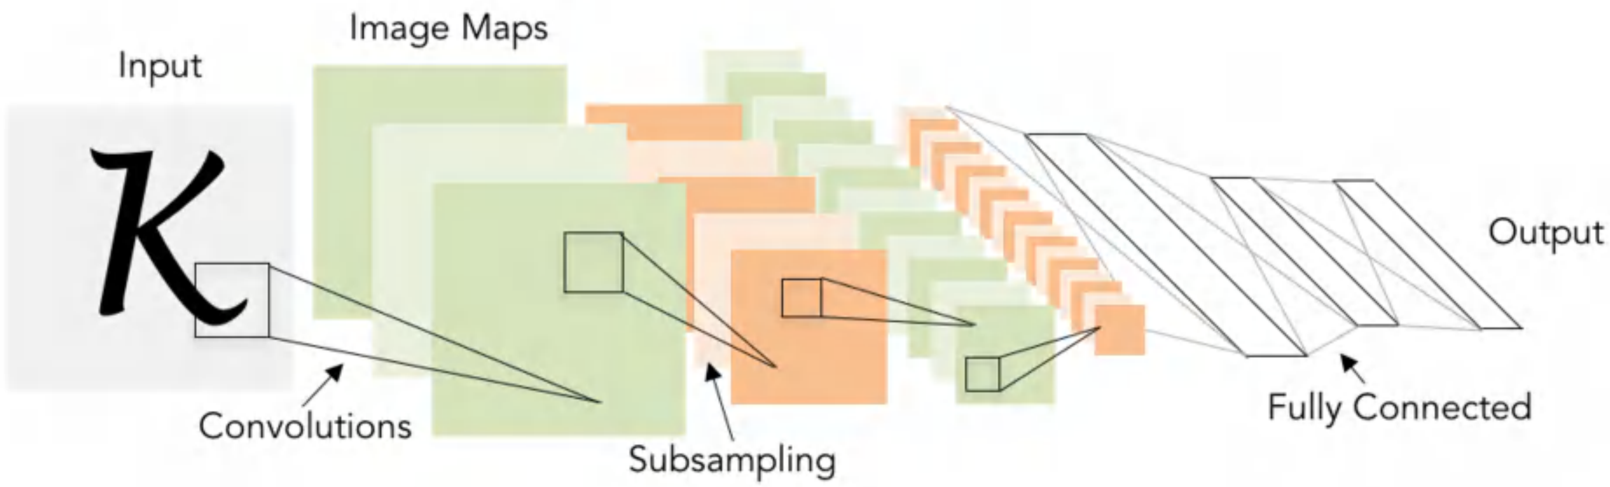

Applied backprop algorithm to a Neocognitron-like architecture
Learned to recognize handwritten digits
Was deployed in a commercial system by NEC, processed handwritten checks
Very similar to our modern convolutional networks!

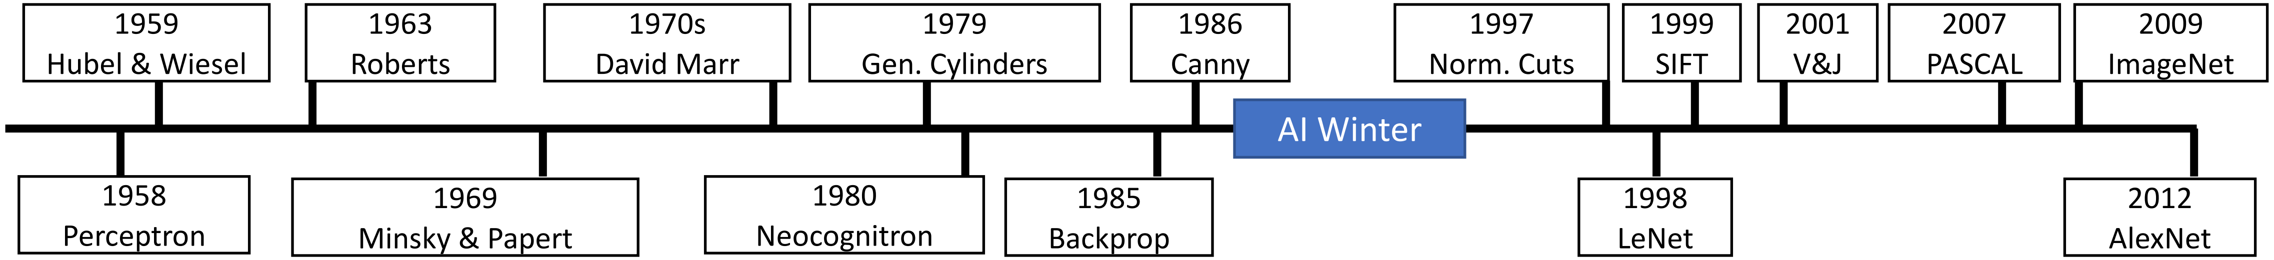

2000s: “Deep Learning”

People tried to train neural networks that were deeper and deeper

Not a mainstream research topic at this time

Hinton and Salakhutdinov, 2006
Bengio et al, 2007
Lee et al, 2009
Glorot and Bengio, 2010

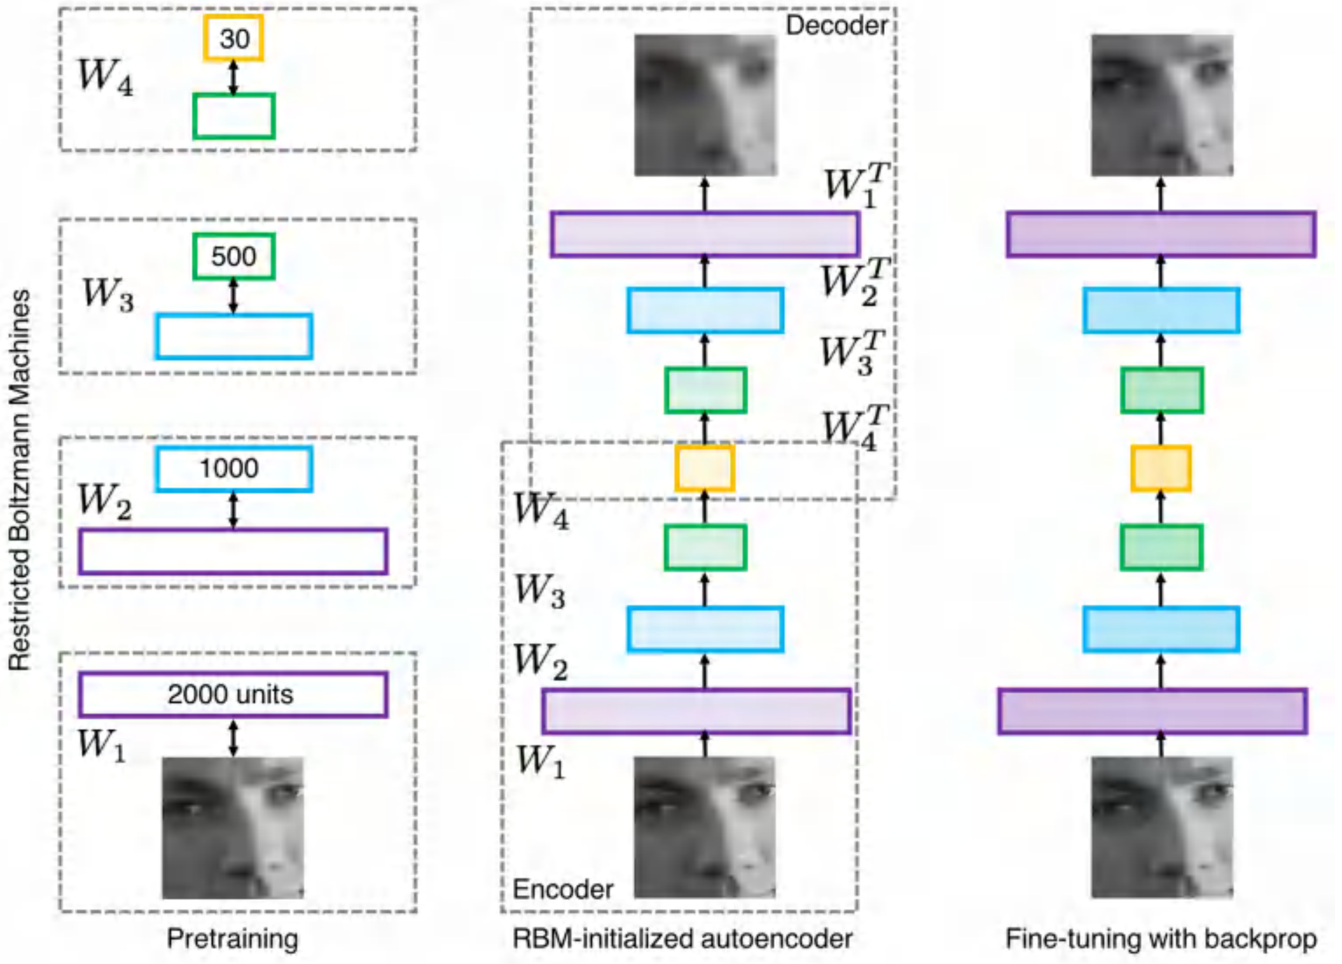

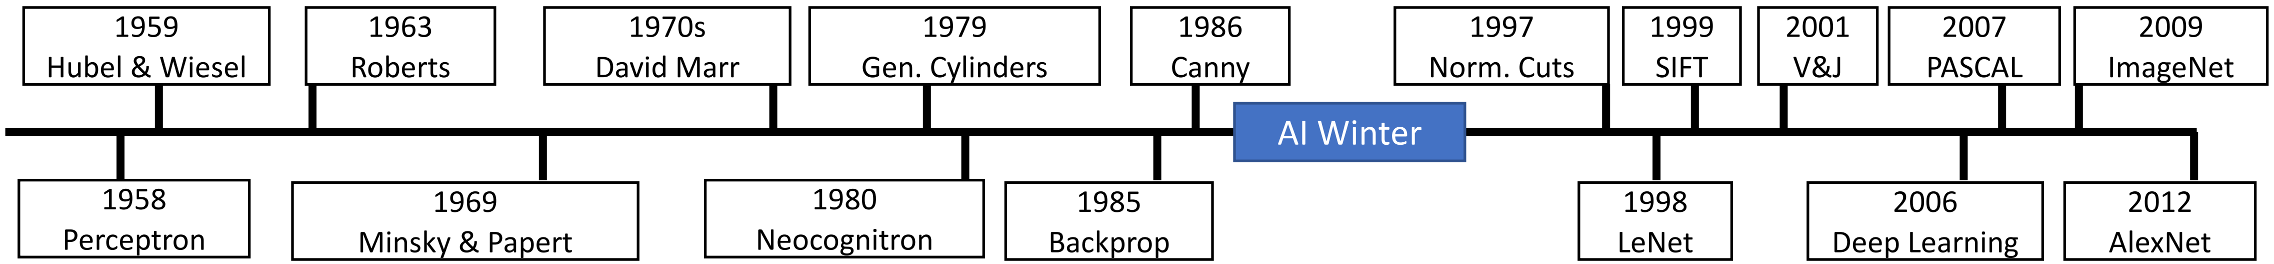

2012 to Present: Deep Learning Explosion

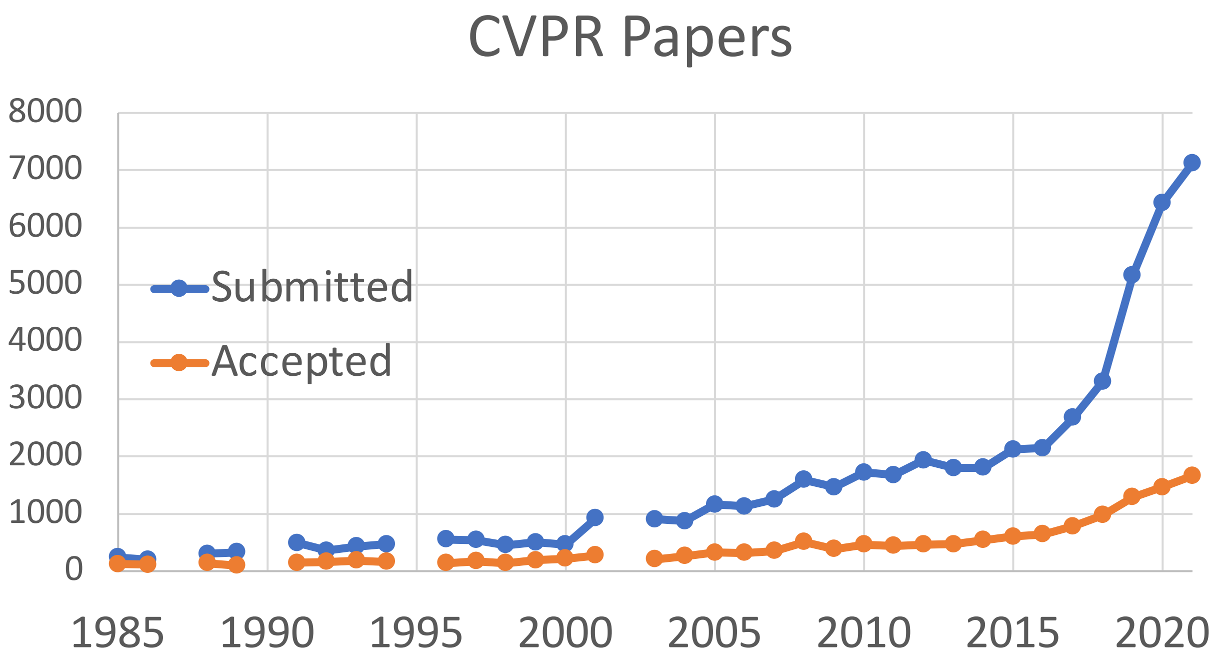

Publications at top Computer Vision conference

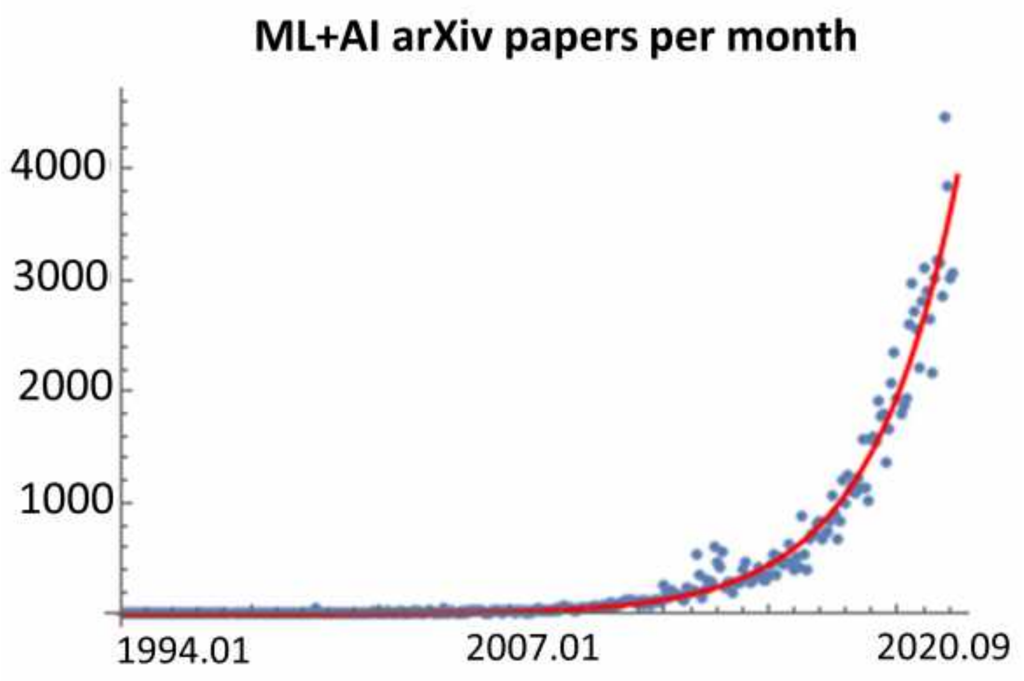

arXiv papers per month [(source)]

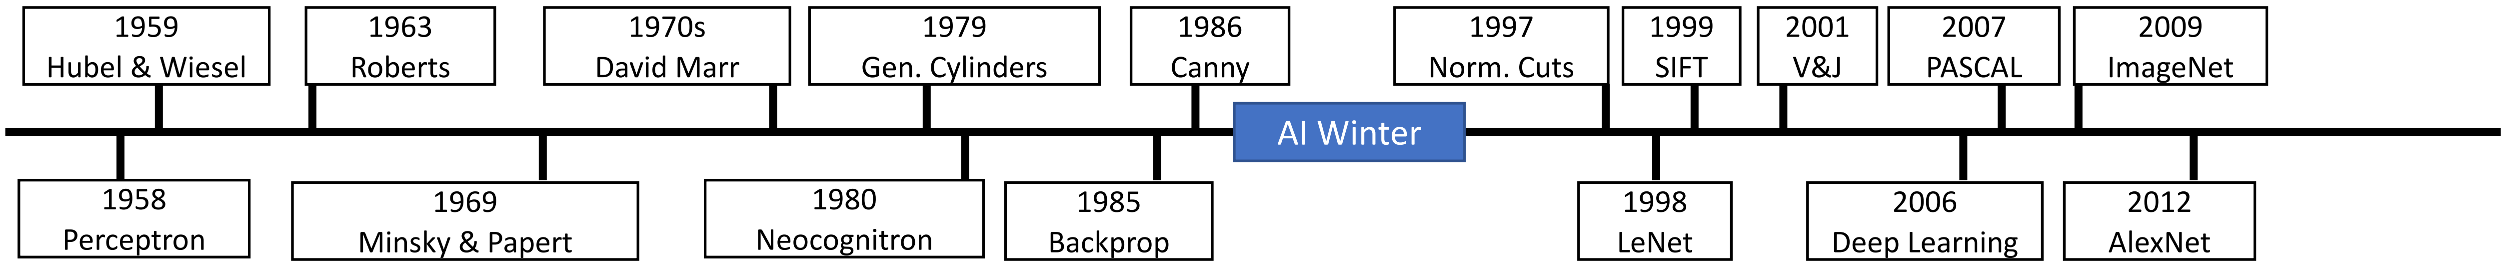

2012 to Present: Deep Learning is Everywhere

Image Classification

Image Retrieval

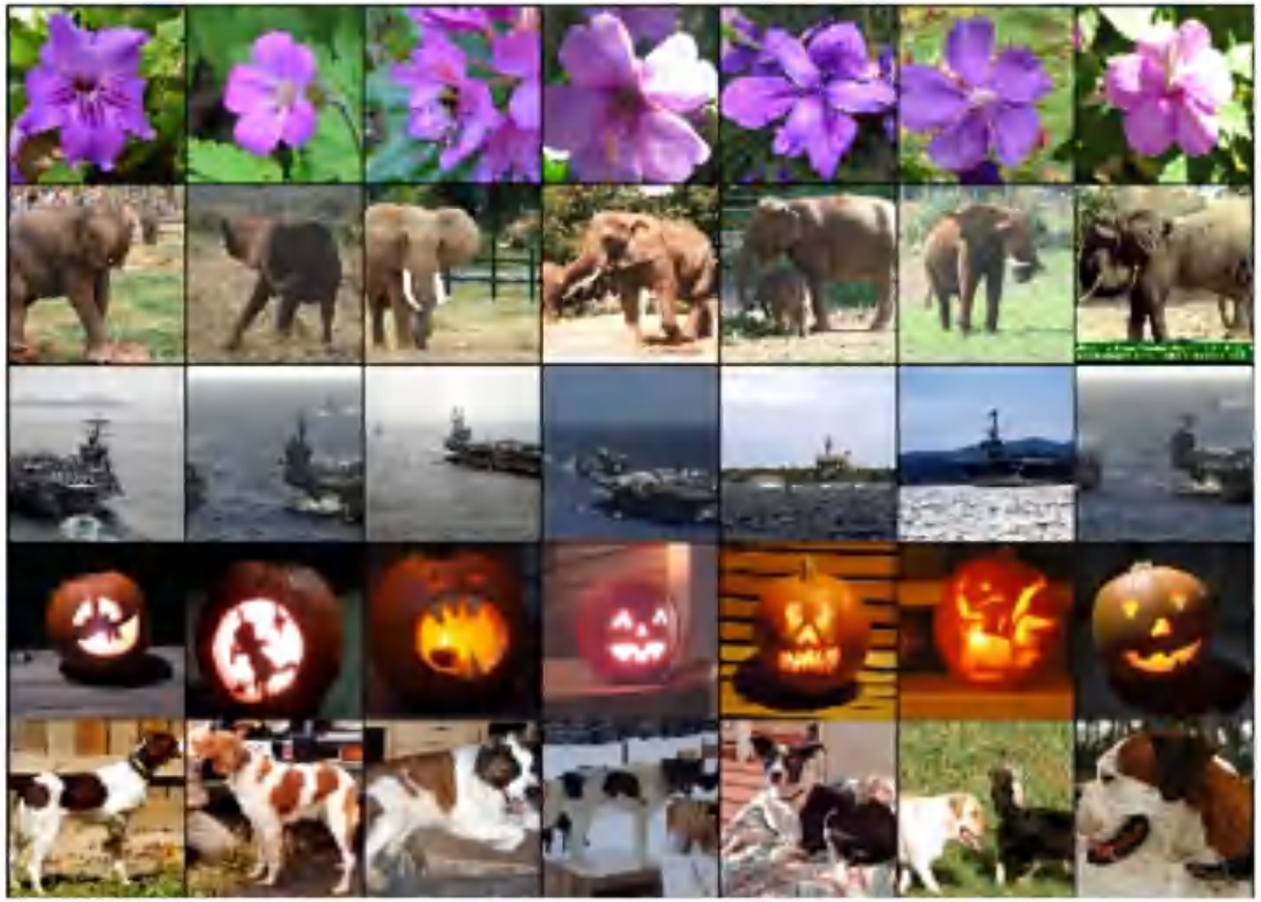

Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

2012 to Present: Deep Learning is Everywhere

Object Detection

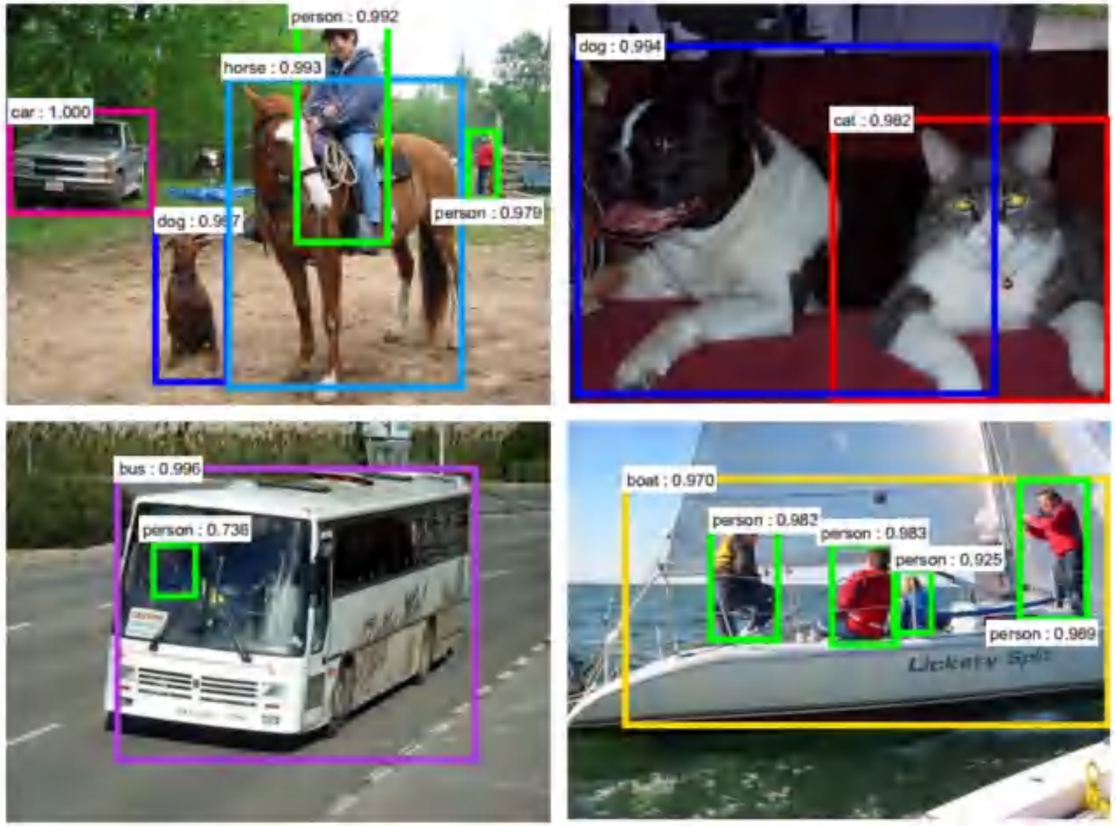

Ren, He, Girshick, and Sun, 2015

Image Segmentation

Fabaret et al, 2012

2012 to Present: Deep Learning is Everywhere

Video Classification

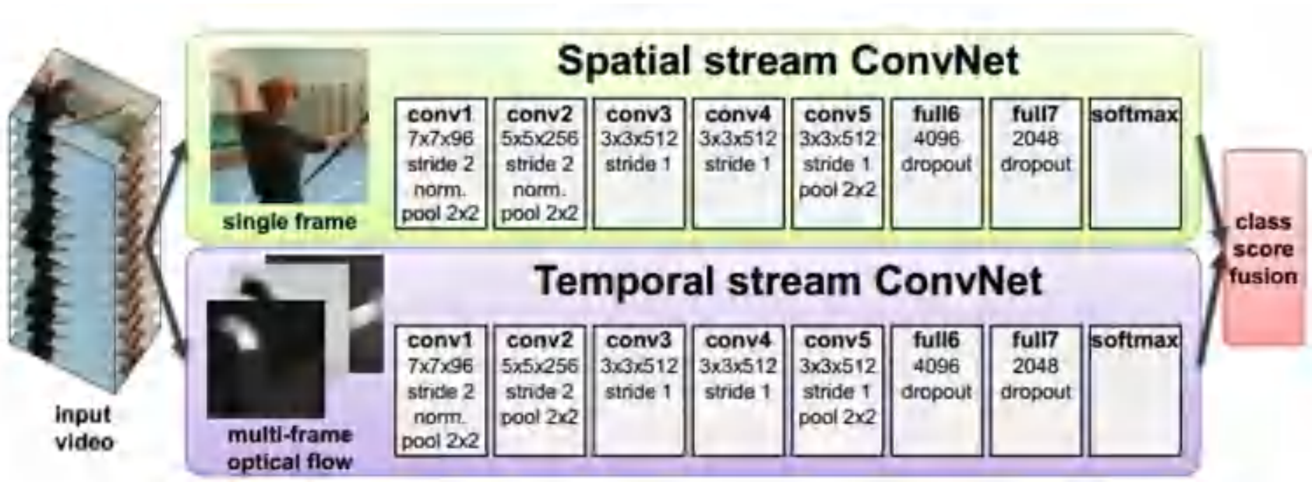

Simonyan et al, 2014

Activity Recognition

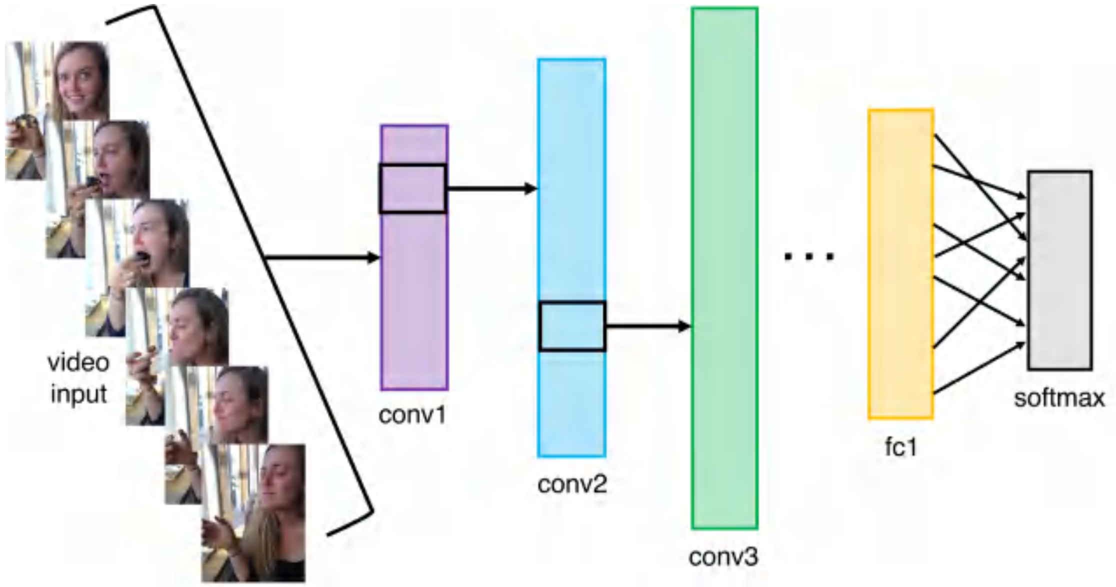

2012 to Present: Deep Learning is Everywhere

Pose Recognition (Toshev and Szegedy, 2014)

Playing Atari games (Guo et al, 2014)

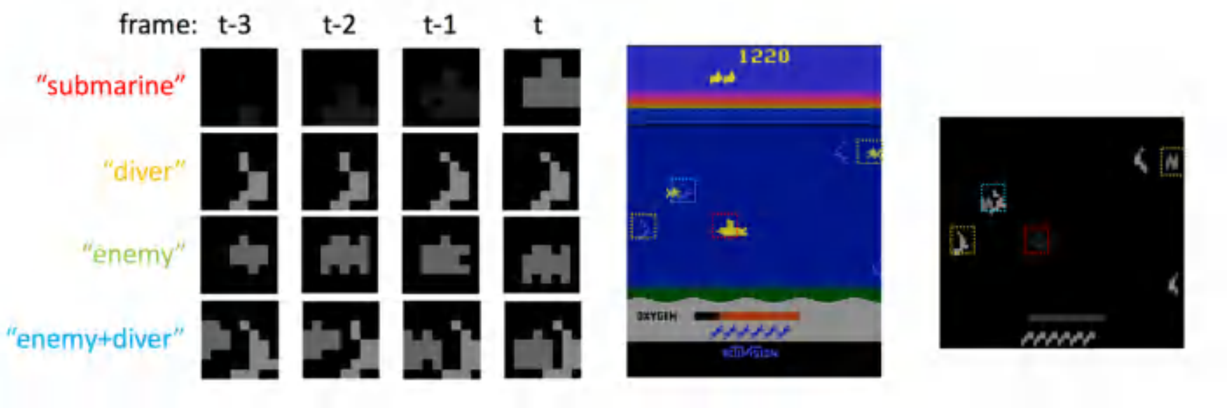

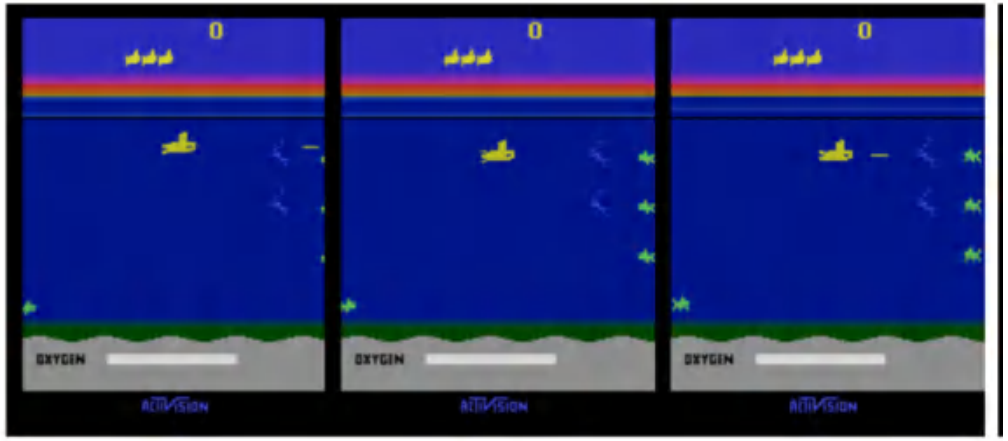

2012 to Present: Deep Learning is Everywhere

Medical Imaging

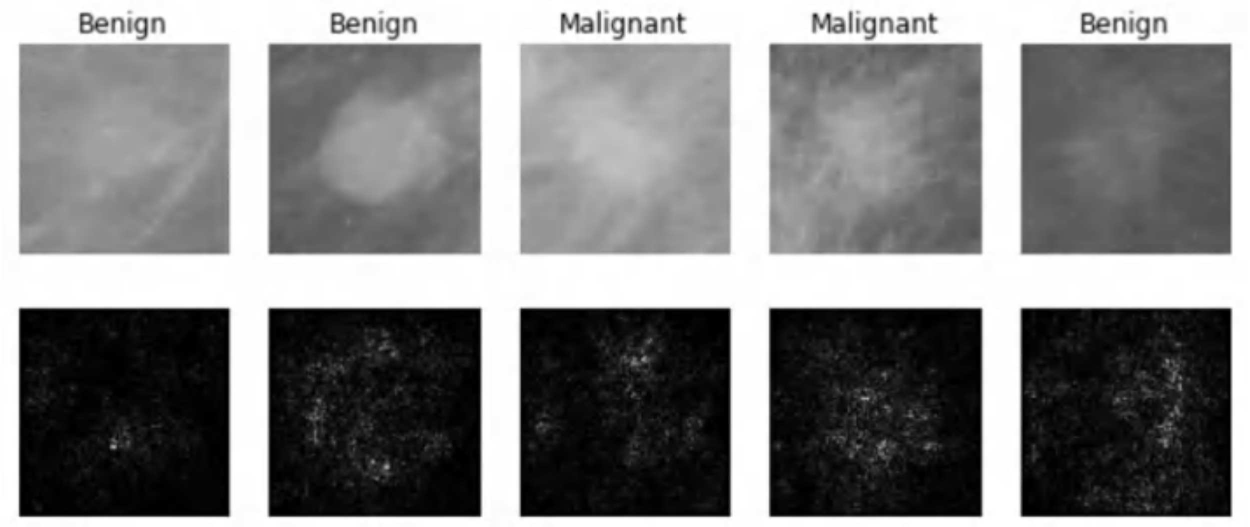

Levy et al, 2016 Figure reproduced with permission

Galaxy Classification

Dieleman et al, 2014

From left to right: public domain by NASA, usage permitted by ESA/Hubble, public domain by NASA, and public domain.

Whale recognition

Kaggle Challenge

This image by Christin Khan is in the public domain and originally came from the U.S. NOAA.

2012 to Present: Deep Learning is Everywhere

Image Captioning

Vinyals et al, 2015
Karpathy and Fei-Fei, 2015

A white teddy bear sitting in the grass

A man in a baseball uniform throwing a ball

A woman is holding a cat in her hand

A man riding a wave on top of a surfboard

A cat sitting on a suitcase on the floor

A woman standing on a beach holding a surfboard

All images are CC0 Public domain:
https://pixabay.com/en/luggage-antique-cat-1643010/
https://pixabay.com/en/teddy-plush-bears-cute-teddy-bear-1623436/
https://pixabay.com/en/surf-wave-summer-sport-litoral-1668716/
https://pixabay.com/en/woman-female-model-portrait-adult-983967/
https://pixabay.com/en/handstand-lake-meditation-496008/
https://pixabay.com/en/baseball-player-shortstop-infield-1045263/

Captions generated by Justin Johnson using Neuraltalk2

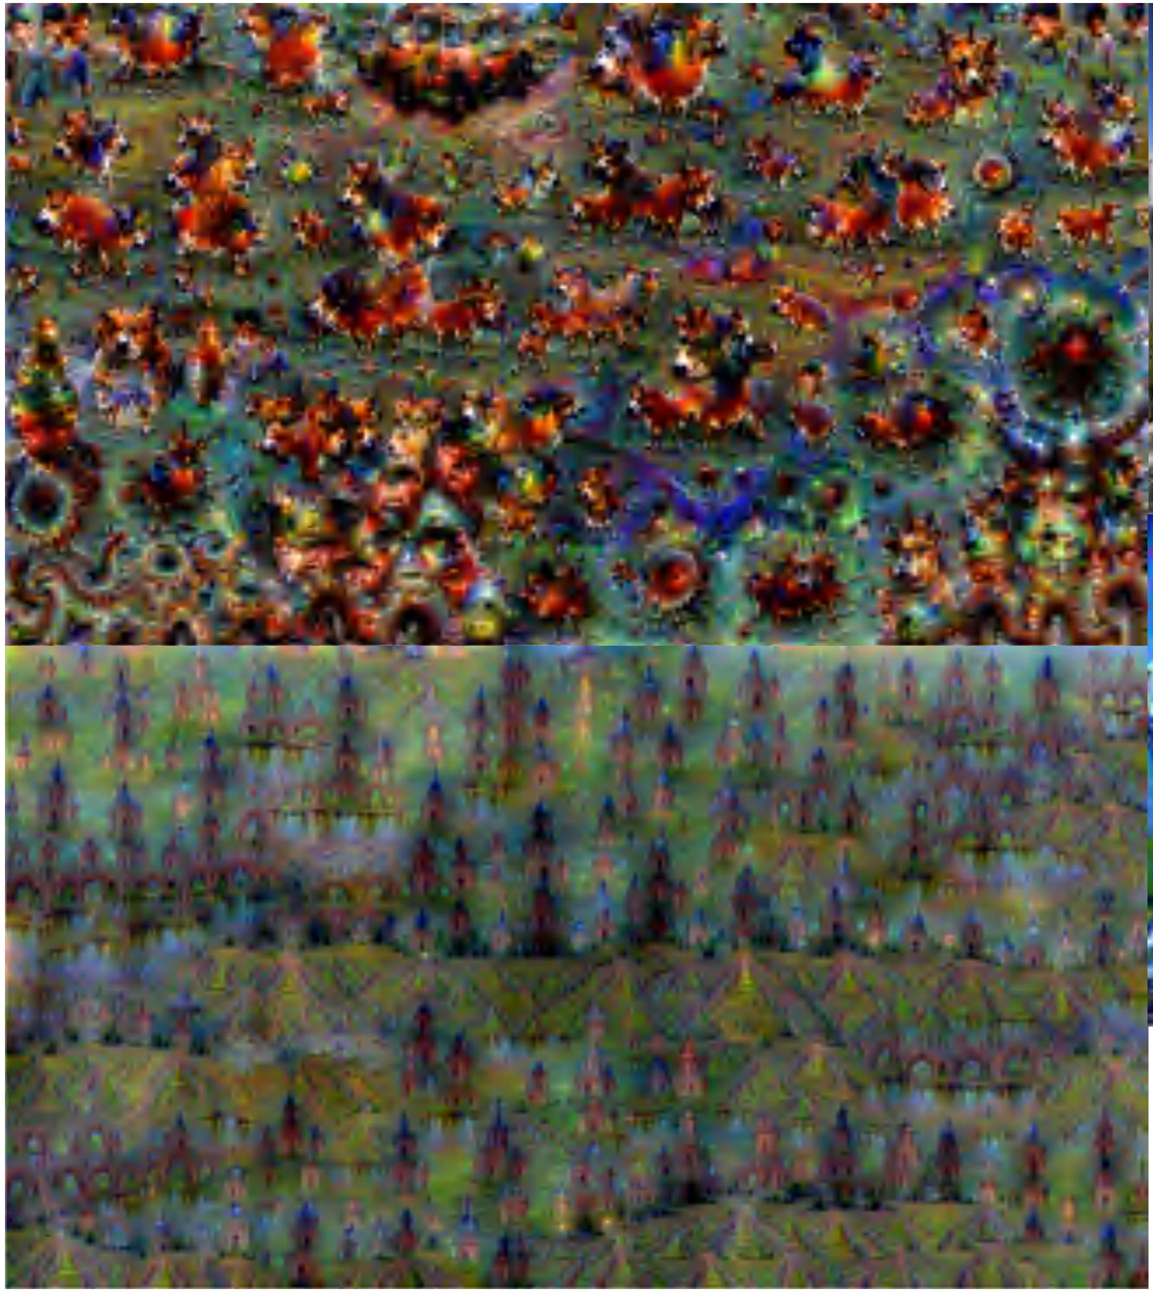

Figures copyright Justin Johnson, 2015. Reproduced with permission. Generated using the Inceptionism approach from a blog post by Google Research.

Original image is CC0 public domain
Starry Night and Tree Roots by Van Gogh are in the public domain
Bokeh image is in the public domain
Stylized images copyright Justin Johnson, 2017;
reproduced with permission

Mordvinsev et al, 2015

Gatys et al, 2016

2012 to Present: Deep Learning is Everywhere

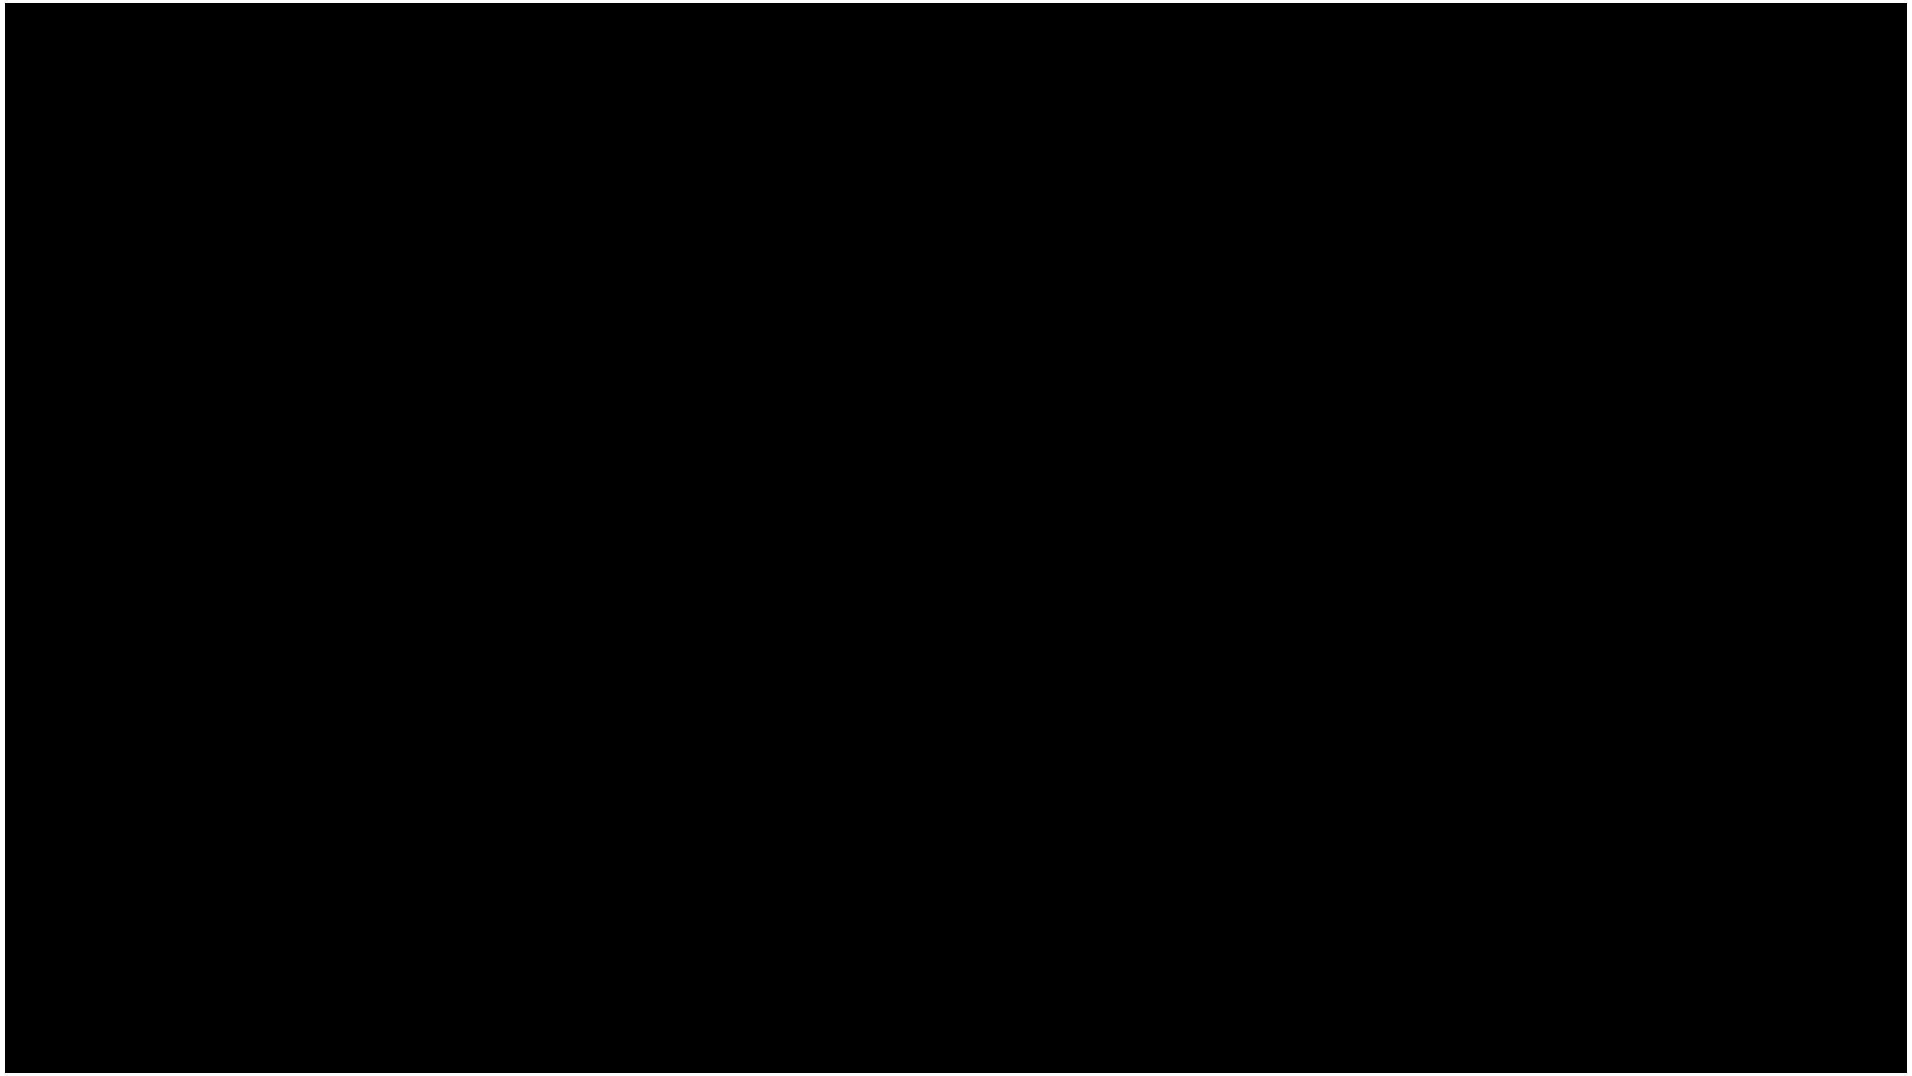

Karras et al, “Progressive Growing of GANs for Improved Quality, Stability, and Variation”, ICLR 2018

2012 to Present: Deep Learning is Everywhere

Ramesh et al, “DALL·E: Creating Images from Text”, 2021. https://openai.com/blog/dall-e/

2012 to Present: Deep Learning is Everywhere

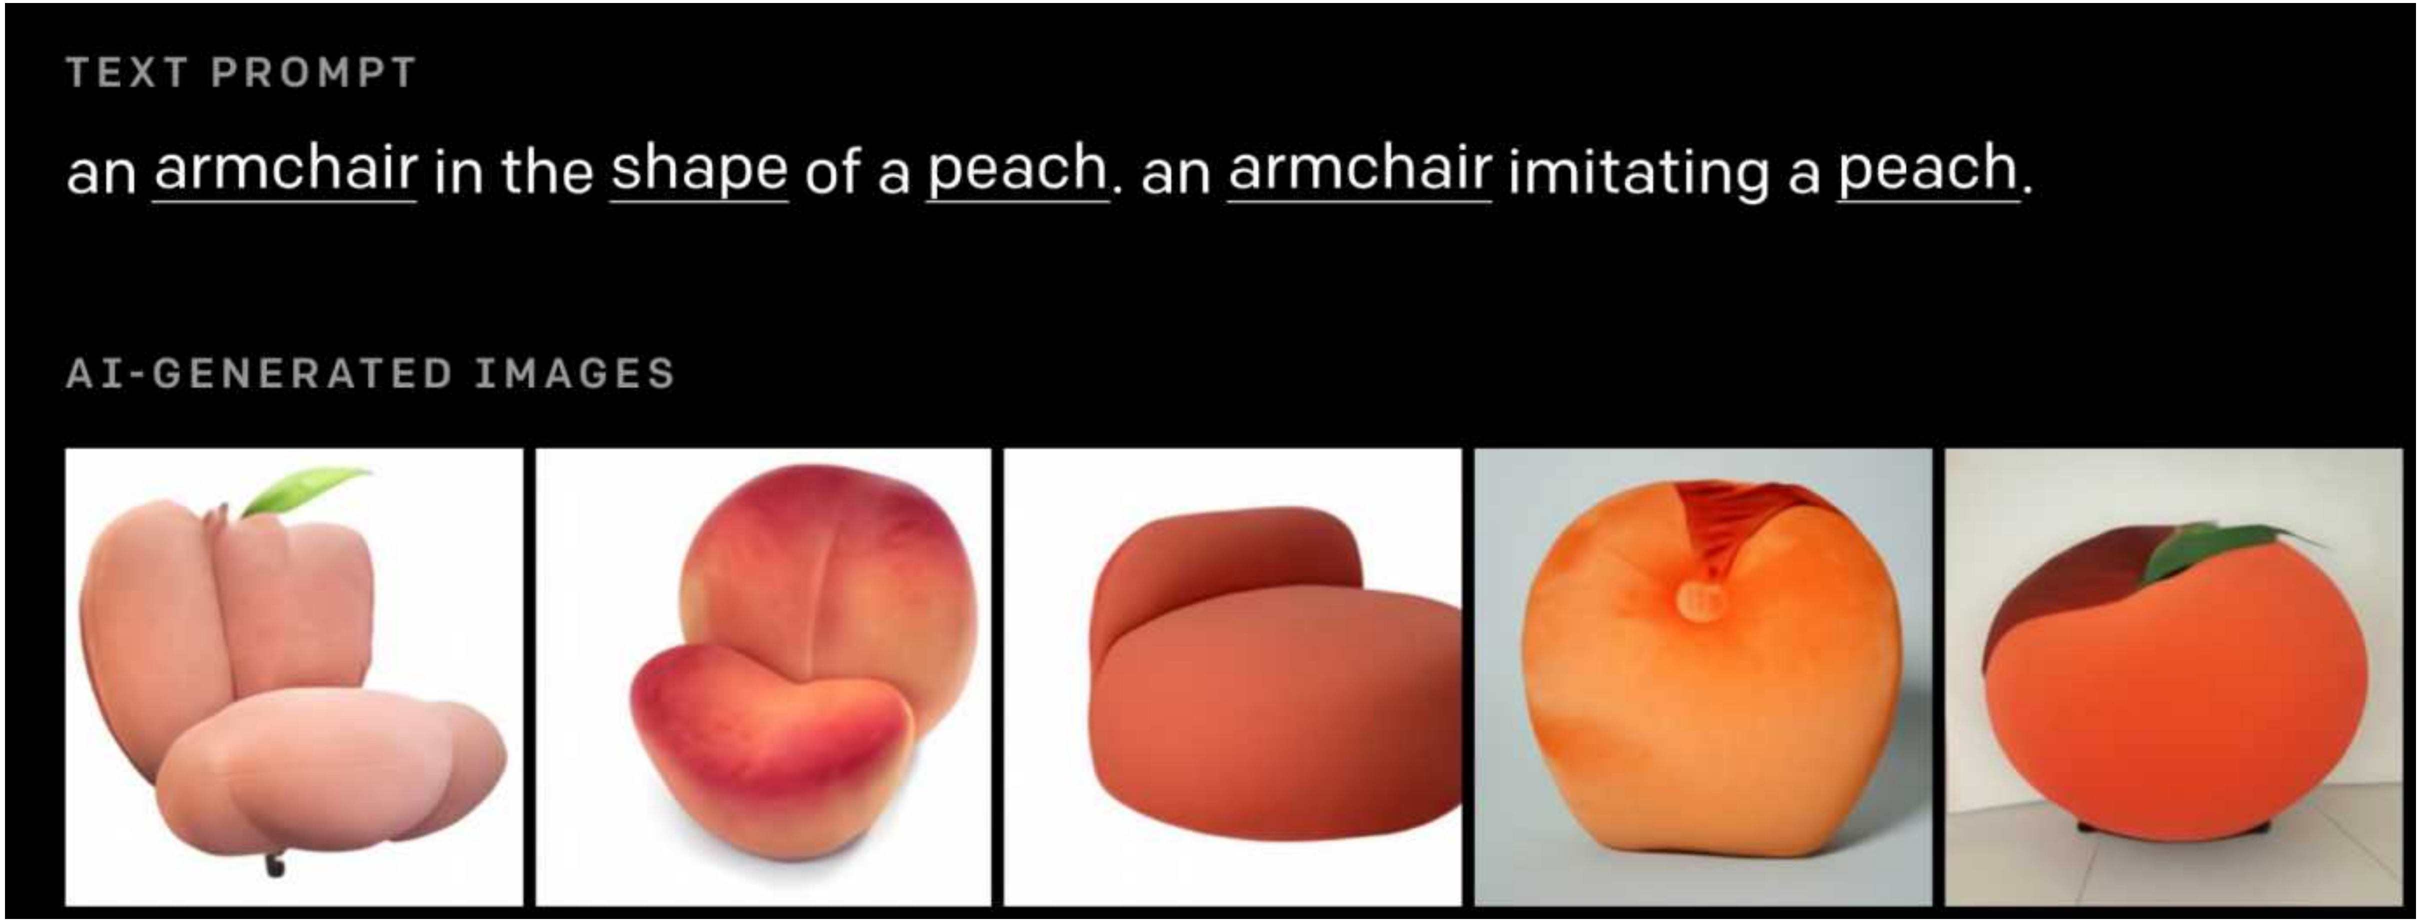

Ramesh et al, “DALL·E: Creating Images from Text”, 2021. https://openai.com/blog/dall-e/


Algorithms
13
192
128
dense
2048
1000
Data
Computation
GTX

GFLOP per Dollar

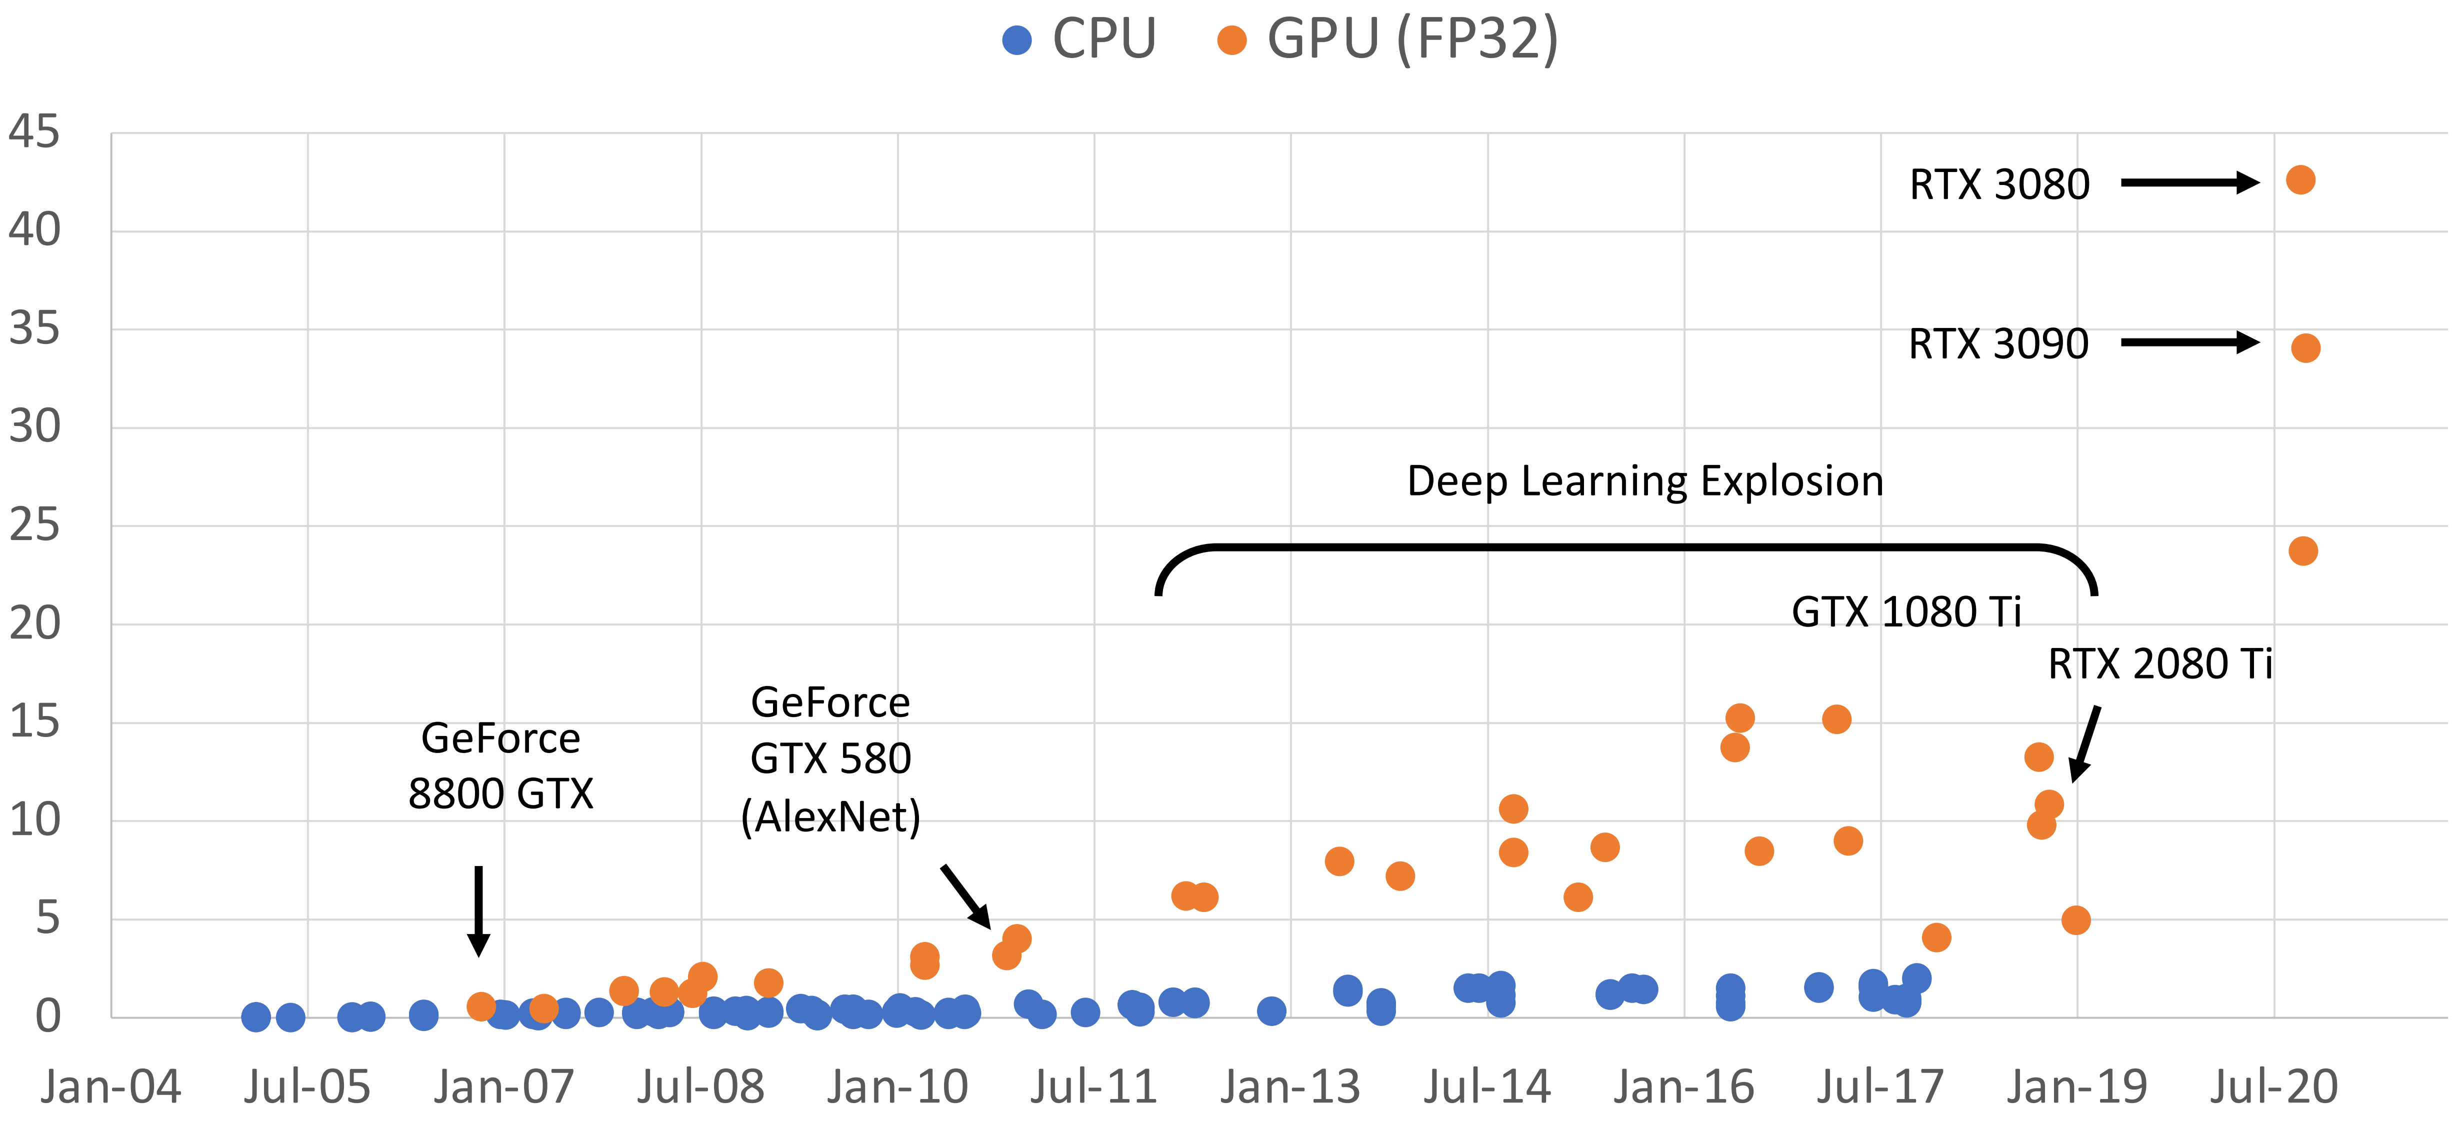

GFLOP per Dollar

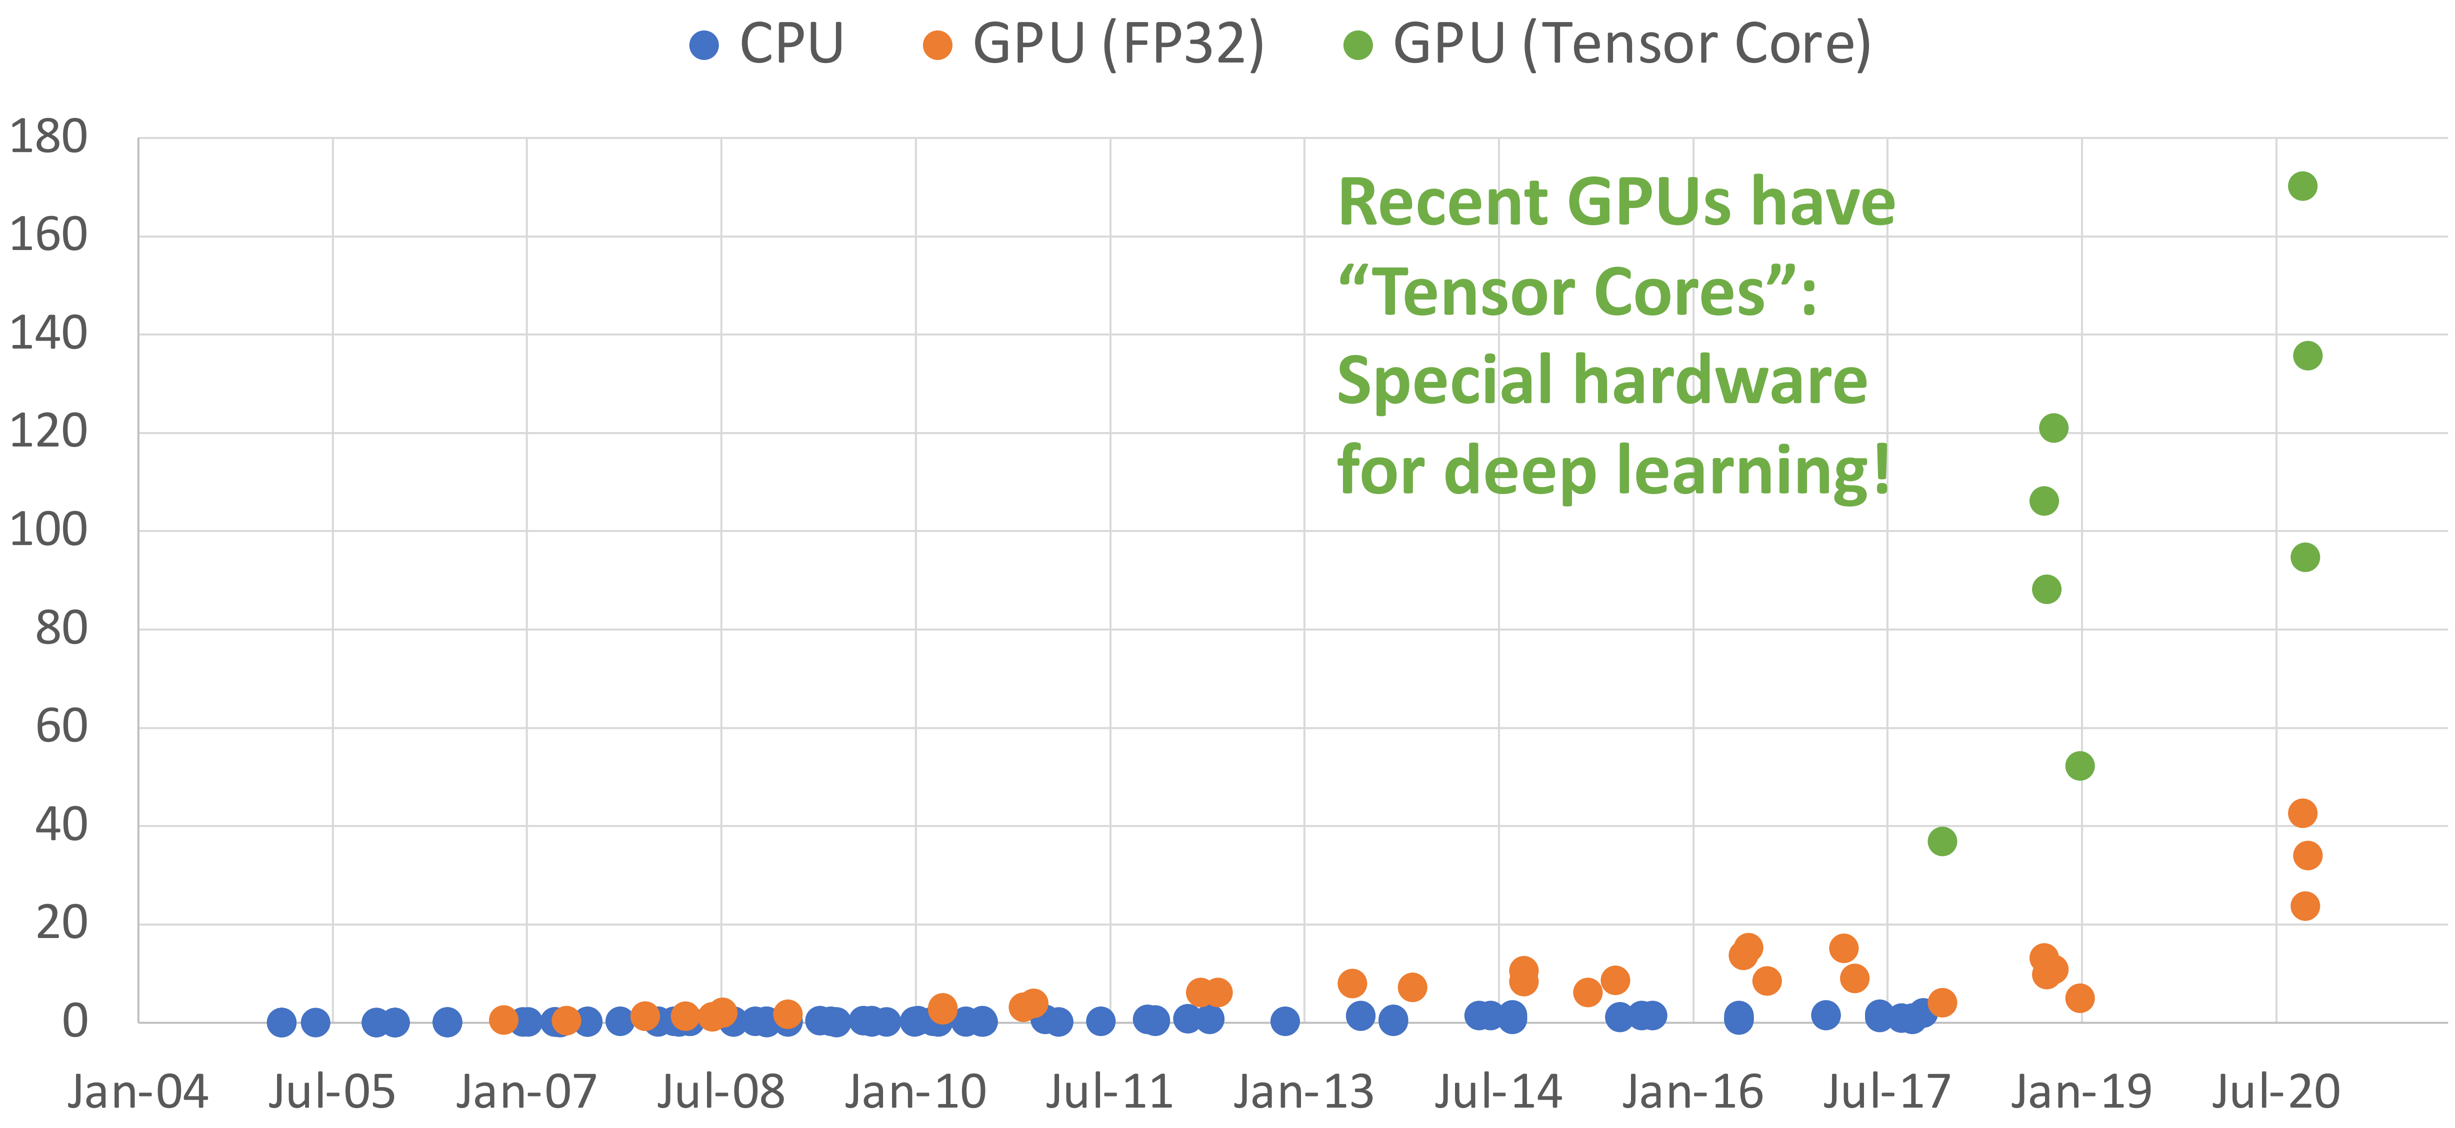

2018 Turing Award


Yoshua Bengio


Geoffrey Hinton


Yann LeCun

1959 Hubel & Wiesel
1963 Roberts
1970s David Marr
1979 Gen. Cylinders
1986 Canny
1997 Norm. Cuts
1999 SIFT
2001 V&J
2007 PASCAL
2009 ImageNet

AI Winter

1958 Perceptron
1969 Minsky & Papert
1980 Neocognitron
1985 Backprop
1998 LeNet
2006 Deep Learning
2012 AlexNet
2018 Turing Award

Despite our success, computer vision still has a long way to go…

Computer Vision can cause harm

Harmful Stereotypes

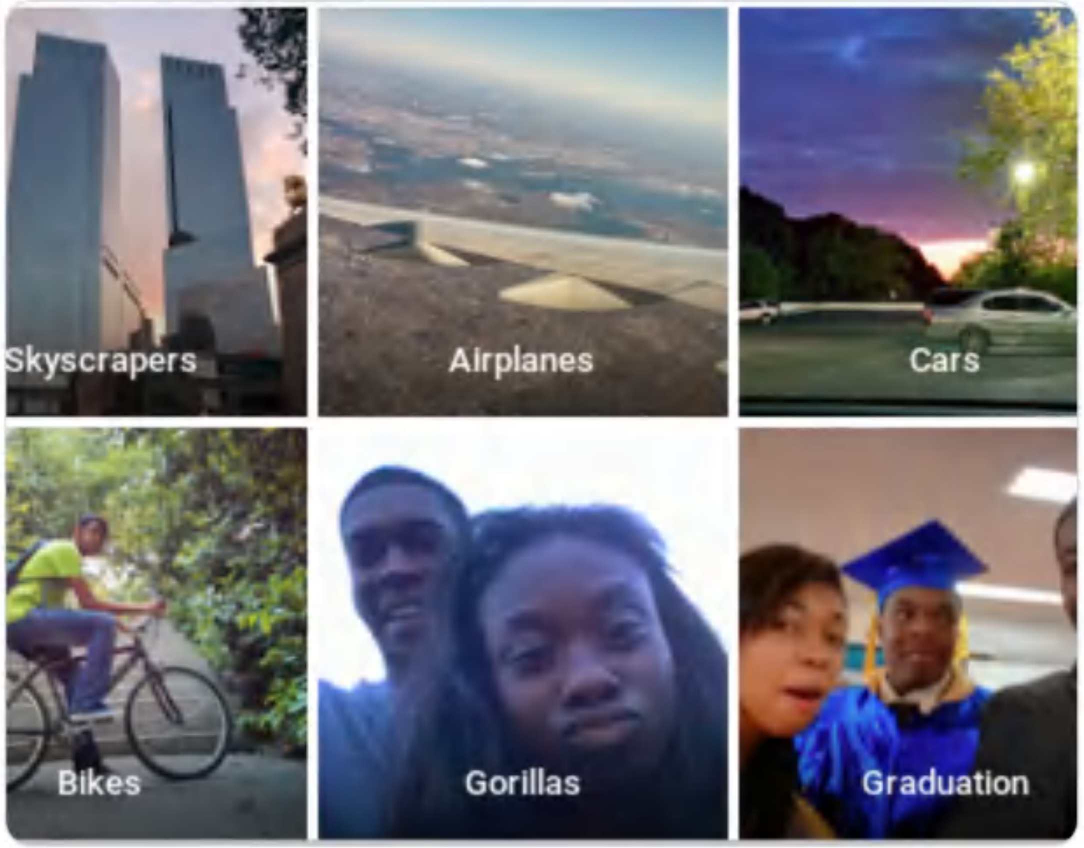

Barocas et al, “The Problem With Bias: Allocative Versus Representational Harms in Machine Learning”, SIGCIS 2017
Kate Crawford, “The Trouble with Bias”, NeurIPS 2017 Keynote
Source: https://twitter.com/jackyalcine/status/615329515909156865 (2015)

Affect people’s lives

Technology

A face-scanning algorithm increasingly decides whether you deserve the job

HireVue claims it uses artificial intelligence to decide who's best for a job. Outside experts call it 'profoundly disturbing.'

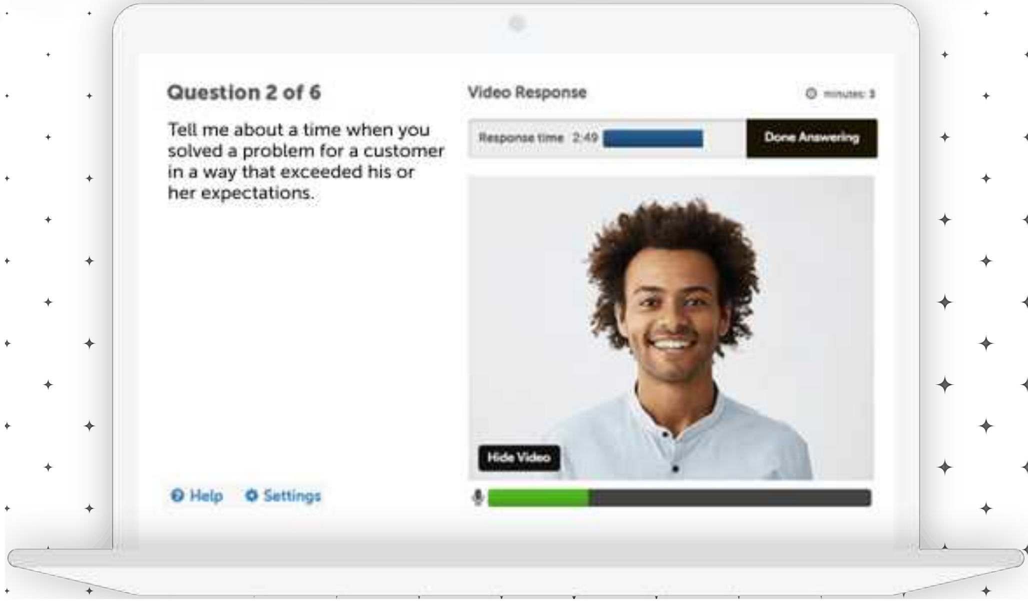

Source: https://www.washingtonpost.com/technology/2019/10/22/ai-hiring-face-scanning-algorithm-increasingly-decides-whether-you-deserve-job/
https://www.hirevue.com/platform/online-video-interviewing-software
Example Credit: Timnit Gebru

This image is copyright-free United States government work

Example credit: Andrej Karpathy

Computer Vision Technology Can Better Our Lives

Today's Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics

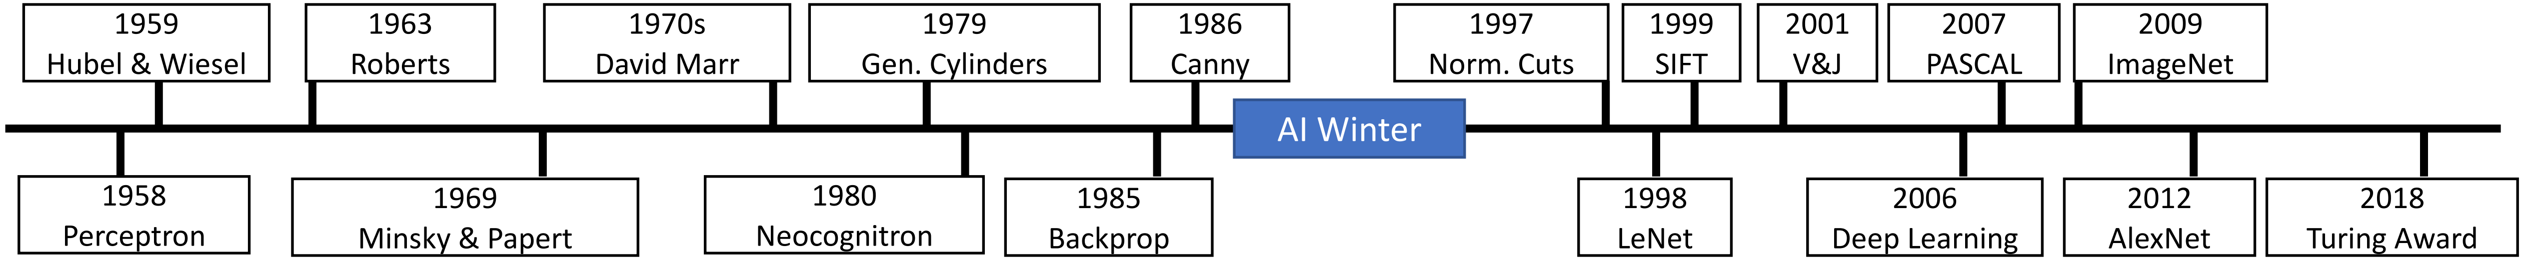

Today’s Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics

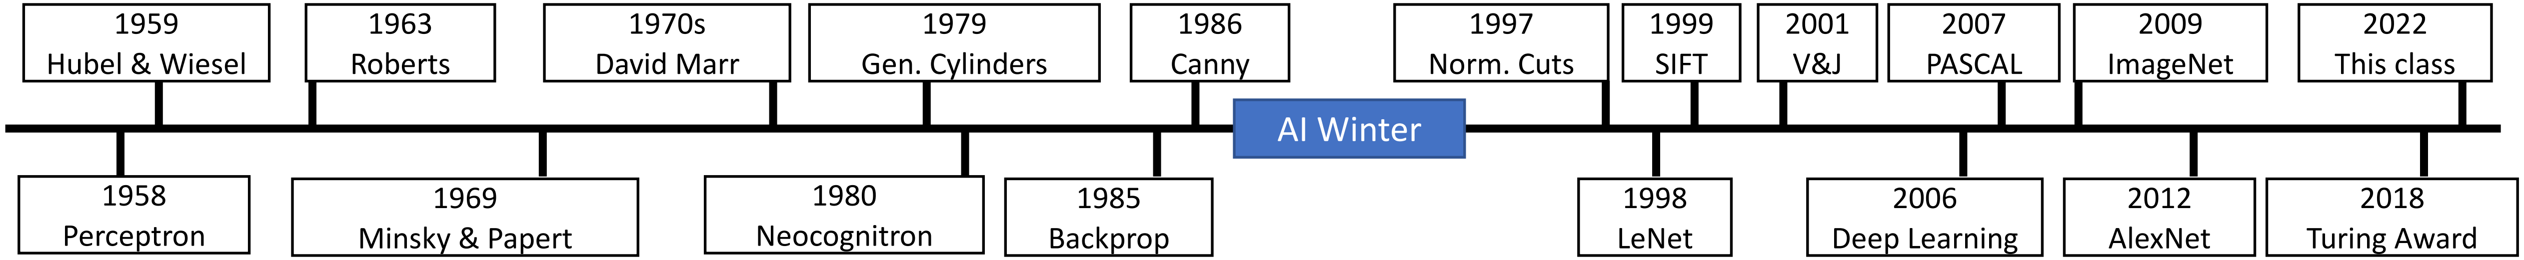

Course Staff

Instructor

Justin Johnson
Assistant Professor, CSE

GSIs / IAs

Karan Desai (KD)

Janpreet Singh (JS)

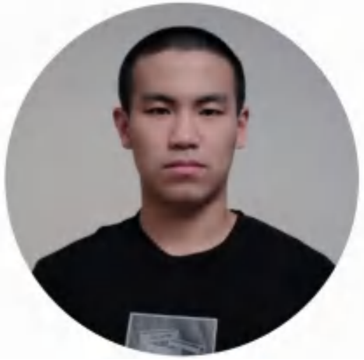

Jim Yang

Wallace Sui (WS)

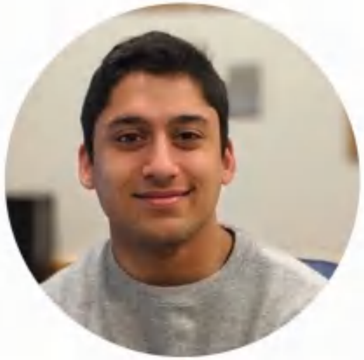

Gaurav Kaul

How to contact us

- Course Website: https://web.eecs.umich.edu/~justincj/teaching/eecs498/
  - Syllabus, schedule, assignments, slides, lecture videos, etc
- Piazza: https://piazza.com/class/kxtai72amx34p0
  - (Almost) all questions about the course should go here!
  - We will also use Piazza to communicate with you
  - Use private questions if you want to post code
- EECS Autograder: https://autograder.io/web/course/151
  - For turning in homework assignments
- Google Calendar: For office hours (starting next week)
- Email: Only for sensitive, confidential issues

Course Website: Check the Schedule!

EECS 498.008 / 598.008
Deep Learning for Computer Vision
Winter 2022

UNIVERSITY OF MICHIGAN

Schedule

Lectures are Mondays and Wednesdays, 4:30pm to 6pm. Attendance is not required. Recordings will be posted after each lecture in case you are unable the attend the scheduled time.

Some lectures have reading drawn from the course notes of Stanford CS 231n, written by Andrej Karpathy.

Some lectures have optional reading from the book *Deep Learning* by Ian Goodfellow, Yoshua Bengio, and Aaron Courville (GBC for short). The entire text of the book is available for free online so you don't need to buy a copy.

Event	Date	Description	Course Materials
Lecture 1	Wednesday January 5	**Course Introduction** Computer vision overview Historical context Course logistics	[slides] [FA2019 video (public)] [Python tutorial] [GBC Sec 1.2] [GBC Sec 6.6]
Lecture 2	Monday January 10	**Image Classification** Data-driven approach K-Nearest Neighbor Hyperparameters Cross-validation	[slides] [FA2019 video (public)] [231n Image Classification]
Lecture 3	Wednesday January 12	**Linear Classifiers** Algebraic / Visual / Geometric viewpoints Softmax / SVM classifiers	[slides] [FA2019 video (public)] [231n Linear Classification]
A1 Due	Friday January 14	**Assignment 1 Due** PyTorch warmup kNN Classifier	[Assignment 1]
	Monday January 17	**No class** MLK Day	

https://web.eecs.umich.edu/~justincj/teaching/eecs498/WI2022/schedule.html

Piazza Etiquette

- Post only short snippets of code (< 20-30 lines)
- Ask a specific, concrete question
- Explain what you have tried so far, and what happened
- See StackOverflow guide on asking good questions: https://stackoverflow.com/help/how-to-ask

Piazza Etiquette

- Post only short snippets of code (< 20-30 lines)
- Ask a specific, concrete question
- Explain what you have tried so far, and what happened
- See StackOverflow guide on asking good questions: https://stackoverflow.com/help/how-to-ask

- Don't expect and answer within 30 minutes of posting
- Monday – Friday, 10am – 6pm EST we'll try to answer within 2 hours
- Other times, we'll try to answer within 12 hours

Optional Textbook

- *Deep Learning* by Goodfellow, Bengio, and Courville
- Free online

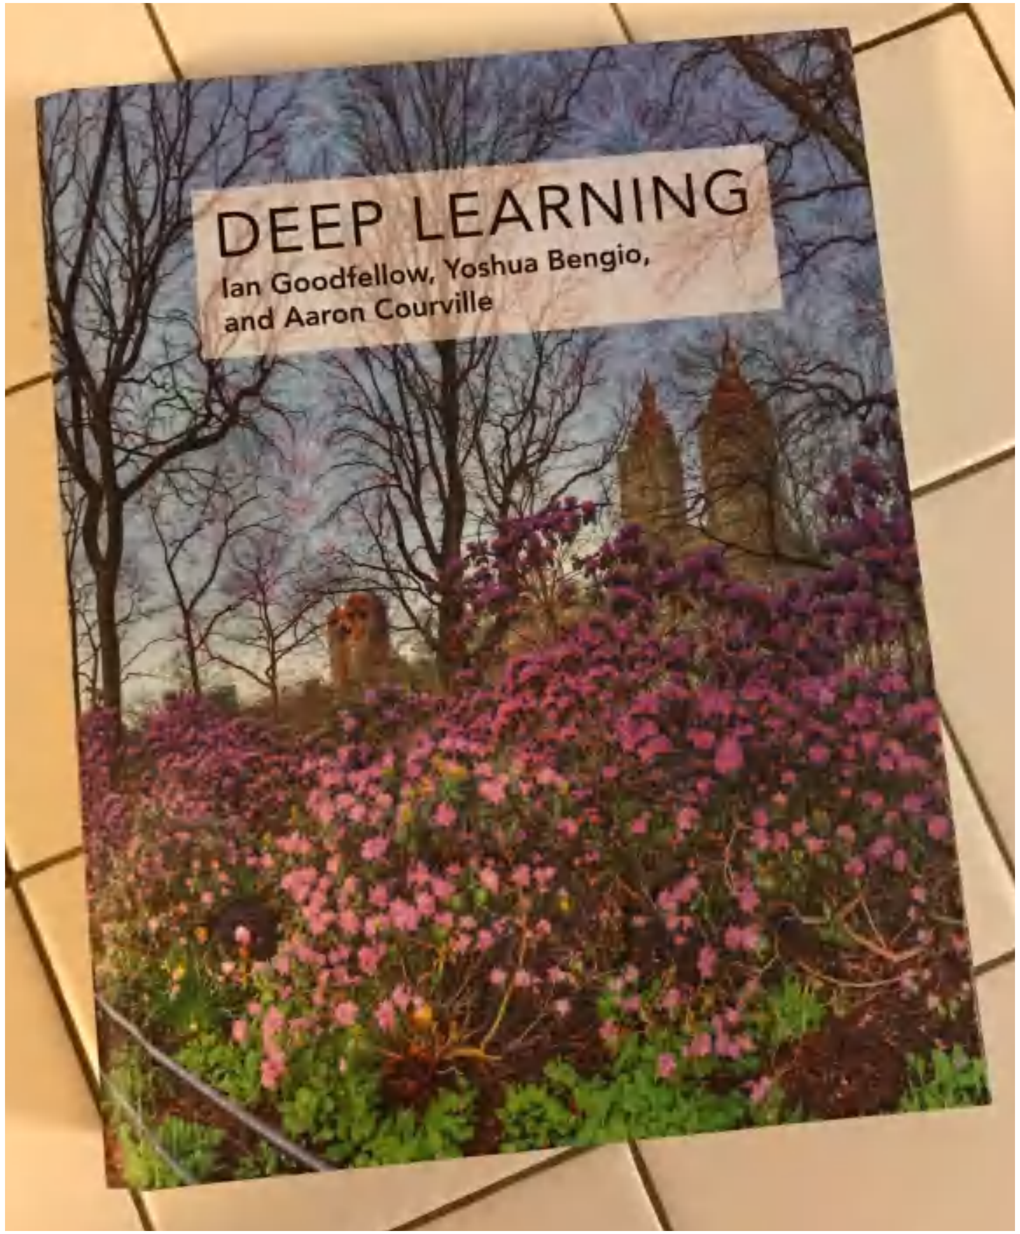

Course Content and Grading

- 6 programming assignments (10% each)
- Midterm Exam (20%)
- Mini-Project (20%)
- Late policy
  - 3 free late days to use on assignments
  - Once free late days are exhausted, 25% penalty per day

Programming Assignments

- Python, PyTorch, will use Google Colab
- “Earn your wings” – implement things “from scratch” in early assignments, then use PyTorch in later assignments
- ”Challenge Questions”
  - Go above and beyond the basic expectations of each assignment
  - Much higher time/points ratio than other parts of the assignment
  - Not necessary to get an A: will be 5% or less of each assignment

Midterm Exam

- Written exam testing basic concepts from first half of course
- True / False, Multiple choice, short answer
- We will provide a practice exam to get a sense of the format of the questions (but not necessarily the length or difficulty)

Mini-Project

- New this year!
- Work in teams of 1 to 3
- Homework assignments: We provide starter code, you “fill in the blank”
- Mini-Project: We provide high-level written description of an algorithm to implement, you do the whole thing “from scratch”
- We will give ~3 project descriptions, you pick one of those
- Deliverable: zip of code, and Colab notebook that walks through your implementation and main results (should be like notebooks from homework!)

Collaboration Policy

- **Rule 1**: Don’t look at solutions or code that are not your own; everything you submit should be your own work
- **Rule 2**: Don’t share your solution code with others; however discussing ideas or general strategies is fine and encouraged
- **Rule 3**: Indicate in your submissions anyone you worked with
- Turning in something late / incomplete is better than violating the honor code

Course Philosophy

- Thorough and Detailed.
  - This not “Learn PyTorch in 90 days”, nor “Deep Learning in 10 lines of code”
  - Understand how to write from scratch, debug, and train convolutional and other types of deep neural networks
  - We prefer to write from scratch, rather than rely on existing implementations
- Practical
  - Focus on practical techniques for training and debugging neural networks
  - Will use state-of-the-art software tools like PyTorch and TensorFlow
- State of the art
  - Most material we cover is research published in the last 5 years

Course Philosophy

- Will also cover some fun topics:
  - Image captioning
  - DeepDream, Artistic Style Transfer

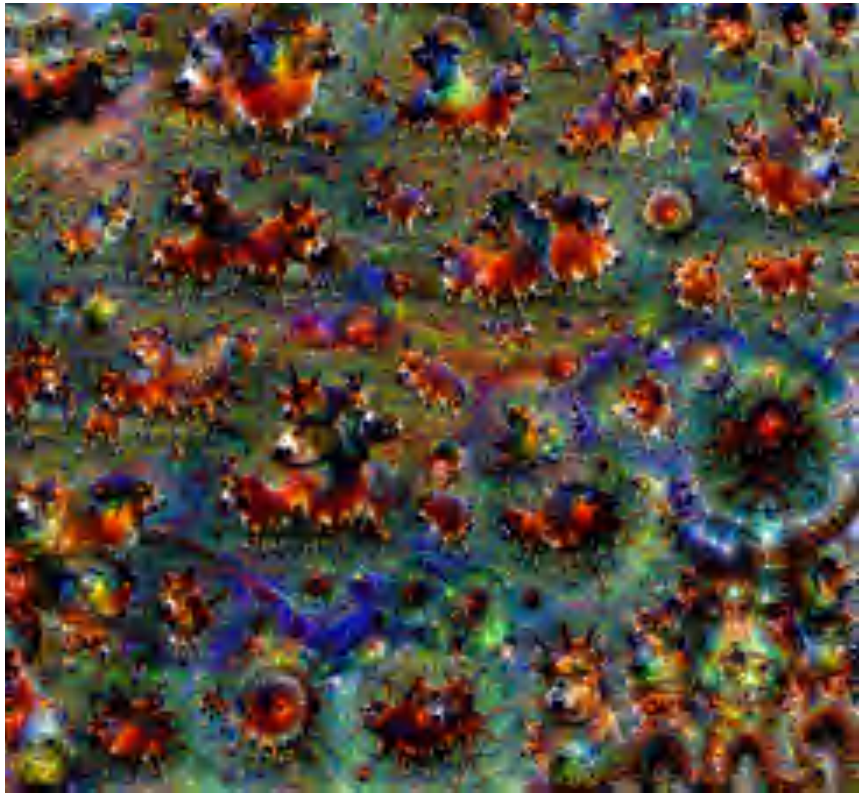

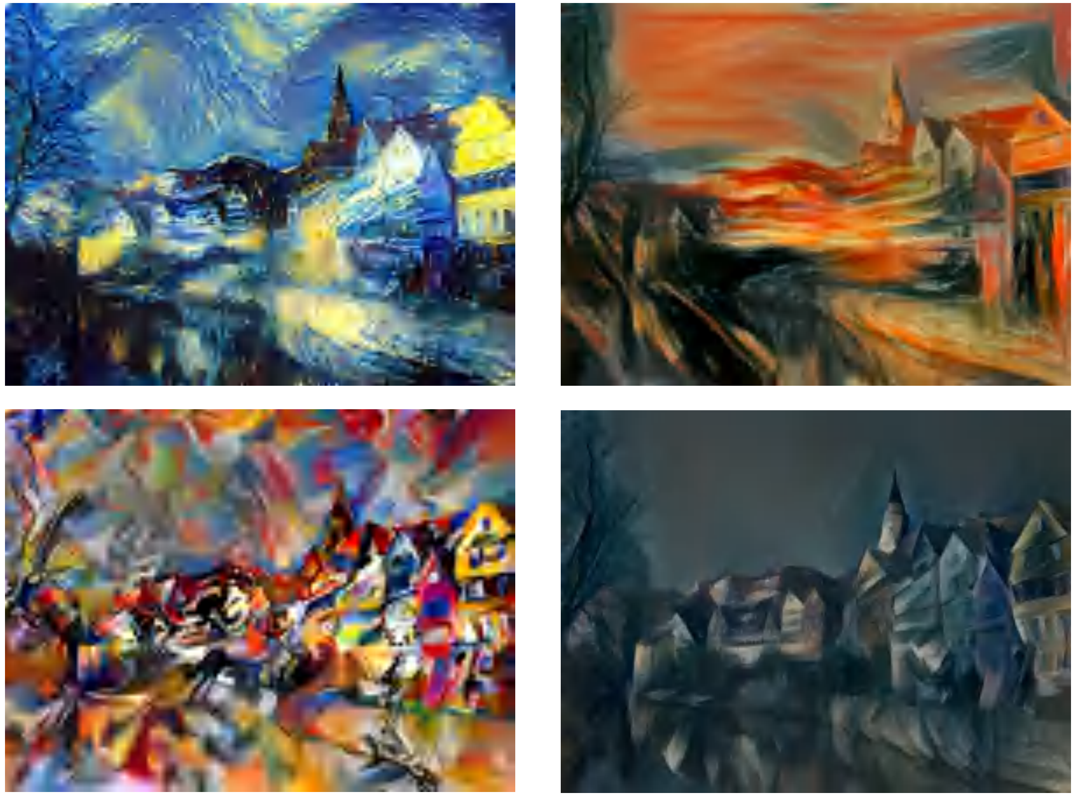

Course Structure

- First half: Fundamentals
  - Details of how to implement and train different types of networks
  - Fully-connected networks, convolutional networks, recurrent networks
  - How to train and debug, very detailed
- Second half: Applications and “Researchy” topics
  - Object detection, image segmentation, 3D vision, videos
  - Attention, Transformers
  - Vision and Language
  - Generative models: GANs, VAEs, etc
  - Less detailed: provide overview and references, but skip some details

New Topics since FA2020

- Modern CNN architectures
  - SENets, MobileNets, NAS
  - EfficientNets, RegNets, NFNets
- Vision Transformers
  - Architectures: ViT, DeiT, Swin, MViT
  - Applications: DETR
  - MLP-like architectures
- Vision + Language
  - Language-based pretraining: CLIP, ALIGN
- Self-Supervised Learning
  - Contrastive learning
  - Masked autoencoding

First homework assignment

- Will be released by today or tomorrow
- Due Friday 1/14/2022
- Next lecture will be enough to complete it

Next time: Image Classification