


Lecture 14:

Object Detectors

Poll Results



- Option 1: Keep mini-project, only 1.5 weeks between each of HW4, HW5, HW6, and project
- Option 2: Cancel mini-project, allowing for 2 weeks between each of HW4, HW5, and HW6

Many comments / suggestions in comments and on Piazza:

- Option 2: Want more weight on HW4-6, less on midterm
- Optional project
- Drop one HW assignment
- Extra late days

Decision

- We will keep 2-week gap between each of HW4-6
- Students can also complete a project if they wish (spec out next week)
- Each student can choose one of the following options:

Decision

- We will keep 2-week gap between each of HW4-6
- Students can also complete a project if they wish (spec out next week)
- Each student can choose one of the following options:

Option A:

Do all assignments,
Do not do project.

Grading scheme:

HW1-3: 12%

Midterm: 22%

HW4-6: 14%

Decision

- We will keep 2-week gap between each of HW4-6
- Students can also complete a project if they wish (spec out next week)
- Each student can choose one of the following options:

Option A:

Do all assignments,
Do not do project.

Grading scheme:

HW1-3: 12%

Midterm: 22%

HW4-6: 14%

Option B:

Do 5 or 6 assignments
Do project

Grading scheme (whichever gives you better grade):

HW1-3: 12%

Midterm: 22%

HW4-6: 14%

Project: Replaces lowest HW

Original grading scheme:

HW1-6: 10%

Midterm: 20%

Project: 20%

In addition: Everyone gets +3 late days
(cannot be applied to A6 or project)

Decision

- We will keep 2-week gap between each of HW4-6
- Students can also complete a project if they wish (spec out next week)
- Each student can choose one of the following options:

Option A:

Do all assignments,
Do not do project.

Grading scheme:

HW1-3: 12%

Midterm: 22%

HW4-6: 14%

Option B:

Do 5 or 6 assignments
Do project

Grading scheme (whichever gives you better grade):

HW1-3: 12%

Midterm: 22%

HW4-6: 14%

Project: Replaces lowest HW

Original grading scheme:

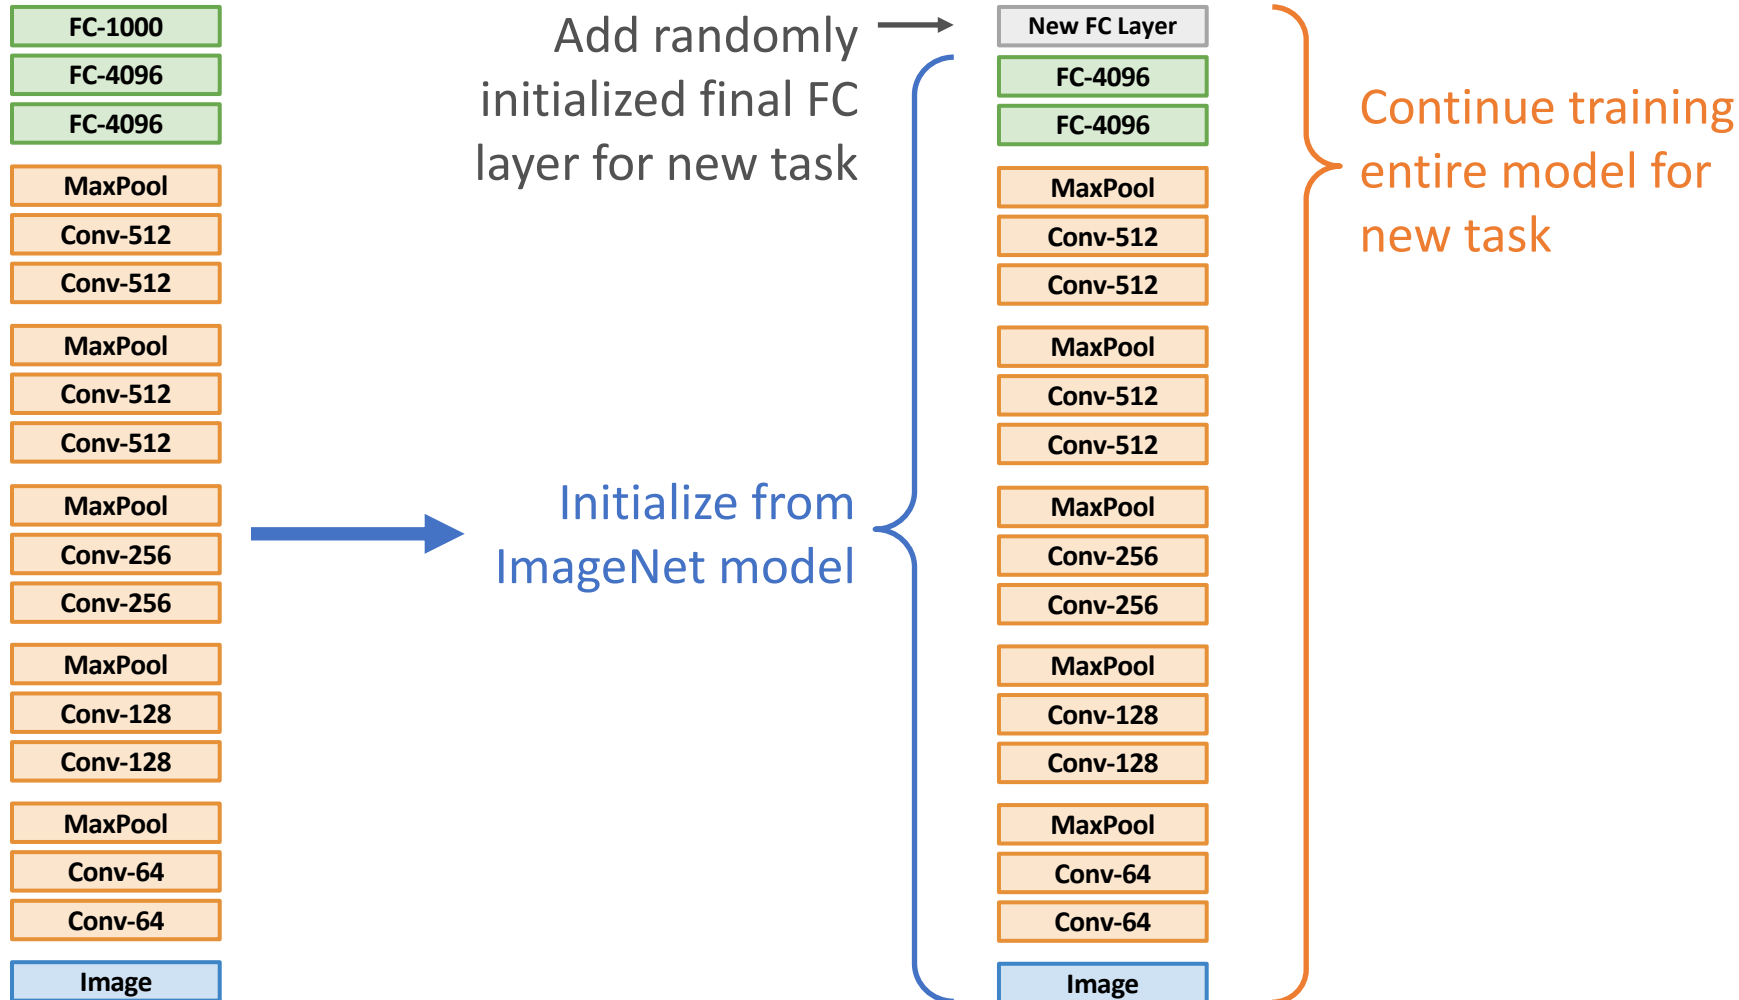
HW1-6: 10%

Midterm: 20%

Project: 20%

Last Time: Transfer Learning

1. Train on ImageNet



Last Time: Localization Tasks


Classification



CAT

No spatial extent


Semantic Segmentation



GRASS, CAT, TREE,
SKY

No objects, just pixels

Object Detection



DOG, DOG, CAT

Multiple Objects

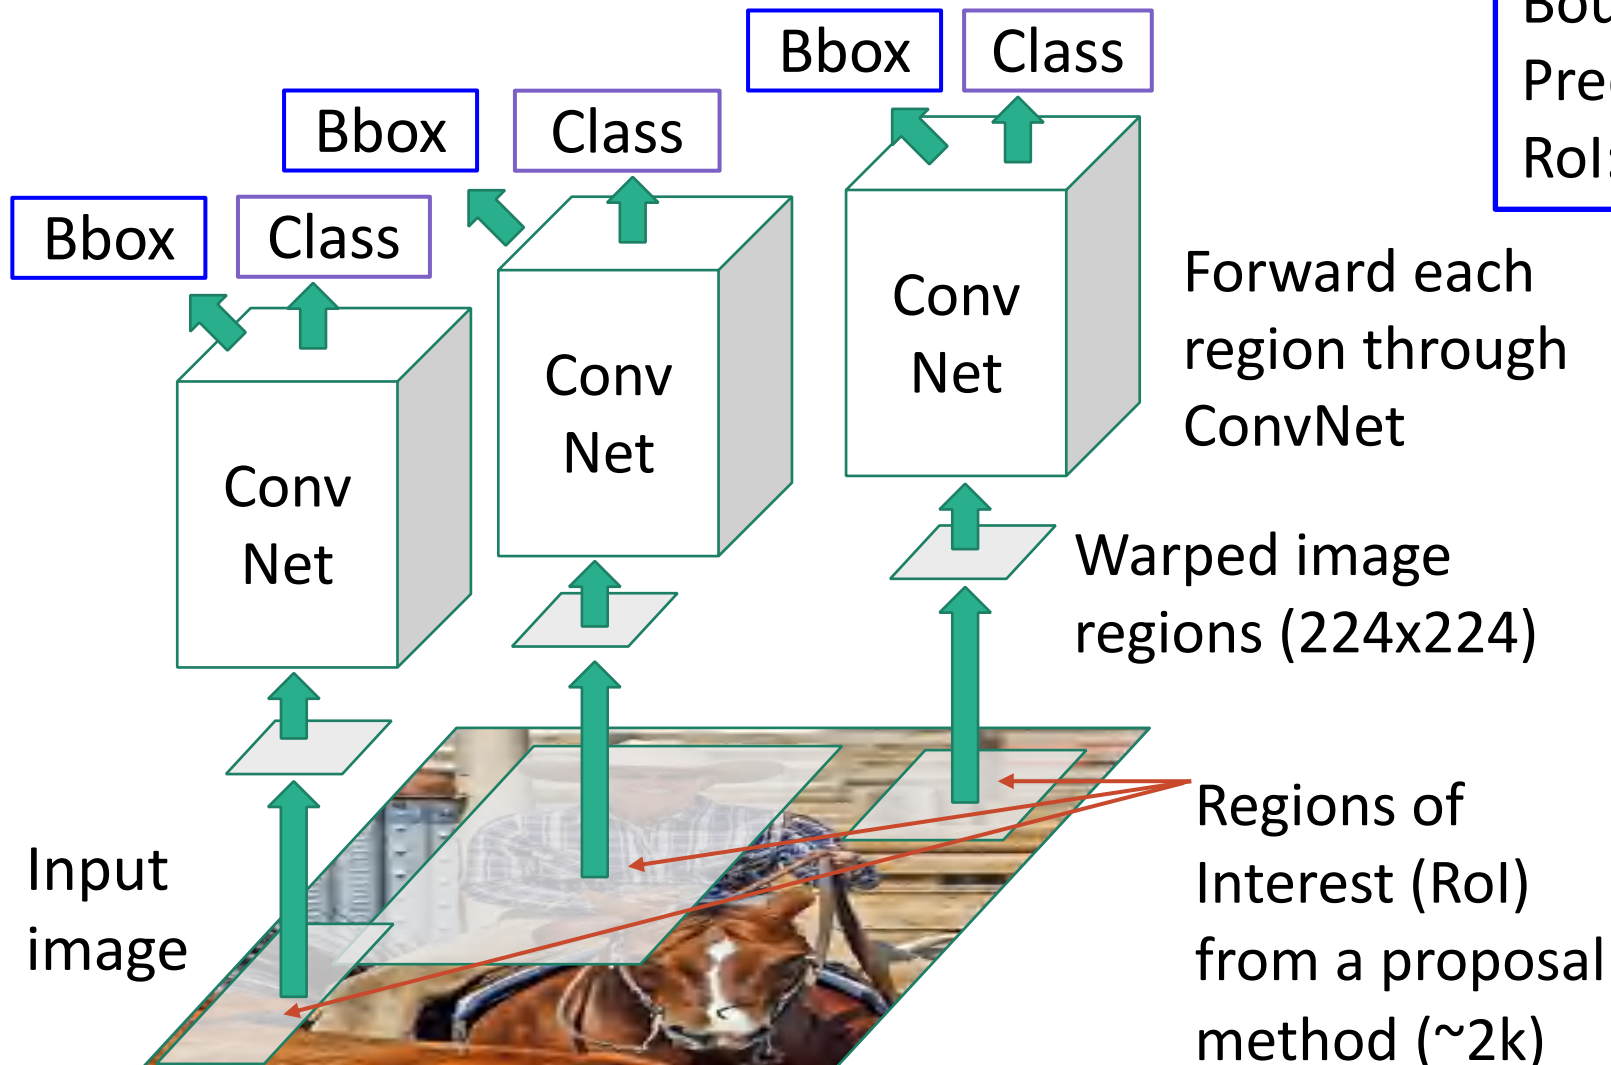
Instance Segmentation



DOG, DOG, CAT

[This image](#) is [CC0 public domain](#)

Last Time: R-CNN




Classify each region

Bounding box regression:
Predict "transform" to correct the
RoI: 4 numbers (t_x , t_y , t_h , t_w)

Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.
Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

Last Time: R-CNN



Classify each region

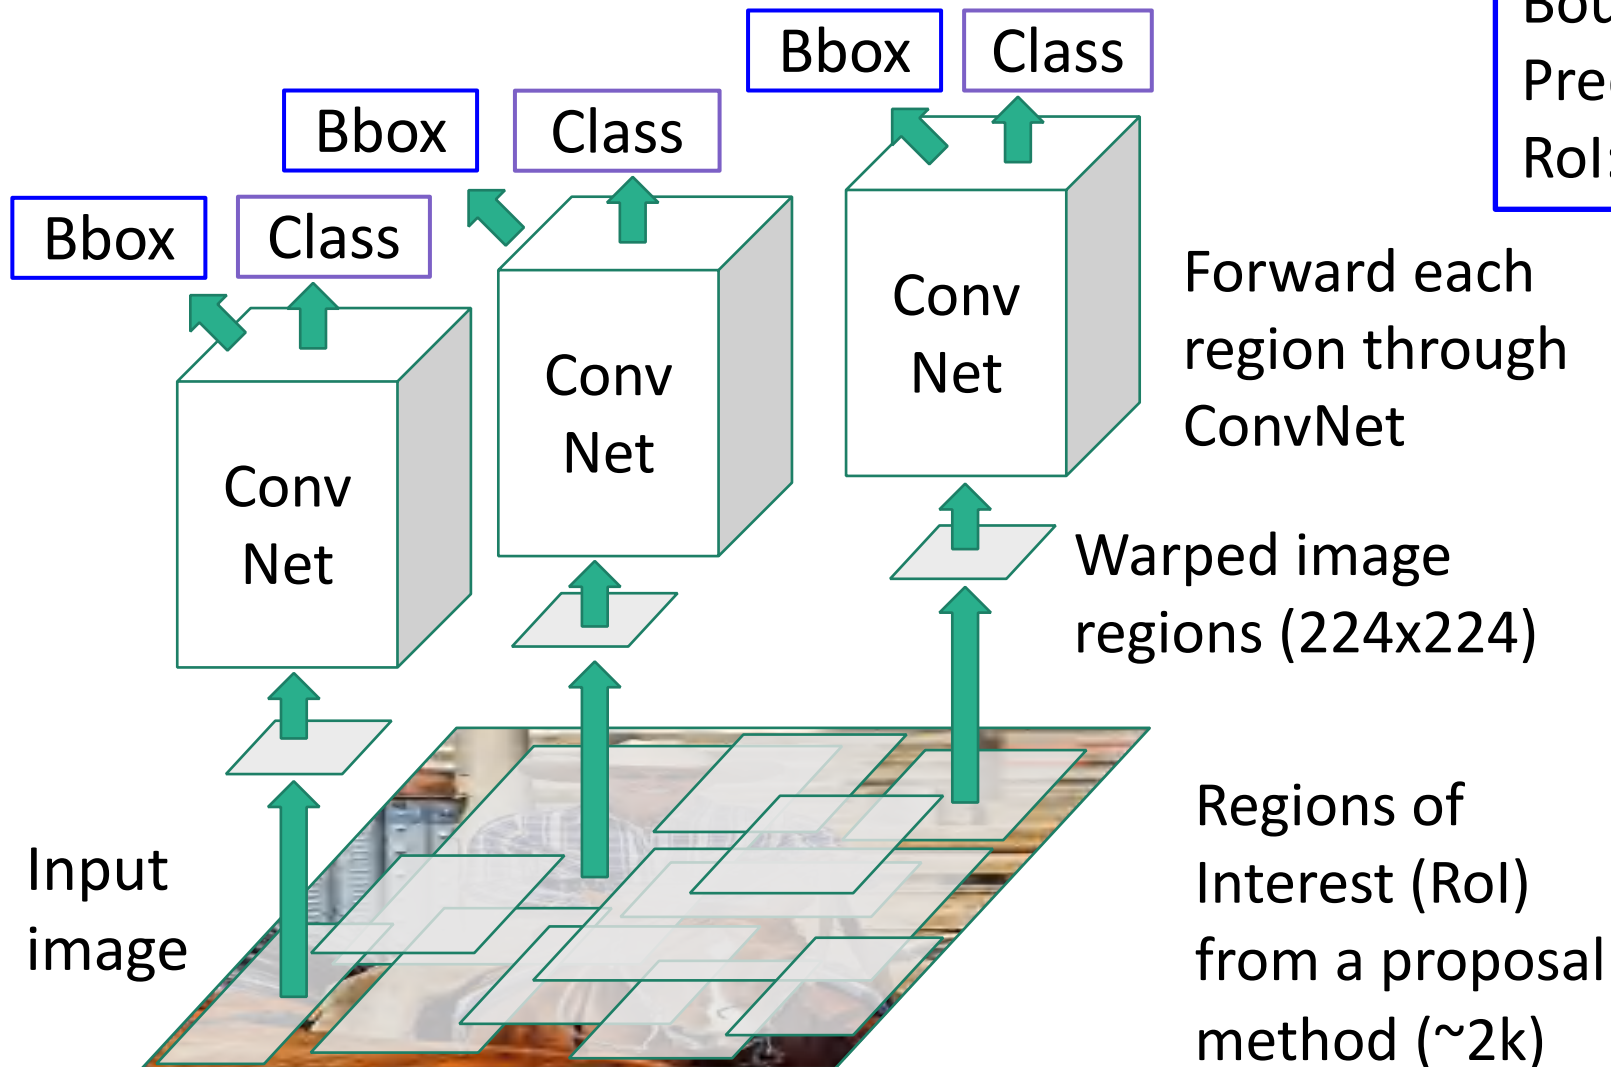
Bounding box regression:
Predict "transform" to correct the RoI: 4 numbers (t_x, t_y, t_h, t_w)

Forward each region through ConvNet

Problem: Very slow! Need to do 2000 forward passes through CNN per image

Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.
Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

Last Time: R-CNN



Classify each region

Bounding box regression:
Predict "transform" to correct the
RoI: 4 numbers (t_x , t_y , t_h , t_w)

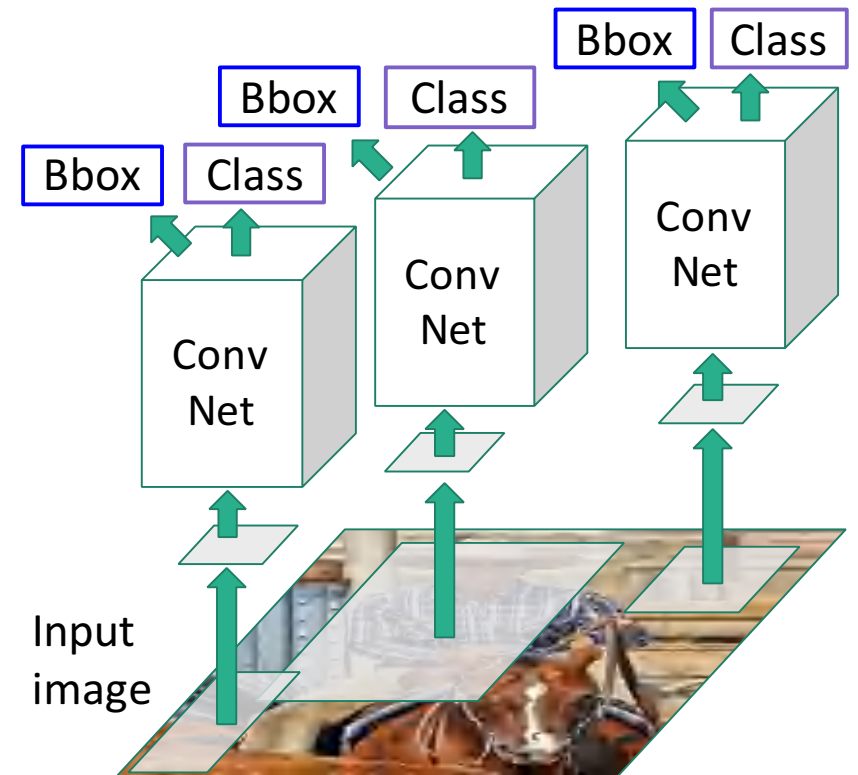
Problem: Very slow! Need to do 2000 forward passes through CNN per image

Idea: Overlapping proposals cause a lot of repeated work: same pixels processed many times. Can we avoid this?

Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.
Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

“Slow” R-CNN

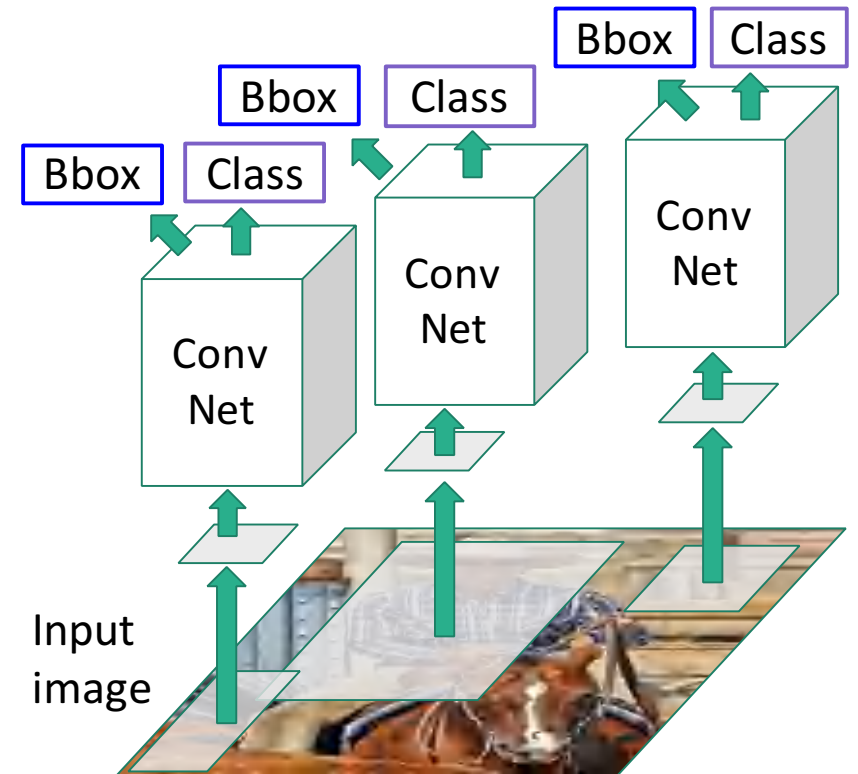
Process each region
independently



Fast R-CNN

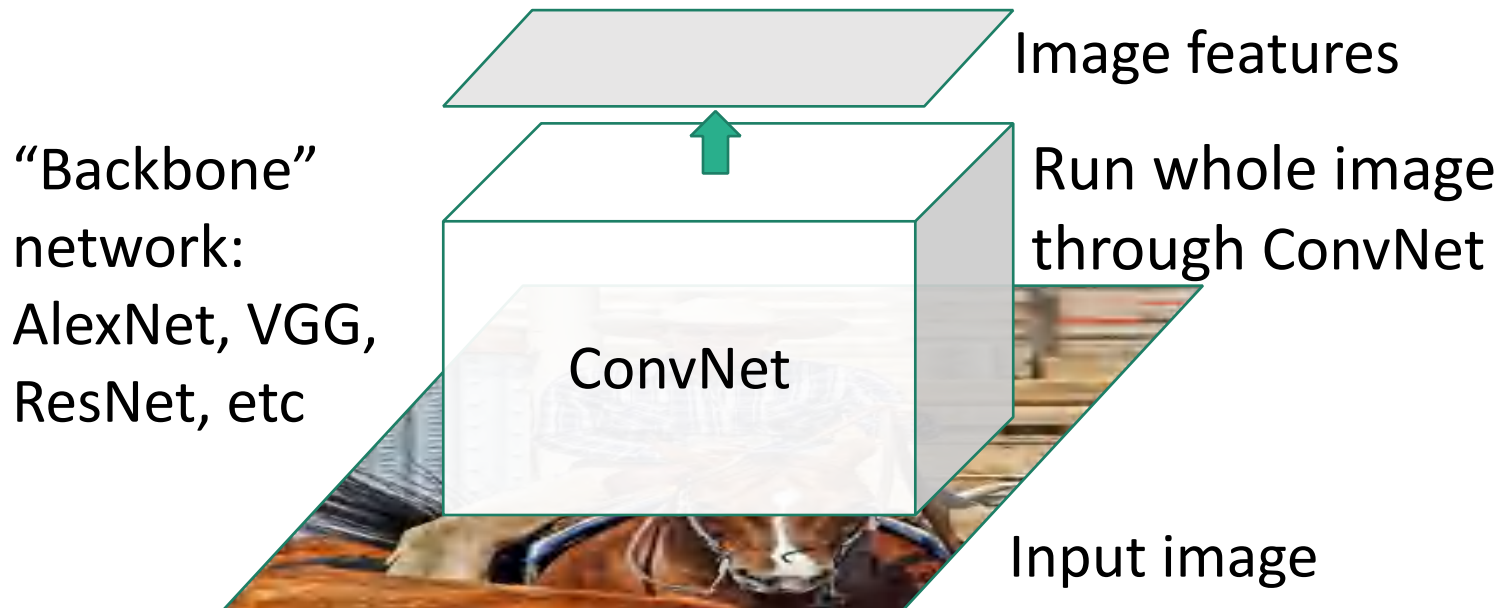


Input image

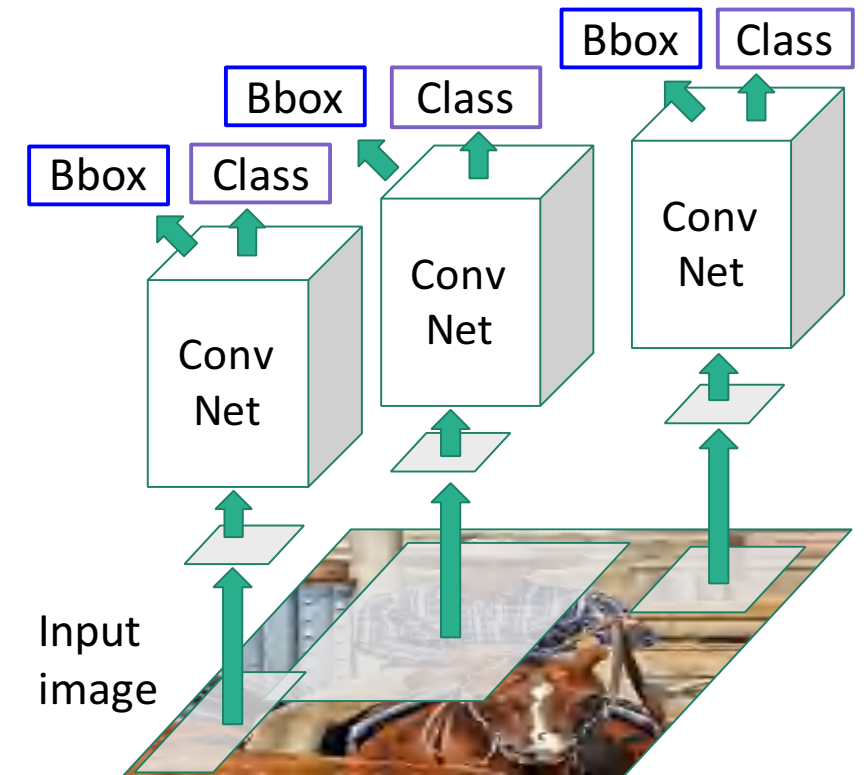


“Slow” R-CNN
Process each region
independently

Fast R-CNN



“Slow” R-CNN
Process each region
independently

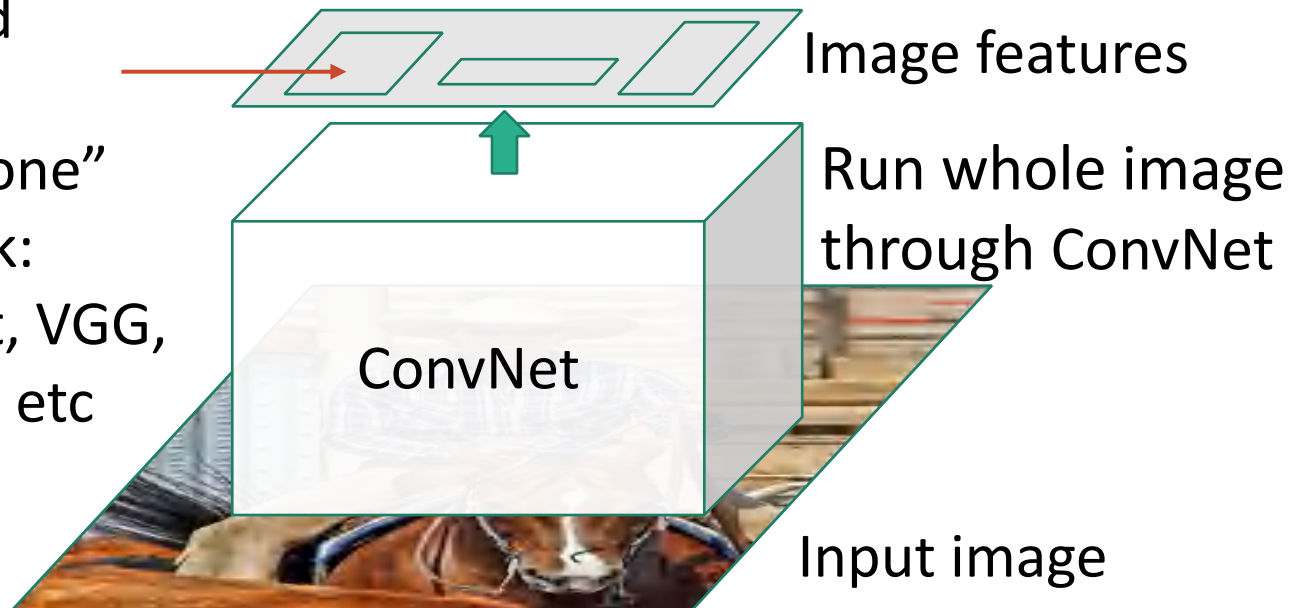


Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

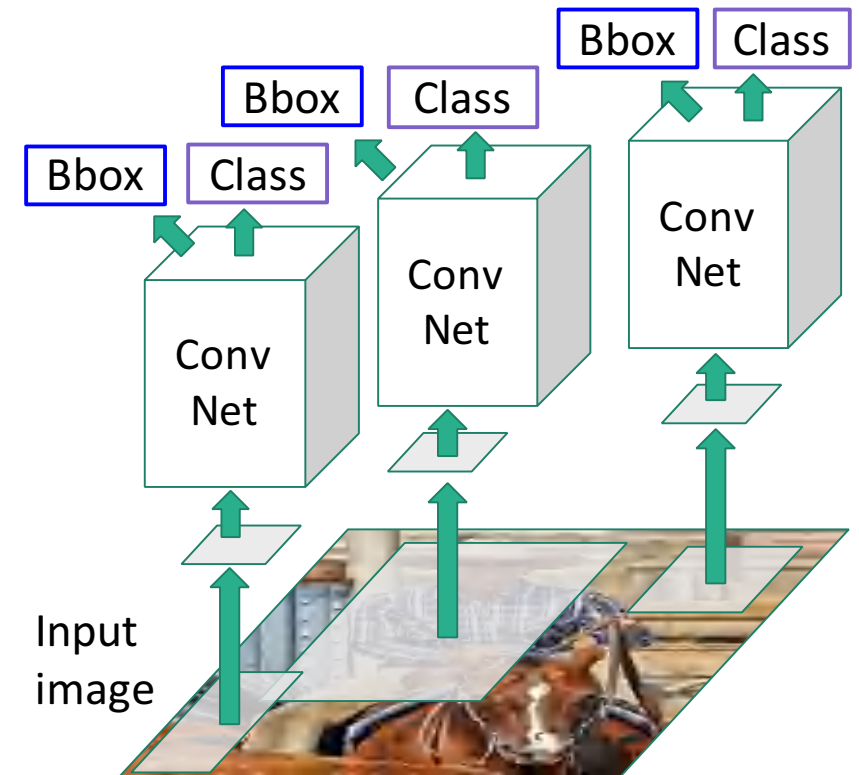
Fast R-CNN

Regions of Interest (RoIs) from a proposal method

“Backbone” network:
AlexNet, VGG, ResNet, etc



“Slow” R-CNN
Process each region independently

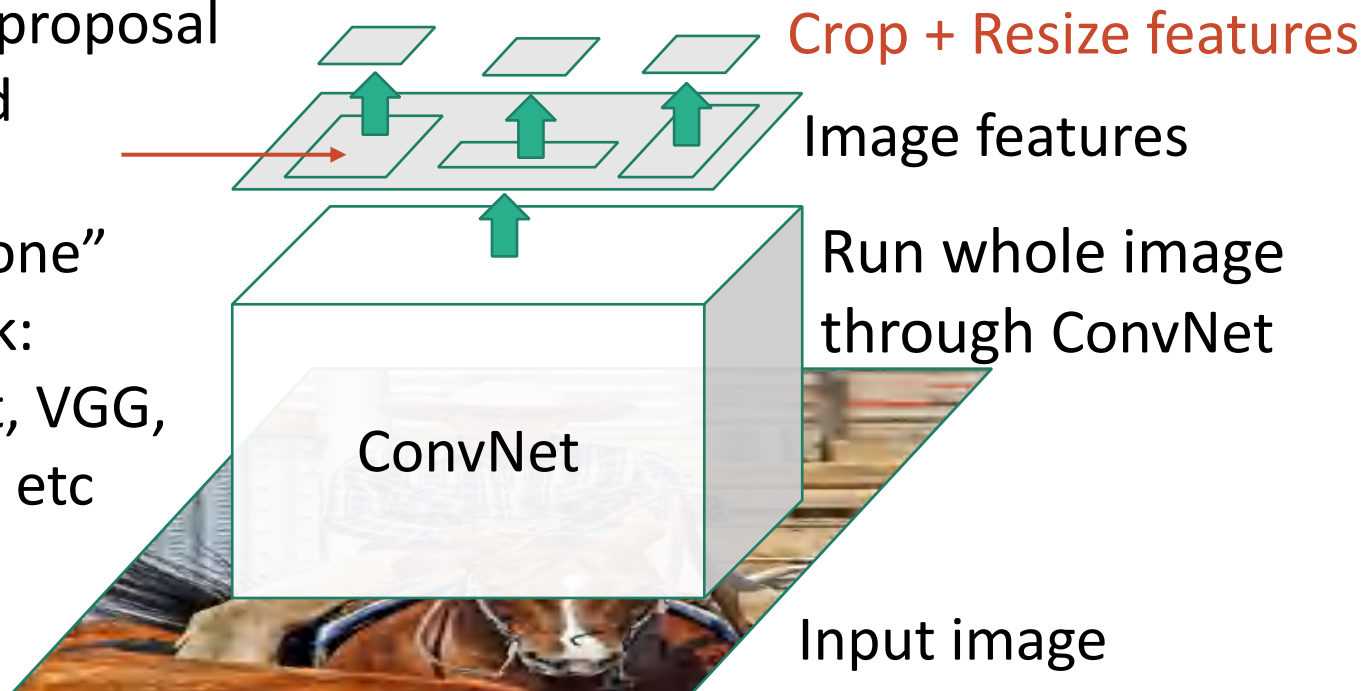


Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

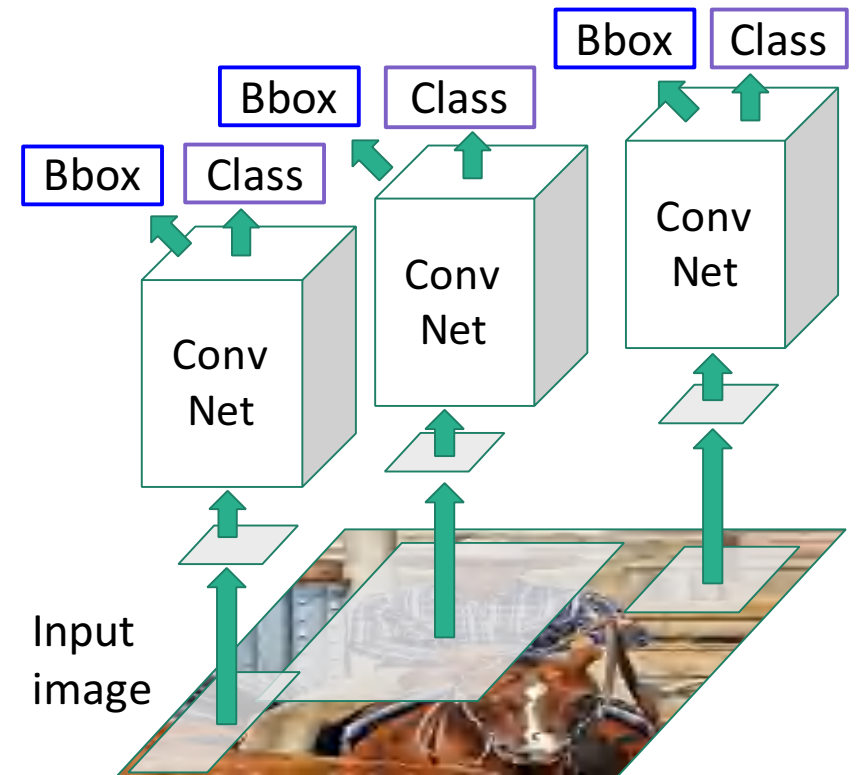
Fast R-CNN

Regions of Interest (RoIs) from a proposal method

“Backbone” network:
AlexNet, VGG, ResNet, etc



“Slow” R-CNN
Process each region independently

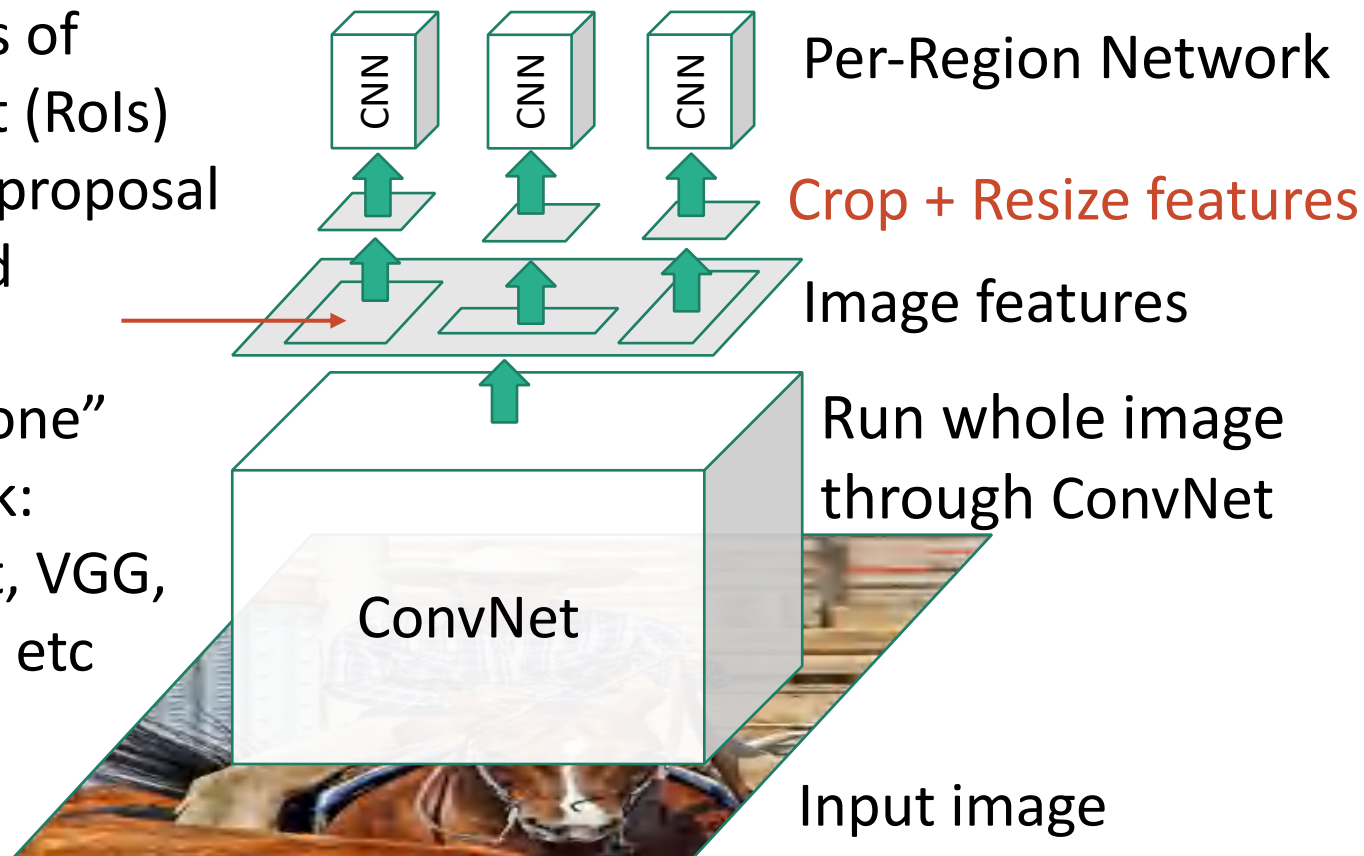


Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

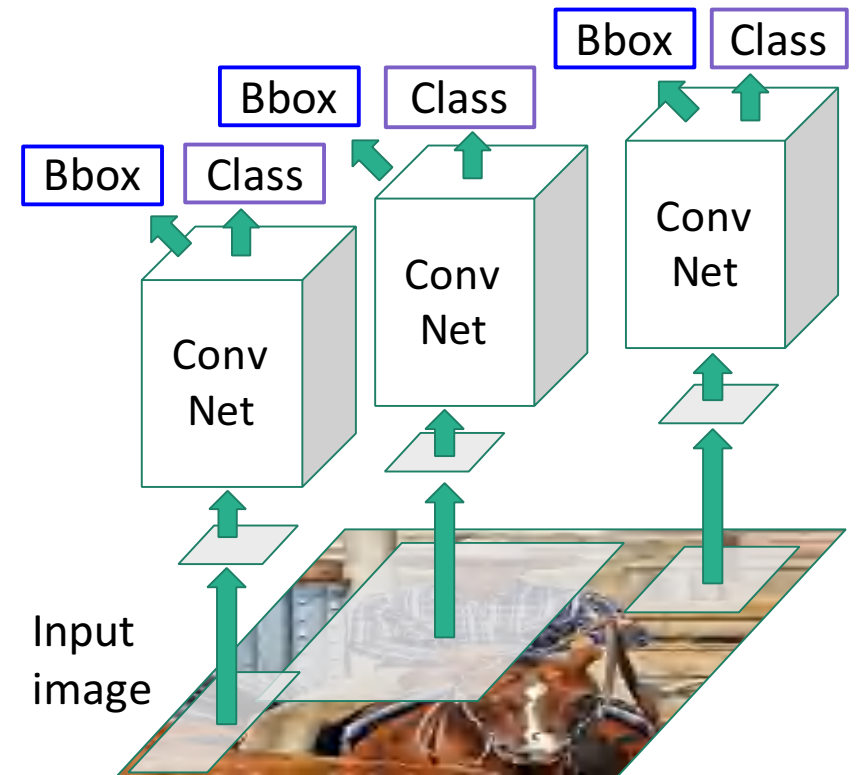
Fast R-CNN

Regions of Interest (RoIs) from a proposal method

“Backbone” network:
AlexNet, VGG, ResNet, etc

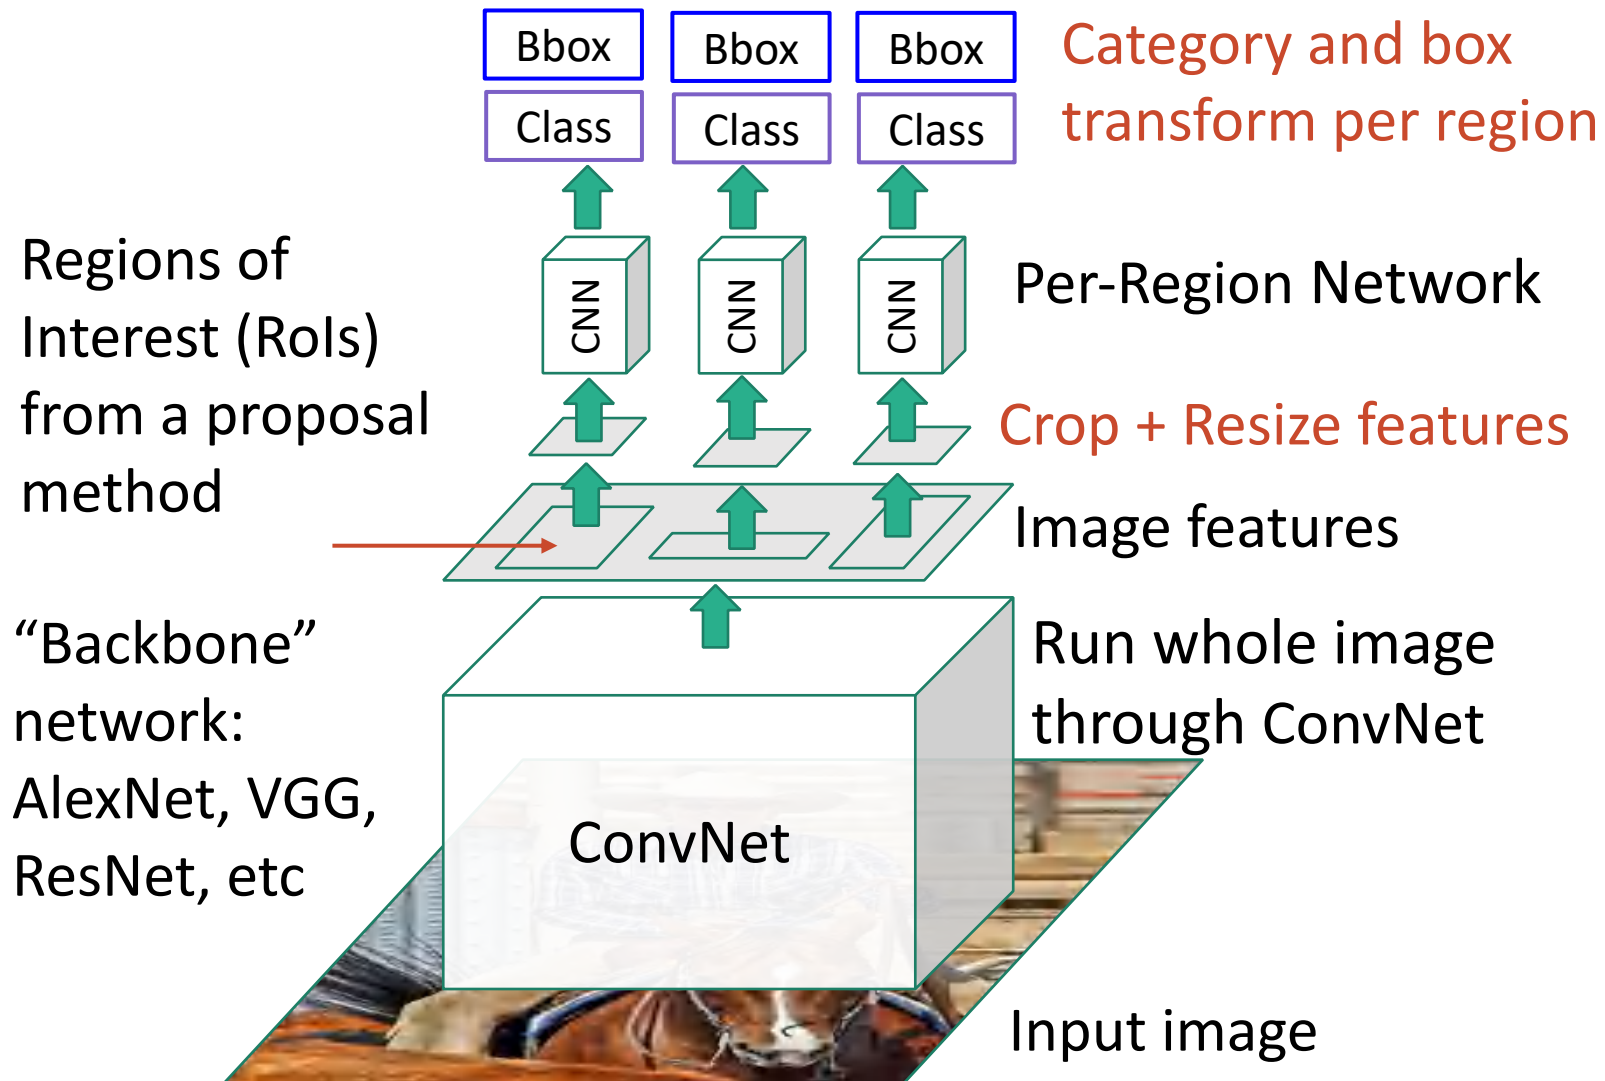


“Slow” R-CNN
Process each region independently




Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

Fast R-CNN

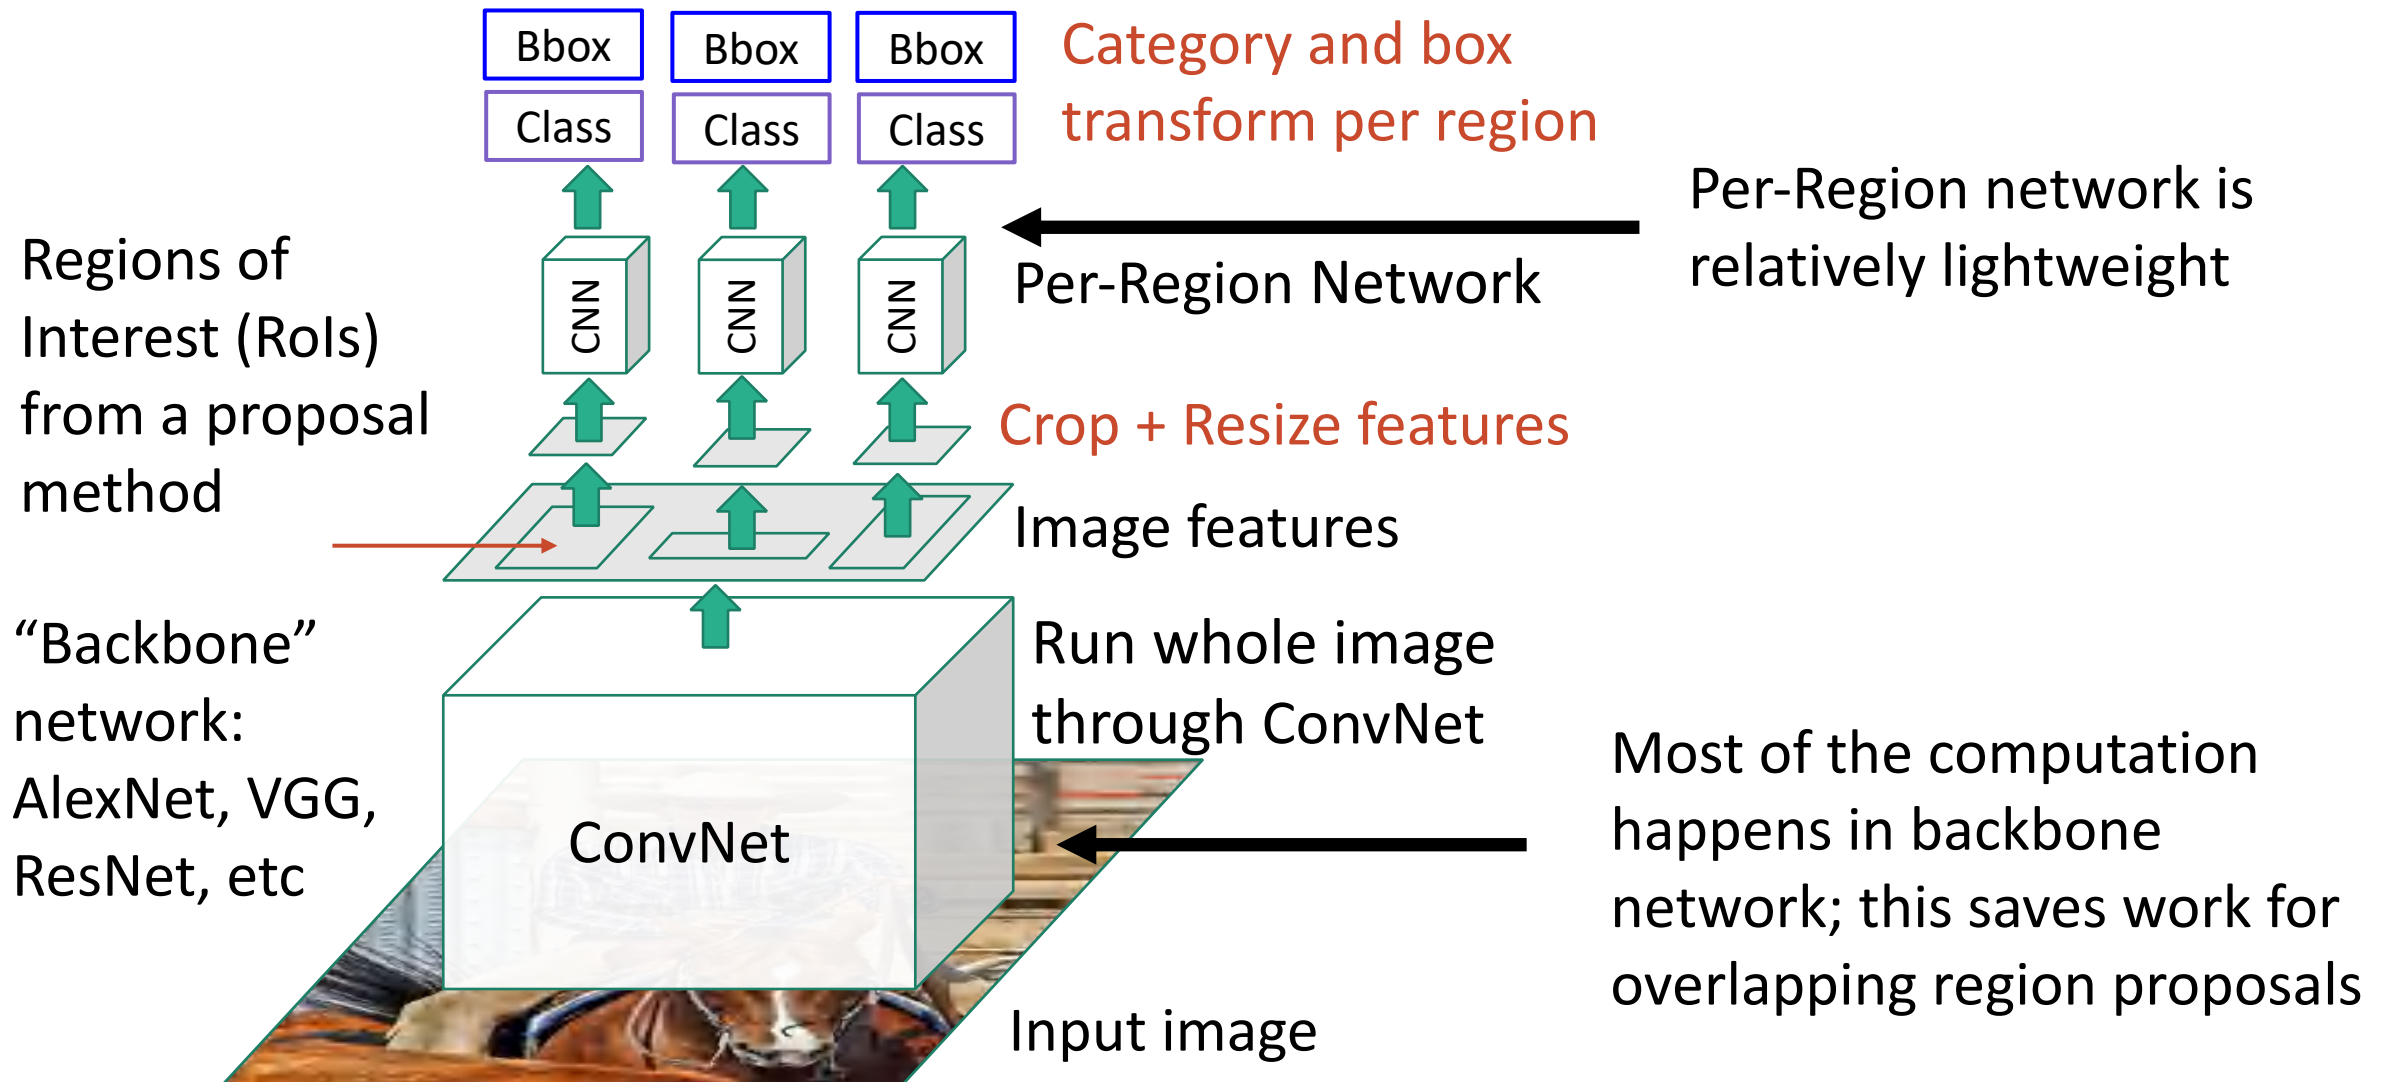


“Slow” R-CNN
Process each region independently




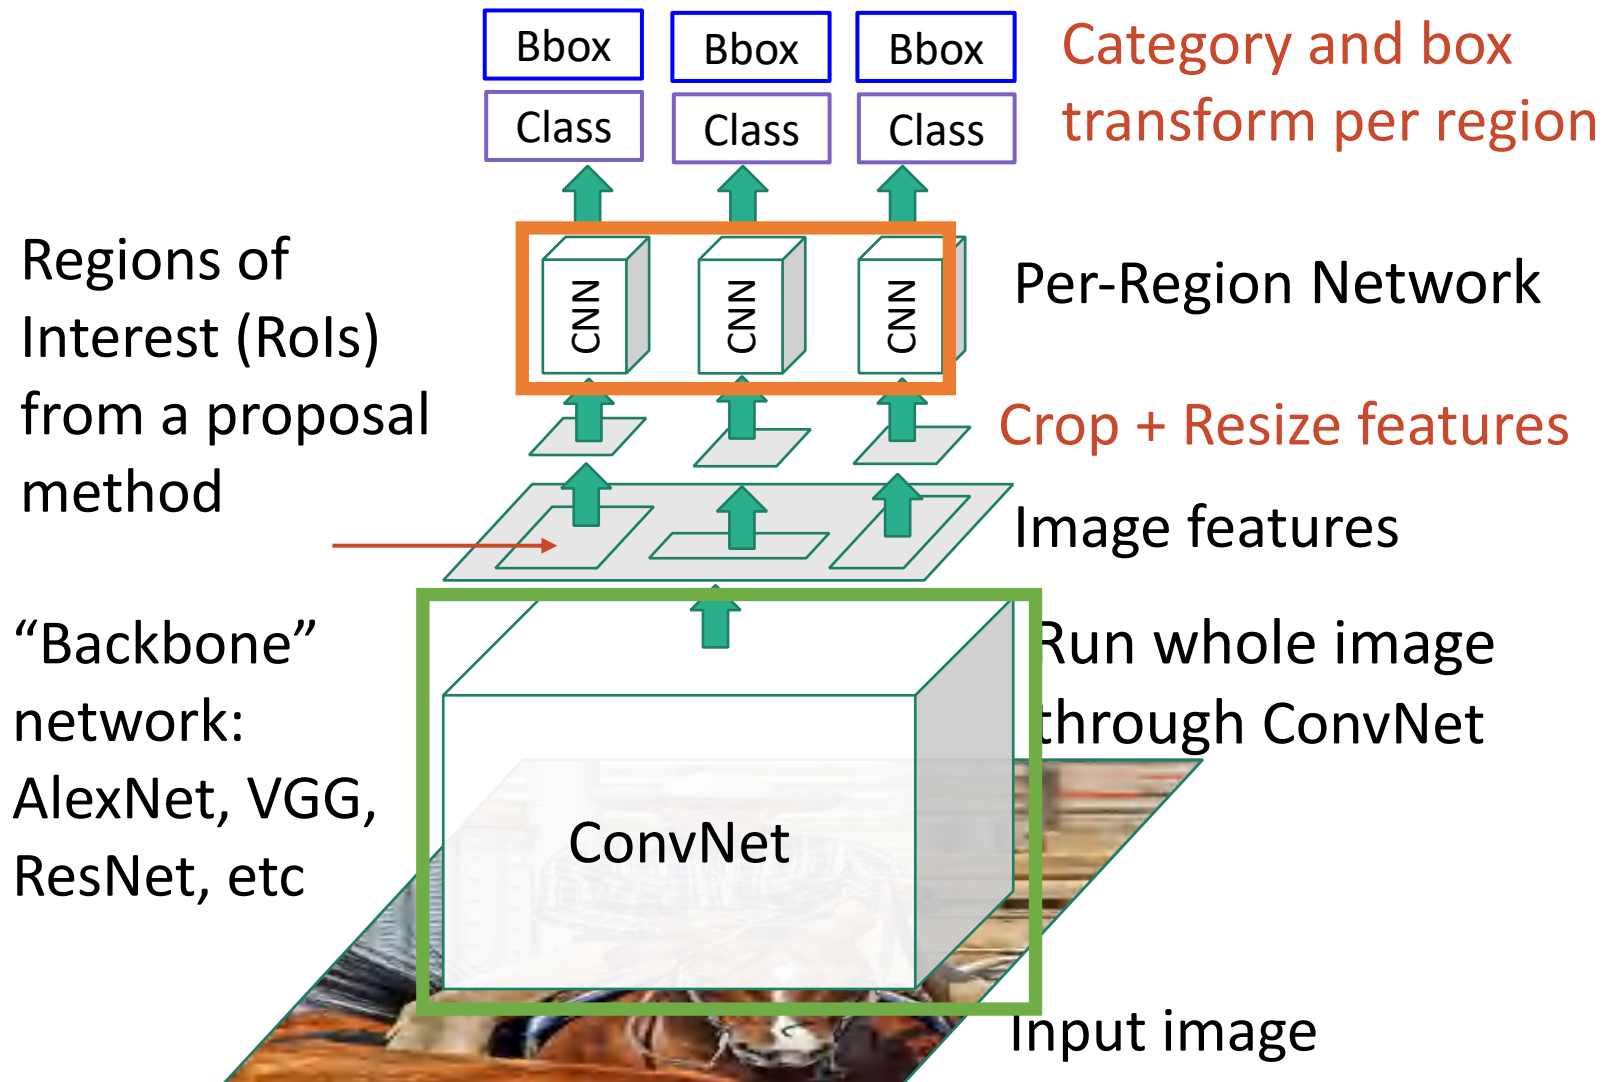
Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

Fast R-CNN



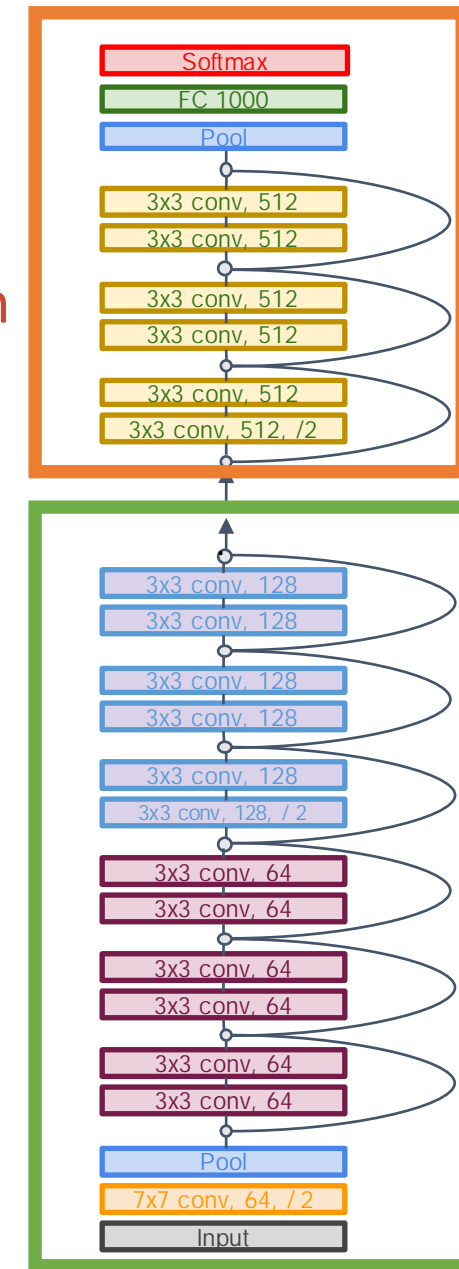
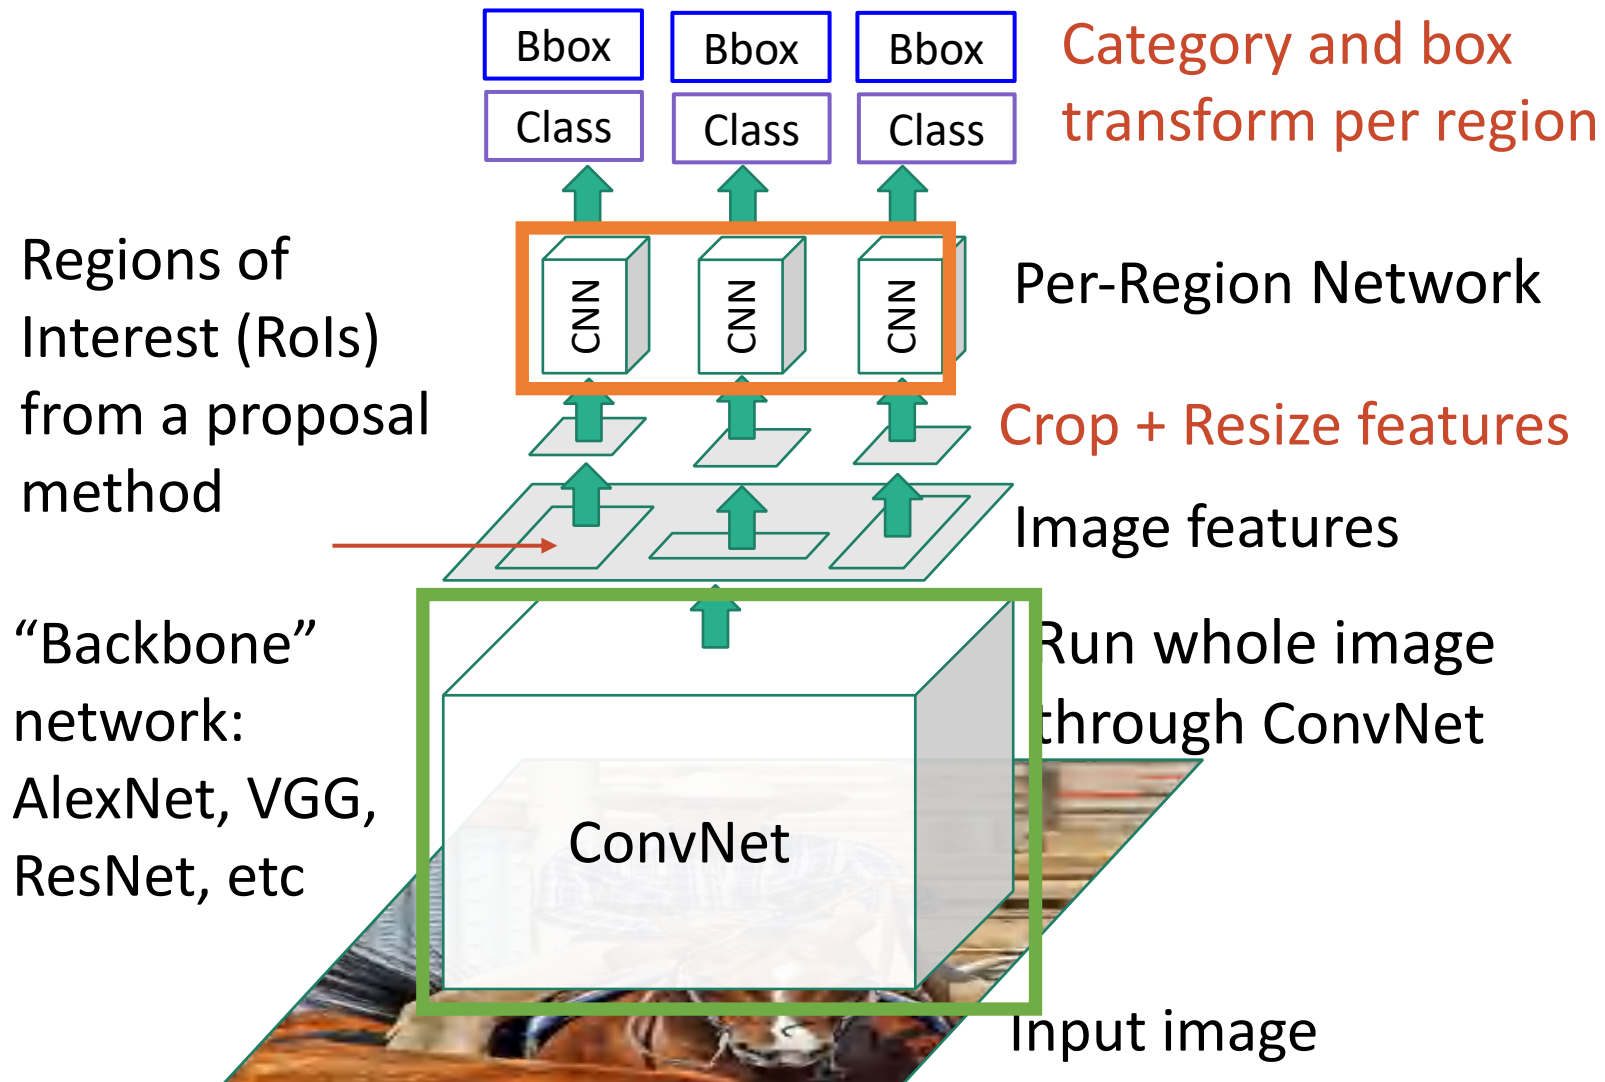
Girshick, "Fast R-CNN", ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

Fast R-CNN



Example:
When using AlexNet for detection, five conv layers are used for backbone and two FC layers are used for per-region network

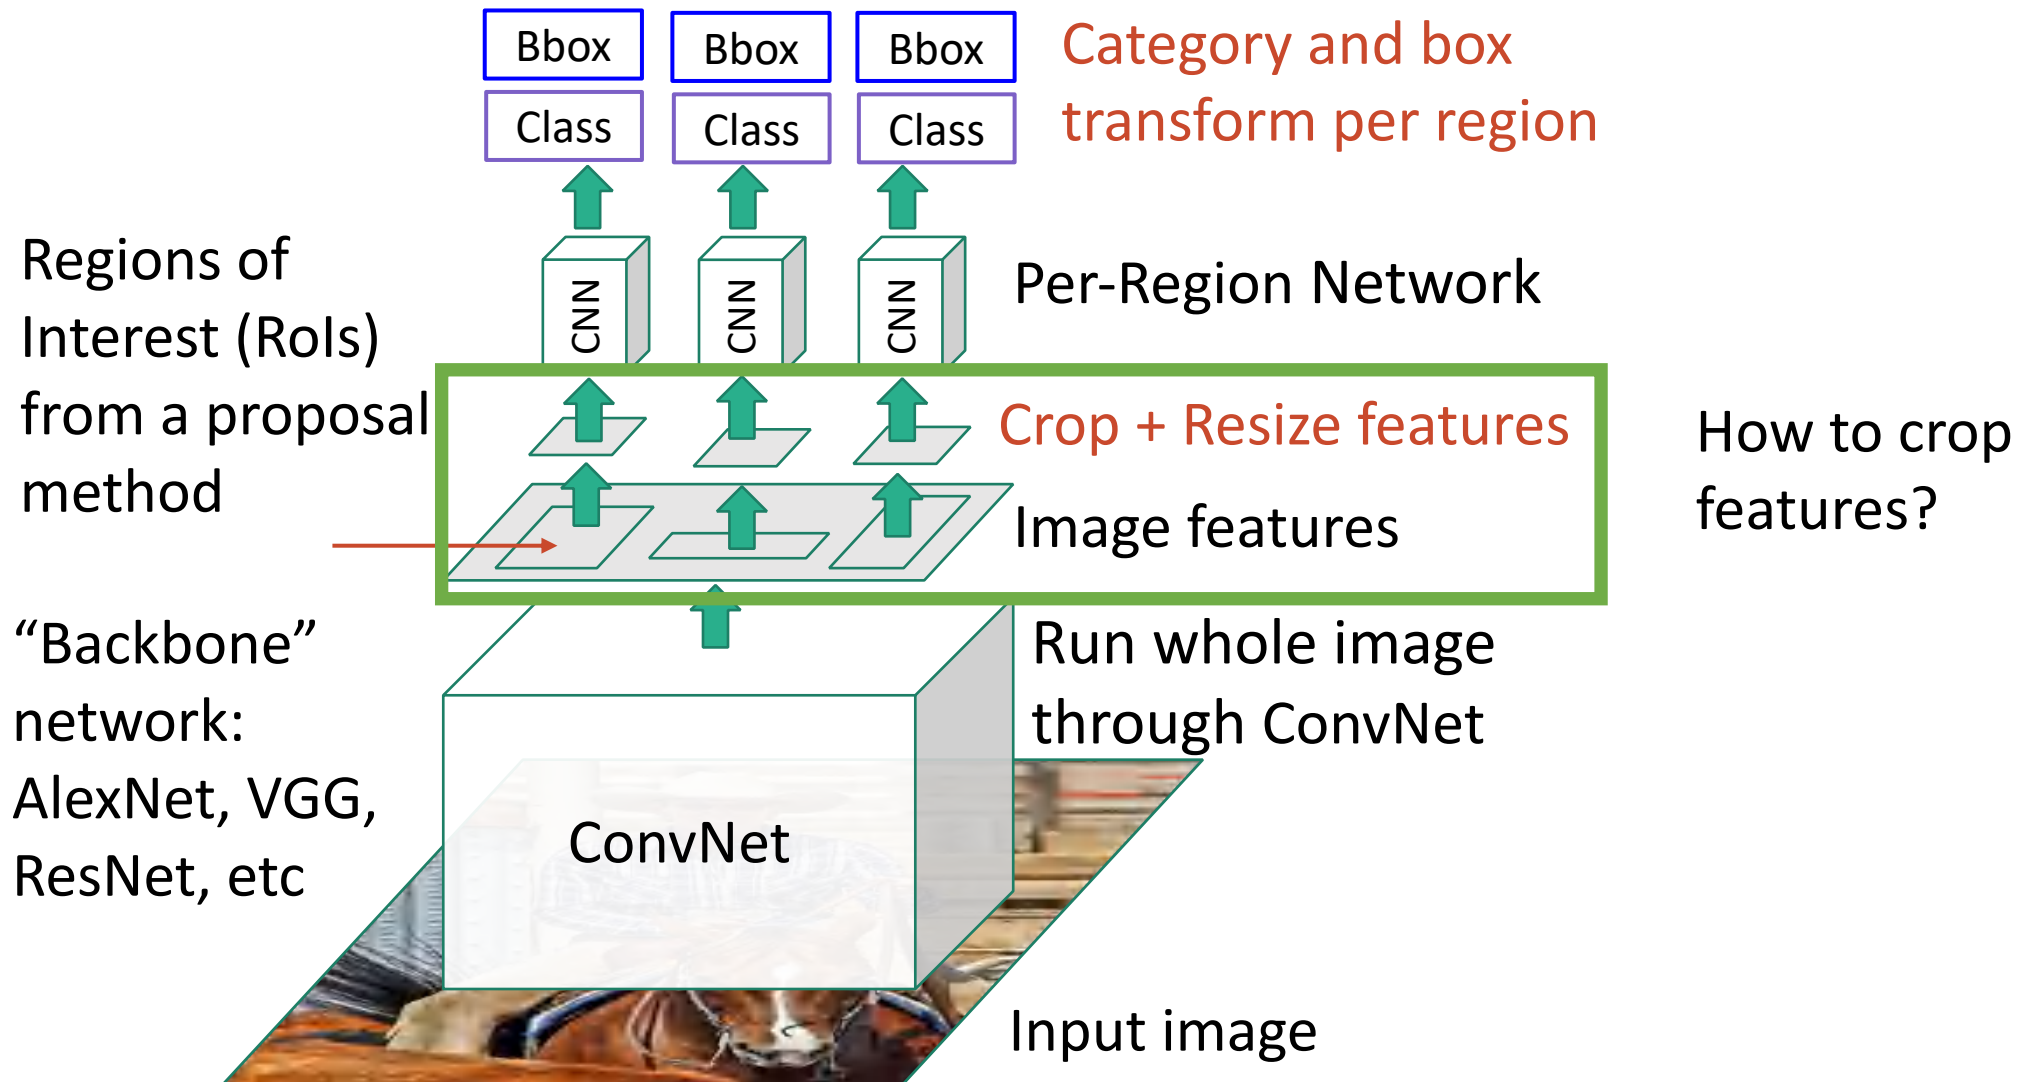
Fast R-CNN



Example:
For ResNet, last stage is used as per-region network; the rest of the network is used as backbone

Girshick, "Fast R-CNN", ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

Fast R-CNN




Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.


Recall: Receptive Fields

Every position in the output feature map depends on a 3x3 receptive field in the input

3x3 Conv
Stride 1, pad 1



Input Image: 8 x 8




Output Image: 8 x 8


Recall: Receptive Fields

Every position in the output feature map depends on a 3x3 receptive field in the input

3x3 Conv
Stride 1, pad 1



Input Image: 8 x 8




Output Image: 8 x 8

Recall: Receptive Fields


Every position in the output feature map depends on a 5x5 receptive field in the input

3x3 Conv
Stride 1, pad 1

3x3 Conv
Stride 1, pad 1




Input Image: 8 x 8



Output Image: 8 x 8

Recall: Receptive Fields




Input Image: 8 x 8

Moving one unit in the output space also moves the receptive field by one

3x3 Conv
Stride 1, pad 1


3x3 Conv
Stride 1, pad 1



Output Image: 8 x 8

Recall: Receptive Fields

(0, 0)



Input Image: 8 x 8


Moving one unit in the output space also moves the receptive field by one

3x3 Conv
Stride 1, pad 1

3x3 Conv
Stride 1, pad 1


There is a correspondence between the **coordinate system of the input** and the **coordinate system of the output**

(0, 0)



Output Image: 8 x 8

Projecting Points




We can align arbitrary points between coordinate system of input and output

3x3 Conv
Stride 1, pad 1


3x3 Conv
Stride 1, pad 1

There is a correspondence between the **coordinate system of the input** and the **coordinate system of the output**



Projecting Points

Same logic holds for more complicated CNNs, even if spatial resolution of input and output are different




We can align arbitrary points between coordinate system of input and output

3x3 Conv
Stride 1, pad 1


2x2 MaxPool
Stride 2

There is a correspondence between the **coordinate system of the input** and the **coordinate system of the output**



Projecting Points

Same logic holds for more complicated CNNs, even if spatial resolution of input and output are different




We can align arbitrary points between coordinate system of input and output


3x3 Conv
Stride 1, pad 1

4x4 MaxPool
Stride 4

There is a correspondence between the **coordinate system of the input** and the **coordinate system of the output**



Projecting Boxes




We can align arbitrary points between coordinate system of input and output

3x3 Conv
Stride 1, pad 1


4x4 MaxPool
Stride 4

There is a correspondence between the **coordinate system of the input** and the **coordinate system of the output**

We can use this idea to project **bounding boxes** between an input image and a feature map




Cropping Features: RoI Pool



Input Image
(e.g. 3 x 640 x 480)

Girshick, "Fast R-CNN", ICCV 2015.

Cropping Features: RoI Pool



Input Image
(e.g. 3 x 640 x 480)


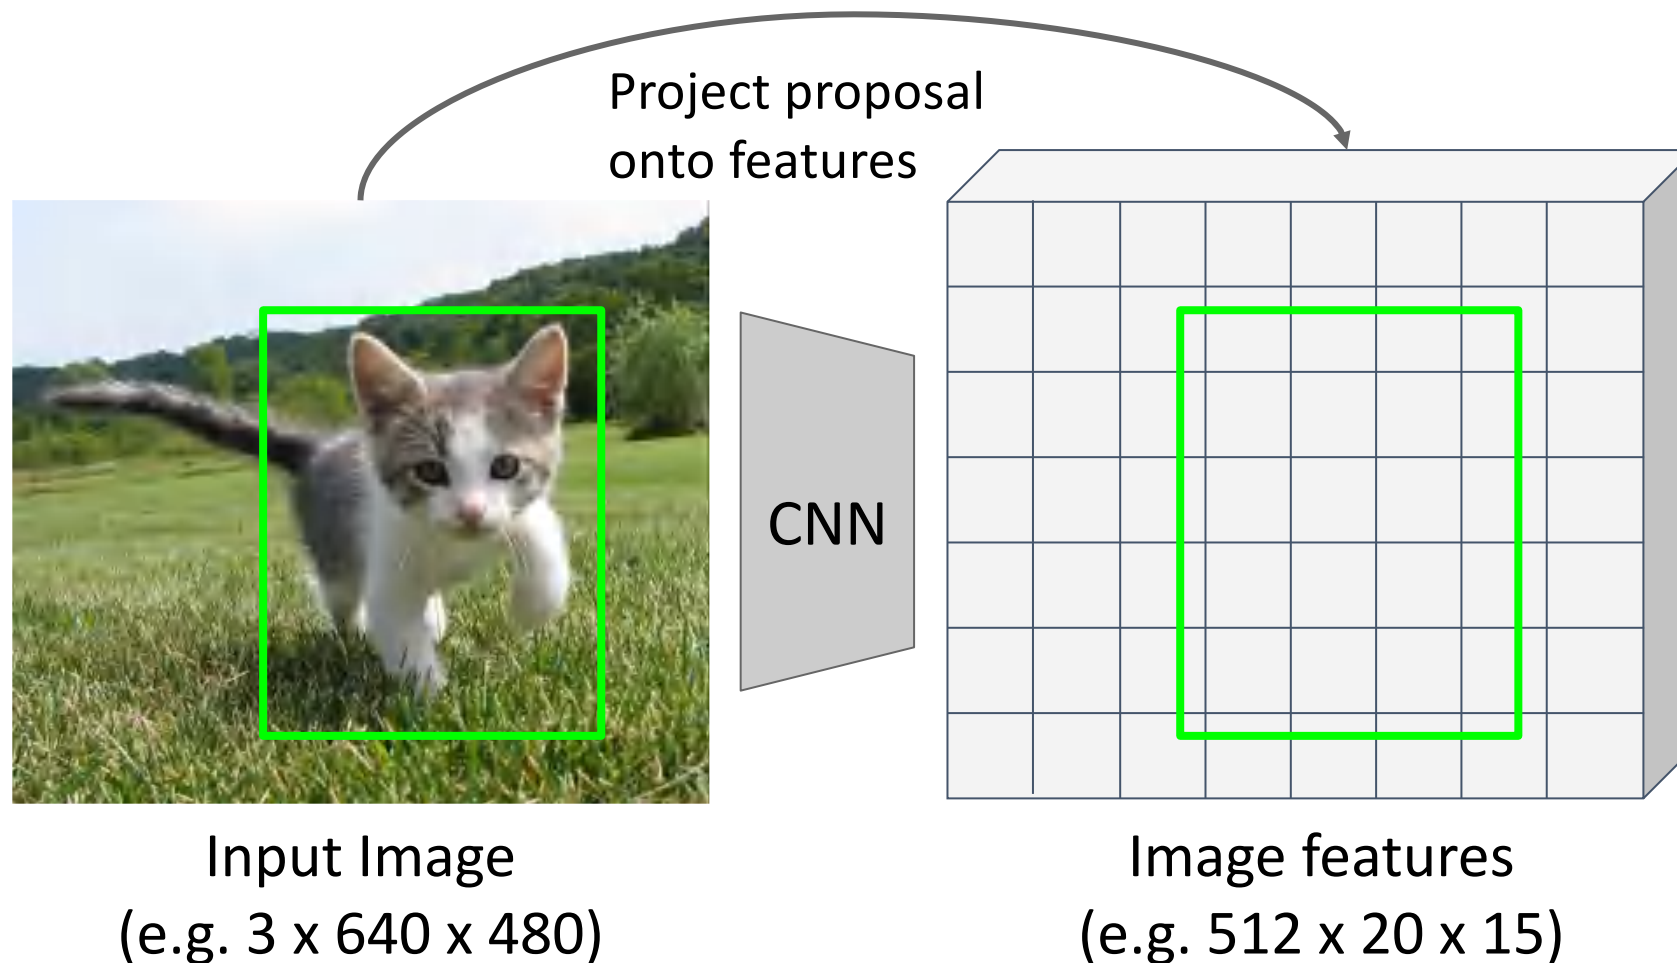


Image features
(e.g. 512 x 20 x 15)

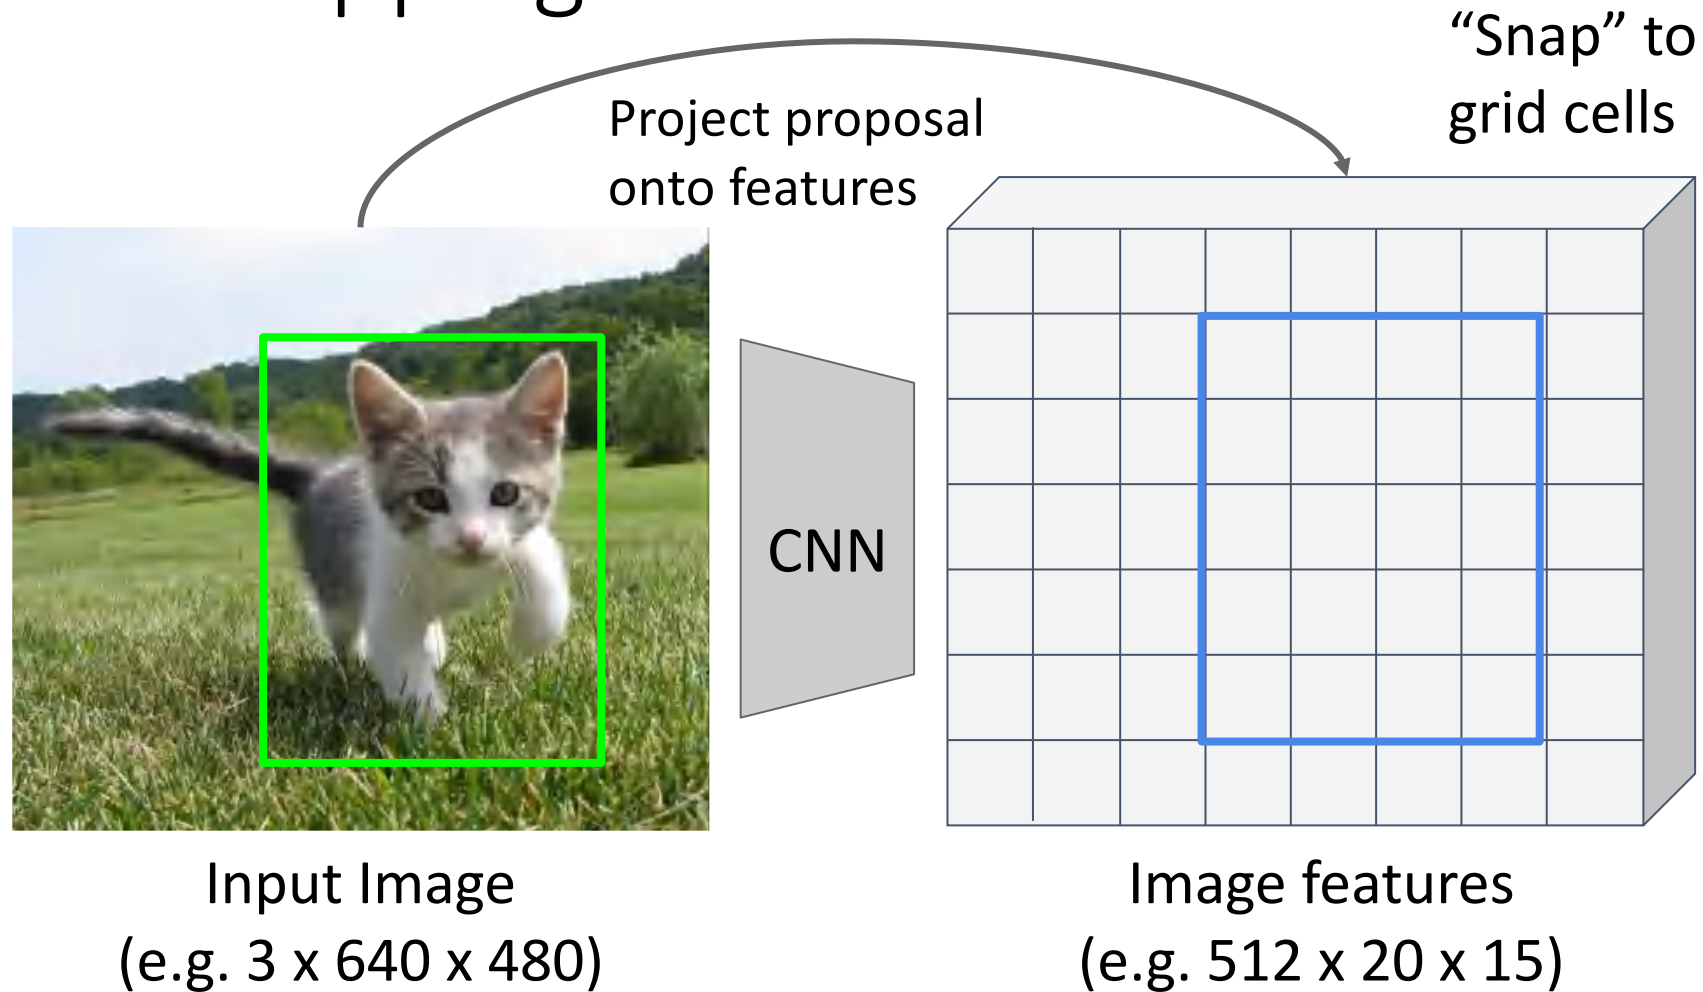
Want features for the
box of a fixed size
(2x2 in this example,
7x7 or 14x14 in practice)

Cropping Features: RoI Pool



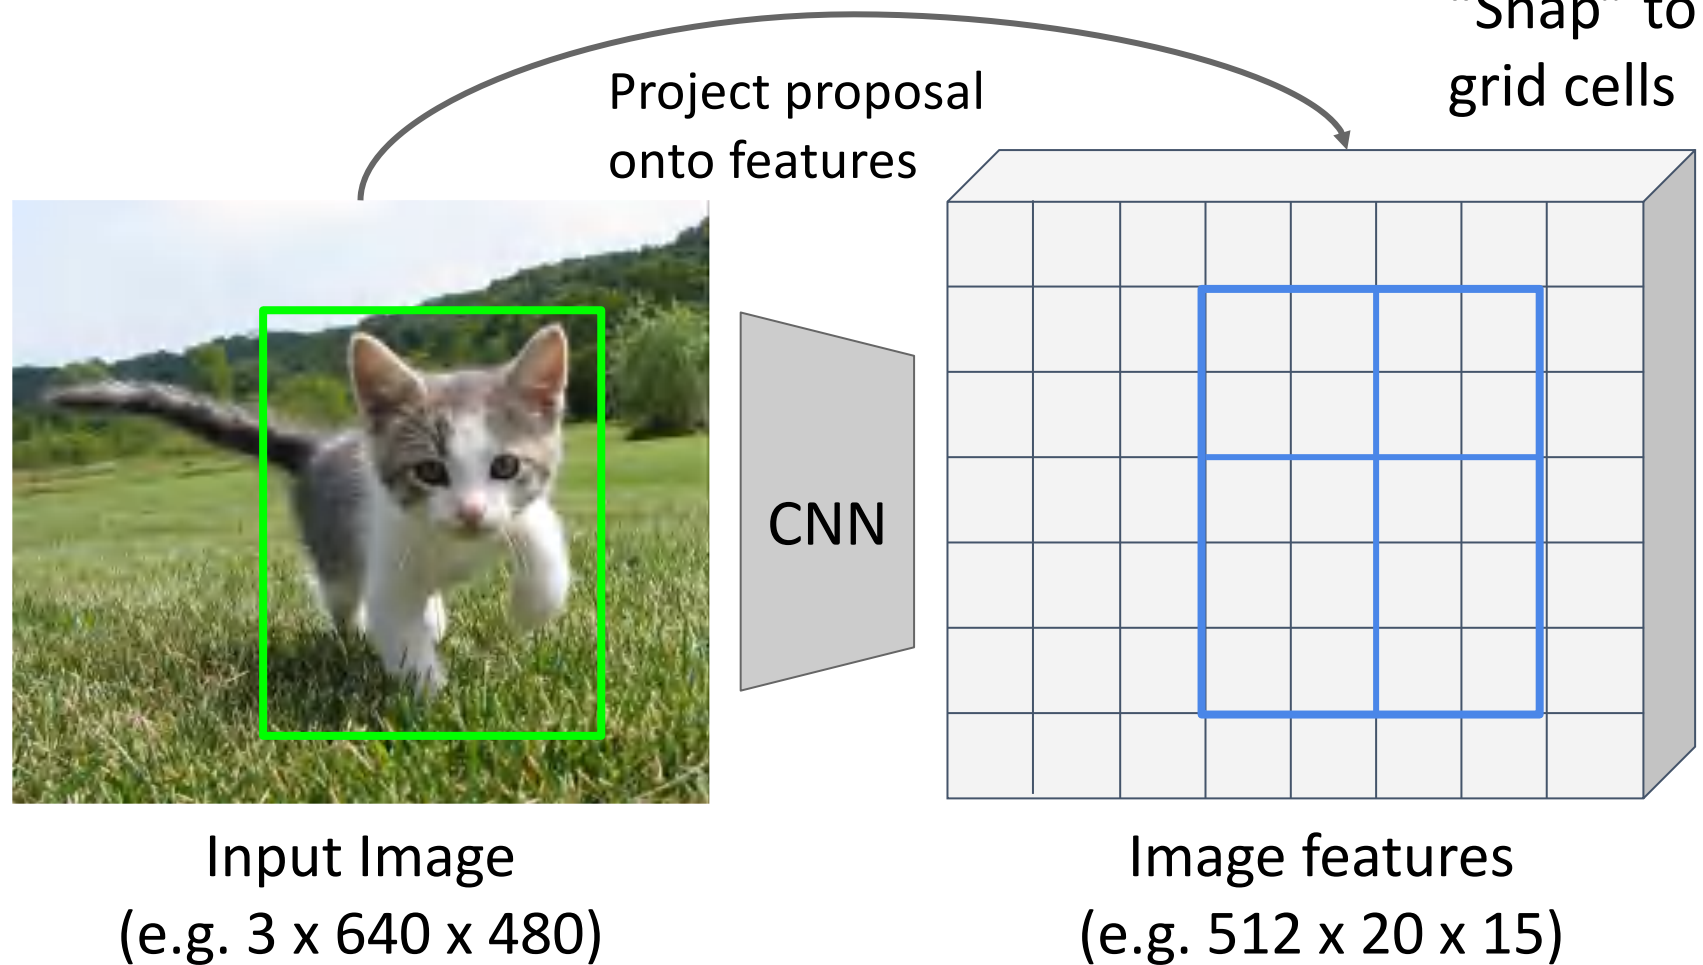
Want features for the box of a fixed size (2x2 in this example, 7x7 or 14x14 in practice)

Cropping Features: RoI Pool



Want features for the box of a fixed size (2x2 in this example, 7x7 or 14x14 in practice)

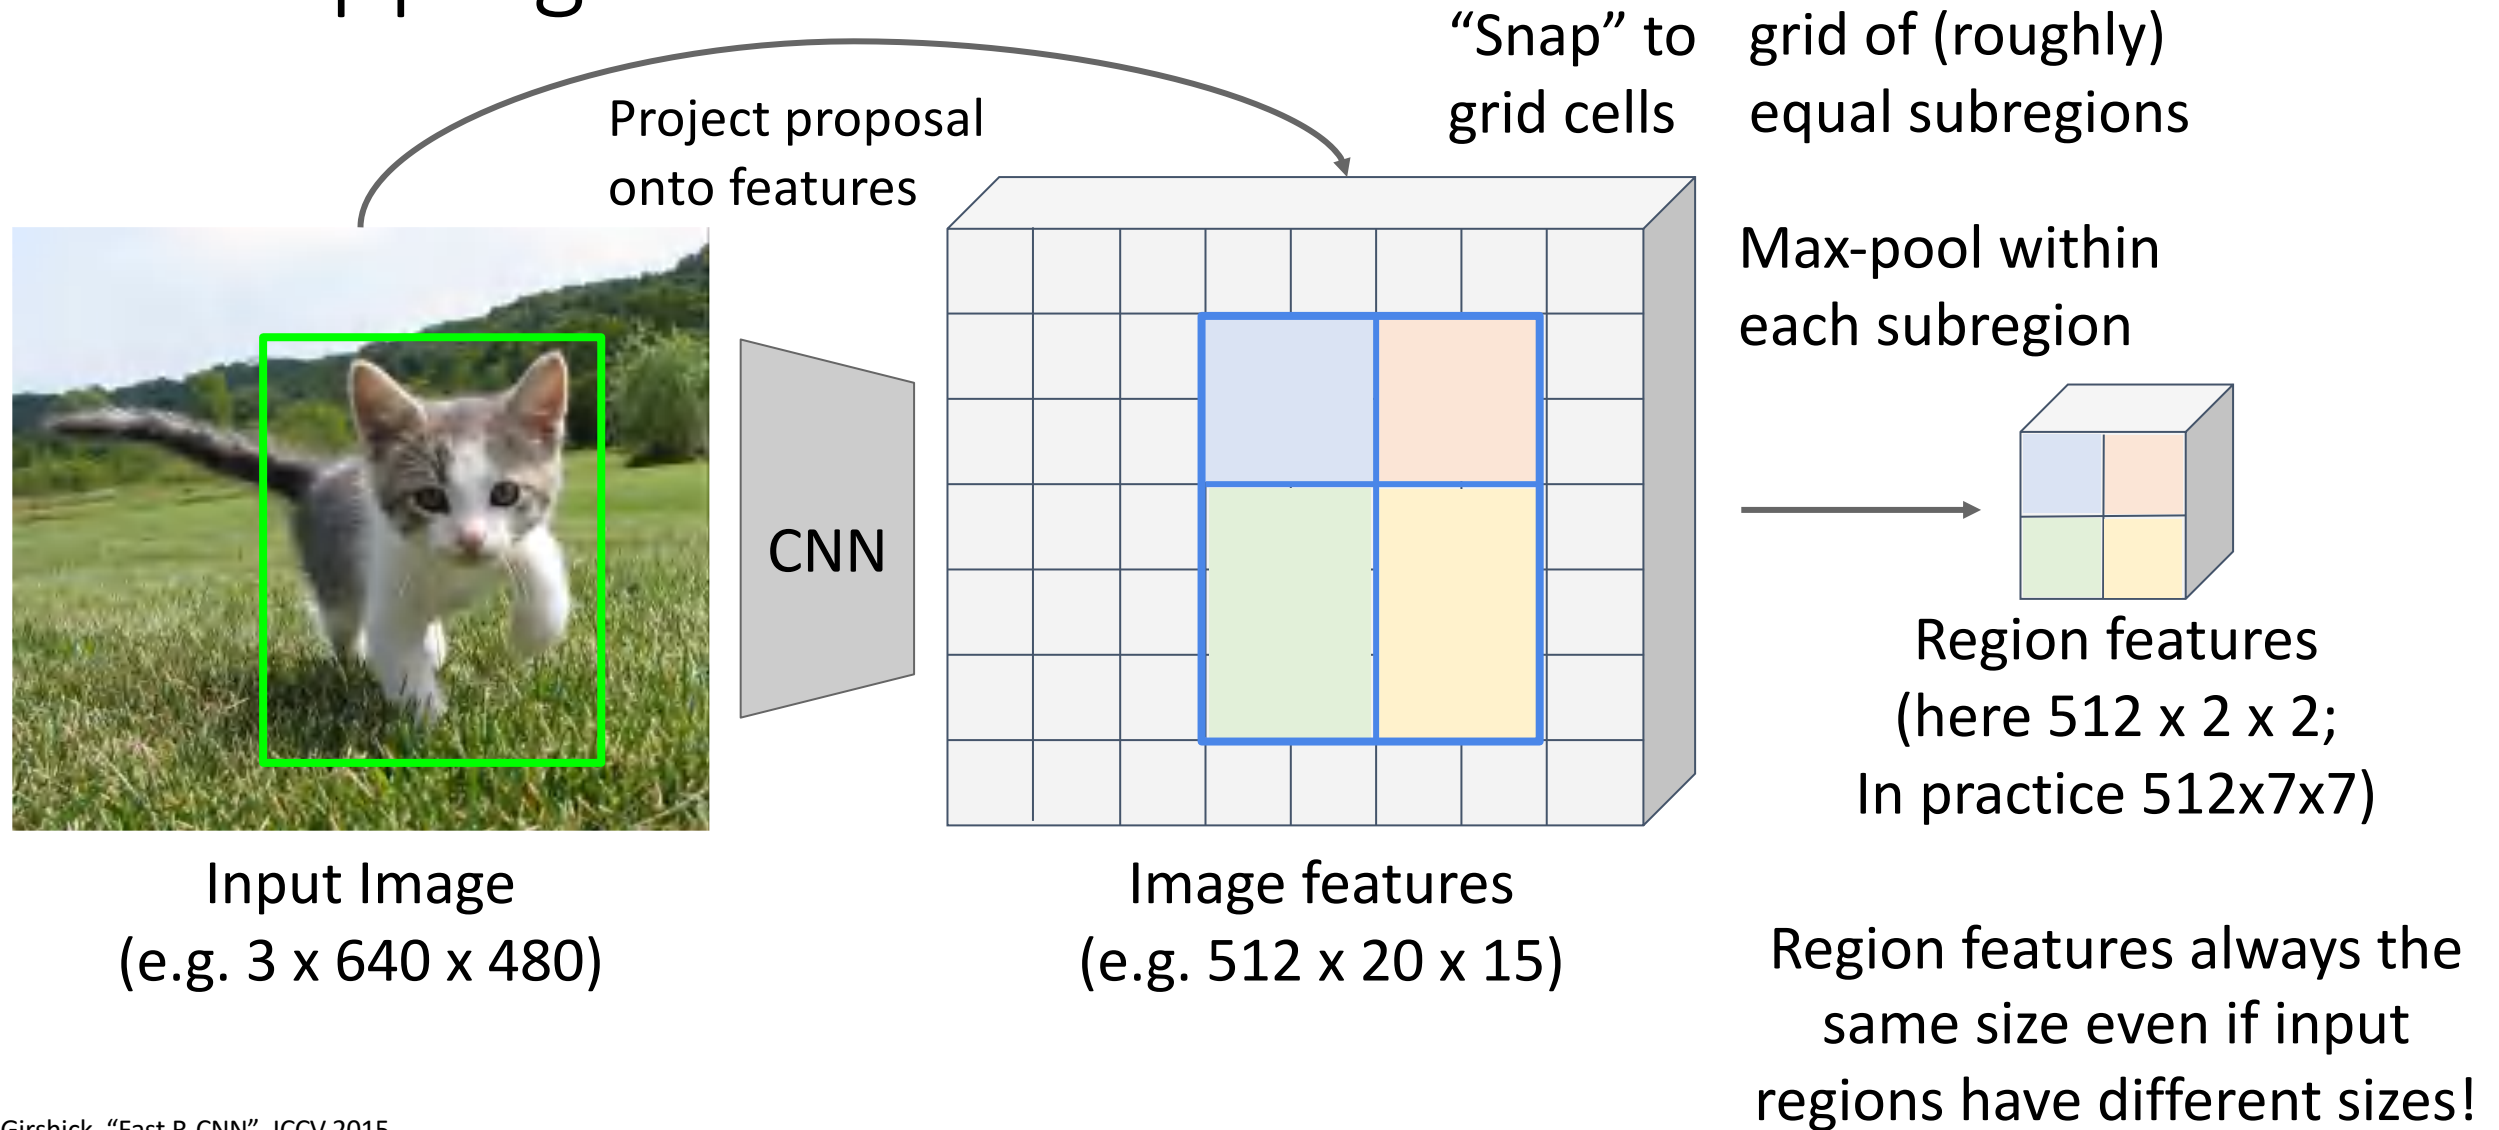
Cropping Features: RoI Pool



Divide into 2x2 grid of (roughly) equal subregions

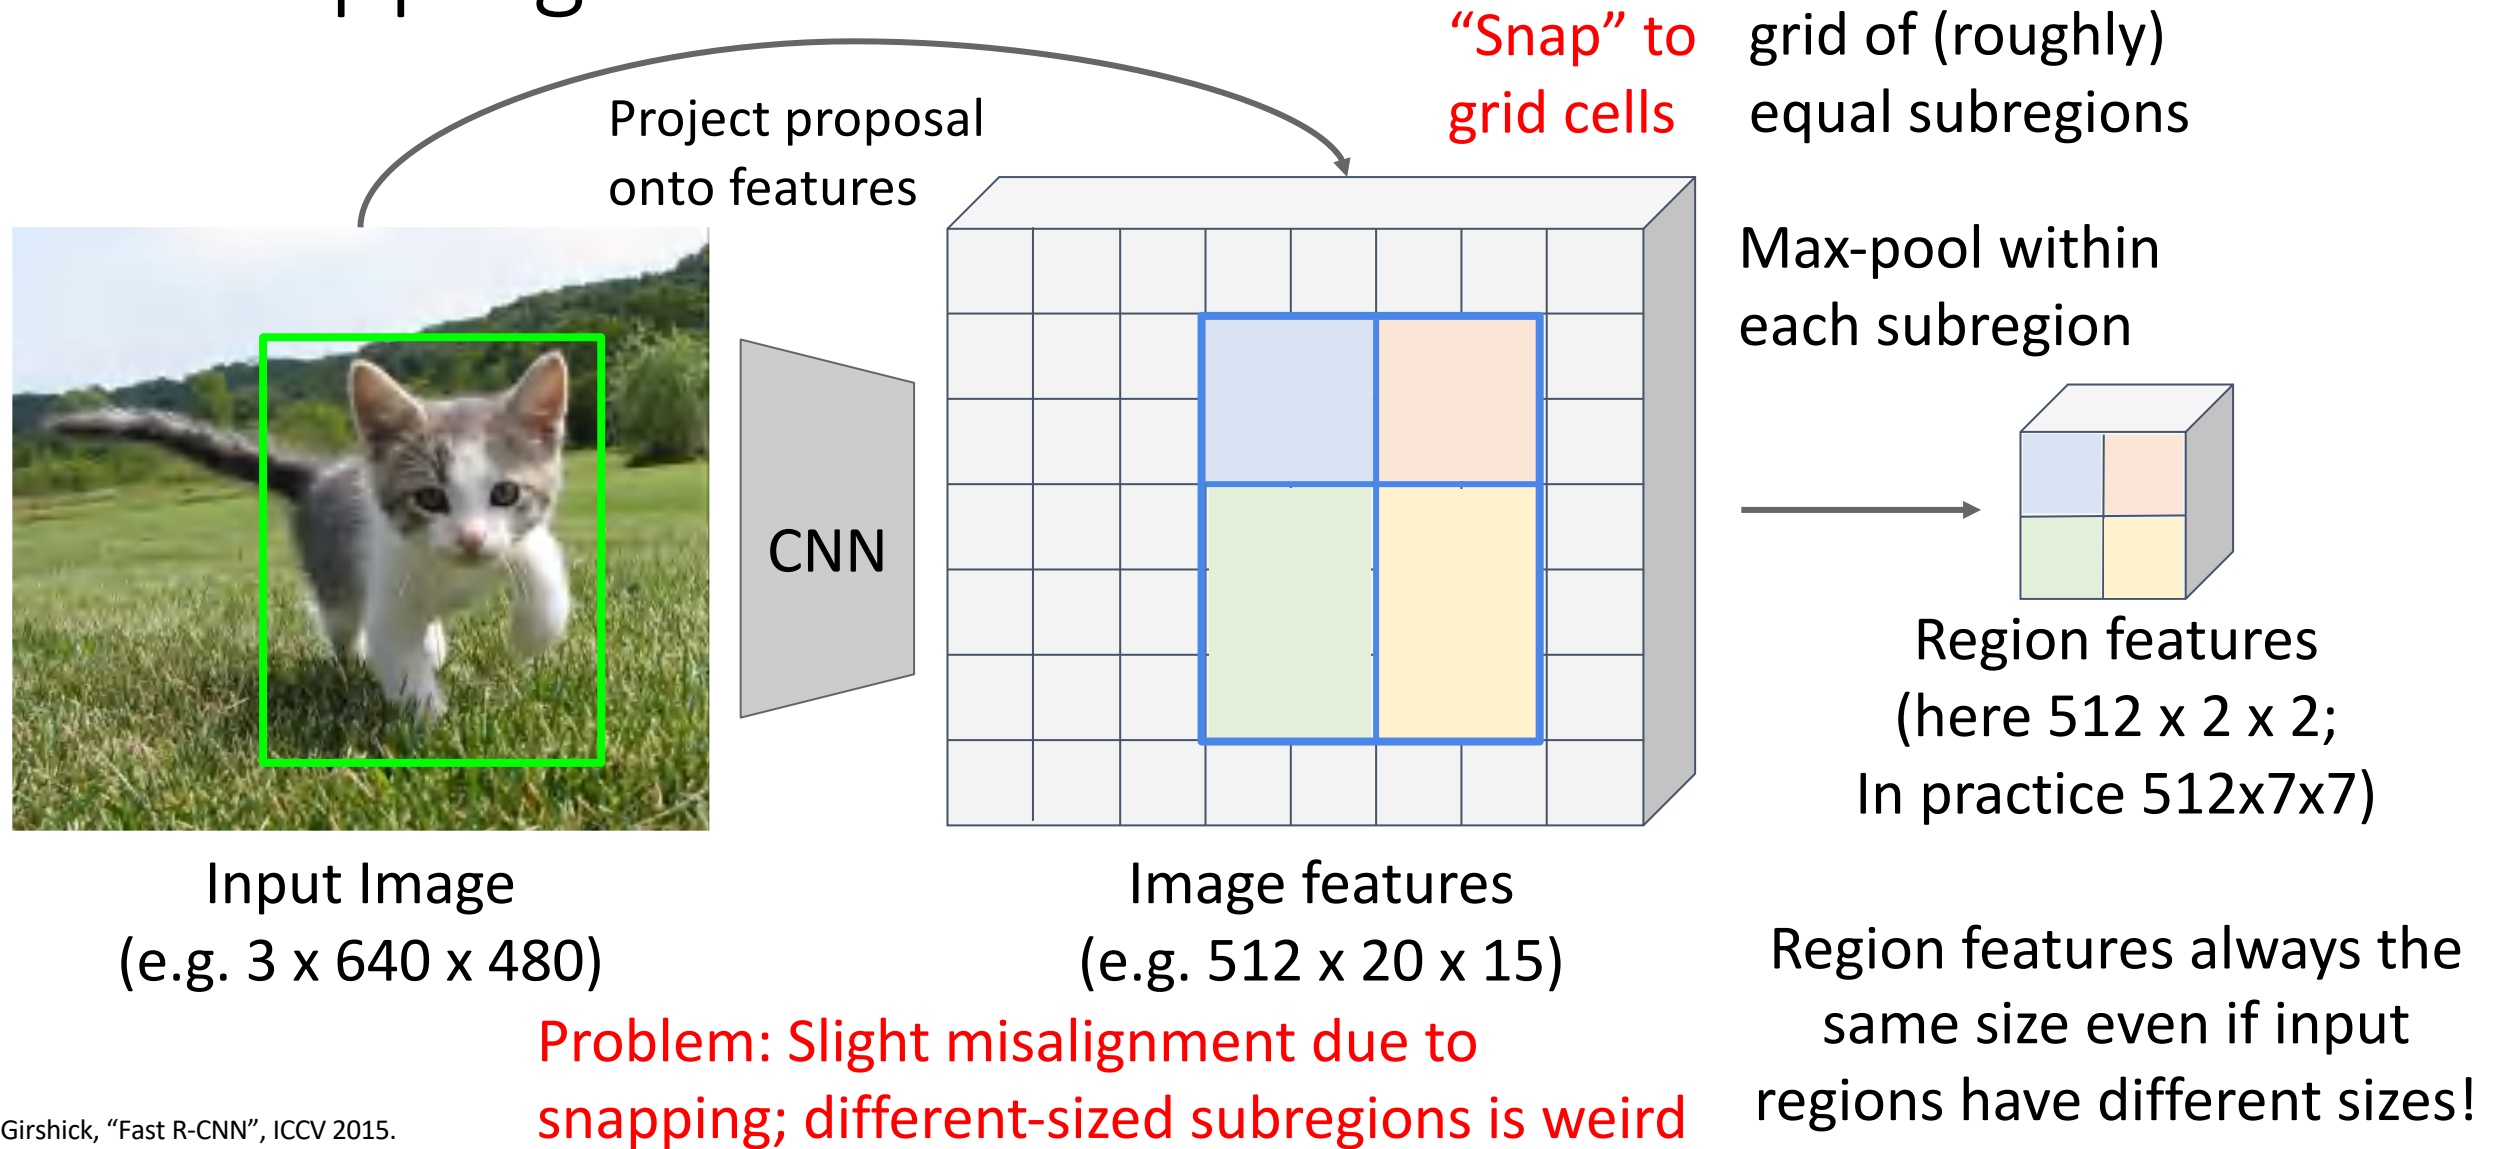
Want features for the box of a fixed size (2x2 in this example, 7x7 or 14x14 in practice)

Cropping Features: RoI Pool



Girshick, “Fast R-CNN”, ICCV 2015.

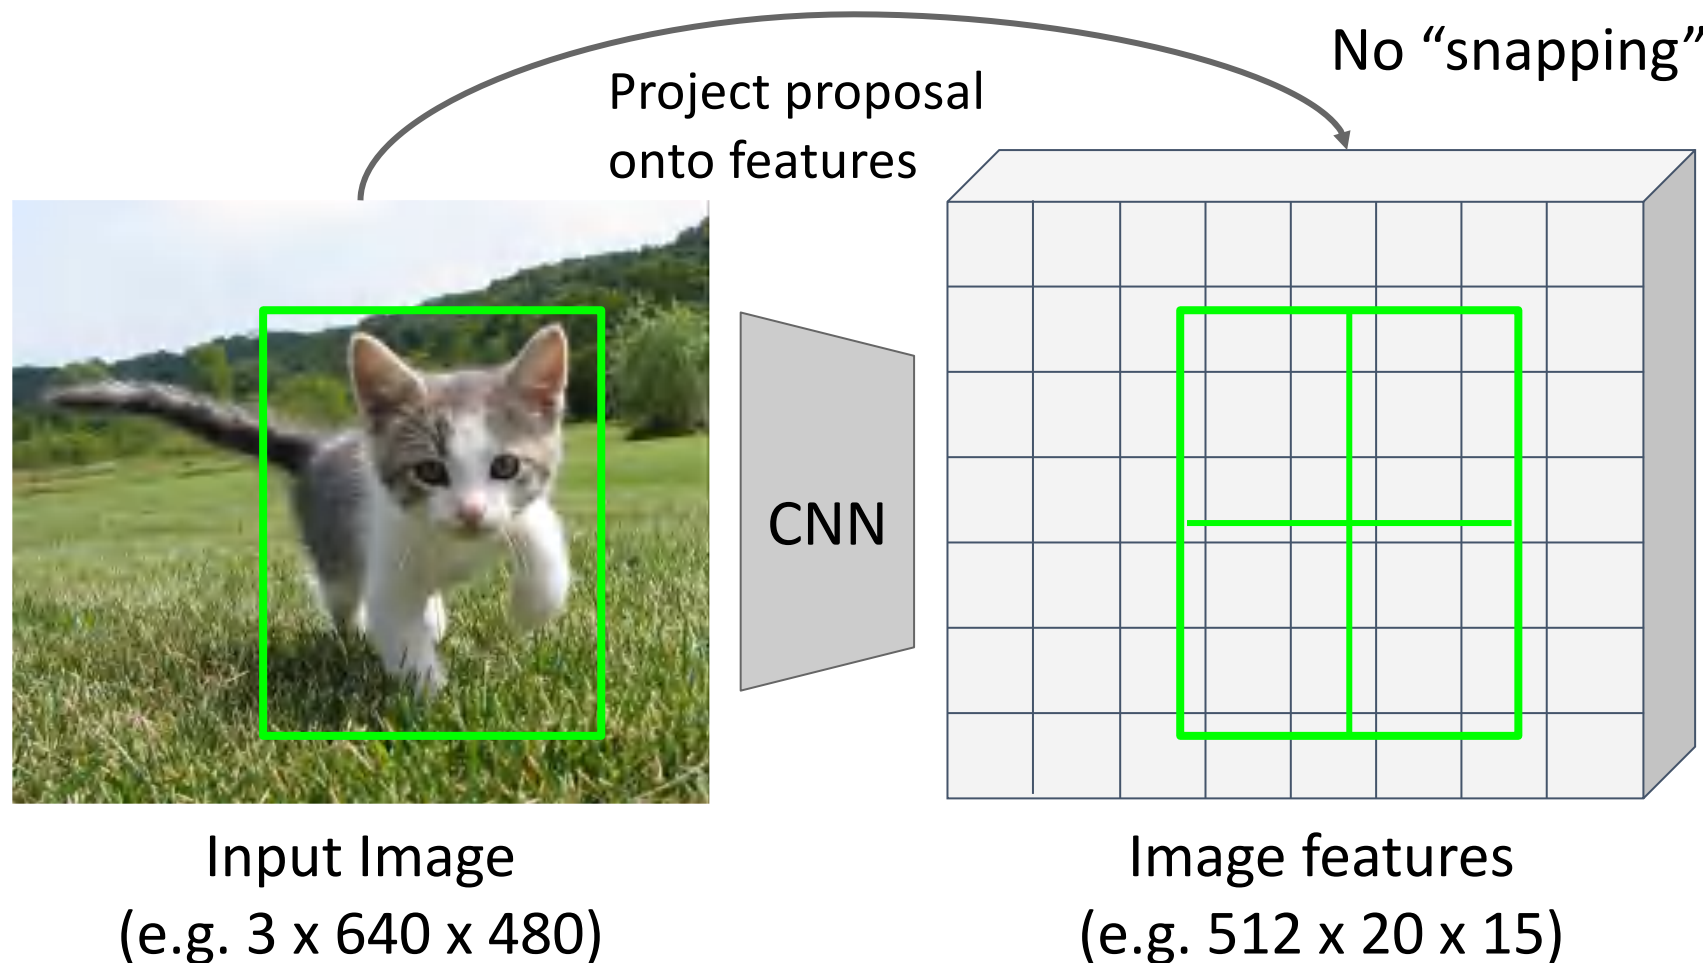
Cropping Features: RoI Pool



Girshick, "Fast R-CNN", ICCV 2015.

Cropping Features: RoI Align

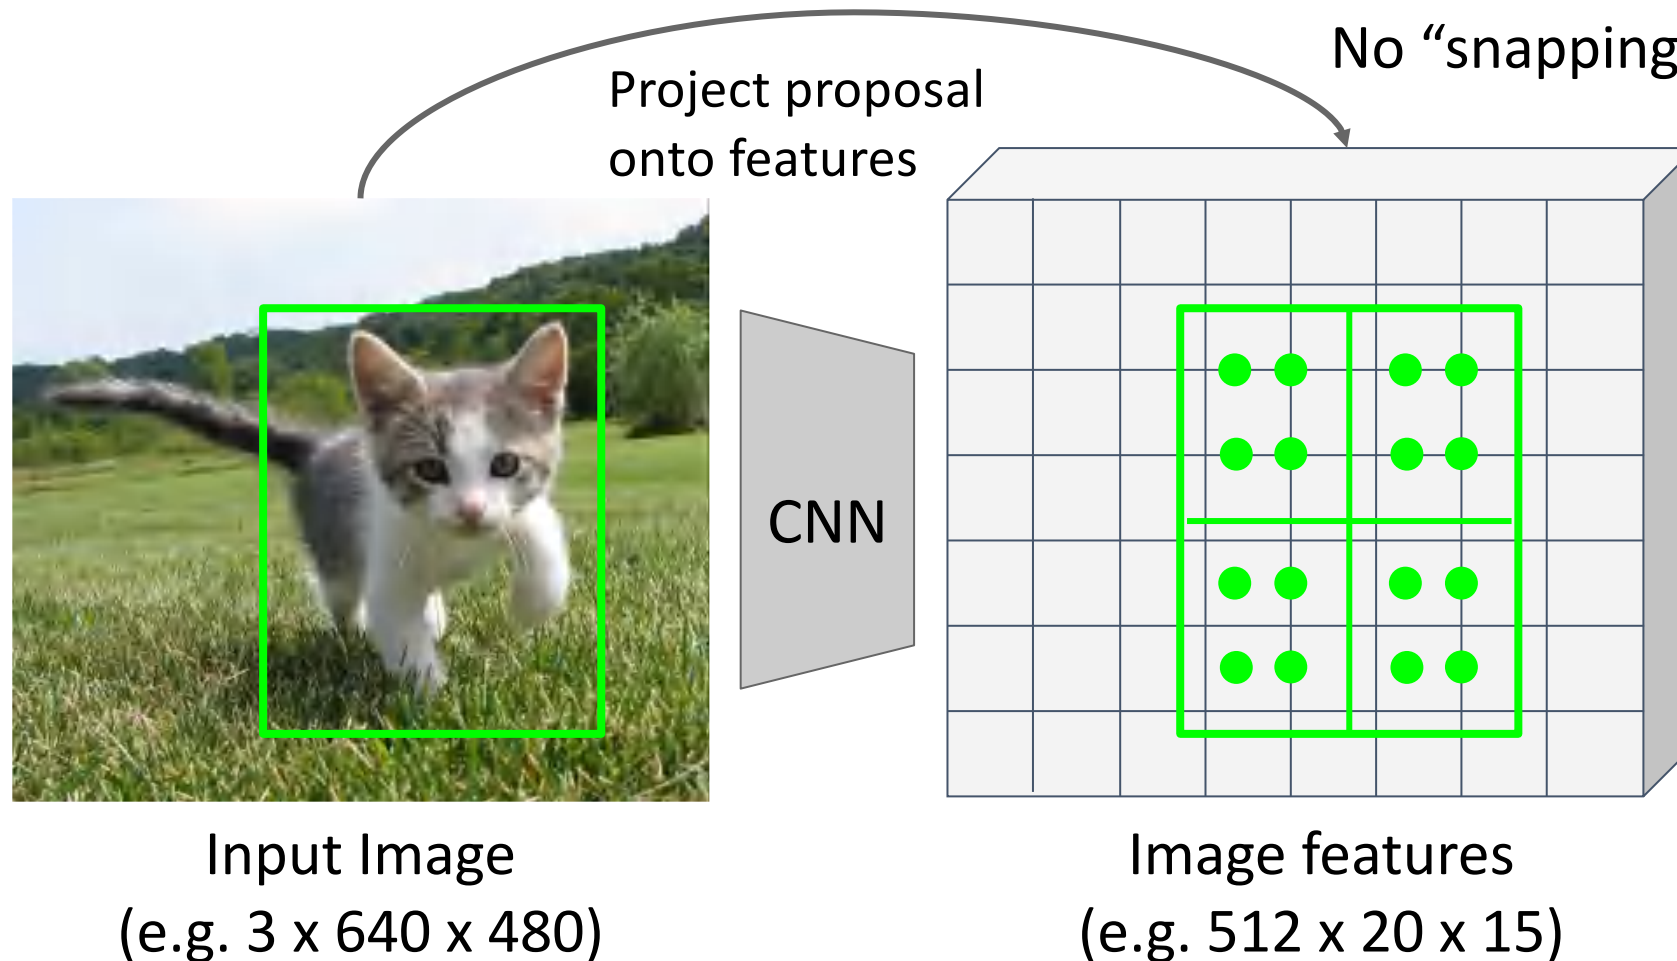
Divide into equal-sized subregions
(may not be aligned to grid!)



Want features for the
box of a fixed size
(2x2 in this example,
7x7 or 14x14 in practice)

Cropping Features: RoI Align

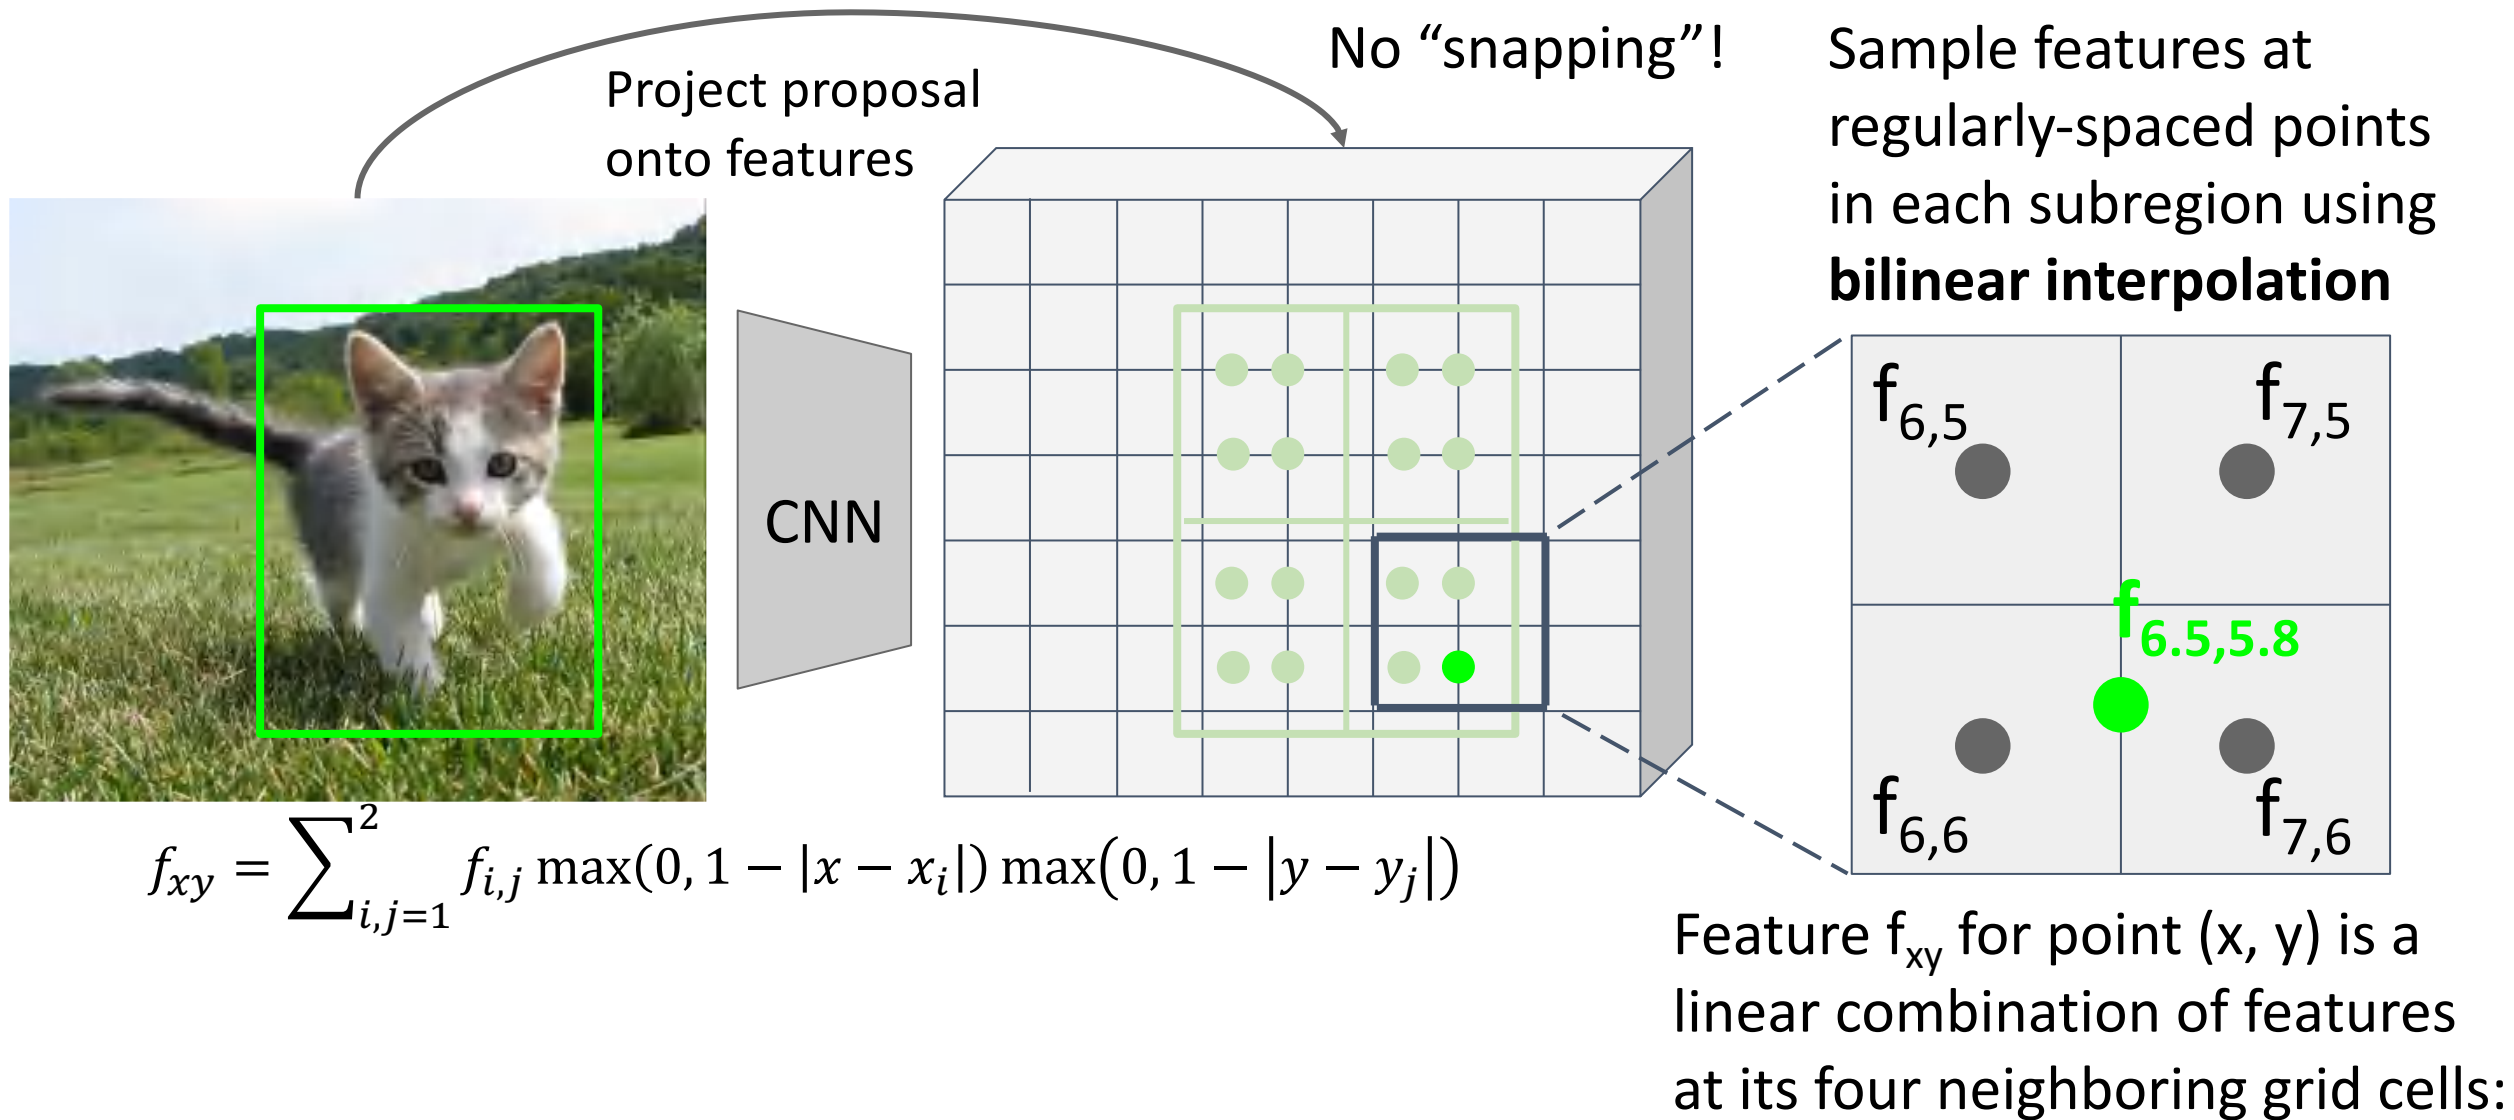
Divide into equal-sized subregions
(may not be aligned to grid!)



Sample features at
regularly-spaced points
in each subregion using
bilinear interpolation

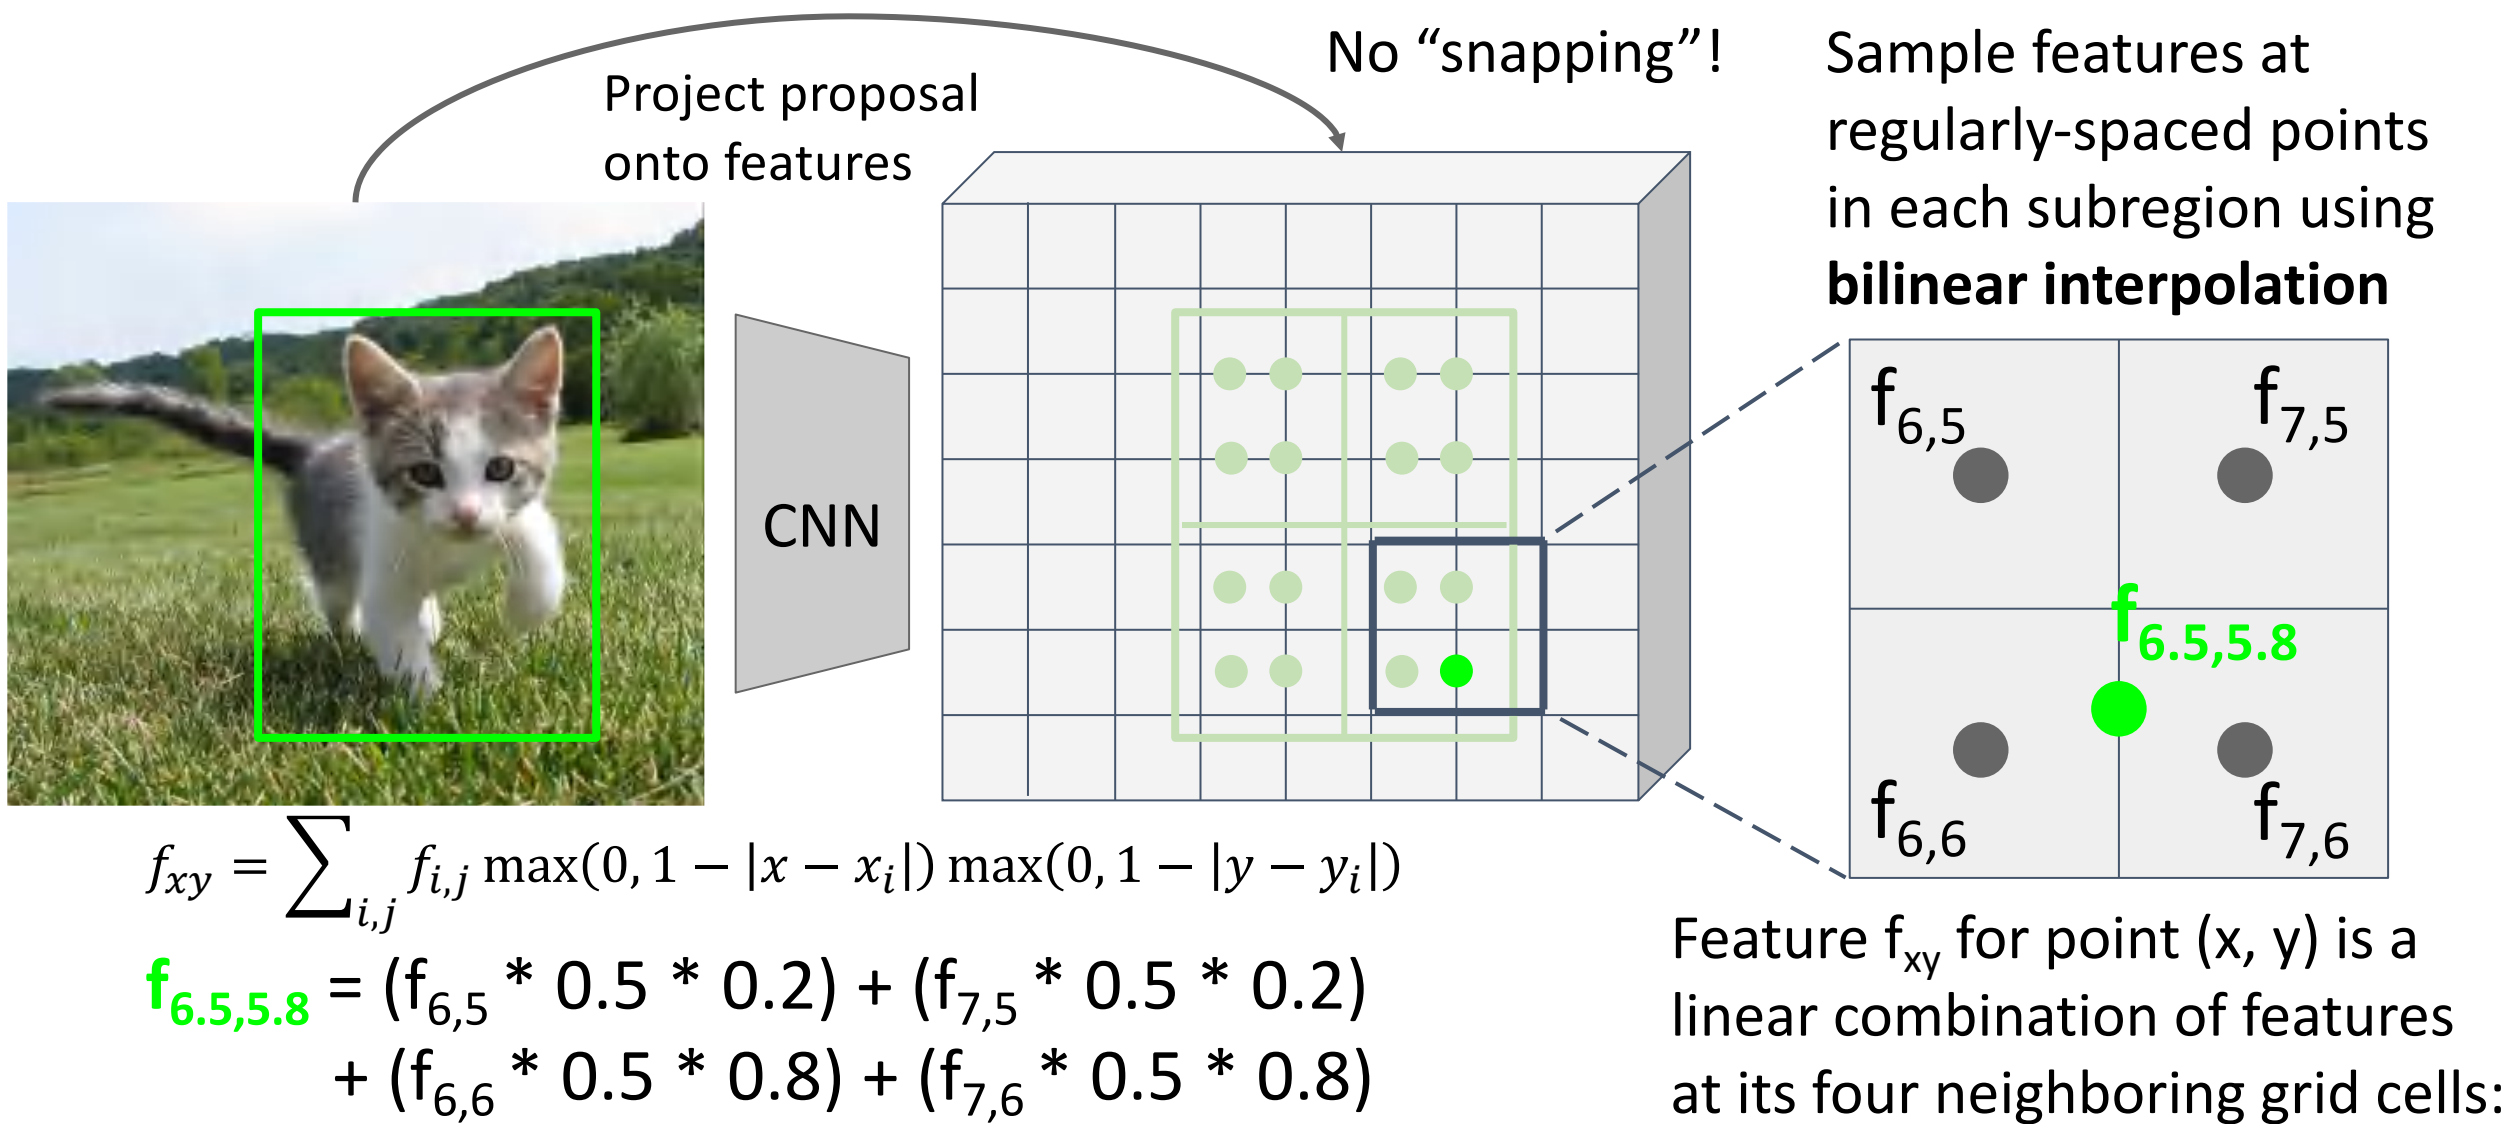
Cropping Features: RoI Align

Divide into equal-sized subregions
(may not be aligned to grid!)

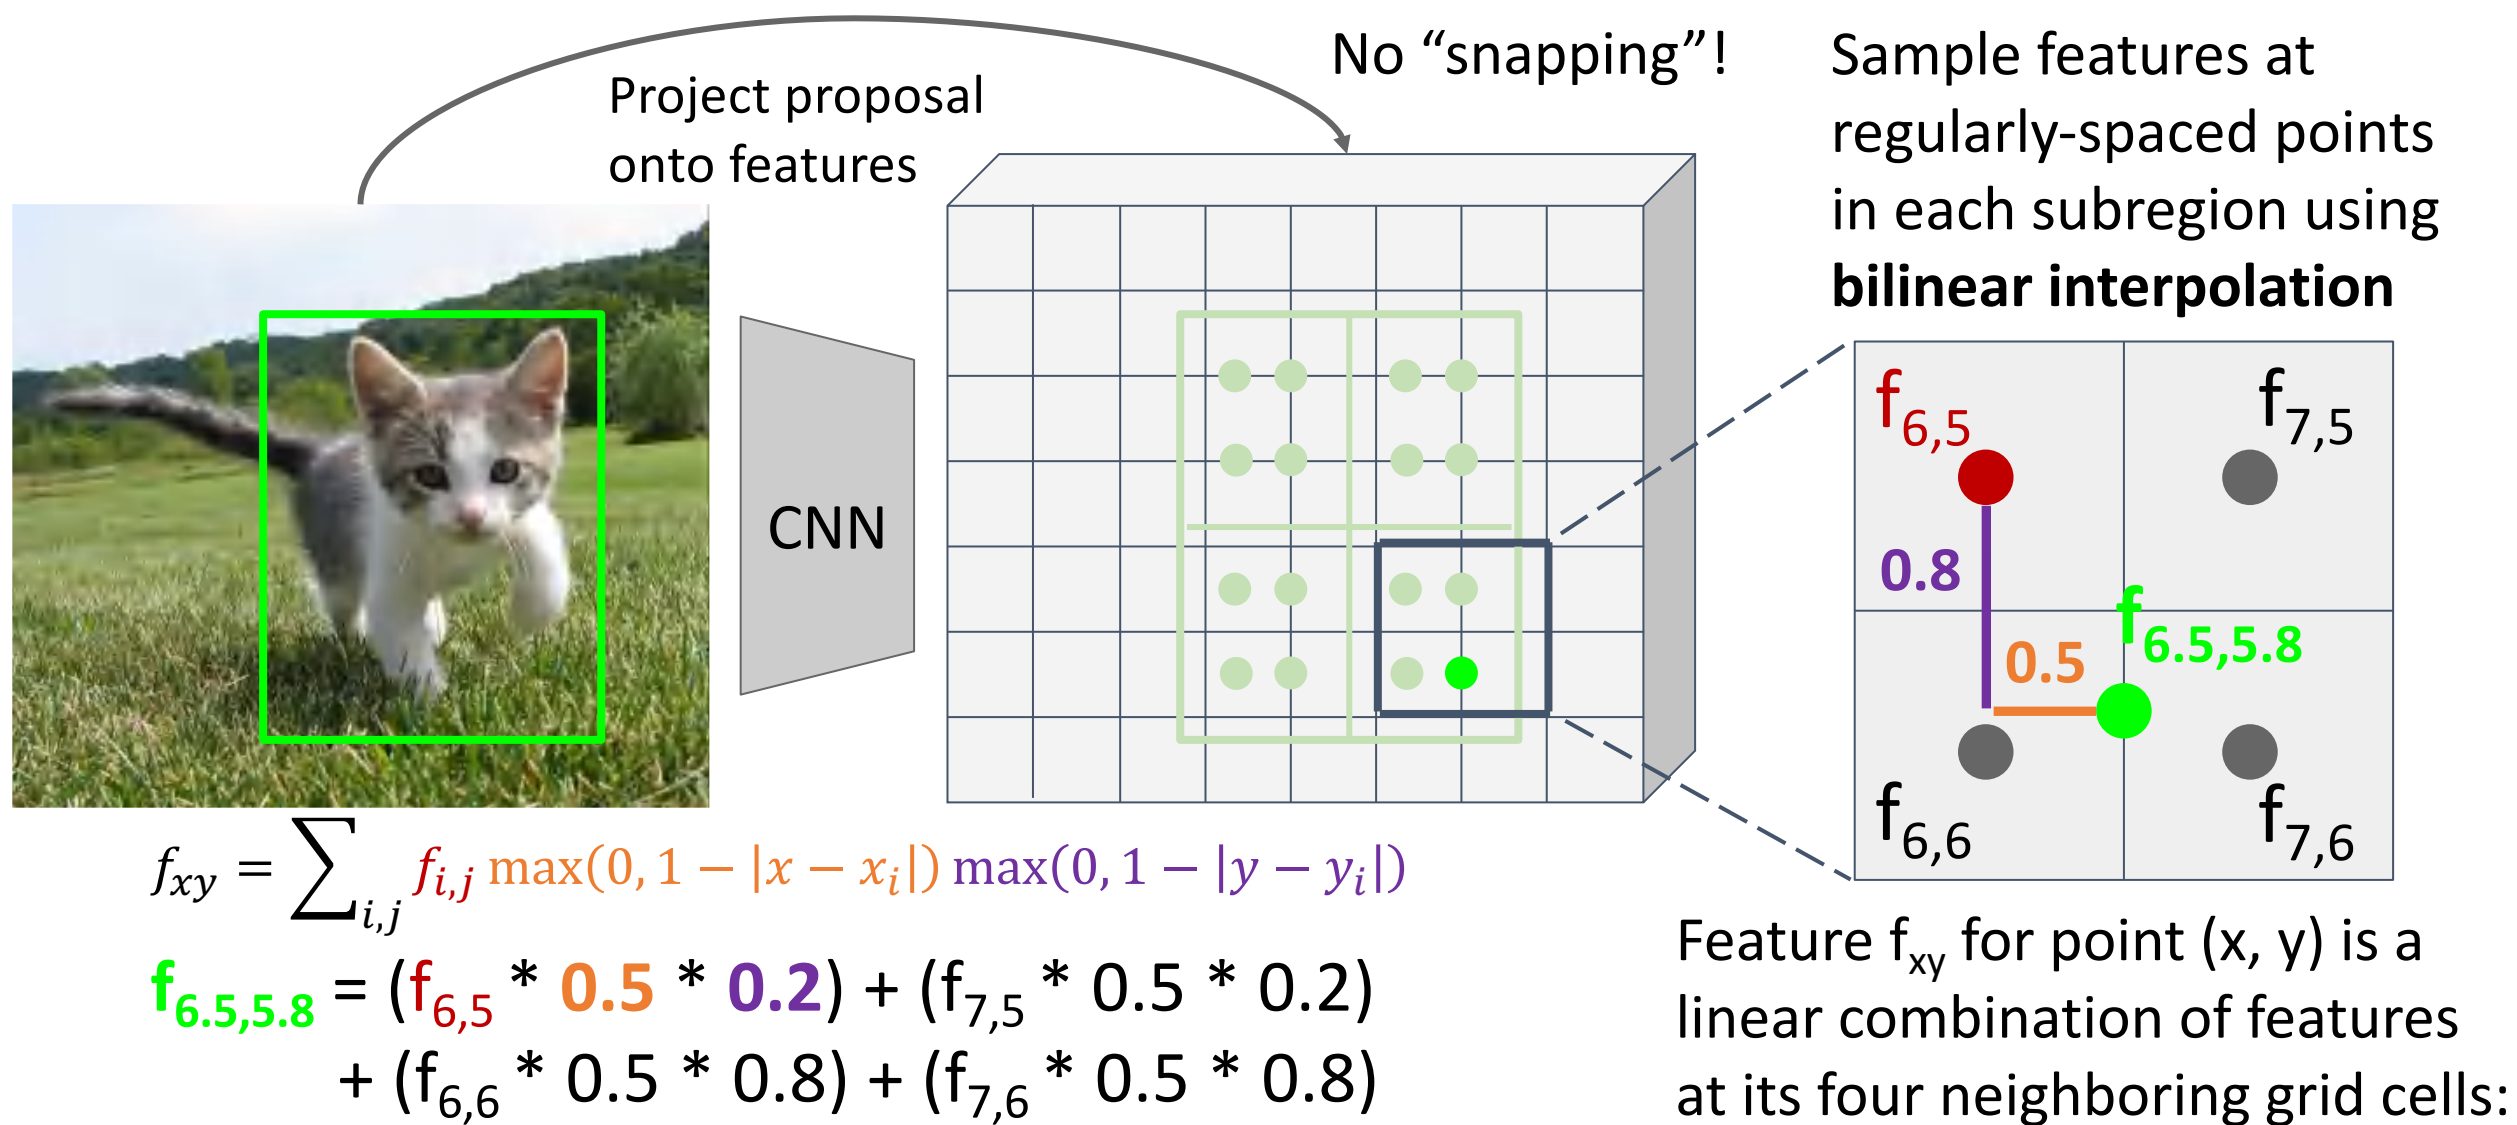


Cropping Features: RoI Align

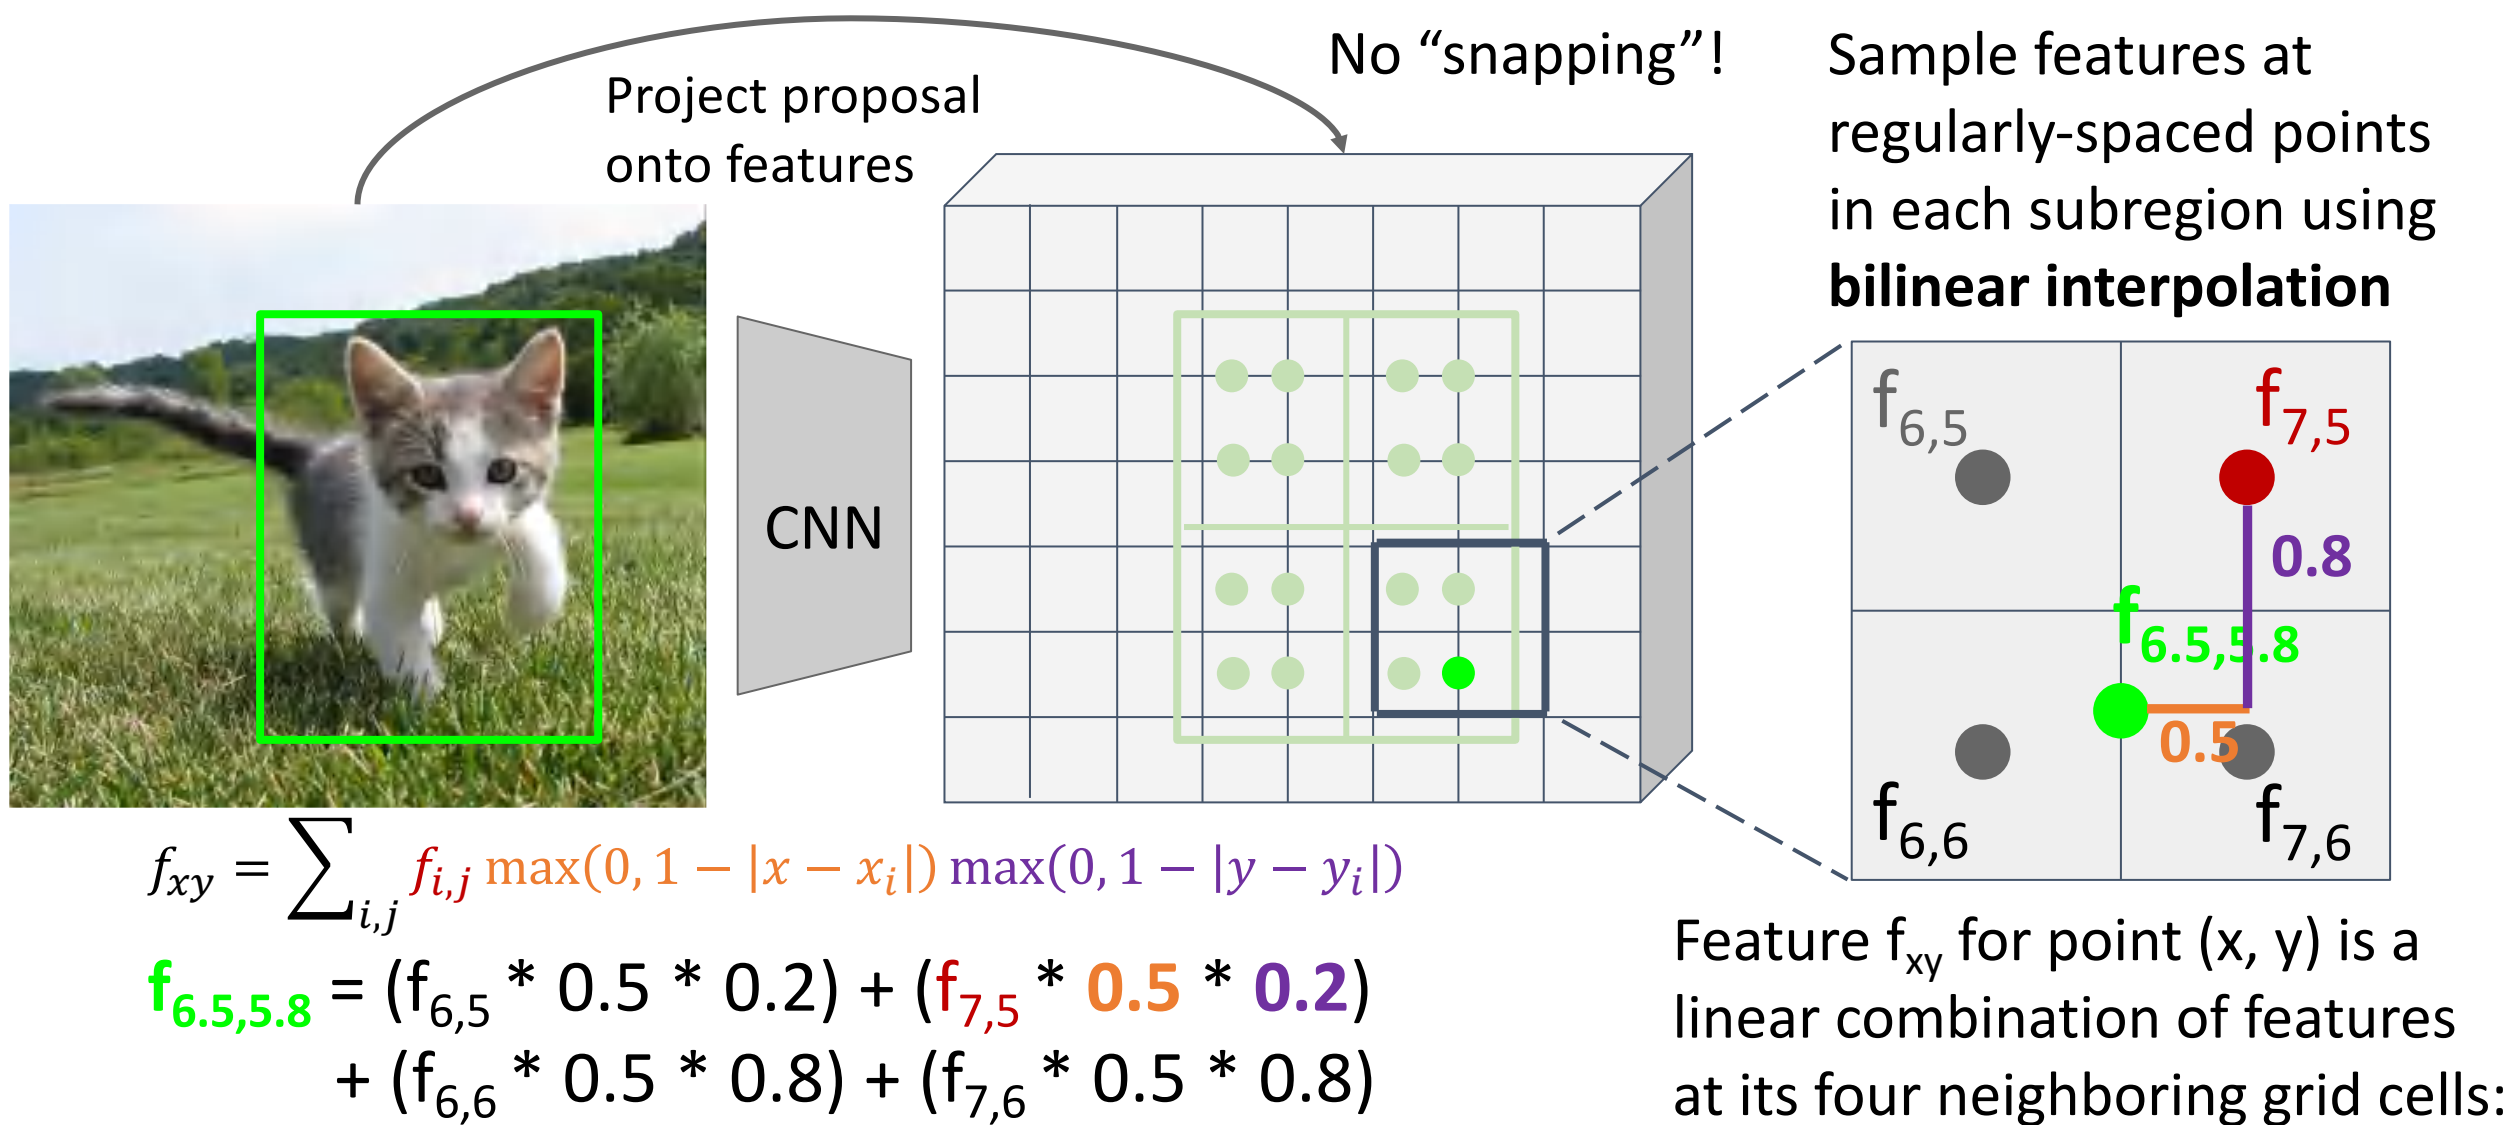
Divide into equal-sized subregions
(may not be aligned to grid!)



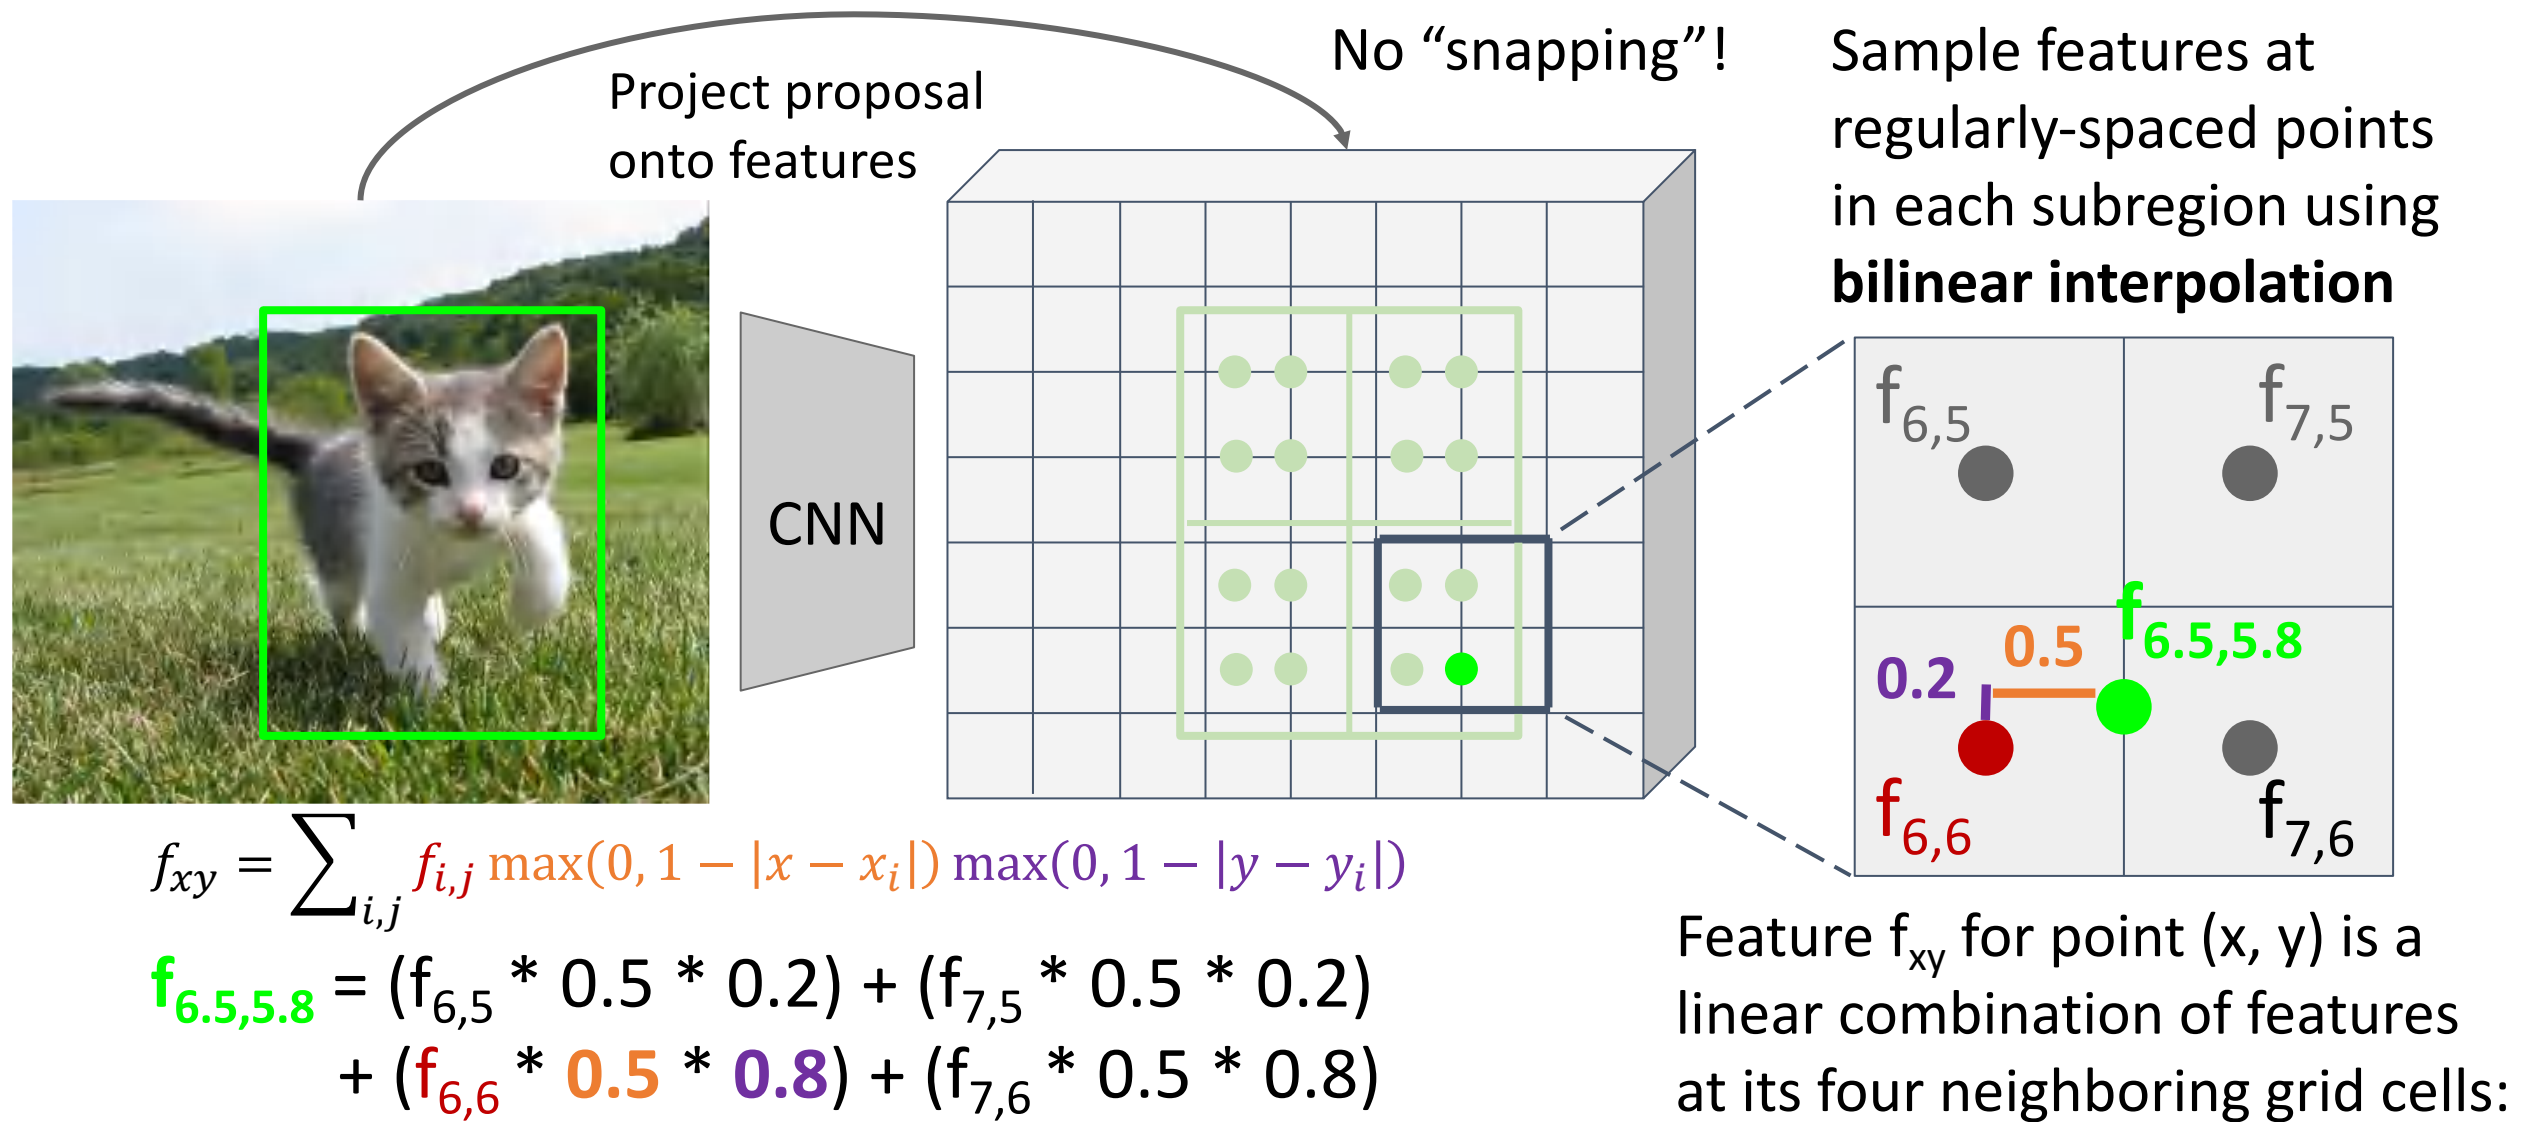
Cropping Features: RoI Align



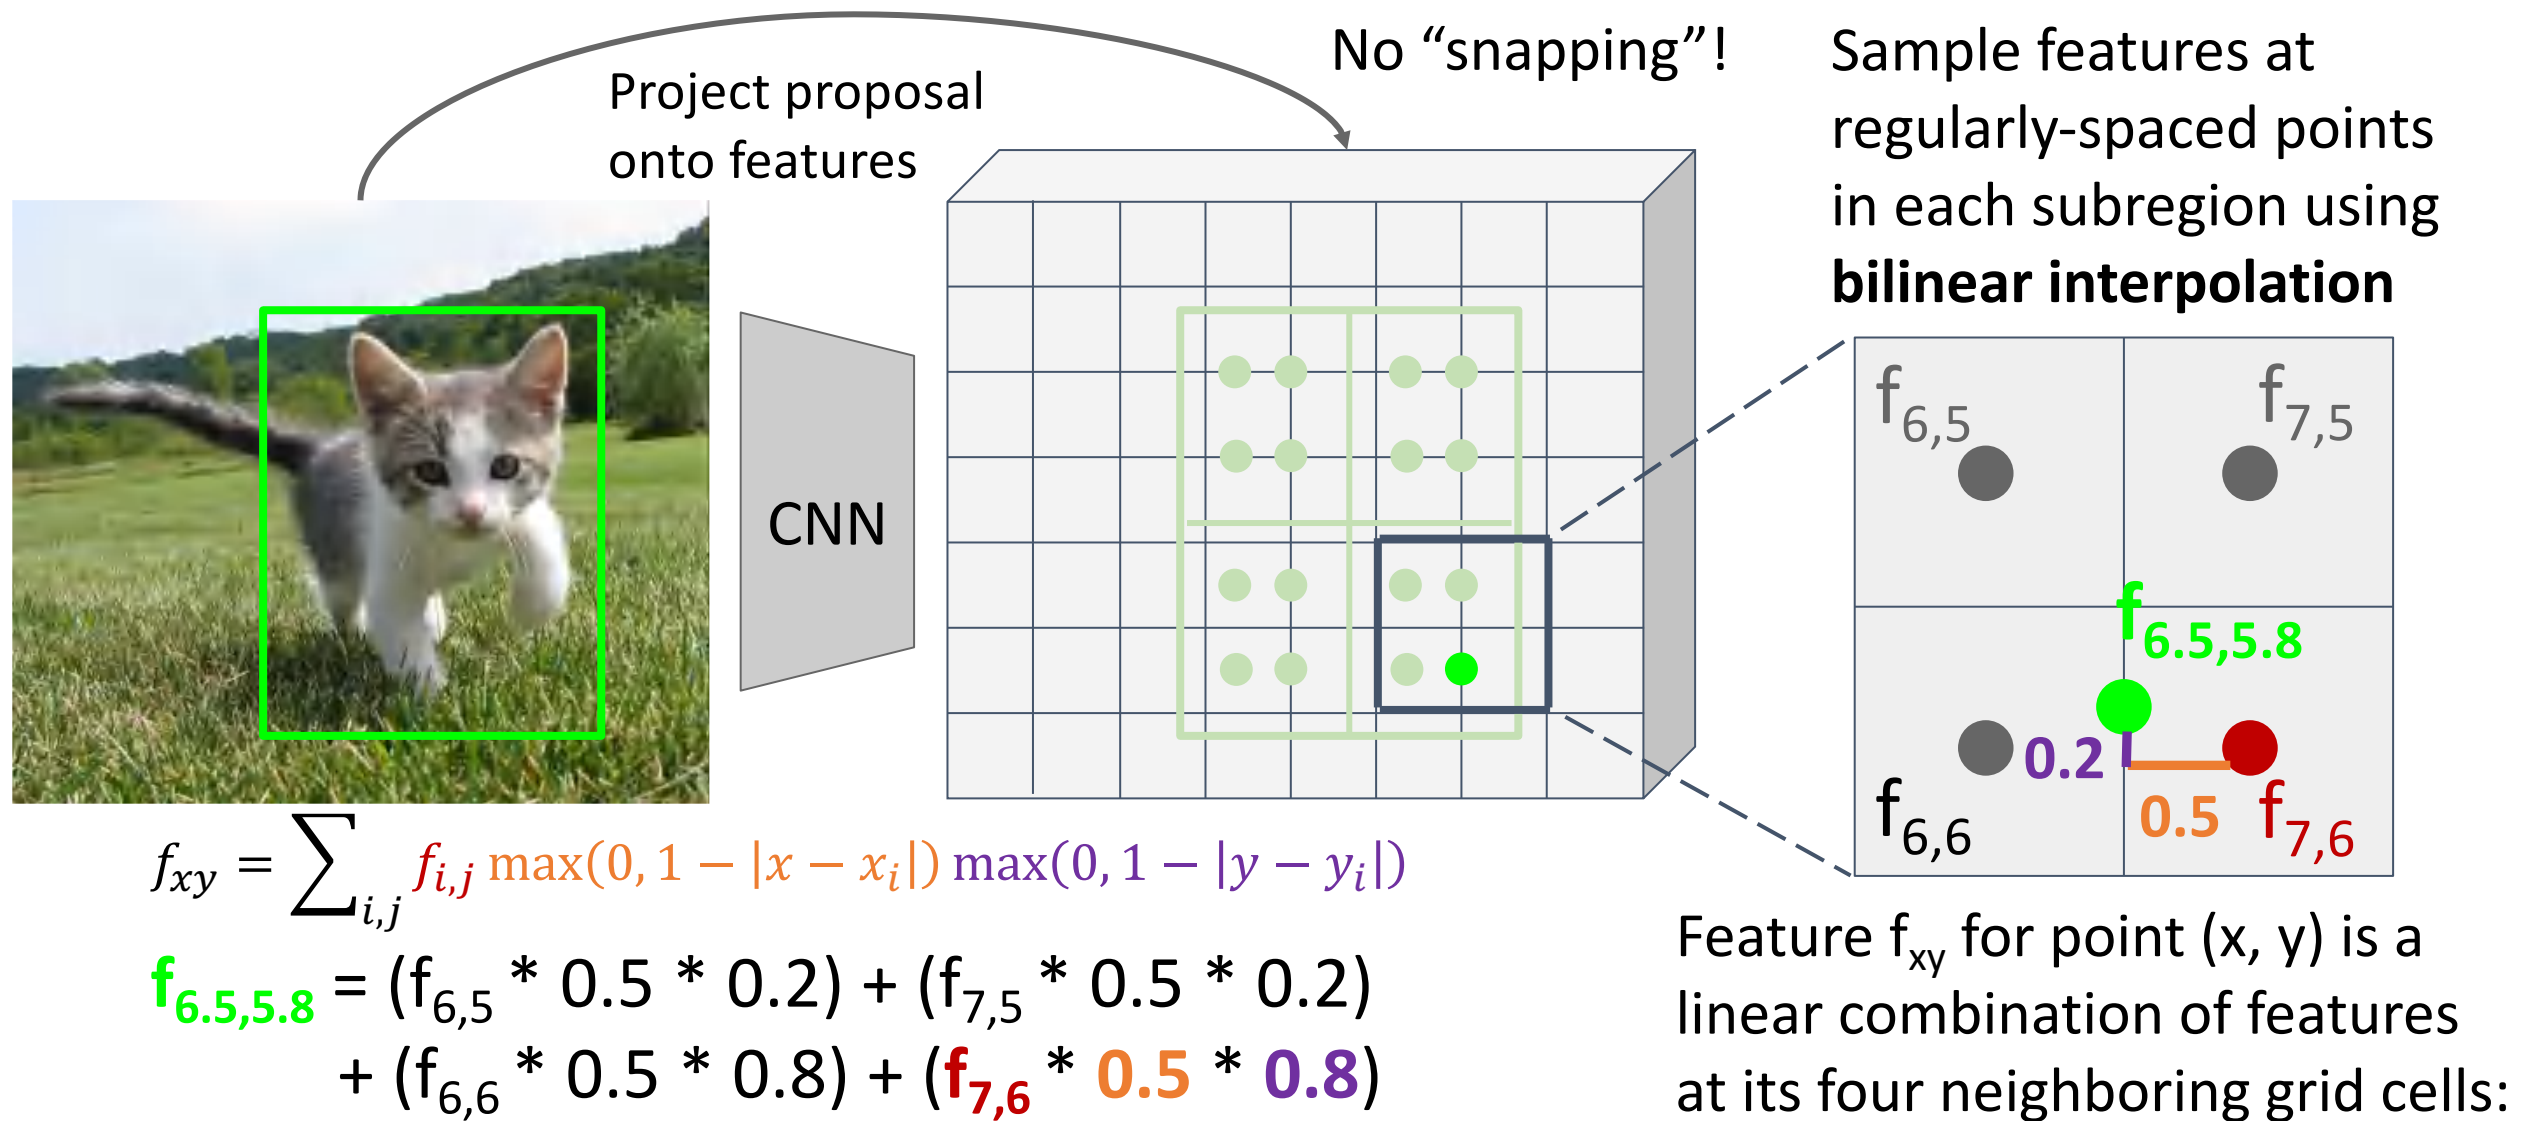
Cropping Features: RoI Align



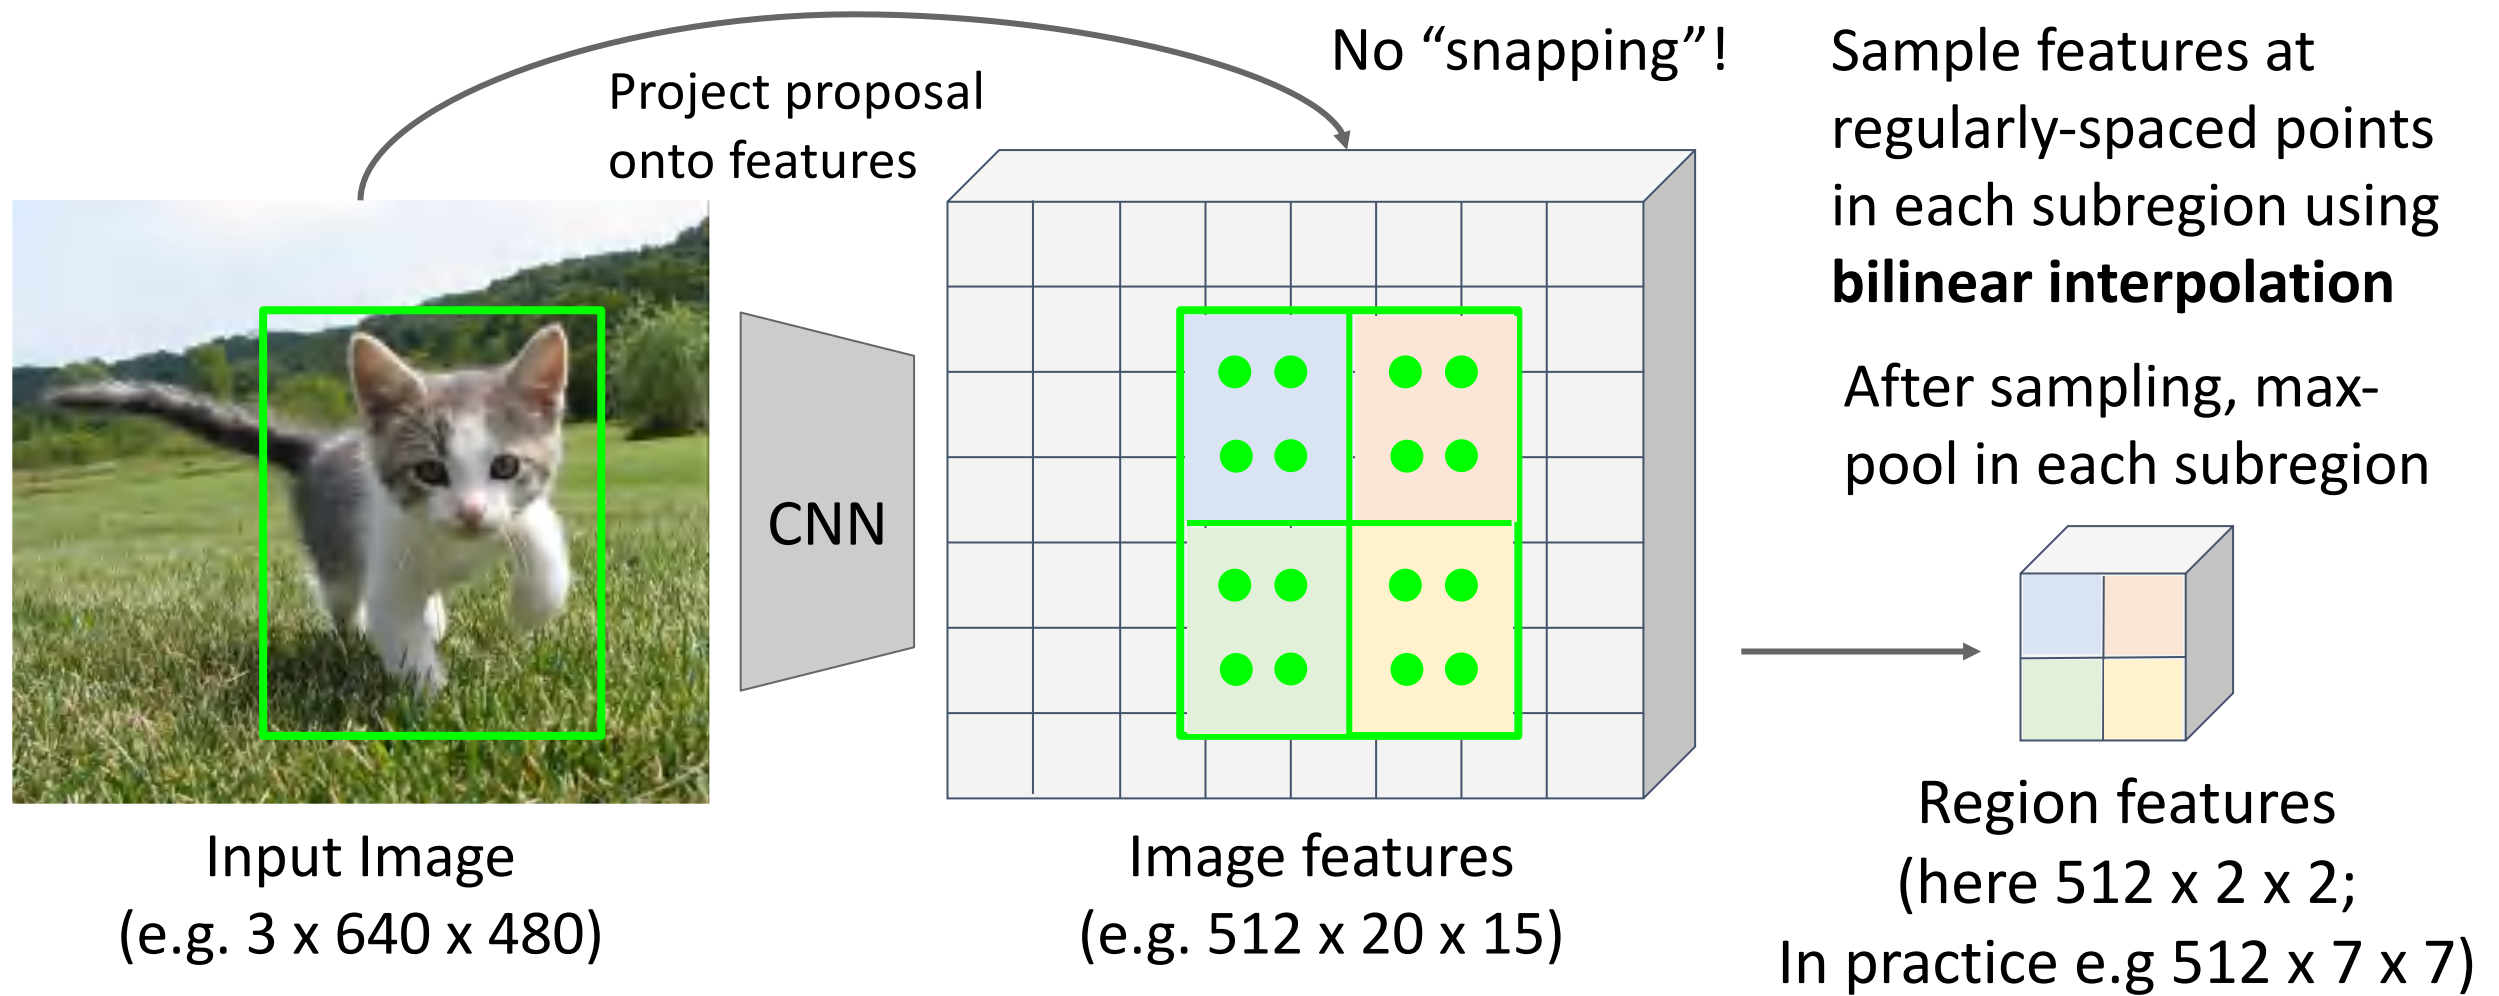
Cropping Features: RoI Align



Cropping Features: RoI Align

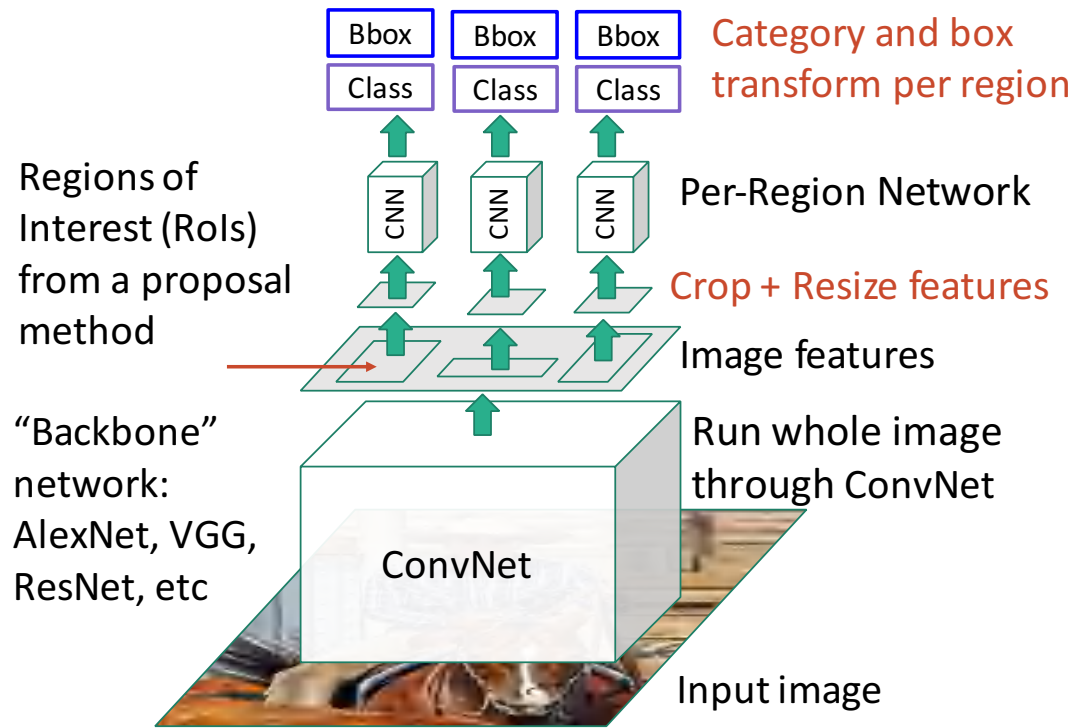


Cropping Features: RoI Align

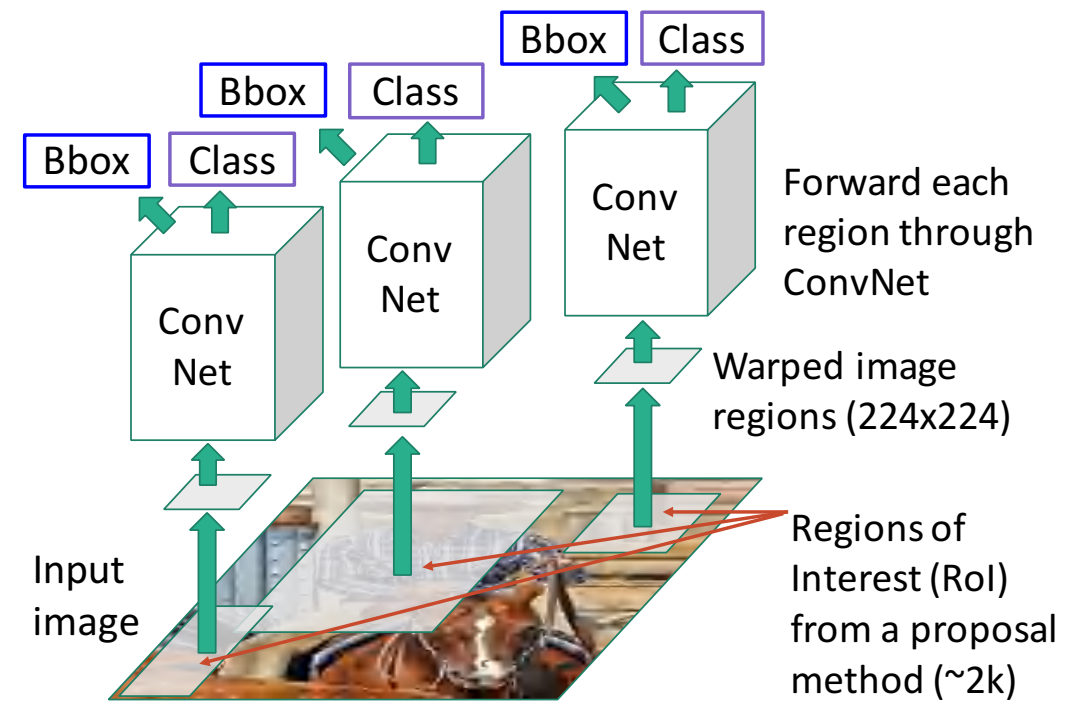


Fast R-CNN vs “Slow” R-CNN

Fast R-CNN: Apply differentiable cropping to shared image features




“Slow” R-CNN: Apply differentiable cropping to shared image features




Fast R-CNN vs “Slow” R-CNN

Training time (Hours)



Test time (seconds)




Girshick et al, “Rich feature hierarchies for accurate object detection and semantic segmentation”, CVPR 2014.

He et al, “Spatial pyramid pooling in deep convolutional networks for visual recognition”, ECCV 2014


Girshick, “Fast R-CNN”, ICCV 2015

Fast R-CNN vs “Slow” R-CNN

Training time (Hours)




Test time (seconds)




Problem: Runtime dominated by region proposals!

Fast R-CNN vs “Slow” R-CNN

Training time (Hours)



Test time (seconds)



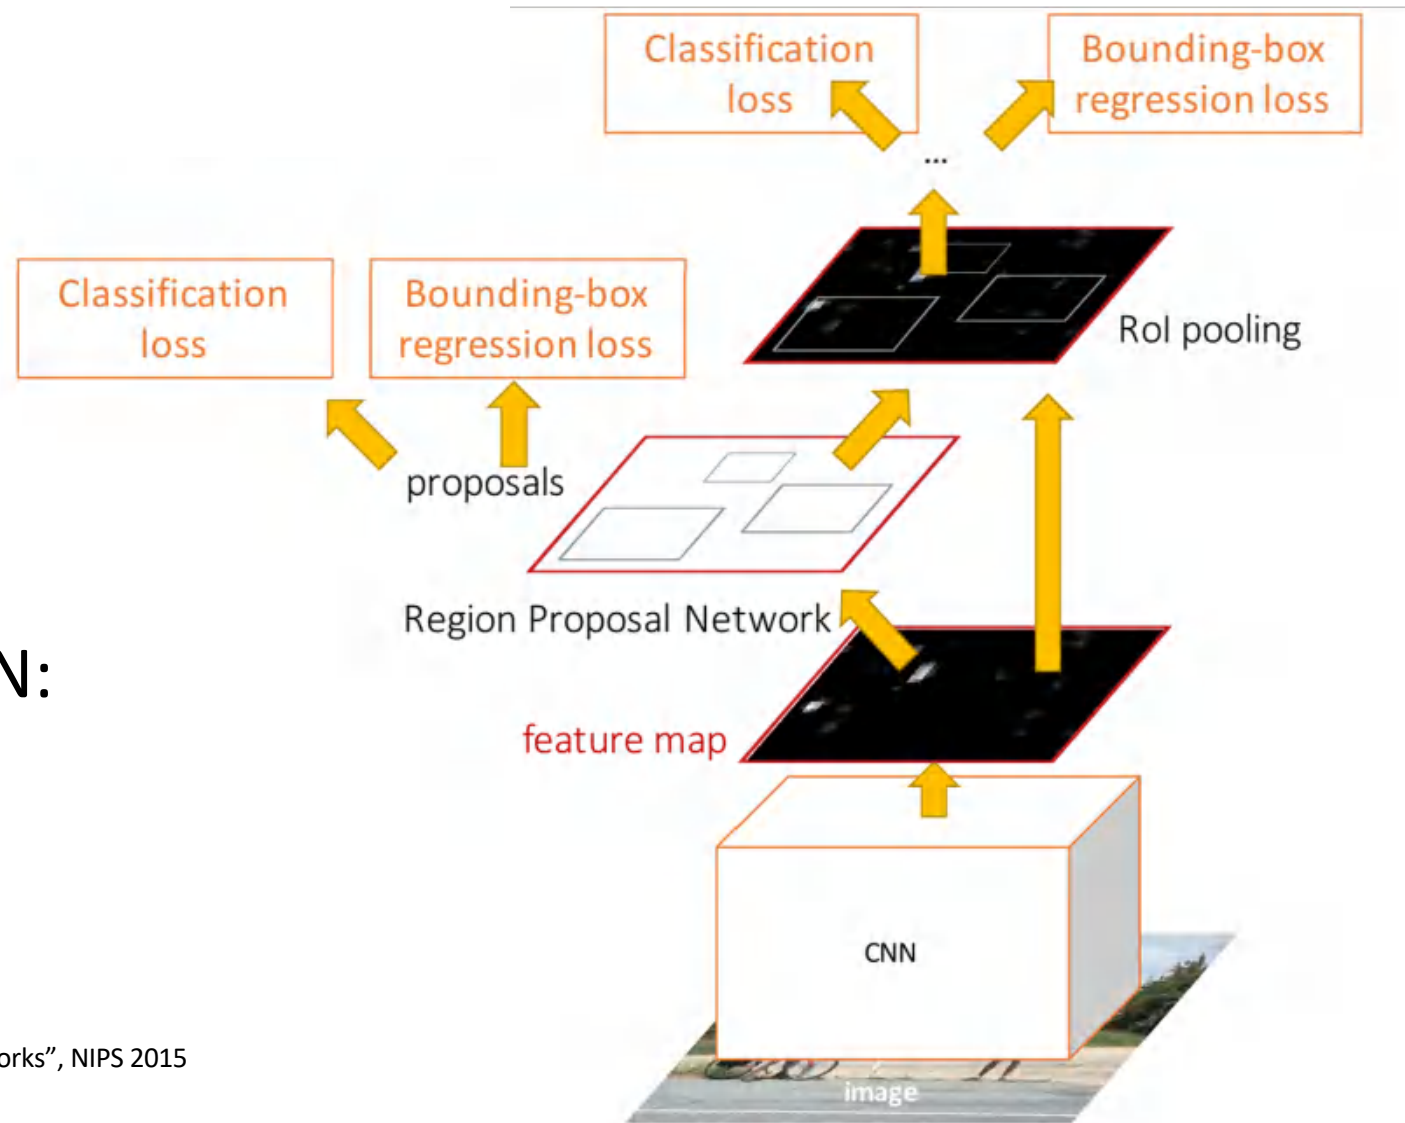
Problem: Runtime dominated by region proposals!

Recall: Region proposals computed by heuristic “Selective Search” algorithm on CPU -- let’s learn them with a CNN instead!

Faster R-CNN: Learnable Region Proposals

Insert **Region Proposal Network (RPN)** to predict proposals from features

Otherwise same as Fast R-CNN:
Crop features for each proposal, classify each one



Ren et al, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015
Figure copyright 2015, Ross Girshick; reproduced with permission

Region Proposal Network (RPN)

Run backbone CNN to get
features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

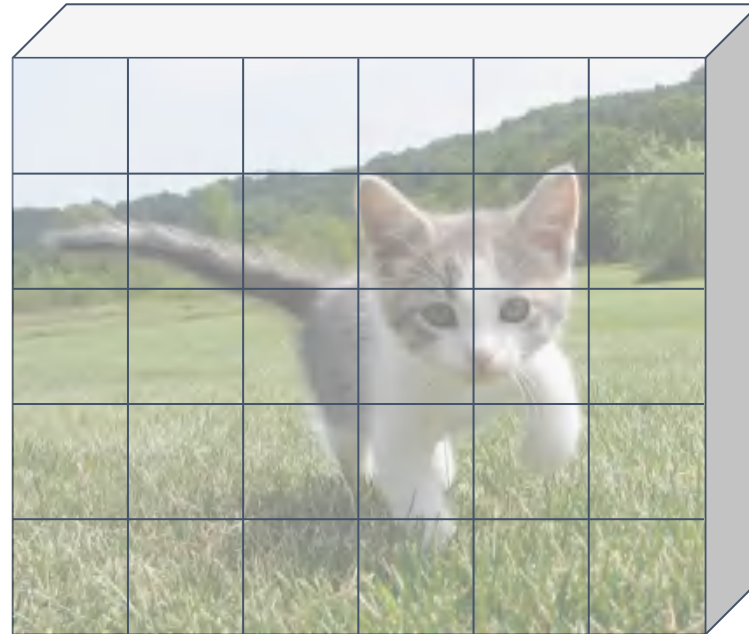


Image features
(e.g. 512 x 5 x 6)

Ren et al, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

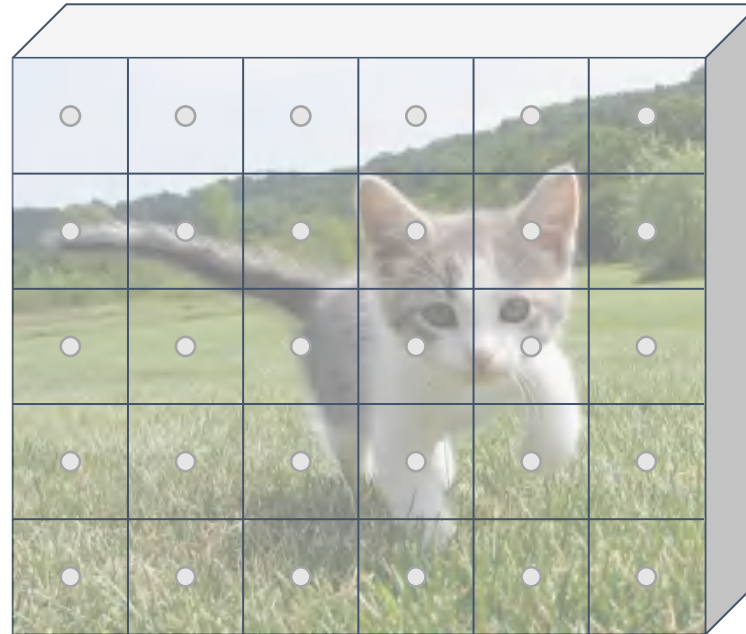



Image features
(e.g. 512 x 5 x 6)

Region Proposal Network (RPN)

Imagine an **anchor box**
of fixed size at each
point in the feature map

Run backbone CNN to get
features aligned to input image

Each feature corresponds
to a point in the input



Input Image
(e.g. 3 x 640 x 480)

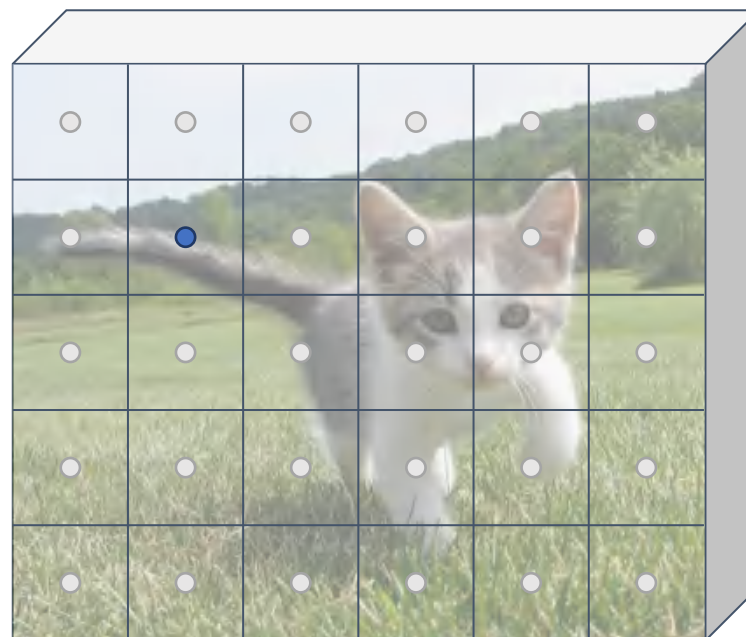



Image features
(e.g. 512 x 5 x 6)

Region Proposal Network (RPN)

Imagine an **anchor box**
of fixed size at each
point in the feature map

Run backbone CNN to get
features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds
to a point in the input

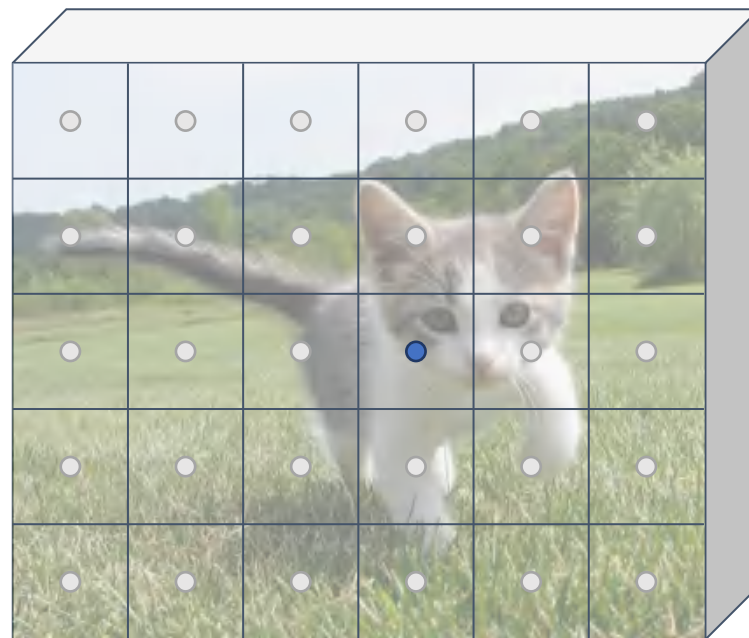



Image features
(e.g. 512 x 5 x 6)

Region Proposal Network (RPN)

Imagine an **anchor box**
of fixed size at each
point in the feature map

Run backbone CNN to get
features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds
to a point in the input

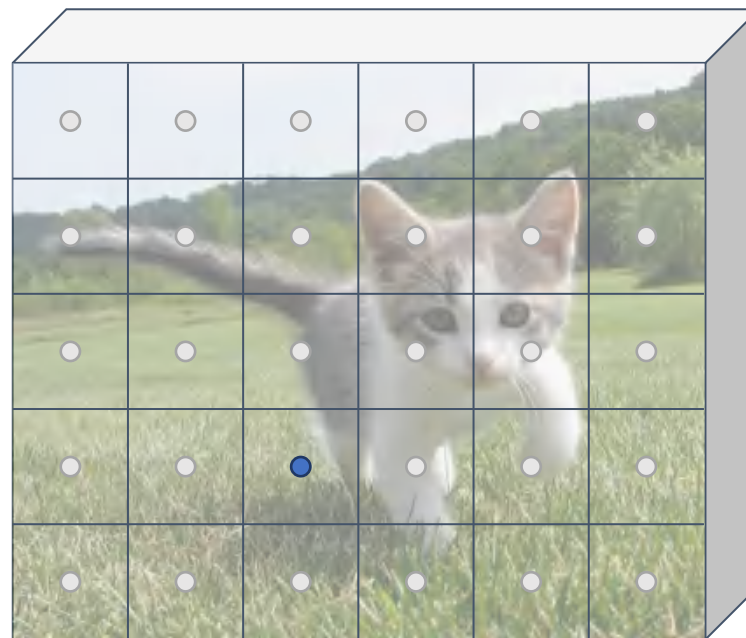



Image features
(e.g. 512 x 5 x 6)

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

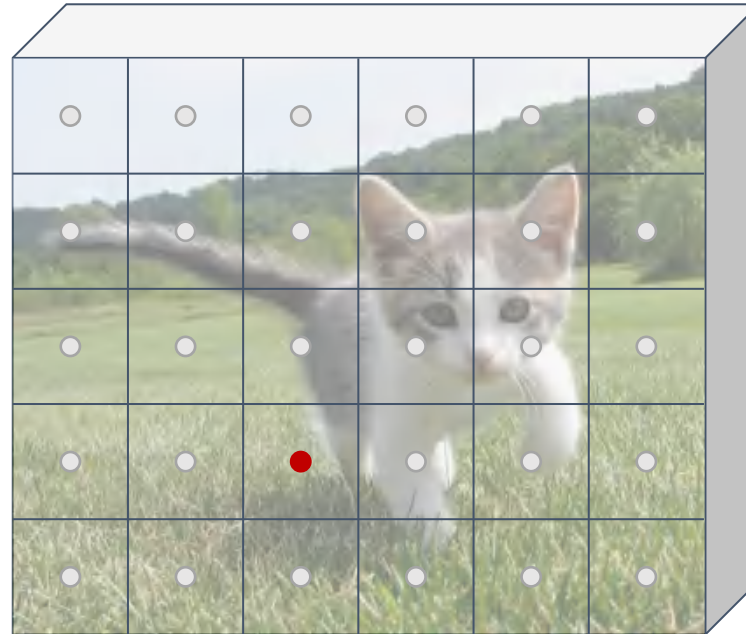


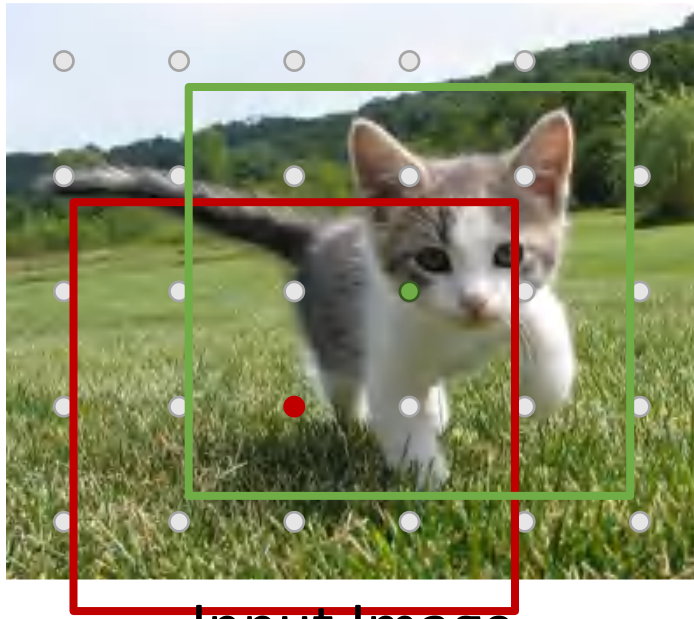
Image features
(e.g. 512 x 5 x 6)

Imagine an **anchor box** of fixed size at each point in the feature map

Classify each anchor as **positive (object)** or **negative (no object)**

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

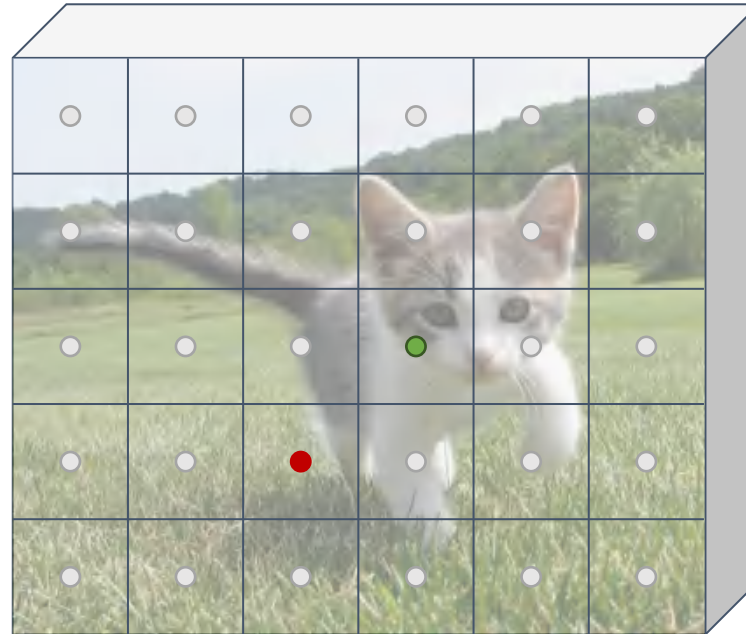



Image features
(e.g. 512 x 5 x 6)

Imagine an **anchor box** of fixed size at each point in the feature map

Classify each anchor as **positive (object)** or **negative (no object)**

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

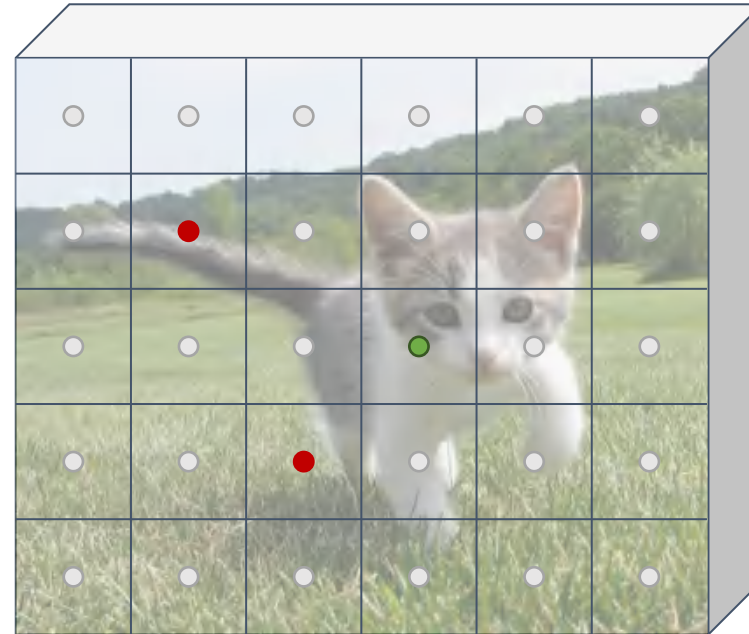


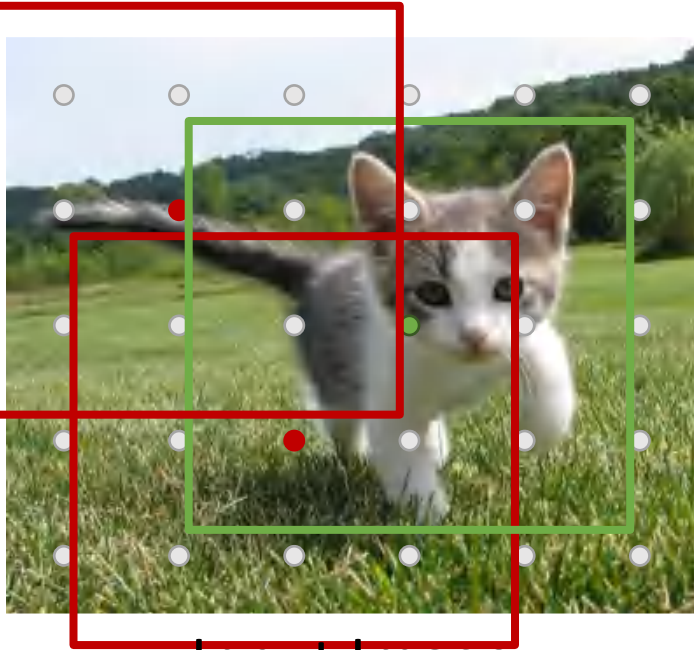
Image features
(e.g. 512 x 5 x 6)

Imagine an **anchor box** of fixed size at each point in the feature map

Classify each anchor as **positive (object)** or **negative (no object)**

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

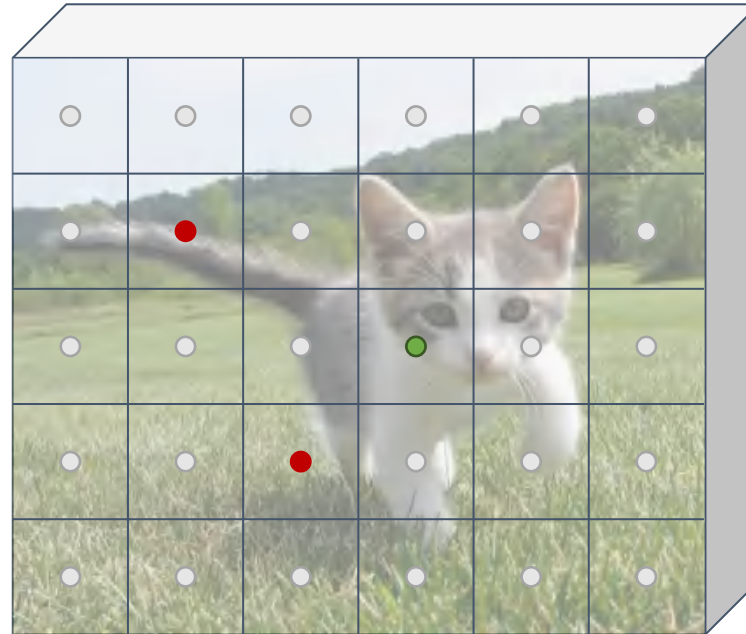


Image features
(e.g. 512 x 5 x 6)

Predict object vs not object scores for all anchors with a conv layer (512 input filters, 2 output filters)




Anchor is
object?
2 x 5 x 6

Classify each anchor as
positive (object) or
negative (no object)

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

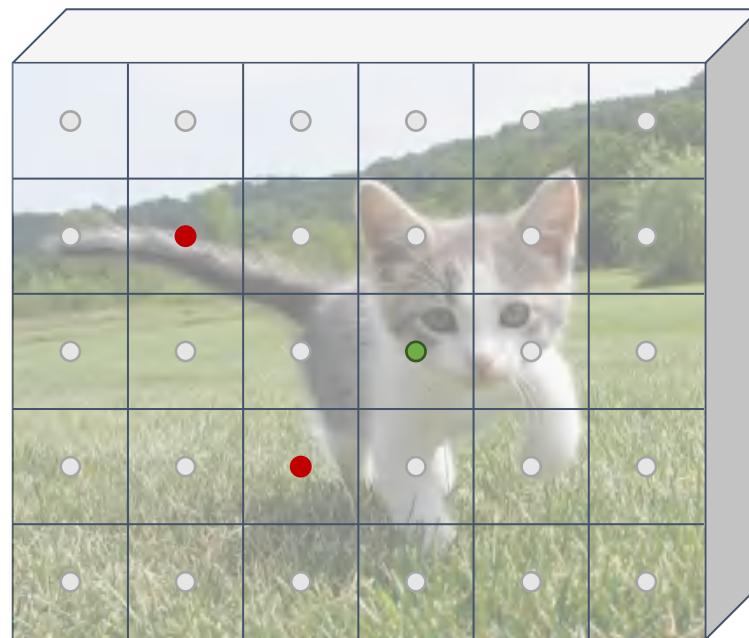


Image features
(e.g. 512 x 5 x 6)



For **positive anchors**, also predict a **transform** that converting the anchor to the **GT box** (like R-CNN)




Anchor is
object?
2 x 5 x 6

Classify each anchor as
positive (object) or
negative (no object)

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

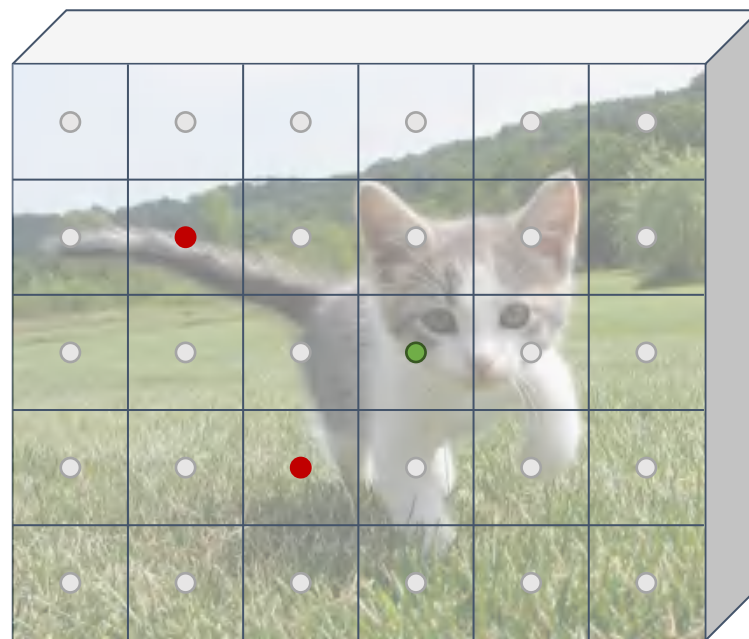


Image features
(e.g. 512 x 5 x 6)



For **positive anchors**, also predict a **transform** that converting the anchor to the **GT box** (like R-CNN)
Predict transforms with conv



Anchor is
object?
2 x 5 x 6

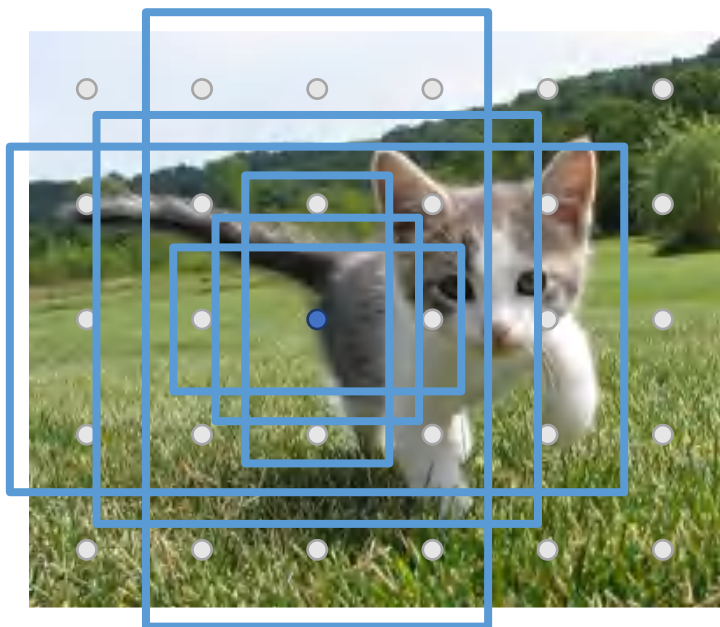


Anchor
transforms
4 x 5 x 6

Classify each anchor as
positive (object) or
negative (no object)

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

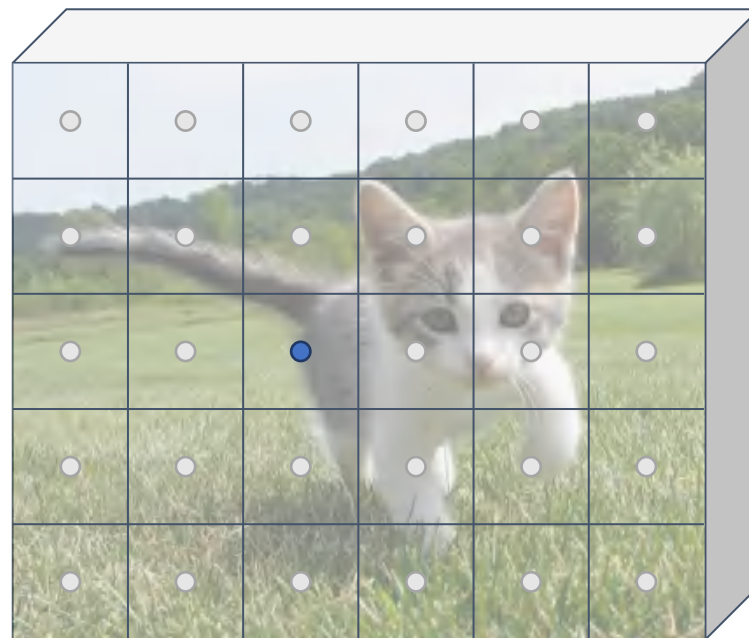


Image features
(e.g. 512 x 5 x 6)



Anchor is
object?
 $2K \times 5 \times 6$




Anchor
transforms
 $4K \times 5 \times 6$

In practice: Rather than using one anchor per point, instead consider K different anchors with different size and scale (here $K = 6$)

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

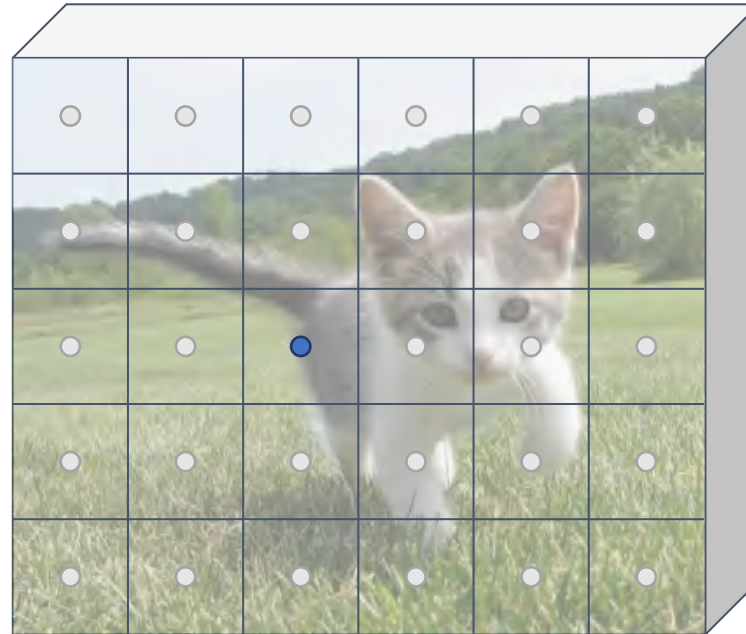


Image features
(e.g. 512 x 5 x 6)

In practice: Rather than using one anchor per point, instead consider K different anchors with different size and scale (here K = 6)




Anchor is
object?
2K x 5 x 6

Anchor
transforms
4K x 5 x 6

During training, supervised positive / negative anchors and box transforms like R-CNN

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input




Image features
(e.g. 512 x 5 x 6)

In practice: Rather than using one anchor per point, instead consider K different anchors with different size and scale (here K = 6)




Anchor is object?
2K x 5 x 6

Anchor transforms
4K x 5 x 6

Positive anchors: ≥ 0.7 IoU with some GT box (plus highest IoU to each GT)

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input




Image features
(e.g. 512 x 5 x 6)

In practice: Rather than using one anchor per point, instead consider K different anchors with different size and scale (here $K = 6$)



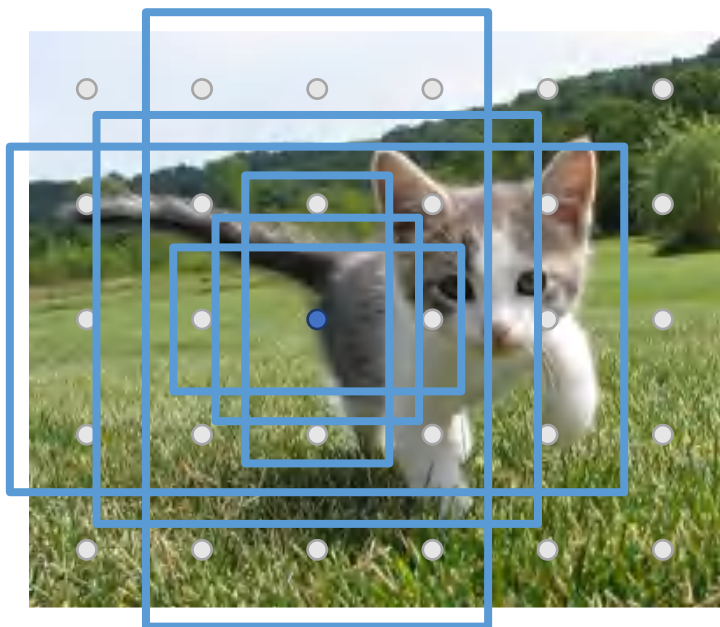
Anchor is
object?
 $2K \times 5 \times 6$

Anchor
transforms
 $4K \times 5 \times 6$

Negative anchors: < 0.3 IoU with all GT boxes. Don't supervised transforms for negative boxes.

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input

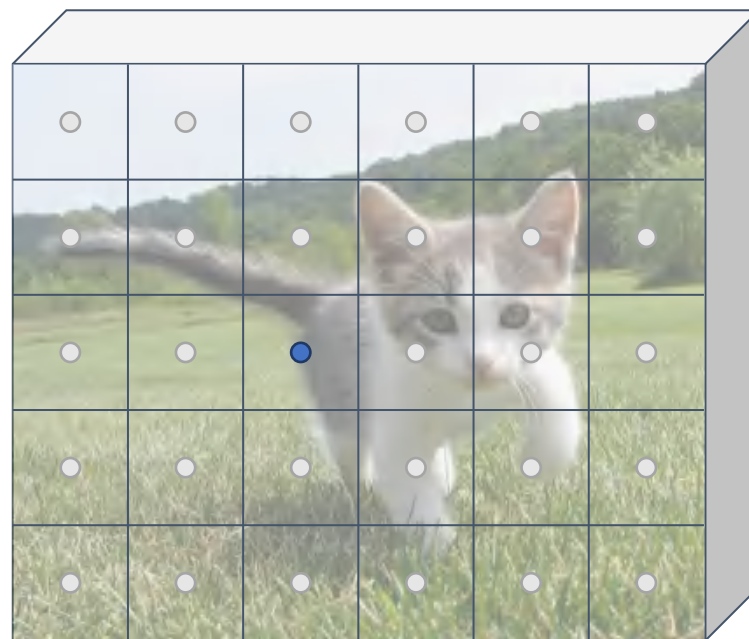


Image features
(e.g. 512 x 5 x 6)

In practice: Rather than using one anchor per point, instead consider K different anchors with different size and scale (here K = 6)




Anchor is
object?
2K x 5 x 6

Anchor
transforms
4K x 5 x 6

Neutral anchors: between 0.3 and 0.7 IoU with all GT boxes; ignored during training

Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)



Each feature corresponds to a point in the input




Image features
(e.g. 512 x 5 x 6)

In practice: Rather than using one anchor per point, instead consider K different anchors with different size and scale (here $K = 6$)



Anchor is
object?
 $2K \times 5 \times 6$


Anchor
transforms
 $4K \times 5 \times 6$

At test-time, sort all $K \times 5 \times 6$ boxes by their positive score, take top 300 as our region proposals

Faster R-CNN: Learnable Region Proposals


Jointly train with 4 losses:

1. **RPN classification:** anchor box is object / not an object
2. **RPN regression:** predict transform from anchor box to proposal box
3. **Object classification:** classify proposals as background / object class
4. **Object regression:** predict transform from proposal box to object box



Ren et al, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015
Figure copyright 2015, Ross Girshick; reproduced with permission

Faster R-CNN: Learnable Region Proposals



Faster R-CNN: Learnable Region Proposals

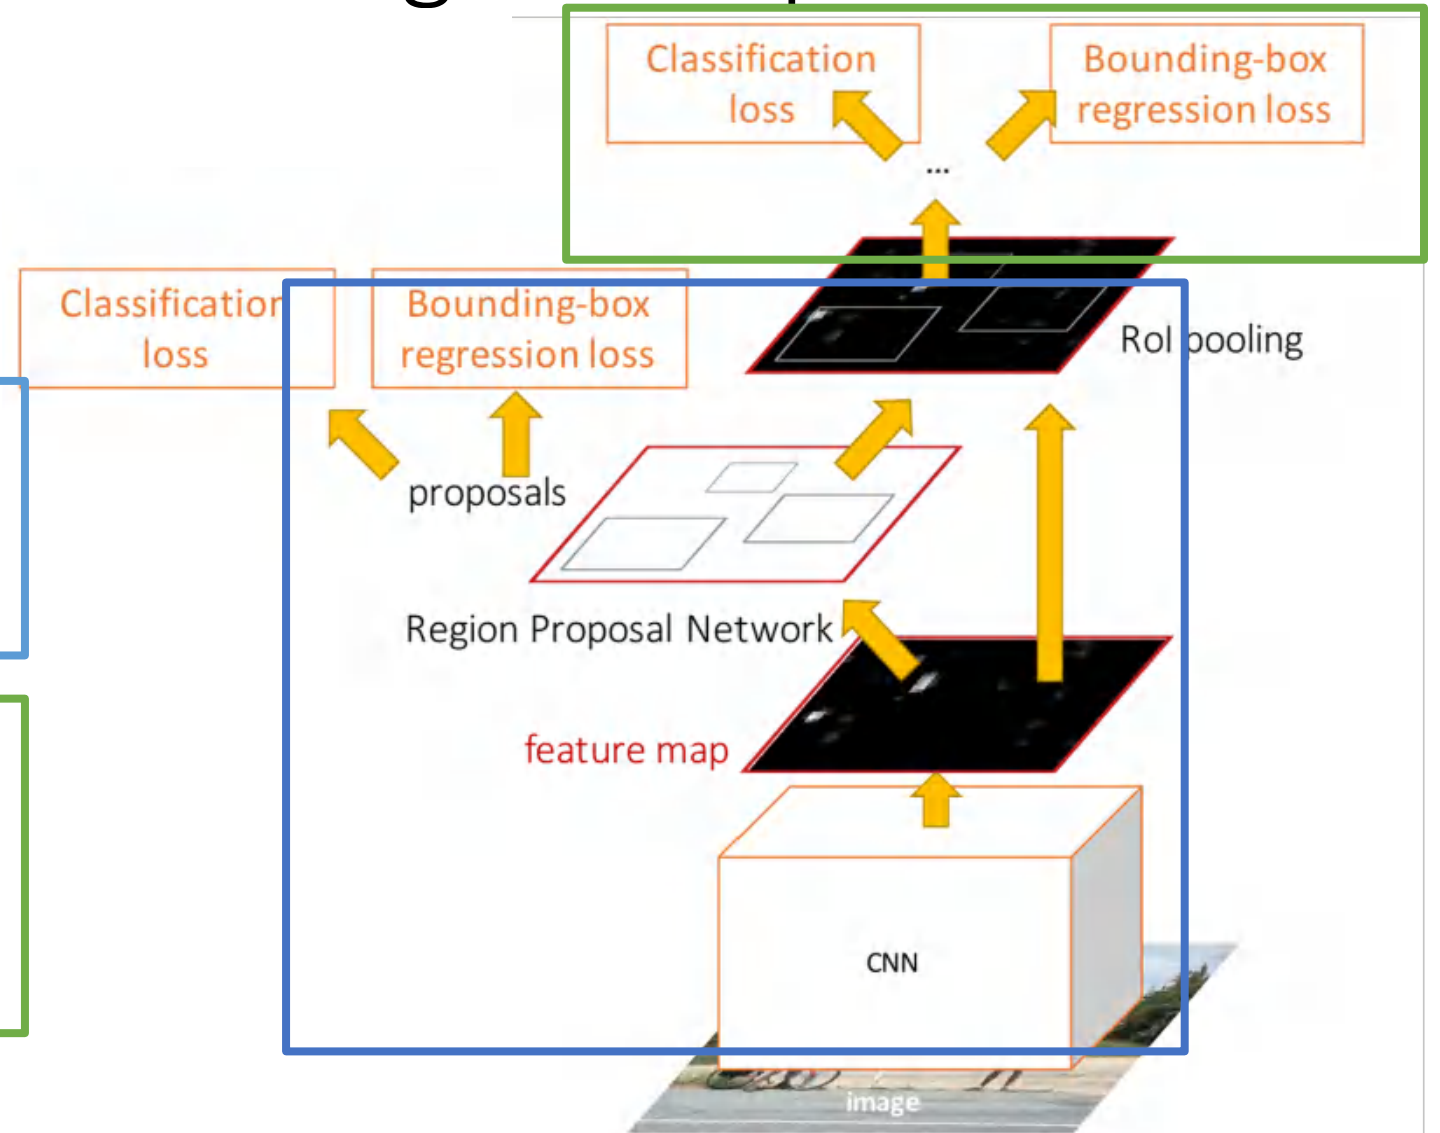
Faster R-CNN is a
Two-stage object detector

First stage: Run once per image

- Backbone network
- Region proposal network


Second stage: Run once per region

- Crop features: RoI pool / align
- Predict object class
- Prediction bbox offset



Dealing with Scale

We need to detect objects of many different scales.
How to improve *scale invariance* of the detector?



[This image](#) is free for commercial use under the [Pixabay license](#)

Dealing with Scale: Image Pyramid

Classic idea: build an *image pyramid* by resizing the image to different scales, then process each image scale independently.



Object
Detector



Object
Detector



Object
Detector

Lin et al, "Feature Pyramid Networks
for Object Detection", ICCV 2017

Dealing with Scale: Image Pyramid

Classic idea: build an *image pyramid* by resizing the image to different scales, then process each image scale independently.

Problem: Expensive! Don't share any computation between scales



Object
Detector



Object
Detector

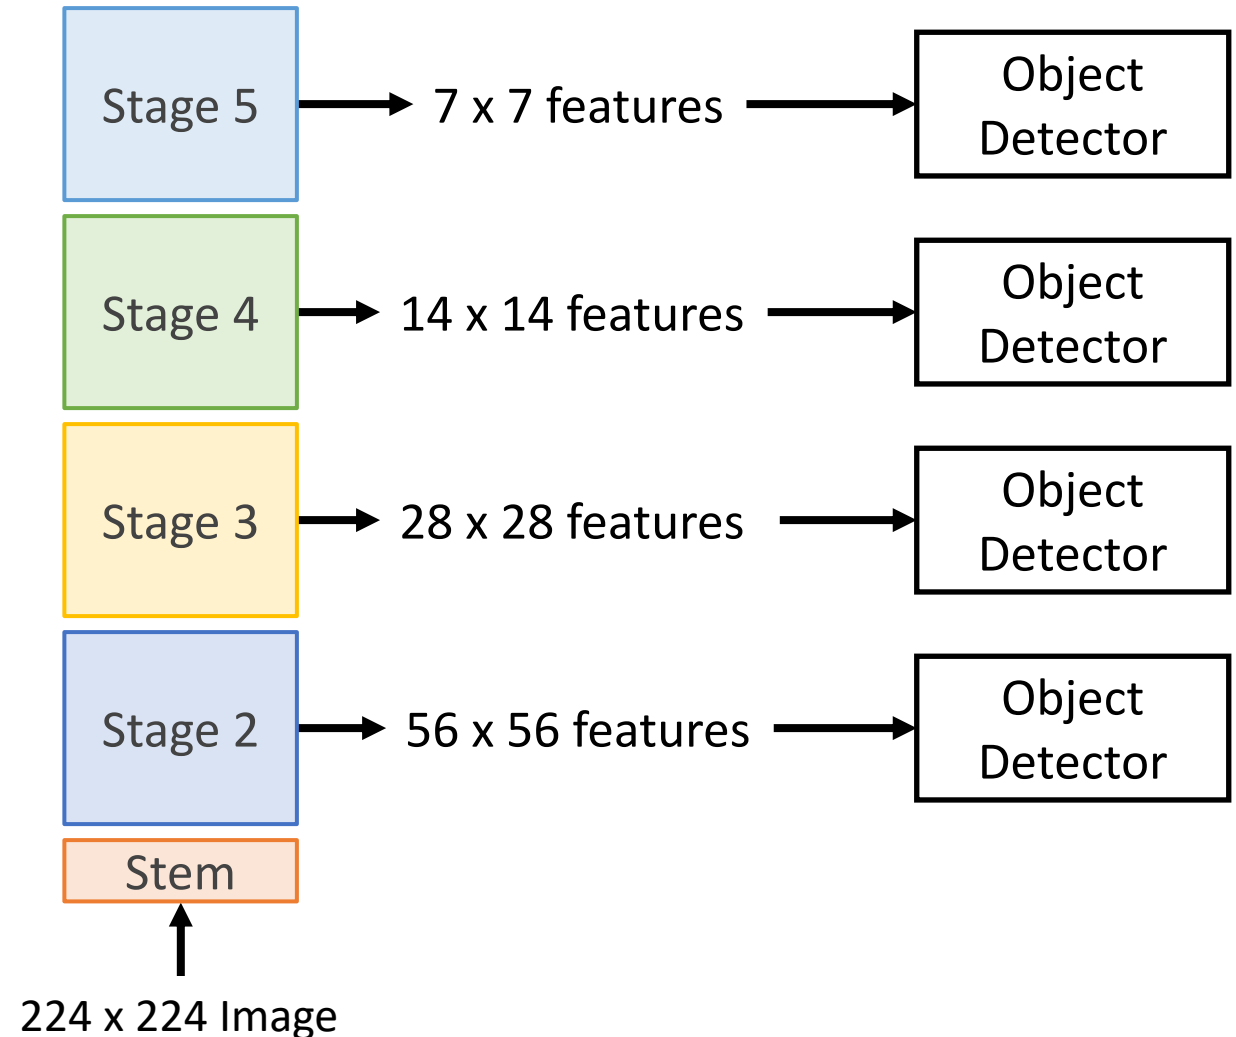


Object
Detector

Lin et al, "Feature Pyramid Networks for Object Detection", ICCV 2017

Dealing with Scale: Multiscale Features

CNNs have multiple *stages* that operate at different resolutions. Attach an independent detector to the features at each level

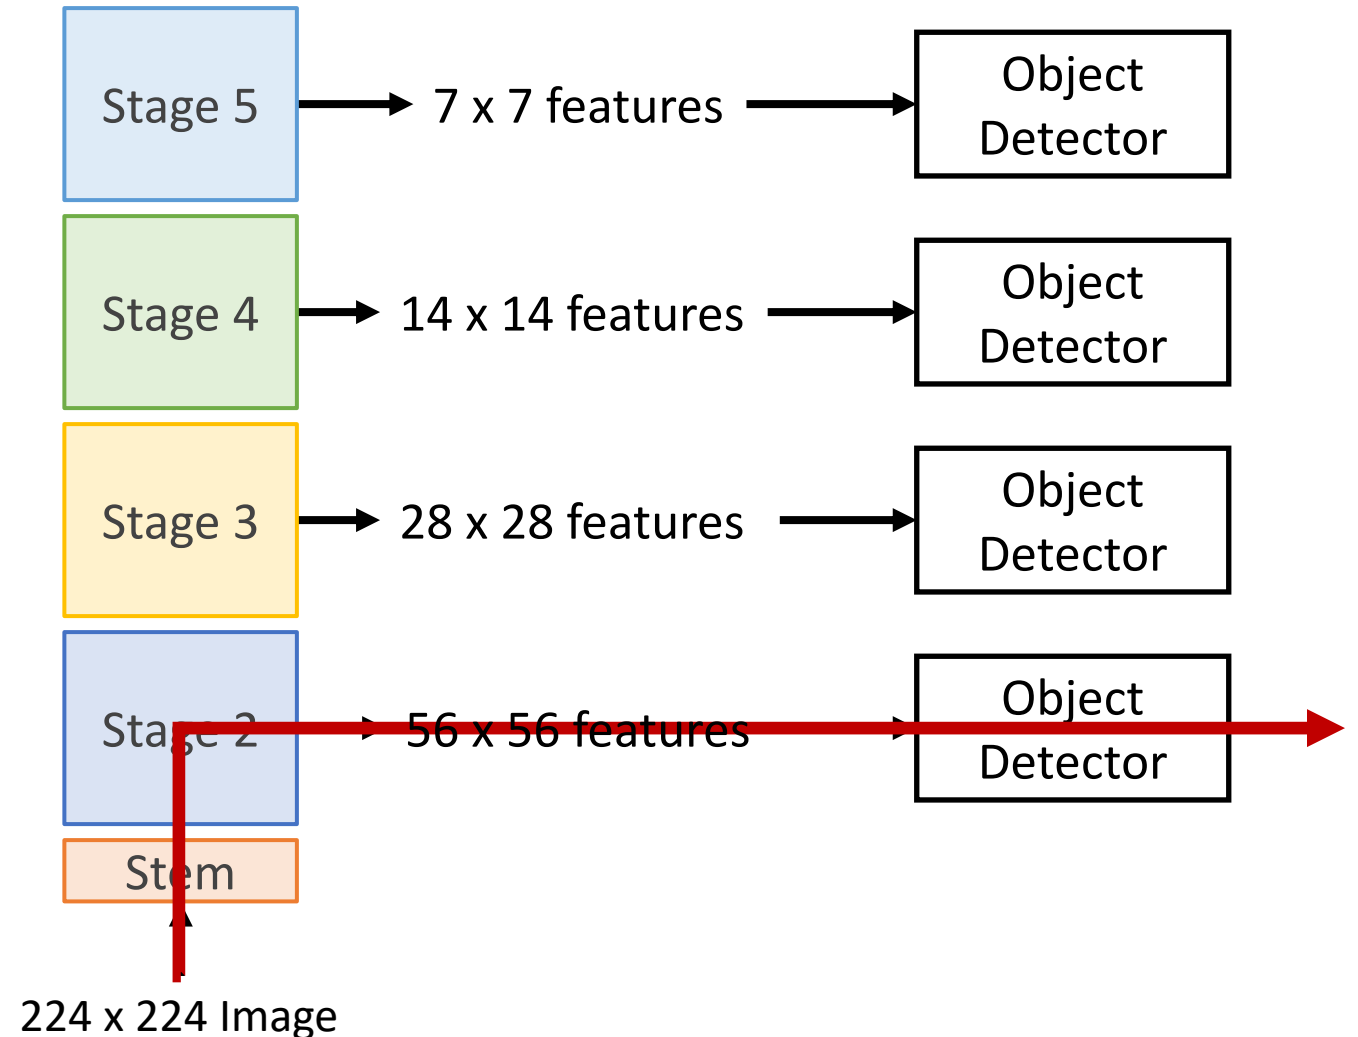


Lin et al, "Feature Pyramid Networks for Object Detection", ICCV 2017

Dealing with Scale: Multiscale Features

CNNs have multiple *stages* that operate at different resolutions. Attach an independent detector to the features at each level


Problem: detector on early features doesn't make use of the entire backbone; doesn't get access to high-level features



Lin et al, "Feature Pyramid Networks for Object Detection", ICCV 2017

Dealing with Scale: Feature Pyramid Network

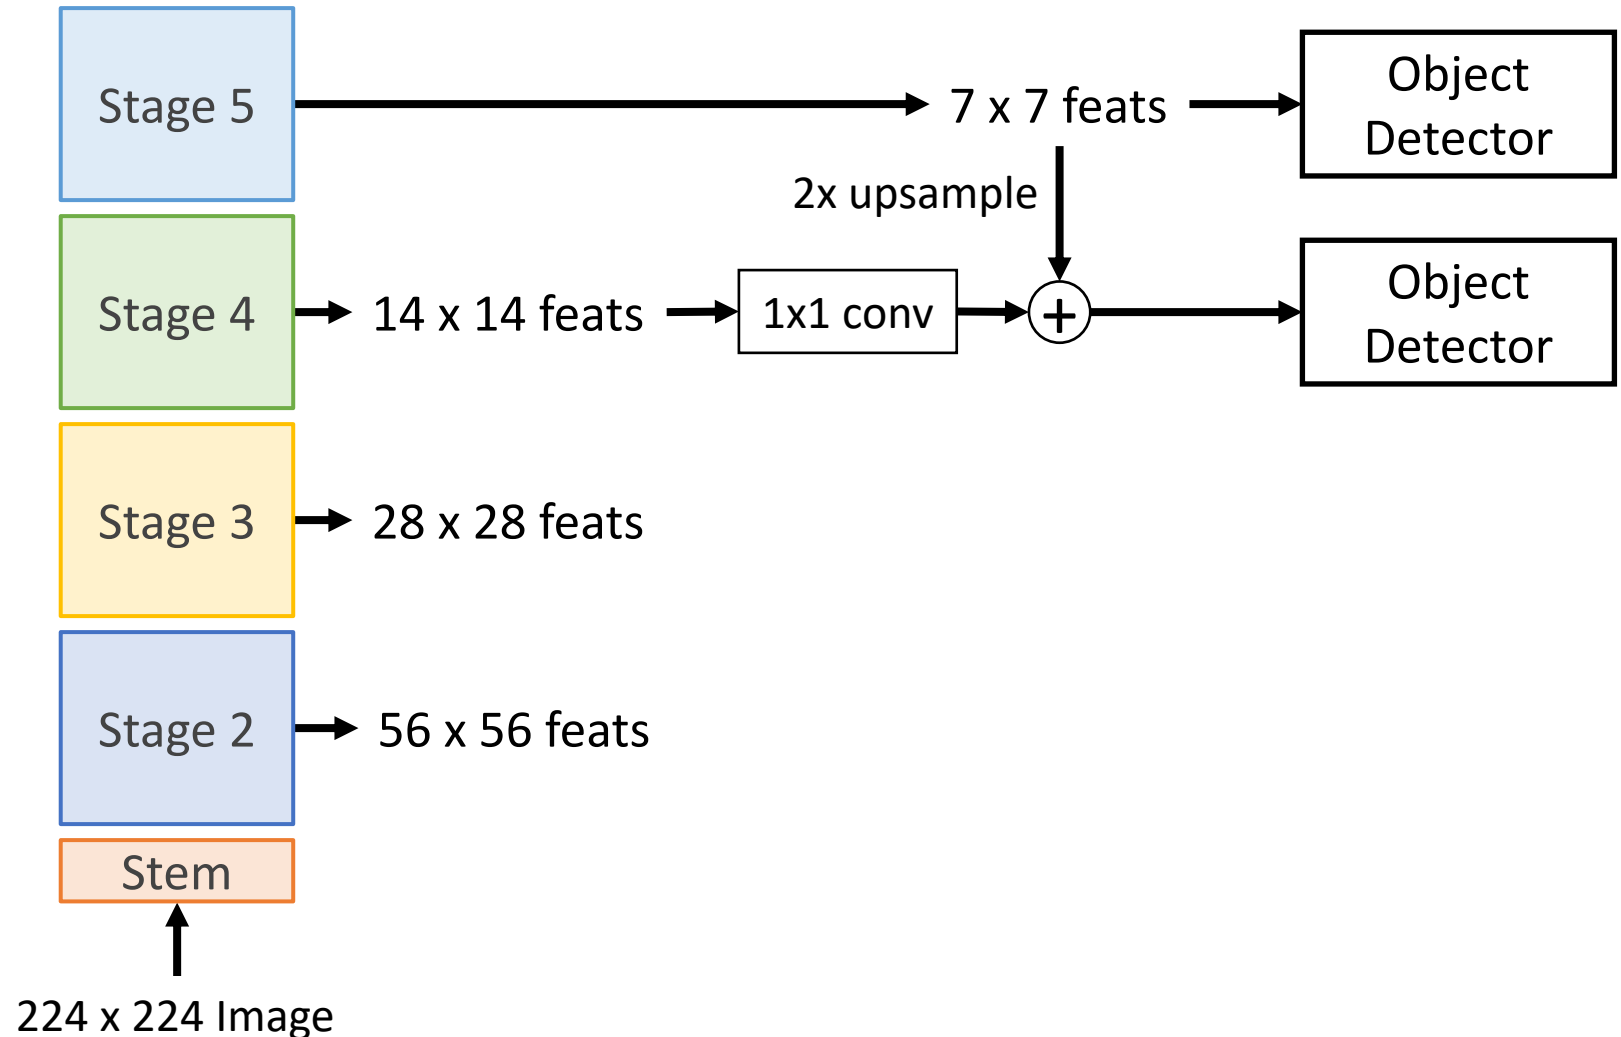
Add *top down* connections that feed information from high level features back down to lower level features



Lin et al, "Feature Pyramid Networks for Object Detection", ICCV 2017

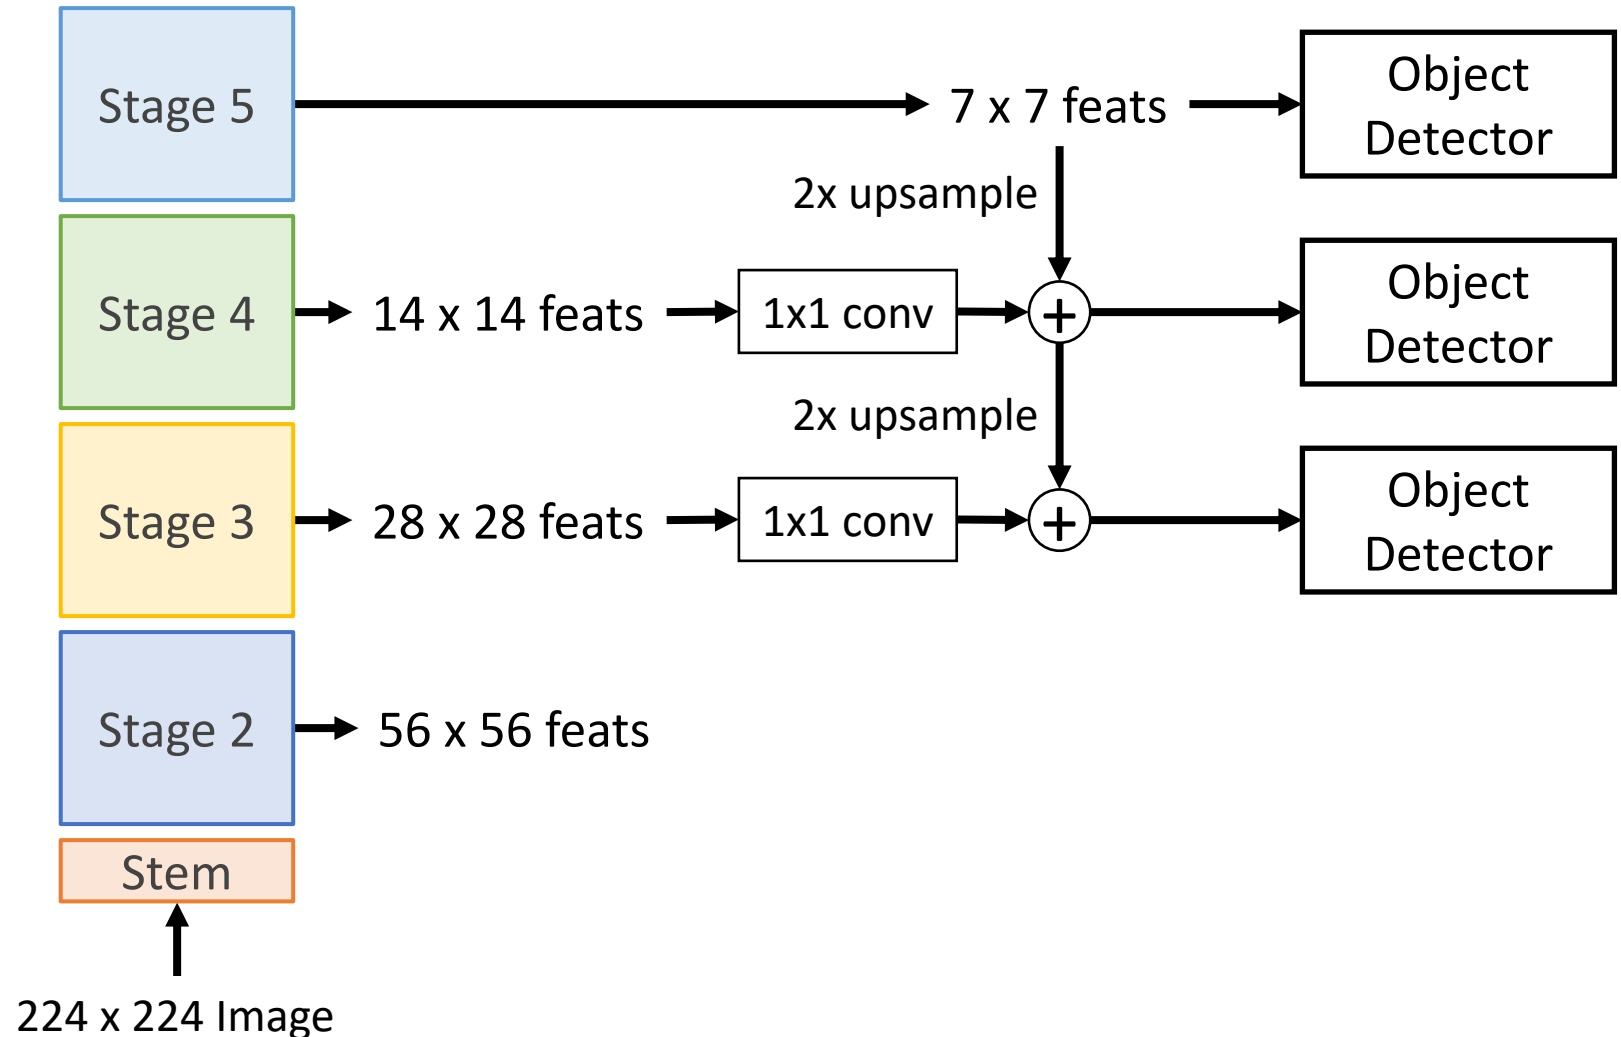
Dealing with Scale: Feature Pyramid Network

Add *top down* connections that feed information from high level features back down to lower level features



Dealing with Scale: Feature Pyramid Network

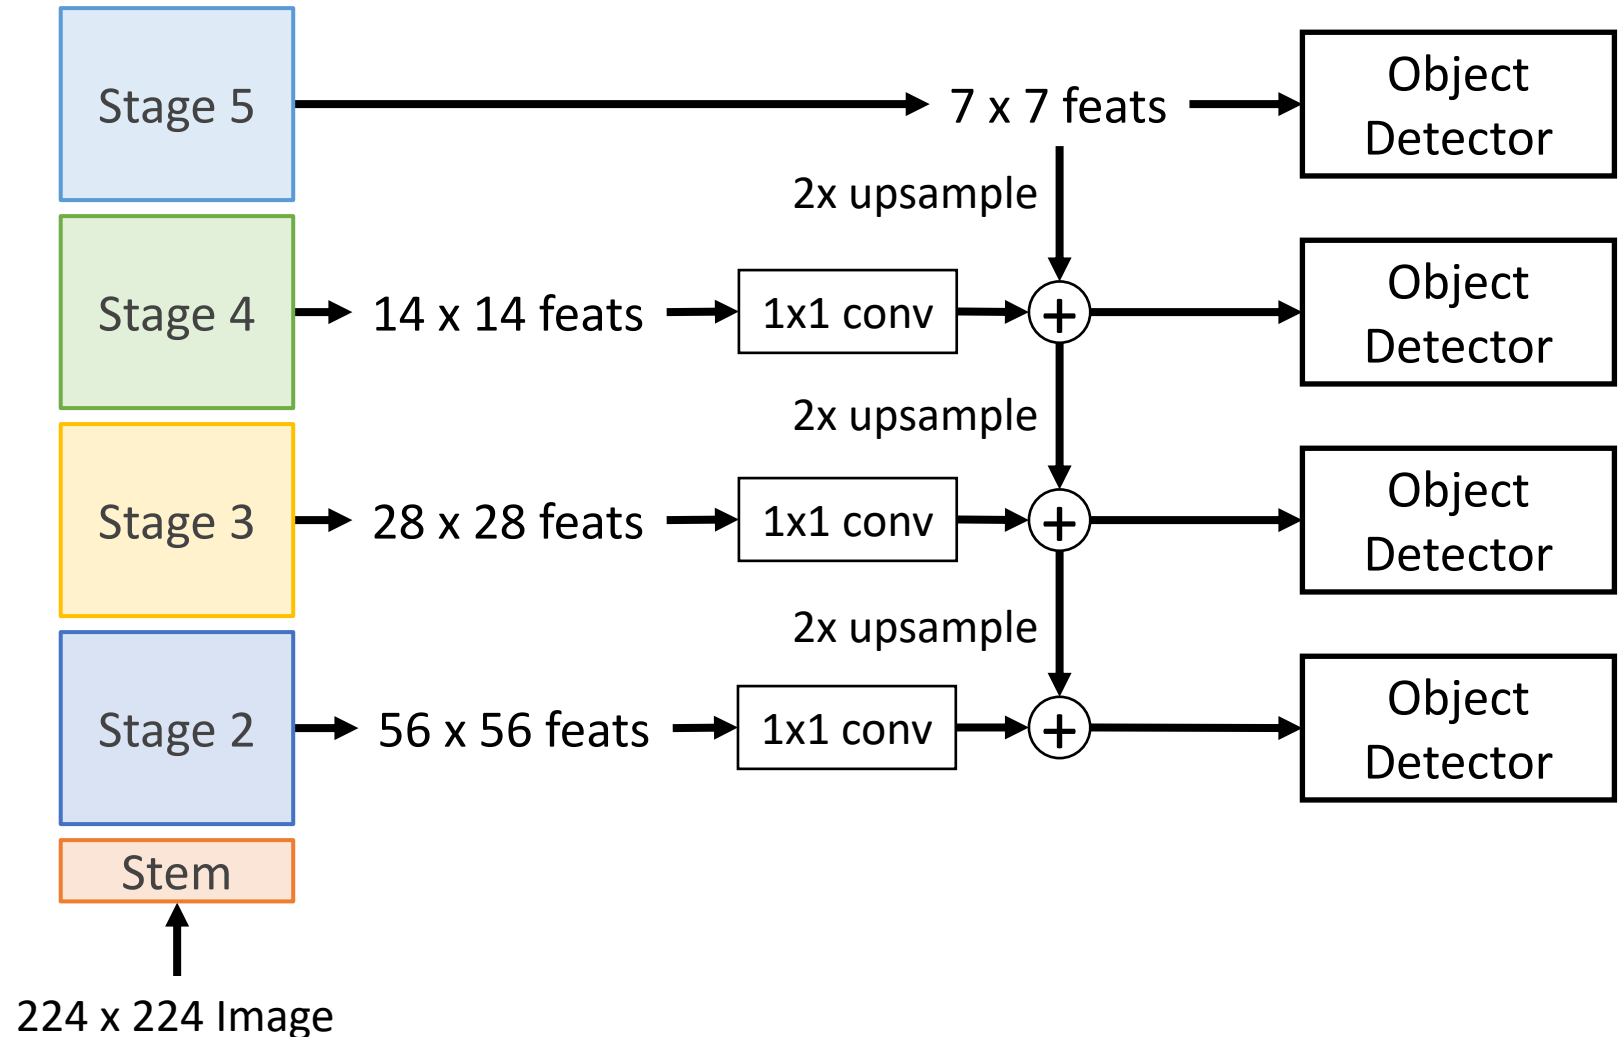
Add top down connections that feed information from high level features back down to lower level features



Lin et al, "Feature Pyramid Networks for Object Detection", ICCV 2017

Dealing with Scale: Feature Pyramid Network


Add *top down* connections that feed information from high level features back down to lower level features



Dealing with Scale: Feature Pyramid Network

Add *top down* connections that feed information from high level features back down to lower level features

Efficient multiscale features where all levels benefit from the whole backbone! Widely used in practice




Lin et al, "Feature Pyramid Networks for Object Detection", ICCV 2017

Dealing with Scale: Feature Pyramid Network

Add *top down* connections that feed information from high level features back down to lower level features

Efficient multiscale features where all levels benefit from the whole backbone! Widely used in practice



Lin et al, "Feature Pyramid Networks for Object Detection", ICCV 2017

Faster R-CNN: Learnable Region Proposals

Question: Do we really need the second stage?

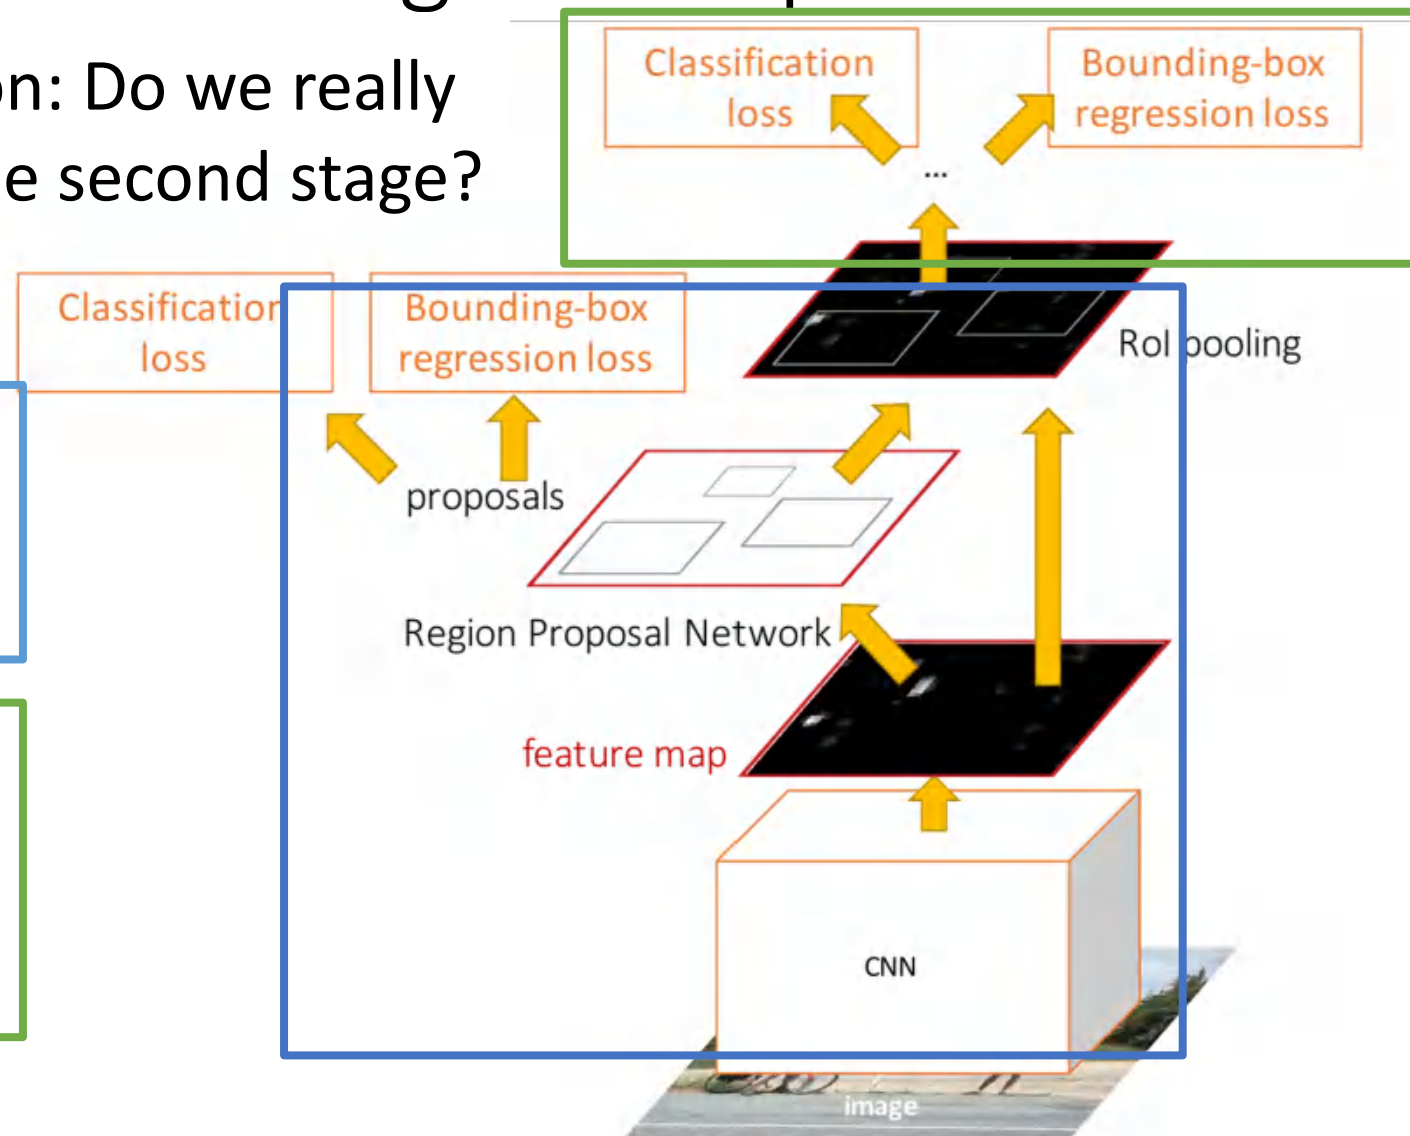
Faster R-CNN is a
Two-stage object detector

First stage: Run once per image

- Backbone network
- Region proposal network


Second stage: Run once per region

- Crop features: RoI pool / align
- Predict object class
- Prediction bbox offset



Single-Stage Detectors: RetinaNet

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input




Image features
(e.g. 512 x 5 x 6)

Similar to RPN – but rather than classify anchors as object/no object, directly predict object category (among C categories) or background

Conv


Anchor
→ classification
 $2K \times (C+1) \times 5 \times 6$

Anchor
→ transforms
 $4K \times 5 \times 6$

Single-Stage Detectors: RetinaNet

Problem: class imbalance –
many more background
anchors vs non-background

Run backbone CNN to get
features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds
to a point in the input




Image features
(e.g. 512 x 5 x 6)




Anchor
classification
 $2K \times (C+1) \times 5 \times 6$

Anchor
transforms
 $4K \times 5 \times 6$

Single-Stage Detectors: RetinaNet

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input





Image features
(e.g. 512 x 5 x 6)

Problem: class imbalance – many more background anchors vs non-background

Solution: new loss function (Focal Loss); see paper

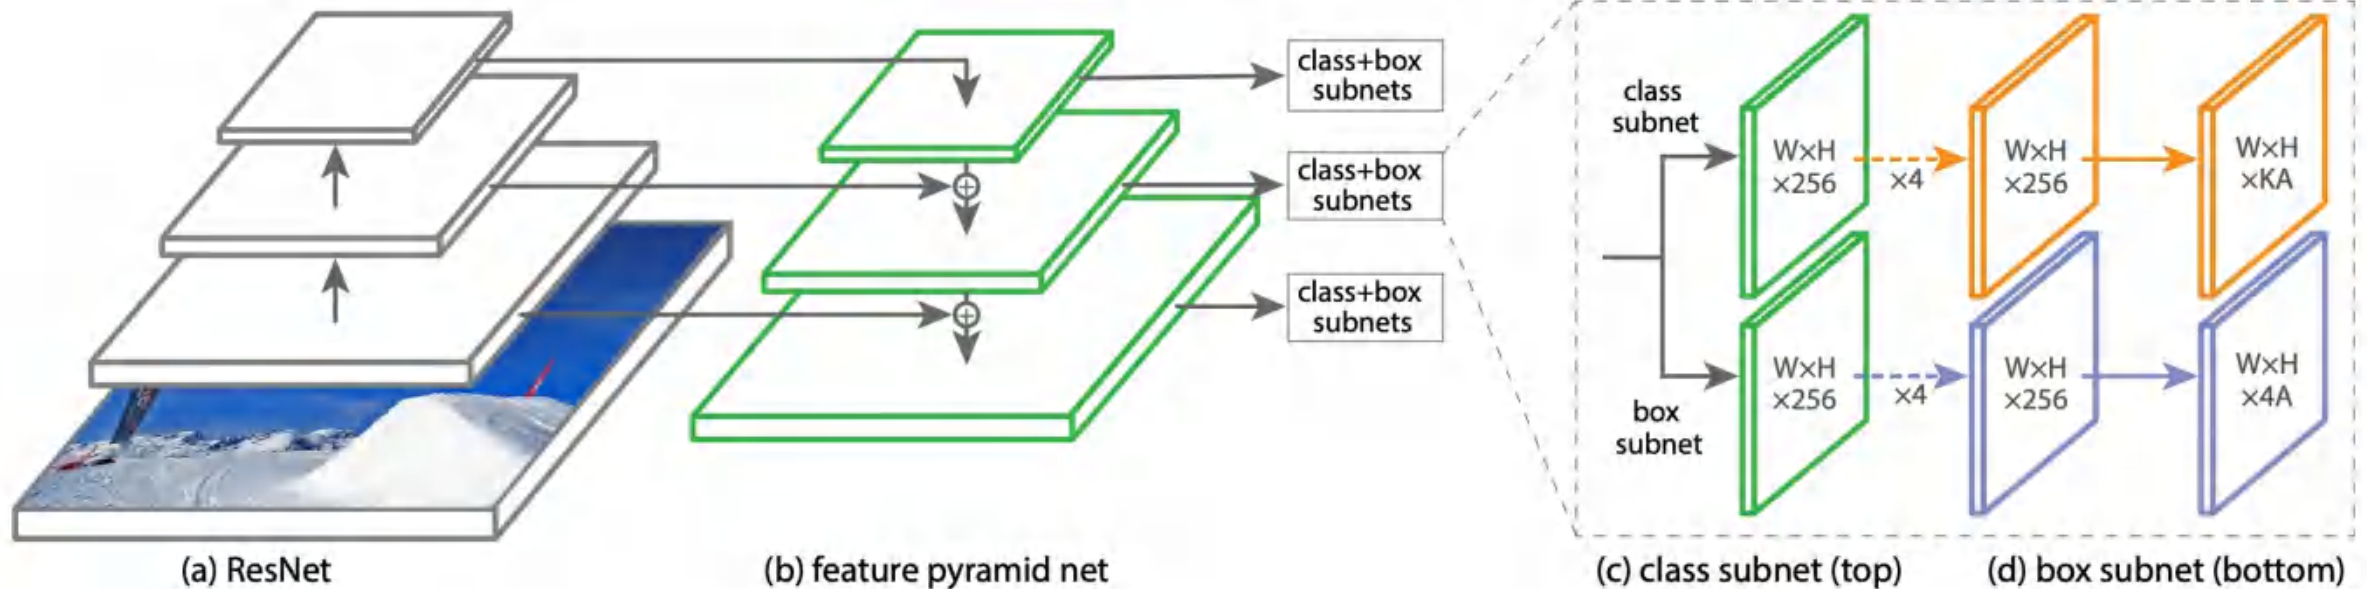


$$\text{CE}(p_t) = -\log(p_t)$$

$$\text{FL}(p_t) = -(1 - p_t)^\gamma \log(p_t)$$


Single-Stage Detectors: RetinaNet

In practice, RetinaNet also uses Feature Pyramid Network to handle multiscale




Single-Stage Detectors: RetinaNet

Single-Stage detectors can be much faster than two-stage detectors



Single-Stage Detectors: RetinaNet

Single-Stage detectors can be much faster than two-stage detectors




Faster R-CNN
with Feature
Pyramid
Network

Single-Stage Detectors: FCOS

“Anchor-free” detector

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input





Image features
(e.g. 512 x 5 x 6)

Single-Stage Detectors: FCOS

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input

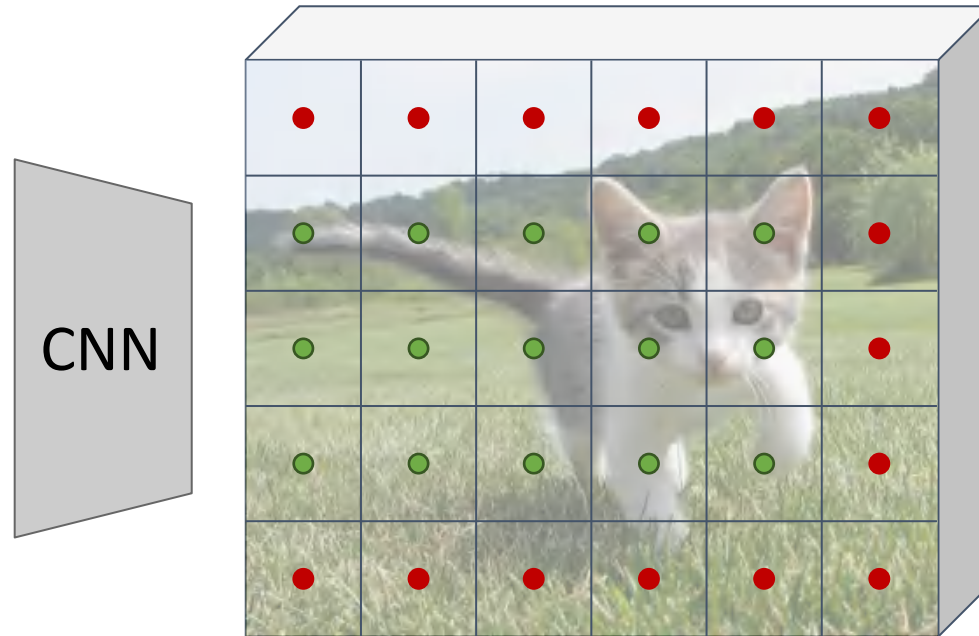


Image features
(e.g. 512 x 5 x 6)

“Anchor-free” detector

Classify points as positive if they fall into a GT box, or negative if they don't

Train independent per-category logistic regressors

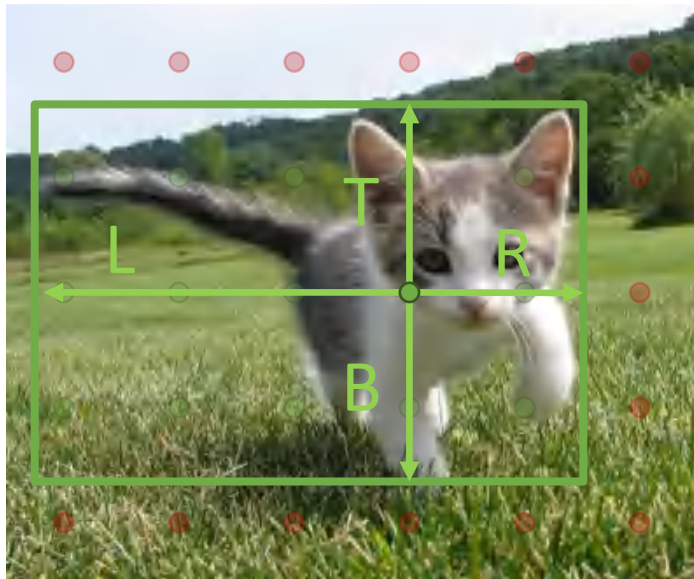
Class scores
 $C \times 5 \times 6$



Single-Stage Detectors: FCOS

“Anchor-free” detector

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input




Image features
(e.g. 512 x 5 x 6)

For positive points, also regress distance to left, right, top, and bottom of ground-truth box (with L2 loss)




Class scores
 $C \times 5 \times 6$
Box edges
 $4 \times 5 \times 6$

Single-Stage Detectors: FCOS

“Anchor-free” detector

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input




Image features
(e.g. 512 x 5 x 6)

For positive points, also regress distance to left, right, top, and bottom of ground-truth box (with L2 loss)




Class scores
 $C \times 5 \times 6$

Box edges
 $4 \times 5 \times 6$

Single-Stage Detectors: FCOS

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input




Image features
(e.g. 512 x 5 x 6)

“Anchor-free” detector

Finally, predict “centerness” for all positive points (using logistic regression loss)




- Class scores
C x 5 x 6
- Box edges
4 x 5 x 6
- Centerness
1 x 5 x 6

$$centerness = \sqrt{\frac{\min(L, R)}{\max(L, R)} \cdot \frac{\min(T, B)}{\max(T, B)}}$$

Ranges from 1 at box center to 0 at box edge

Single-Stage Detectors: FCOS

Run backbone CNN to get features aligned to input image



Input Image
(e.g. 3 x 640 x 480)

Each feature corresponds to a point in the input




Image features
(e.g. 512 x 5 x 6)

“Anchor-free” detector

Test-time: predicted
“confidence” for the box from
each point is product of its
class score and centerness



Class scores
→ C x 5 x 6

Box edges
→ 4 x 5 x 6

Centerness
→ 1 x 5 x 6

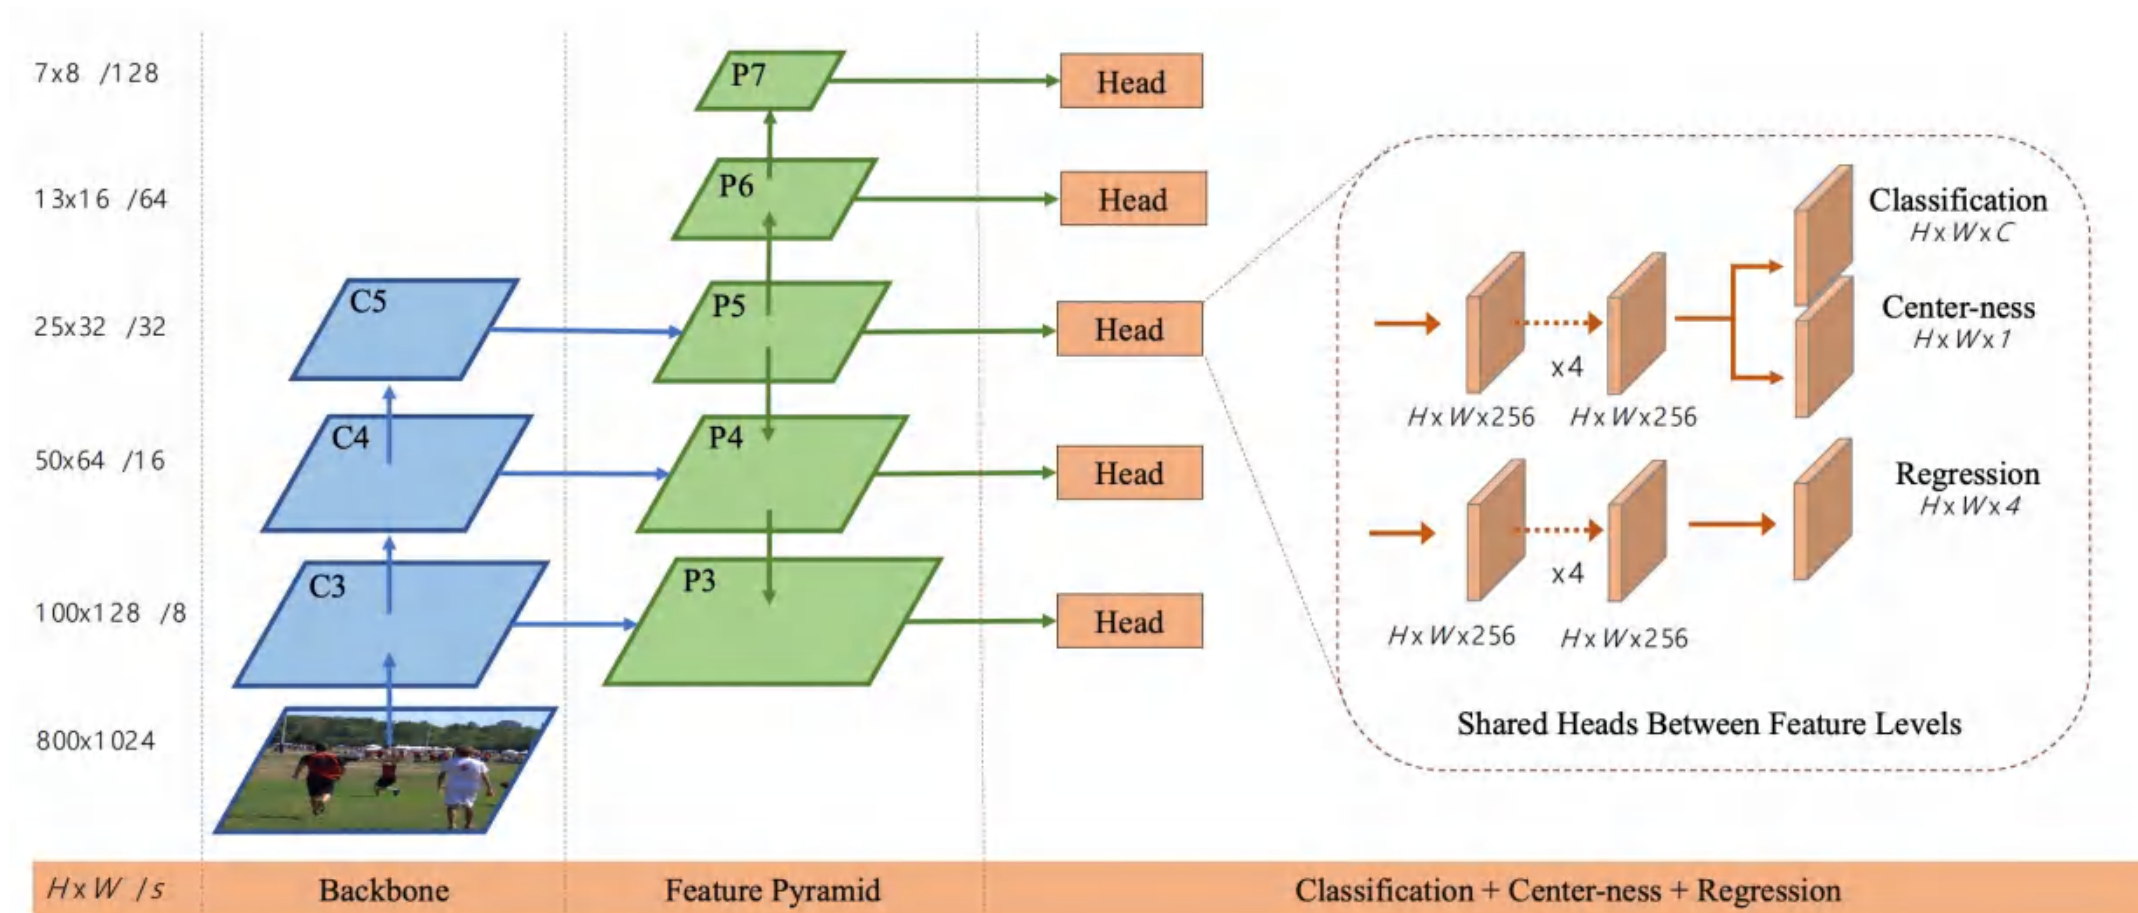
$$centerness = \sqrt{\frac{\min(L, R)}{\max(L, R)} \cdot \frac{\min(T, B)}{\max(T, B)}}$$

Ranges from 1 at box center to 0 at box edge

Single-Stage Detectors: FCOS

“Anchor-free” detector

FCOS also uses a Feature Pyramid Network with heads shared across stages




Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) =
area under Precision vs Recall Curve


Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)




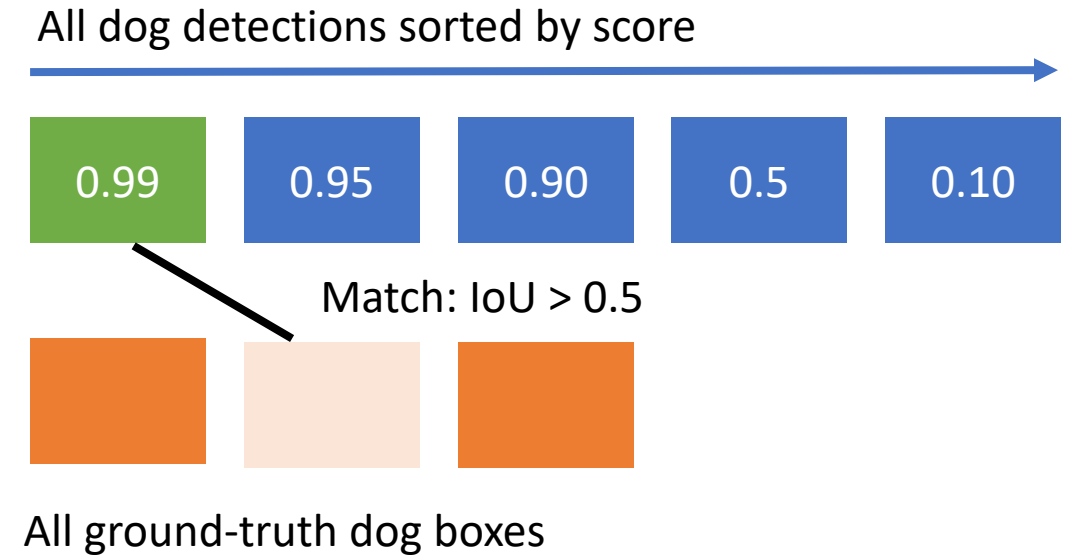
Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with $\text{IoU} > 0.5$, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative





Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with $\text{IoU} > 0.5$, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve





Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with $\text{IoU} > 0.5$, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve





Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve





Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve





Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve



Evaluating Object Detectors: Mean Average Precision (mAP)



1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve
 2. Average Precision (AP) = area under PR curve



Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with $\text{IoU} > 0.5$, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve
 2. Average Precision (AP) = area under PR curve

How to get AP = 1.0: Hit all GT boxes with $\text{IoU} > 0.5$, and have no “false positive” detections ranked above any “true positives”



Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve
 2. Average Precision (AP) = area under PR curve
3. Mean Average Precision (mAP) = average of AP for each category

Car AP = 0.65

Cat AP = 0.80

Dog AP = 0.86

mAP@0.5 = 0.77

Evaluating Object Detectors: Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve
 2. Average Precision (AP) = area under PR curve
3. Mean Average Precision (mAP) = average of AP for each category
4. For “COCO mAP”: Compute mAP@thresh for each IoU threshold (0.5, 0.55, 0.6, ..., 0.95) and take average

$$\text{mAP}@0.5 = 0.77$$

$$\text{mAP}@0.55 = 0.71$$

$$\text{mAP}@0.60 = 0.65$$

...

$$\text{mAP}@0.95 = 0.2$$

$$\text{COCO mAP} = 0.4$$

Summary: Beyond Image Classification

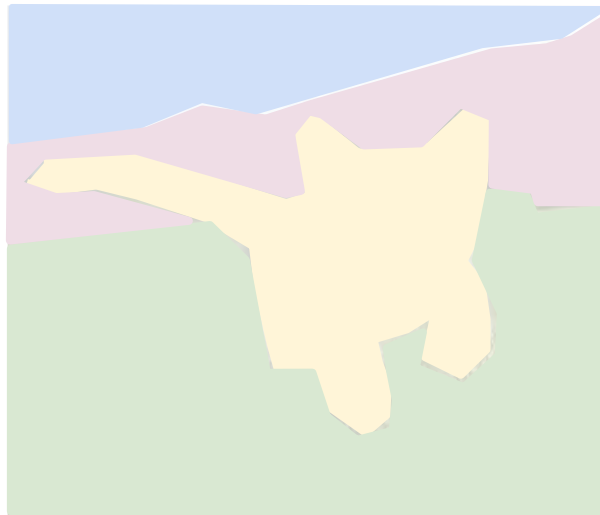
Classification



CAT

No spatial extent


Semantic Segmentation



GRASS, CAT, TREE,
SKY

No objects, just pixels

Object Detection



DOG, DOG, CAT

Multiple Objects

Instance Segmentation

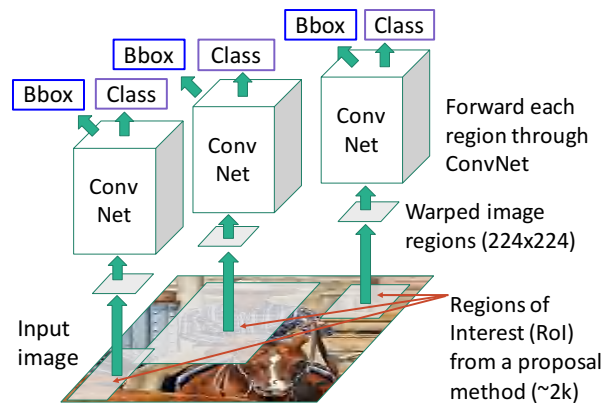


DOG, DOG, CAT

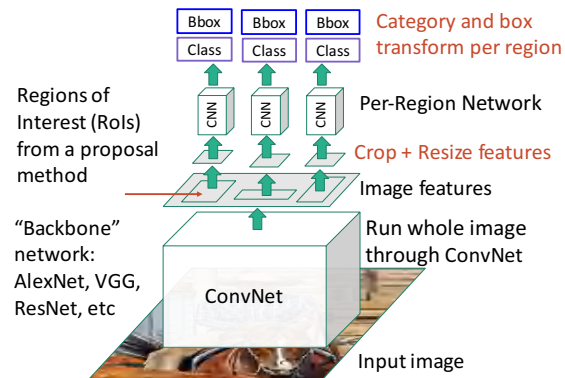
[This image](#) is [CC0 public domain](#)

Summary


“Slow” R-CNN: Run CNN independently for each region




Fast R-CNN: Apply differentiable cropping to shared image features



Faster R-CNN: Compute proposals with CNN



Single-Stage: Fully convolutional detector



With anchors: RetinaNet
Anchor-Free: FCOS

Next time:
Image and Instance Segmentation